
Auditing Correlated Failures in Frozen-Feature Pretrained-Encoder Pools for Medical Segmentation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Medical segmentation ensembles are often treated as providing complementary
2 failure coverage, yet many contemporary pools are built from repeated uses of
3 the same pretrained encoder bundles, input statistics, and training recipes. We
4 provide an E&D audit protocol for this setting: measure case-level co-failure
5 by per-case Dice-score correlation, separate same-bundle/recipe pairs from cross-
6 bundle/family pairs, and release the traces needed to recompute the audit. Across
7 five primary task rows (Kvasir polyp, ACDC LV, BraTS foreground, RIGA Cup,
8 RIGA Disc), a source pool of up to 11 pretrained encoder bundles $\times$ 4 decoder
9 seeds yields same-vs-cross Pearson gaps of $\Delta = 0.261\text{--}0.638$ after a Dice ≥ 0.30
10 functional floor; all paired case and hierarchical bootstrap intervals exclude zero.
11 The result remains positive under floor sweeps, calibration-selected pools, binary
12 failure-event thresholds, ICC(A,1), subject-level ACDC/BraTS checks, quality-
13 controlled pair regressions, and cross-fitted item-difficulty residualisation. The
14 practical implication is that extra decoder seeds or near-duplicate frozen-encoder
15 bundles should not be counted as independent failure evidence without an in-
16 domain co-failure audit. Secondary pixel-error, lift, same-architecture-family ViT
17 / different-pretraining, same-family cross-checkpoint, and same-backbone diag-
18 nostics locate where the dependence appears, while leaving causal separation of
19 architecture, pretraining, and recipe outside the supported claim. The artifact
20 includes per-case Dice traces, summary JSONs, metadata, hashes, scorer code, and
21 selected diagnostic JSONs allowed by upstream redistribution constraints.

22 1 Introduction

23 Ensemble-based safety arguments for medical AI rest on members making different mistakes. Deep
24 ensembles [Lakshminarayanan et al., 2017], nnU-Net 5-fold validation [Isensee et al., 2021], MC
25 dropout [Gal and Ghahramani, 2016], hyperparameter ensembles [Wenzel et al., 2020], and major-
26 ity voting are most useful when errors are complementary rather than perfectly shared [Tumer and
27 Ghosh, 1996]. Deep-ensemble theory motivates diversity but does not require literal statistical inde-
28 pendence; we likewise do not test literal independence. We audit the benchmark-level dependence
29 that determines whether extra pool members add new failure evidence: in frozen or lightly adapted
30 pretrained-encoder segmentation pipelines, do independently trained decoders actually make differ-
31 ent errors on the same cases?

32 The medical-imaging supply chain now reuses a small set of foundation-model encoders and conven-
33 tional pretrained backbones: DINOv2, CLIP/BiomedCLIP, SAM/MedSAM, ConvNeXt, Efficient-
34 Net, ResNet, and domain-specific pretrainings such as RETFound. A pool of “different” segmenters
35 can therefore share backbone family, pretraining bundle, input statistics, and training recipe. This
36 creates an evaluation gap: aggregate Dice and uncertainty scores can show whether members are
37 accurate, but not whether the pool supplies independent failure coverage. The intuition that near-
38 duplicate pipelines should share errors is not new; the missing piece is a released, case-aligned audit

that quantifies how large the dependence is, tests whether it survives simple alternative explanations, and states which evaluation claims the evidence can and cannot support. We use **regime-specific monoculture** for this measured dependence pattern in a tested frozen/light-adaptation regime, not for every segmentation ensemble, prompt-conditioned MedSAM, full nnU-Net pipeline, or clinical outcome.

Our primary contribution is a reusable E&D audit protocol and released primary evidence bundle: a symmetric multi-seed same-bundle/recipe vs. cross-bundle/family Dice-correlation estimator with paired case and hierarchical bootstrap CIs, applied to five primary task rows from an up-to-11-bundle $\times$ 4-seed source pool. The protocol asks a concrete reporting question: before a pool is credited with complementary failure coverage, have same-bundle/recipe and cross-bundle/family co-failure correlations been measured on aligned cases? We also release quality-controlled pair-regression, same-architecture-family ViT / different-pretraining, same-family cross-checkpoint, cross-fitted item-difficulty residualisation, pixel/lift, M_{eff} , same-backbone, adaptation, estimator-sensitivity, distribution-shift, and held-out probes that locate the dependence pattern without turning those probes into causal decomposition or standalone pool-selection rules. The appendix is organized by evidence role so that primary claims, diagnostics, robustness checks, and scope boundaries remain separated.

2 Related work

Relation to prior work and scope of novelty. Our risk framing follows algorithmic-monoculture and foundation-model ecosystem work [Kleinberg and Raghavan, 2021, Bommasani et al., 2021, 2022, Toups et al., 2023]. The correlation gap we measure has modelling antecedents: trial-by-trial error consistency asks whether systems fail on the same inputs [Geirhos et al., 2020], error correlations arise when deployed models share algorithms, datasets, or foundation models [Li et al., 2025], shared training trajectories and weight-space basins can constrain checkpoint diversity [Fort et al., 2019, Wortsman et al., 2022], transfer/representation work studies what is reused across pretrained models [Neyshabur et al., 2020, Gontijo-Lopes et al., 2022], pretrained-model ensemble work treats pretraining as a possible diversity source [Mustafa et al., 2020], same-architecture models can align strongly across initializations [Kornblith et al., 2019], and predictive-diversity pathologies can limit deep-ensemble gains [Abe et al., 2024]. LLM monoculture audits [Jiang et al., 2025, Kim et al., 2025a, Gorecki and Hardt, 2025, Goel et al., 2025] and null-model critiques [Jo et al., 2026] motivate our estimator choices. Pathology-FM and precision-oncology studies provide adjacent evidence about center robustness, benchmark behavior, and complementary domain-specialised features [de Jong et al., 2025, Campanella et al., 2025, Luo et al., 2025, Neidlinger et al., 2025]. Our contribution is not the abstract claim that correlated errors exist; rather, we provide a released medical-segmentation audit of frozen-feature pretrained-encoder pools that distinguishes primary from diagnostic evidence and reports where same-checkpoint/recipe, same-architecture-family, and broader cross-family correlations appear or contract under a shared protocol.

Medical segmentation, pretrained/foundation-model benchmarks, and ensemble baselines. nnU-Net, deep ensembles, MC dropout, adjacent uncertainty or parameter-efficient ensembles, and hyperparameter ensembles are usually evaluated through Dice, calibration, or uncertainty quality [Isensee et al., 2021, 2024, Lakshminarayanan et al., 2017, Gal and Ghahramani, 2016, Mehrtash et al., 2020, Jungo and Reyes, 2019, Mühlematter et al., 2026, Gu et al., 2024]. General ensemble-diversity theory [Kuncheva and Whitaker, 2003, Ganaie et al., 2022, Wood et al., 2023] and medical-segmentation diversity or failure-detection studies [Georgescu et al., 2023, Zenk et al., 2025] are closest methodologically, but they study adjacent diversity settings rather than the released frozen-encoder per-case co-failure audit studied here. Recent medical-segmentation benchmarks and robustness-specification work evaluate or discuss generic-vision probes, SAM medical-image evaluations, standardized segmentation datasets, promptable SAM adaptations, universal 3D/medical segmenters, task-specific robustness requirements, transferability estimation, and VFM segmentation benchmarking [Baharoon et al., 2023, Huang et al., 2024, Kuş and Aydin, 2024, Kirillov et al., 2023, Ma et al., 2024, Cheng et al., 2023, Archit et al., 2025, Zhu et al., 2024, Ma et al., 2025, Gu et al., 2025, He et al., 2025, Kim et al., 2025b, Xian et al., 2025, Yang et al., 2023, Kerssies et al., 2024]. Our protocol is not a VFM performance benchmark or full fine-tuning comparison: it deliberately fixes frozen/light-adaptation probes to audit co-failure, while generic and medical FMs

plus conventional pretrained backbones define the candidate universe and adjacent model classes. Promptable/universal models are adjacent to, but outside, our frozen-vision-encoder audit.

3 Formal setup and benchmark estimands

Pool and estimands. Let $\mathcal{F}=\{F_1, \dots, F_K\}$ be the operational pretrained-encoder bundles used by the Table 1 scorer. In the primary 11-bundle source pool, a bundle is the concrete released encoder / pretraining / input-statistic wrapper used for frozen-feature extraction; same-bundle decoder seeds form the main same group, and DINOv2-B $\times$ DINOv2-S is the only recurring cross-checkpoint member of that group because both variants share the same upstream DINOv2 release and pretraining recipe. Other broad-family groupings (for example OpenAI CLIP sizes, ResNet sizes, and expanded DINOv2/ConvNeXt variants) are not folded into the Table 1 primary estimator; they are separated in the released same-family cross-checkpoint diagnostic in Appendix A31. BiomedCLIP (ViT-B/16, biomedical image-text pretraining) is treated as a different bundle from DINOv2 despite being a ViT. A model is a tuple (F_k, σ) (bundle, decoder seed) producing per-case Dice $D_{(k,\sigma)}(i, t)$. Appendix A30 and Appendix A31 report released diagnostics that separate Table-1 product bundles, same-architecture-family ViT / different-pretraining, and broader same-family cross-checkpoint readouts; these clarify that a “cross” column can still contain broad architectural structure. For a pair of models $a \neq b$ we define $r_{a,b}(t) = \text{Pearson}_i[D_a(i, t), D_b(i, t)]$ (and a McGraw-Wong ICC(A,1) variant $\iota_{a,b}$ [McGraw and Wong, 1996], chosen to penalise per-bundle Dice offsets that Pearson absorbs). The Table-1 same/cross means at task t are the expectations of $r_{a,b}$ over pairs in the same group (same bundle, plus DINOv2-B/S) and the remaining cross group under this scorer, and $\Delta(t) = \bar{r}_{\text{same}}(t) - \bar{r}_{\text{cross}}(t)$ is the *monoculture gap*. Secondary estimands: pixel-level error-mask Dice (computed the same way over binarised error masks) and joint-failure *lift* over independence $L_{\text{mono}}(\tau) = \Pr[\text{all fail} \mid \text{mono pool}] / \prod_m \Pr[\text{fail}_m \mid \tau]$.

Symmetric multi-seed averaging. For each bundle pair we enumerate all $|\Sigma|^2=16$ seed combinations (4 per model) rather than seed-matched pairs, removing the one-seed-per-bundle noise that inflates ad-hoc gap estimates. Same-bundle averaging uses $\binom{|\Sigma|}{2}=6$ same-bundle seed pairs per F_k plus the full $|\Sigma|^2$ on DINOv2-B $\times$ DINOv2-S; the cross group uses all remaining Table-1 bundle pairs.

CI and scope. A 3-level nested hierarchical bootstrap (Table-1 groups $\rightarrow$ seeds $\rightarrow$ cases, $B=2000$) yields 95% CIs as the 2.5th/97.5th percentiles of bootstrap-sampled $\bar{r}_{\text{same}}, \bar{r}_{\text{cross}}, \Delta$. Zero-exclusion is reported as a sign-check rather than as population-level proof over all possible encoder bundles or families. The estimands assume a shared-decoder, shared-task pool and measure *observed* correlation within such pools; a released CNN random-init diagnostic (Appendix A33.2) shows positive same-vs-cross structure without pretrained weights, so we report the pretrained-bundle gap descriptively rather than as a causal isolation of pretrained-weight effects. With a finite panel and task-specific post-floor pools, our hierarchical bootstrap has few-cluster bias in the Cameron et al. [2008] sense; we therefore also report paired case-only bootstrap CIs ($B=2000$) throughout and take agreement across the two as the informative CI.

Primary-group structural caveat. The primary rows retain task-specific functional pools of 11/10/10/9/11 bundles after the Dice floor. In these pools, the Table 1 “same” column is dominated by same-bundle, different-seed decoder-head pairs, with DINOv2-B $\times$ DINOv2-S providing the only recurring same-family cross-checkpoint pair when both variants are functional. Thus the high same-group $\bar{r}$ values in Table 1 (0.834–0.959) are best read as same-frozen-checkpoint/recipe-dominated dependence readouts, not a broad sample of many same-family cross-checkpoint comparisons. Section 5.1 and Appendix A30 report a 5-variant same-architecture-family ViT / different-pretraining diagnostic (DINOv2-B, BiomedCLIP, MAE-B, DeiT-B, CLIP-B) that measures cross-pretraining correlation under a shared broad architecture family and is the closest “within-architecture-family, cross-pretraining” control in this benchmark.

4 Experimental design

Datasets & targets. The five primary rows use predefined splits recorded in each JSON’s `test_ids`: Kvasir-SEG polyp [Jha et al., 2020] (loader-defined 880/120 image split; the released IDs are image-level file identifiers, with no patient/video identifiers available for patient-

level duplicate clustering), RIGA+ cup/disc [Almazroa et al., 2018, Almazroa, 2018] (Bin-Rushed+MESSIDOR train, Magrabia (Magrabi Eye Center) source test, $n=95$; the released `test_ids` enumerate the local loader-retained Magrabia source files, and source/subfolder labels are retained only for deterministic alignment), ACDC LV-positive ED/ES slices [Bernard et al., 2018] from public patients 001–100 (fold-0, 20 patients, $n=361$ slices), and a BraTS 2024 GLI 2D FLAIR non-background foreground derivative [BraTS 2024 Organizers, 2024, de Verdier et al., 2024] (release key `brats_wt`) that keeps up to the top-10 axial slices per held-out subject by foreground area, with a minimum 64 positive pixels and a subject-level 80/20 split using NumPy seed 42 ($n=3,228$ test slices from 324 subjects). The BraTS count is 12 below the 324×10 maximum because three held-out subjects have fewer than ten eligible foreground-positive slices after the 64-pixel filter (3, 6, and 9 slices). Because this target is `seg>0` and includes the BraTS-GLI resection-cavity label, it is not the official 3D BraTS WT challenge target.

Encoder bundles. *Generic pretrained-encoder bundles* (11, mixing foundation-model encoders and conventional pretrained backbones): DINOv2-B/S, BiomedCLIP, CLIP ViT-L/14, CLIP ViT-B/16, MAE ViT-B/16, DeiT ViT-B/16, ConvNeXt-T, EfficientNet-B0, ResNet-50/18. Here “CLIP ViT-L/14” denotes the bundle’s preserved OpenCLIP ViT-L-14/openai dense-token loader, not OpenAI `clip.load` inference parity; strict QuickGELU parity would be a separate sensitivity loader. *Medical-domain encoder probe:* MedSAM vision-encoder-only (SAM ViT-B; prompt encoder and mask decoder discarded) is reported as a secondary boundary probe on RIGA/Kvasir/ACDC. RETFound support is documented as a gated-checkpoint path, but RETFound cells are not reported because the checkpoint/authentication was unavailable for the reported artifact run.

Protocol. The primary protocol resizes images and masks to 224×224 , applies encoder-specific upstream normalisation (ImageNet-style statistics for DINOv2/MAE/DeiT/torchvision CNNs and OpenAI CLIP statistics for CLIP/BiomedCLIP), freezes the encoder, and trains a 1×1 -conv decoder with Adam 10^{-3} for 30 epochs, batch 16, and BCE loss on the binary primary rows. Table 1 starts from the 11-bundle generic source pool where cells are available and applies the estimator’s $\text{Dice} \geq 0.30$ functional floor to remove non-segmenting, near-constant traces: Kvasir retains 11 bundles, ACDC 10, BraTS foreground 10, RIGA Cup 9, and RIGA Disc 11. This floor is a test-trace-based analytic convention chosen for functional-pool definition, not a training filter, ingestion filter, or pre-registered hypothesis threshold; Table 1 is therefore conditional on the post-floor analytic pool, and Appendix A5 reports no-floor and stricter-floor sensitivity with retained pool counts. The MedSAM sensitivity is a secondary encoder-only boundary probe, not part of the primary pool or a prompt-conditioned MedSAM evaluation.

5 Main result: same-frozen-checkpoint pools show higher co-failure correlation

Table 1: **Same-bundle/recipe vs. Table-1 cross-bundle/family per-case Dice correlation across the five primary medical segmentation task rows.** Values are Pearson $\bar{r}$ outputs from `results/_merged/paper_table.json`. Symmetric multi-seed averaging uses completed decoder-seed combinations per bundle pair. The same column is dominated by same-bundle seed pairs plus the DINOv2-B/S pair; broader same-family cross-checkpoint groupings are diagnostic rows in Appendix A31. The Bundles column reports the functional bundle count after the $\text{Dice} \geq 0.30$ floor, not the nominal source pool size.

Task row	Unit N	Bundles	ρ_{same}	ρ_{cross}	Δ	Case CI [§]	Hier CI [¶]
Kvasir polyp	image 120	11	0.958	0.697	0.261	[0.211,0.320]	[0.174,0.378]
ACDC LV	slice 361 [‡]	10	0.959	0.639	0.321	[0.289,0.355]	[0.211,0.435]
BraTS foreground [†]	slice 3,228*	10	0.938	0.623	0.315	[0.302,0.329]	[0.224,0.425]
RIGA Cup	fundus 95	9	0.834	0.361	0.473	[0.395,0.556]	[0.280,0.780]
RIGA Disc	fundus 95	11	0.870	0.232	0.638	[0.569,0.687]	[0.475,0.831]

[§] Case-level paired bootstrap from `gap_ci95_case` ($B=2,000$), conditional on the functional pool shown. [¶] Hierarchical bootstrap over Table-1 grouping and seed clusters from `gap_ci95_hier`; it is the more conservative finite-pool interval, and all hierarchical CIs exclude zero. $\Delta = \rho_{\text{same}} - \rho_{\text{cross}}$ (Pearson). [‡]20 ACDC patients. ^{*}324 BraTS subjects. [†]BraTS uses the completed 11-bundle source pool for the 2D FLAIR `seg>0` derivative; ResNet-18 is below the Dice floor, leaving 10 post-floor bundles.




Figure 1: **Spatial and lift diagnostics that motivate the case-level audit.** Table 1 supplies the primary paired per-case Dice evidence; this figure shows the same dependence problem in complementary spatial and threshold views. (a) Single-case illustrations (RIGA Cup, upper; BraTS foreground, lower) compare four DINOv2-B decoder seeds with four diverse backbones; red=false positives, blue=false negatives, green=GT boundary. (b) Pixel-level error-mask Dice on secondary, non-case-identical aggregate rows (RIGA Cup/Disc $N=224$, ACDC LV $N=90$, Kvasir $N=300$). (c) Selected matched-size joint-failure lift ratios; sparse rows are descriptive. Panels (b–c) are released aggregate diagnostics rather than the primary per-case-Dice rows in Table 1; the submission artifact distributes summary statistics rather than raw masks or full prediction caches.

181 Across four imaging domains and five task rows, same-group per-case Dice correlation is high
 182 (0.834–0.959) while the Table-1 cross group is lower (0.232–0.697). Every case-level CI excludes
 183 zero, and the more conservative hierarchical bootstrap remains positive. Table 2 summarizes four
 184 main-text sensitivity checks: a calibration-selected functional pool evaluated on disjoint cases, a
 185 binary failure-event readout at the empirical P_{33} threshold, quality-controlled pair regression, and
 186 distinct-checkpoint same-family decomposition.

187 **Estimator robustness.** A released CPU-only sensitivity file
 188 (results/_merged/diagnostics/audit_sensitivity.json) recomputes the primary
 189 Pearson gap under a functional-floor sweep, an ICC(A,1) alternative, leave-one-bundle/task-
 190 out summaries, and a random family-label permutation stress test. The split-half calibration
 191 diagnostic selects the functional pool on held-out calibration units and evaluates on disjoint
 192 units; the five evaluation gaps remain positive (0.250–0.630). A binary failure-event diagnostic
 193 (failure_event_summary.json) reduces the concern that the result is only a continuous-
 194 Dice-trace correlation: at the empirical P_{33} failure threshold, same-cross event-correlation gaps
 195 are 0.329–0.578, and all case-bootstrap intervals exclude zero. A third released diagnostic
 196 (quality_controlled_pairs.json) fits the pair-level OLS check in Table 2; the same-group

Table 2: **Main robustness checks for the primary rows.** Cal.-split Δ selects functional bundles on a deterministic calibration subset and evaluates the same/cross gap on disjoint cases. P_{33} fail ϕ gap is the same/cross gap in binary failure-event correlation at the empirical 33rd percentile of post-floor model-case Dice. Quality-control β_{same} is the same-group coefficient in a pair-level OLS diagnostic controlling for marginal quality and variance. Expanded-grid xckpt-cross is the same-family cross-checkpoint diagnostic from Appendix A31, not the Table 1 analytic-pool cross-family stratum.

Task	Cal.-split Δ	P_{33} fail ϕ gap	QC β_{same}	Xckpt-cross
Kvasir polyp	0.250	0.329	0.169	0.080
ACDC LV	0.350	0.386	0.208	0.134
BraTS foreground	0.307	0.413	0.224	0.093
RIGA Cup	0.478	0.453	0.372	0.267
RIGA Disc	0.630	0.578	0.582	0.227

coefficient stays positive on all five rows (0.169–0.582). At the paper’s Dice $\geq$ 0.30 floor the five-task mean gap is 0.401 (leave-one-task-out means 0.342–0.437); every task remains positive at the stricter Dice $\geq$ 0.40 floor where the pool is non-empty, ICC gaps are larger than Pearson gaps, and in 1000 random-label permutations the observed row gap is exceeded with rank fraction 0.001 (Appendix A5). Because family labels are not a clean exchangeable sampling design, we treat these as estimator-stability evidence, not as independent experiments or population-level p -values.

Spatial and lift readouts. Figure 1 explains why a case-level audit is needed: visually similar failures and joint-threshold failures can persist even when members have separate decoder seeds. In the released aggregate diagnostics, same-family pairwise error-mask Dice is 0.80–0.87 versus 0.32–0.40 cross-family on RIGA Cup/Disc ($N=224$), ACDC LV ($N=90$), and Kvasir ($N=300$) rows (Fig. 1b, Appendix A3, and Appendix A8), and matched-size diverse pools have lower lift than same-bundle pools in selected joint-failure diagnostics (Fig. 1c; Appendix A9). These readouts are released only as aggregate JSONs, not as raw masks or full prediction caches, so sparse rare-event rows are descriptive rather than headline statistics.

Item-difficulty diagnostic. Following the warning that monoculture claims depend on the null model [Jo et al., 2026], we residualise per-case Dice against held-out item-difficulty estimates. The naive all-bundle difficulty regressor leaks the response and inflates Δ . A released all-row cross-fitted diagnostic estimates each trace’s item-difficulty regressor from traces outside that trace’s broad upstream family, then recomputes the Table-1 same/cross rule on residuals; residualised gaps remain positive on all five rows ($\Delta = 0.830$ –0.976, all case-bootstrap intervals exclude zero; Appendix A27). This is a robustness diagnostic, not a symmetric causal adjustment: under this particular held-out residualisation, item difficulty does not absorb the measured separation.

5.1 Limited ViT diagnostic: cross-pretraining correlations remain intermediate

A natural Kornblith-type concern [Kornblith et al., 2019] is that two ViTs produce similar representations *because they share architecture family*, not because they share pretraining. We partially control broad architecture family with a 5-variant ViT pool—DINOv2-B (SSL on LVD-142M), Biomed-CLIP (CLIP on biomed image-text), MAE-B (ImageNet MAE), DeiT-B (ImageNet supervised), OpenAI CLIP-B—trained with the identical frozen 1×1 -conv probe (Table 3). The released aggregate covers RIGA Cup and ACDC LV. In the functional-filtered readouts, same-architecture-family ViT / different-pretraining pairs sit below the same-bundle decoder-seed floor but above the primary cross-family reference. In this limited diagnostic, that pattern is consistent with both architecture family and pretraining bundle affecting correlation, but it does not identify their separate causal contributions.

Because the primary same column is dominated by same-bundle decoder-seed pairs, we also release a CPU-only same-family cross-checkpoint decomposition over the five primary rows (Appendix A31). It separates the same-bundle seed floor, distinct checkpoints within broad upstream families (DINOv2-S/B, OpenAI CLIP-B/L, ResNet-18/50 when functional), and cross-family pairs. The distinct-checkpoint same-family minus cross-family gap remains positive on all five rows (0.080–0.267), while staying well below the same-bundle seed floor. Thus the primary effect is strongest for shared checkpoints and recipes, but the released traces do not collapse to a same-decoder-seed artifact.

Table 3: **Same-architecture-family ViT / different-pretraining diagnostic.** Same-architecture-family ViT cross-pretraining lies below the same-bundle seed floor; after functional filtering where needed, it remains above the primary cross-family reference. This is consistent with architecture family and pretraining bundle leaving measurable signal in this grid, not a causal decomposition.

Dataset	ViT subset	Func. ViTs	Seed floor	ViT cross-pretrain	Primary cross ref.	Seed-floor gap
RIGA Cup	all 5 ViTs	5/5	0.968	0.680	0.361	0.288
ACDC LV	CLIP-B dropped [†]	4/5	0.988	0.807	0.639	0.182

Values are from `results/_merged/diagnostics/r7_arch_confound_analysis.json` and `paper_table.json`.

[†]The all-5 ACDC same-architecture-family ViT cross-pretraining value is 0.631 because CLIP-B is nonfunctional on ACDC (mean Dice 0.205); the functional-filtered readout is 0.807. ViT×CNN subcells are kept as appendix diagnostics and are not used as main-table numeric evidence.

Table 4: **Released same-backbone diagnostics.** These rows ask whether common same-backbone or same-bundle levers reduce the measured dependence; they are diagnostic comparisons, not dose-matched causal effects.

Control	Released source	Main readout	Cautious interpretation
Head-capacity sweep (M3), RIGA Cup	<code>per_case_dice_m3</code> , App. A33.1	4 heads, 769–3.94M params; Dice 0.725–0.848; cross-head same-bundle $\bar{r}=0.690$ (range 0.510–0.896)	Head design changes accuracy and some pairwise r , but stays above primary RIGA cross-family $\bar{r}=0.361$.
UNet-skip decoder replication	<code>m4_unet_summary.json</code> , App. A16	20 cells; Dice 0.806–0.902; seed-pair $r=0.774$ –0.979	Stronger skip head improves accuracy but does not make same-backbone pools look cross-family.
MC dropout, B/RIGA Cup	DINOv2- <code>per_case_dice_m22</code> , App. A33.7	MC $\bar{r}=0.947, 0.925, 0.893$ for $p=0.1, 0.2, 0.4$	Weak decorrelation lever; largest tested drop is about 0.10.

6 Same-backbone controls and mechanism limits

(C1) Head architecture/capacity. The released head-capacity and UNet-skip decoder controls change accuracy and reduce some pairwise correlations, but their same-backbone correlations remain above the primary RIGA cross-family reference. Thus decoder capacity does not eliminate the measured gap in these released diagnostics.

(C2) Stochastic/adaptation levers. MC dropout is a weak same-bundle lever in the released RIGA Cup diagnostic. Bootstrap-bagging and M15 fold-reslicing are appendix diagnostics with different sampling surfaces; M15 is a single-seed split-stability check, not a within-fold vs cross-fold same-backbone estimator. Unbundled full-fine-tuning and LoRA mechanism probes are kept only as transparency notes and are not used as bundle-supported evidence.

(C3) Medical-domain encoder-only probe. A MedSAM image-encoder-only sensitivity is described in Appendix A10 and released as an aggregate diagnostic summary rather than primary evidence. The probe strips MedSAM’s prompt encoder and mask decoder, so it is not a fair evaluation of prompt-conditioned MedSAM as released. RETFound is documented only as a gated-checkpoint path; without the gated checkpoint/authentication, we do not report RETFound rows or use them to support the claims.

7 Evidence taxonomy and scope boundaries

Same-backbone perturbations are useful but incomplete in our measurements. Data bootstrap-bagging reduces within-pool $\bar{r}$ by only 0.05–0.12, the released M15 fold-reslicing diagnostic keeps cross-bundle $\bar{r}$ stable across held-out folds but has only one seed per fold, and the released MC-dropout diagnostic is a weak same-bundle lever (Appendices A14, A17, A33.7). Released and explicitly labelled appendix diagnostics suggest that shared pretrained weights are not the only source of dependence, while the same-architecture-family ViT and CNN random-init readouts indicate that pretraining, architecture, and shared recipe effects remain entangled in this benchmark rather than collapsing to a single identified mechanism (Appendices A33.2, A30).

The strongest caution is the full-pipeline boundary. The nnU-Net v2 2D/3D complements are seeded task-level runs using custom trainer subclasses and task-level splits; the RIGA rows are not case-

identical to the Magrabia-source held-out frozen-encoder row. They are descriptive scope checks rather than official nnU-Net benchmark claims, case-identical baselines, or causal decompositions. Appendix A25 reports an unbundled 3D segmentation-architecture pilot as appendix scope context only, not as primary evidence.

Table 5: **What is primary, what is diagnostic, and what is a scope boundary.** Detailed appendix sections preserve the distinction between primary evidentiary results, mechanism diagnostics, and scope checks.

Evidence role	Bundle support	Interpretation
Primary frozen-encoder audit	Per-case Dice, <code>paper_table.json</code> , <code>within_cross/*.json</code>	Primary main claim: five task rows give $\Delta = 0.261 - 0.638$ for frozen/light-adaptation 2D pretrained-encoder pools under the Table-1 grouping.
Estimator and event robustness	<code>audit_sensitivity.json</code> ; <code>quality_controlled_pairs.json</code> ; <code>failure_event_summary.json</code> ; <code>item_difficulty_robustness.json</code>	Stability checks only: floor/ICC/leave-one-out/permutation, pair-quality controls, binary failure-event readouts, and cross-fitted item-difficulty residualisation; no new training evidence.
Architecture/pretraining and cross-checkpoint diagnostics	<code>r7_arch_confound_analysis.json</code> ; <code>cross_checkpoint_family.json</code>	Same-architecture-family ViT / different-pretraining and same-family cross-checkpoint rows lie between seed floor and primary cross-family reference; unbundled cross-architecture probes are labelled separately and not used as support.
Same-backbone levers	Released M3/M4/M22 traces plus selected aggregate diagnostic summaries; unbundled probes are explicitly labelled and not used as support	These weaken simple alternatives but are heterogeneous and not dose-matched causal effects.
Full-pipeline complements	Released nnU-Net JSONs; unbundled 3D segmentation-architecture pilot reported only as appendix scope context	Boundary evidence: nnU-Net complements give much smaller gaps, while the 3D pilot is not included in the submission artifact; do not generalize the 2D frozen-encoder result to all segmentation ensembles.
Held-out stress test	Aggregate held-out JSONs only; no case CIs for the held-out aggregator	External-validity diagnostic: the 11-bundle stress test gives large direction-positive Hippocampus/ISIC rows, a tiny- N positive Prostate row, a near-zero Spleen row, and a Heart floor failure; held-out rows remain secondary scope evidence.

8 Discussion

Practical reading. Treat a same-bundle 2D frozen-feature pool as potentially $M_{\text{eff}} \ll M$ until an in-domain audit shows otherwise. In our released rows, diverse $M=4$ pools yield M_{eff} from 2.05 (Kvasir) to 2.97 (RIGA Disc) under the paper’s equicorrelation readout at $\tau = 0.85$, rather than 4; same-bundle seed pools yield M_{eff} near 1 under the same readout (Appendix A28). M_{eff} is only a dependence readout, so it should be paired with marginal-Dice floors and matched-quality comparisons before any pool-selection decision.

Limitations. The evidence is retrospective and benchmark-based, and should be read at the level of the audit it supports. We do not measure clinician behavior, diagnoses, patient outcomes, prompt-conditioned MedSAM, broad cross-demographic shifts, or case-identical full-pipeline baselines for every primary row. The BraTS foreground row is a 2D derivative of the 2024 GLI release rather than the official 3D BraTS challenge target, and all primary rows are finite 11-bundle-panel audits after a task-specific functional floor applied to the evaluation traces. This is not a VFM performance benchmark, SOTA segmentation benchmark, or full fine-tuning study; the frozen/light-adaptation design is chosen to isolate the co-failure reporting problem under a controlled readout. Table 1 is the confirmatory evidence surface; secondary threshold, lift, pool-selection, and held-out rows are exploratory diagnostics rather than multiplicity-adjusted confirmatory tests. CKA and the random-effects representation are descriptive tools, not causal proofs.

What the evaluation supports. The evaluation supports a practical rule for E&D reporting: in the tested 2D frozen/light-adaptation regime, same-bundle decoder-seed pools should not be credited as independent failure sources unless an in-domain co-failure audit shows that they are. Cross-bundle/family pools reduce but do not eliminate dependence, so pool reports should include aligned-case co-failure correlations in addition to aggregate Dice. The artifact is intended for reproducing the reported dependence estimators and auditing fixed derived traces; it is not evidence of clinical safety, demographic fairness, diagnostic accuracy, or deployable pool-selection performance without separate in-domain validation.

Release. The artifact bundle includes primary per-case Dice JSONs, summary tables, code, artifact metadata/manifests, hashes, and selected appendix/diagnostic aggregate JSONs. It does not introduce a new raw dataset; the released JSONs are derived trace artifacts that retain upstream de-identified alignment identifiers and source labels, including subject/patient/slice IDs and RIGA source/gender-coded path labels, solely for deterministic alignment and subject-level resampling. Users should not attempt re-identification or patient-level profiling. The bundle does not include raw medical images, raw masks, direct identifiers, full prediction caches, or held-out raw prediction arrays; raw data and checkpoints remain governed by upstream licenses and DUAs.

9 Artifact and reproducibility scope

The artifact supports auditing the numerical claims from fixed derived traces without including raw medical images, raw masks, or prediction caches. Table 6 maps each class of main-text evidence to its released files and recomputation path. The mapping is intentionally conservative: primary claims are reproducible from per-case Dice traces, while mechanism, pixel, lift, held-out, and full-pipeline rows are labelled as diagnostics or scope boundaries when their raw masks, prediction arrays, or case-identical comparators are not included. This separation is part of the E&D contribution: the artifact makes the evidentiary status of each result explicit without requiring readers to reconstruct it from appendix provenance.

Table 6: **Main claim-to-artifact map.** Every primary main-text number is tied to a released JSON or a CPU-only scorer. Aggregate-only diagnostics are marked as such rather than promoted to primary evidence.

Claim/readout	Released source	Recompute / limitation
Primary same-vs-cross gaps, CIs, and post-floor pools	<code>paper_table.json</code> , <code>within_cross/*.json</code> , <code>per_case_dice/*.json</code>	<code>compute_all.py</code> ; case IDs and per-case Dice are included, raw images and masks are not.
Subject, floor, ICC, permutation, event, and quality-control robustness	<code>subject_level/*.json</code> , <code>audit_sensitivity.json</code> , <code>failure_event_summary.json</code> , <code>quality_controlled_pairs.json</code>	<code>compute_all.py</code> , <code>compute_audit_sensitivity.py</code> , <code>compute_failure_event_summary.py</code> , <code>compute_quality_controlled_pairs.py</code> ; these are estimator checks, not new training runs.
Item-difficulty and cross-checkpoint diagnostics	<code>item_difficulty_robustness.json</code> , <code>cross_checkpoint_family.json</code> , <code>r7_arch_confound_analysis.json</code>	<code>compute_item_difficulty_robustness.py</code> and <code>compute_cross_checkpoint_family.py</code> ; diagnostic evidence for alternative explanations, not causal identification.
Pixel/lift, same-backbone, held-out, and nnU-Net scope rows	<code>diagnostics/pixel_error/*.json</code> , <code>diagnostics/matched_lift/*.json</code> , <code>appendix trace trees, nnunet/*.json</code>	Aggregate or task-level support only where raw masks/prediction caches or case-identical full-pipeline comparators are not included.

The artifact reproduction script recomputes the primary scorer and CPU-only appendix scorers used by the main-text robustness checks. The SHA-256 manifest covers the released files, and the meta-data describes the trace/code artifact rather than a standalone raw dataset. These mechanics do not make the study fully reproducible from raw data; they make the released evidence auditable at the level at which the claims are made.

References

- Taiga Abe, E. Kelly Buchanan, Geoff Pleiss, and John Patrick Cunningham. Pathologies of predictive diversity in deep ensembles. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=TQfQUksaC8>.
- Ahmed Almazroa. Retinal fundus images for glaucoma analysis: The riga dataset. University of Michigan - Deep Blue Data, 2018. URL <https://doi.org/10.7302/Z23R0R29>.
- Ahmed Almazroa, Sami Alodhayb, Essameldin Osman, Eslam Ramadan, Mohammed Hummadi, Mohammed Dlaim, Muhannad Alkatee, Kaamran Raahemifar, and Vasudevan Lakshminarayanan. Retinal fundus images for glaucoma analysis: The riga dataset. In *Medical Imaging 2018: Imaging Informatics for Healthcare, Research, and Applications*, volume 10579, page 105790B. SPIE, 2018. doi: 10.1117/12.2293584.
- Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *arXiv:2107.07511*, 2021.
- Michela Antonelli, Annika Reinke, Spyridon Bakas, et al. The medical segmentation decathlon. *Nature Communications*, 13:4128, 2022. doi: 10.1038/s41467-022-30695-9. URL <http://medicaldecathlon.com/>.
- Anwai Archit, Luca Freckmann, and Constantin Pape. Medicosam: Robust improvement of sam for medical imaging. *arXiv:2501.11734*, 2025.
- Mohammed Baharoon et al. Evaluating general purpose vision foundation models for medical image analysis: An experimental study of dinov2 on radiology benchmarks. *arXiv:2312.02366*, 2023.
- Olivier Bernard, Alain Lalande, Clement Zotti, et al. Deep learning techniques for automatic mri cardiac multi-structures segmentation and diagnosis: Is the problem solved? *IEEE Transactions on Medical Imaging*, 37(11):2514–2525, 2018. doi: 10.1109/TMI.2018.2837502.
- Rishi Bommasani, Drew A Hudson, Ehsan Adeli, et al. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021.
- Rishi Bommasani, Kathleen A. Creel, Ananya Kumar, Dan Jurafsky, and Percy Liang. Picking on the same person: Does algorithmic monoculture lead to outcome homogenization? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- BraTS 2024 Organizers. Brats 2024 glioma challenge data. Synapse: syn53708249, 2024. URL <https://www.synapse.org/Synapse:syn53708249/wiki/628815>.
- A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *Review of Economics and Statistics*, 90(3):414–427, 2008.
- Gabriele Campanella, Shengjia Chen, Manbir Singh, et al. A clinical benchmark of public self-supervised pathology foundation models. *Nature Communications*, 16:3640, 2025. doi: 10.1038/s41467-025-58796-1.
- Junlong Cheng, Jin Ye, Zhongying Deng, Jianpin Chen, Tianbin Li, Haoyu Wang, Yanzhou Su, Ziyang Huang, Jilong Chen, Lei Jiang, Hui Sun, Junjun He, Shaoting Zhang, Min Zhu, and Yu Qiao. Sam-med2d. *arXiv:2308.16184*, 2023.
- Noel Codella, Veronica Rotemberg, Philipp Tschandl, M. Emre Celebi, Stephen Dusza, David Gutman, Brian Helba, Aadi Kalloo, Konstantinos Liopyris, Michael Marchetti, Harald Kittler, and Allan Halpern. Skin lesion analysis toward melanoma detection 2018: A challenge hosted by the international skin imaging collaboration (isic). *arXiv:1902.03368*, 2019.
- Edwin de Jong, Eric Marcus, and Jonas Teuwen. Current pathology foundation models are unrobust to medical center differences. *arXiv:2501.18055*, 2025.
- Maria Correia de Verdier, Rachit Saluja, Louis Gagnon, Dominic LaBella, Ujjwall Baid, Nourel Hoda Tahon, Martha Foltyn-Dumitru, Jikai Zhang, Maram Alafif, Saif Baig, et al. The 2024 brain tumor segmentation (brats) challenge: Glioma segmentation on post-treatment mri. *arXiv preprint arXiv:2405.18368*, 2024. doi: 10.48550/arXiv.2405.18368.
- Stanislav Fort, Huiyi Hu, and Balaji Lakshminarayanan. Deep ensembles: A loss landscape perspective. *arXiv preprint arXiv:1912.02757*, 2019.
- Yarin Gal and Zoubin Ghahramani. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. In *ICML*, 2016.

366 M A Ganaie, Minghui Hu, A K Malik, M Tanveer, and P N Suganthan. Ensemble deep learning: A review.
367 *Engineering Applications of Artificial Intelligence*, 115, 2022.

368 Robert Geirhos, Kristof Meding, and Felix A. Wichmann. Beyond accuracy: Quantifying trial-by-trial be-
369 haviour of cnns and humans by measuring error consistency. In *Advances in Neural Information Processing*
370 *Systems*, 2020.

371 Mariana-Iuliana Georgescu, Radu Tudor Ionescu, and Andreea-Iuliana Miron. Diversity-promoting ensemble
372 for medical image segmentation. In *Proceedings of the 38th ACM/SIGAPP Symposium on Applied Comput-*
373 *ing*, pages 599–606, 2023. doi: 10.1145/3555776.3577822.

374 Shashwat Goel, Joschka Struber, Ilze Amanda Auzina, Karuna K. Chandra, Ponnuram Kumaraguru, Douwe
375 Kiela, Ameya Prabhu, Matthias Bethge, and Jonas Geiping. Great models think alike and this undermines
376 AI oversight. arXiv:2502.04313, 2025.

377 Raphael Gontijo-Lopes, Yann Dauphin, and Ekin Dogus Cubuk. No one representation to rule them all: Over-
378 lapping features of training methods. In *International Conference on Learning Representations (ICLR)*,
379 2022.

380 Mila Gorecki and Moritz Hardt. Monoculture or multiplicity: Which is it? In *NeurIPS*, 2025. URL <https://openreview.net/forum?id=D05LtJc80w>.

382 Hanxue Gu, Haoyu Dong, Jichen Yang, and Maciej A. Mazurowski. How to build the best medical image
383 segmentation algorithm using foundation models: A comprehensive empirical study with segment anything
384 model. *MELBA*, 2025. arXiv:2404.09957.

385 Yi Gu, Yi Lin, Kwang-Ting Cheng, and Hao Chen. Revisiting deep ensemble uncertainty for enhanced medical
386 anomaly detection. In *MICCAI*, 2024.

387 Yufan He et al. Vista3d: A unified segmentation foundation model for 3d medical imaging. In *CVPR*, 2025.

388 Yuhao Huang, Xin Yang, Lian Liu, Han Zhou, Ao Chang, Xinrui Zhou, Rusi Chen, Junxuan Yu, Jiongquan
389 Chen, Chaoyu Chen, Sijing Liu, Haozhe Chi, Xindi Hu, Kejuan Yue, Lei Li, Vicente Grau, Deng-Ping Fan,
390 Fajin Dong, and Dong Ni. Segment anything model for medical images? *Medical Image Analysis*, 92:
391 103061, 2024. doi: 10.1016/j.media.2023.103061.

392 Fabian Isensee, Paul F Jaeger, Simon A A Kohl, Jens Petersen, and Klaus H Maier-Hein. nnU-Net: a self-
393 configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18:203–211,
394 2021.

395 Fabian Isensee et al. nnu-net revisited: A call for rigorous validation in 3d medical image segmentation. In
396 *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, 2024.

397 Debesh Jha, Pia H. Smedsrud, Michael A. Riegler, Pål Halvorsen, Thomas de Lange, Dag Johansen, and
398 Håvard D. Johansen. Kvasir-SEG: A segmented polyp dataset. In *International Conference on Multimedia*
399 *Modeling (MMM)*, 2020.

400 Liwei Jiang, Yejin Choi, et al. Artificial hivemind: The open-ended homogeneity of language models (and
401 beyond). arXiv:2510.22954; NeurIPS 2025 Datasets & Benchmarks Track oral, 2025.

402 Nathanael Jo, Nikhil Garg, and Manish Raghavan. The subjectivity of monoculture. arXiv:2602.24086, 2026.

403 Alain Jungo and Mauricio Reyes. Assessing reliability and challenges of uncertainty estimations for medical
404 image segmentation. In *MICCAI*, 2019.

405 Tommie Kerssies et al. How to benchmark vision foundation models for semantic segmentation? In *CVPR*
406 *Workshop*, 2024.

407 Elliot Kim, Avi Garg, Kenny Peng, and Nikhil Garg. Correlated errors in large language models. In *Inter-*
408 *national Conference on Machine Learning (ICML)*, 2025a. arXiv:2506.07962.

409 Yunhe Kim et al. Show and segment: Universal medical image segmentation via in-context learning.
410 arXiv:2503.19359, 2025b.

411 Alexander Kirillov, Eric Mintun, Nikhila Ravi, et al. Segment anything. In *2023 IEEE/CVF International*
412 *Conference on Computer Vision (ICCV)*, pages 3992–4003, 2023. doi: 10.1109/ICCV51070.2023.00371.

413 Jon Kleinberg and Manish Raghavan. Algorithmic monoculture and social welfare. *Proceedings of the National*
414 *Academy of Sciences*, 118(22), 2021.

415 Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network
416 representations revisited. In *ICML*, 2019.

417 Ludmila I Kuncheva and Christopher J Whitaker. Measures of diversity in classifier ensembles and their rela-
418 tionship with the ensemble accuracy. *Machine Learning*, 51(2):181–207, 2003.

419 Zeki Kuş and Musa Aydin. Medsegbench: A comprehensive benchmark for medical image segmentation in
420 diverse data modalities. *Scientific Data*, 11(1):1283, 2024. doi: 10.1038/s41597-024-04159-2.

421 Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty
422 estimation using deep ensembles. In *NeurIPS*, 2017.

423 Yuanyuan Li, Neeraj Sarna, and Yang Lin. Quantifying correlations of machine learning models. *arXiv preprint*
424 *arXiv:2502.03937*, 2025.

425 Xiangde Luo, Xiyue Wang, Feyisope Eweje, Xiaoming Zhang, et al. Ensemble learning of foundation models
426 for precision oncology. *arXiv:2508.16085*, 2025. URL <https://arxiv.org/abs/2508.16085>.

427 Jun Ma, Yuting He, Feifei Li, Lin Han, Chenyu You, and Bo Wang. Segment anything in medical images.
428 *Nature Communications*, 15, 2024.

429 Jun Ma, Zongxin Yang, Sumin Kim, Bihui Chen, Mohammed Baharoon, Adibvafa Fallahpour, Reza Asakereh,
430 Hongwei Lyu, and Bo Wang. Medsam2: Segment anything in 3d medical images and videos. *arXiv preprint*
431 *arXiv:2504.03600*, 2025.

432 Kenneth O. McGraw and S. P. Wong. Forming inferences about some intraclass correlation coefficients. *Psy-*
433 *chological Methods*, 1(1):30–46, 1996.

434 Alireza Mehrtash, William M Wells, Clare M Tempny, Purang Abolmaesumi, and Tina Kapur. Confidence
435 calibration and predictive uncertainty estimation for deep medical image segmentation. *IEEE Transactions*
436 *on Medical Imaging*, 39(12):3868–3878, 2020.

437 Dominik J. Mühlematter, Michelle Halbheer, Alexander Becker, Dominik Narnhofer, Helge Aasen, Konrad
438 Schindler, and Mehmet Ozgur Turkoglu. Lora-ensemble: Efficient uncertainty modelling for self-attention
439 networks. *Transactions on Machine Learning Research*, 2026. URL [https://openreview.net/forum?](https://openreview.net/forum?id=yhXXm0MpSQ)
440 [id=yhXXm0MpSQ](https://openreview.net/forum?id=yhXXm0MpSQ).

441 Basil Mustafa, Carlos Riquelme, Joan Puigcerver, André Susano Pinto, Daniel Keysers, and Neil Houlsby.
442 Deep ensembles for low-data transfer learning. *arXiv preprint arXiv:2010.06866*, 2020.

443 Peter Neidlinger, Omar S. M. El Nahhas, Hannah Sophie Muti, et al. Benchmarking foundation models as
444 feature extractors for weakly-supervised computational pathology. *Nature Biomedical Engineering*, 2025.
445 doi: 10.1038/s41551-025-01516-3.

446 Behnam Neyshabur, Hanie Sedghi, and Chiyuan Zhang. What is being transferred in transfer learning? In
447 *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

448 Connor Toups, Rishi Bommasani, Kathleen Creel, Sarah H. Bana, Dan Jurafsky, and Percy Liang. Ecosystem-
449 level analysis of deployed machine learning reveals homogeneous outcomes. In *Advances in Neural Infor-*
450 *mation Processing Systems (NeurIPS)*, 2023.

451 Kagan Tumer and Joydeep Ghosh. Error correlation and error reduction in ensemble classifiers. *Connection*
452 *Science*, 8(3–4):385–404, 1996. doi: 10.1080/095400996116839.

453 Florian Wenzel, Jasper Snoek, Dustin Tran, and Rodolphe Jenatton. Hyperparameter ensembles for robustness
454 and uncertainty quantification. In *NeurIPS*, 2020.

455 Danny Wood, Tingting Mu, Andrew M. Webb, Henry Reeve, Mikel Lujan, and Gavin Brown. A unified theory
456 of diversity in ensemble learning. *arXiv:2301.03962*, 2023.

457 Mitchell Wortsman, Gabriel Ilharco, Samir Ya Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes, Ari S. Morcos,
458 Hongseok Namkoong, Ali Farhadi, Yair Carmon, Simon Kornblith, and Ludwig Schmidt. Model soups:
459 Averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In
460 *International Conference on Machine Learning (ICML)*, 2022.

461 R. Patrick Xian, Noah R. Baker, Tom David, Qiming Cui, A. Jay Holmgren, Stefan Bauer, Madhumita Sushil,
462 et al. Robustness tests for biomedical foundation models should tailor to specifications. *npj Digital Medicine*,
463 8:557, 2025. doi: 10.1038/s41746-025-01926-2.

- 464 Yuncheng Yang, Meng Wei, Junjun He, Jie Yang, Jin Ye, and Yun Gu. Pick the best pre-trained model: Towards
 465 transferability estimation for medical image segmentation. In *Medical Image Computing and Computer*
 466 *Assisted Intervention (MICCAI)*, 2023.
- 467 Maximilian Zenk, Klaus H. Maier-Hein, Jens Petersen, et al. Comparative benchmarking of failure detection
 468 methods in medical image segmentation: Unveiling the role of confidence aggregation. *Medical Image*
 469 *Analysis*, 101:103392, 2025. doi: 10.1016/j.media.2024.103392.
- 470 Yukun Zhou, Mark A. Chia, Siegfried K. Wagner, et al. A foundation model for generalizable disease detection
 471 from retinal images. *Nature*, 622:156–163, 2023.
- 472 Jiayuan Zhu et al. Medical sam 2: Segment medical images as video via segment anything model 2.
 473 arXiv:2408.00874, 2024.

474 A1 Appendix roadmap and evidence inventory

475 The appendix is organised by evidentiary role rather than by experiment chronology. It starts with
 476 companion readouts, then documents primary estimators, uncertainty, task-level sources, and esti-
 477 mator sensitivity. Middle sections collect mechanism diagnostics and same-backbone controls; later
 478 sections cover adaptation, data swaps, segmentation-pipeline variants, distribution shift, operational
 479 metrics, held-out stress tests, and M_{eff} . The final sections record artifact metadata and additional
 480 matrix diagnostics. Secondary rows are retained only when they qualify or explain a paper claim,
 481 and their captions state when they use non-primary pools or splits. Unless explicitly marked primary
 482 or sourced from the primary `results/_merged` summaries, appendix values are diagnostic rather
 483 than primary support for Table 1.

484 **CLIP-L notation.** Throughout the appendix, CLIP-L/14 refers to the preserved OpenCLIP
 485 ViT-L-14/openai dense-token wrapper used to generate the released traces. It is not a claim of
 486 OpenAI clip.load parity; a strict QuickGELU-parity CLIP-L loader would be a separate sensitiv-
 487 ity with a distinct FM identifier.

Table 7: **Appendix guide.** Each block has a single role in the argument; cross-references point to the detailed tables rather than repeating them in the main paper.

Block	Evidence role
Estimators and provenance	Case/hierarchical bootstrap, subject clustering, quality-controlled pair regres- sion, functional-floor/ICC/permutation sensitivity, pixel overlap, annotation- source note, and primary dataset details.
Mechanism diagnostics	Same-architecture-family ViT / different-pretraining, CKA, random-effects, and item-difficulty null analyses.
Same-backbone controls	Matched-size/quality-gated pool checks, decoder capacity, UNet skips, fixed- backbone folds, MC dropout, released CNN random-init, and labelled scope- limited probes.
Scope boundaries	nnU-Net complements, unbundled 3D pilot, BraTS sparse-target sensitivities, standard segmentation-recipe, distribution shift, and held-out failures.
Operational metrics	Conditional recovery, matched- M lift, lift ceiling, pool-selection probes, and M_{eff} .
Artifact record	Bundle manifest, artifact metadata, requirements, and trace inventory.

488 A2 Companion readouts and scope

489 This section records companion readouts that clarify scope without adding new headline claims.

490 **CKA.** We computed CKA only as a descriptive feature-similarity readout. The saved partial-
 491 regression coefficient artifact is not included in the submission artifact, so we do not use partial-
 492 regression p -values or coefficients as formal evidence. The main paper relies on pixel-level failure
 493 overlap and per-case Dice correlation for mechanism-adjacent interpretation, and treats CKA as a
 494 qualitative companion variable rather than as a standalone causal mediator.

495 **RIGA annotation source.** Table 1 uses the six-rater majority consensus as ground truth for both
 496 RIGA Cup and RIGA Disc, matching the released RIGA per-case Dice traces. RIGA also distributes

individual ophthalmologist masks, but individual-rater replay artifacts are not included in the submission artifact. The consensus Cup and Disc rows are therefore the only RIGA annotation-source evidence used by the manuscript.

Probability-map diagnostics. The artifact releases per-case Dice traces and selected aggregate pixel/lift diagnostics, but not raw probability maps. We therefore do not use mean-confidence, probability-calibration, or failure-flagging AUROC values as evidence for any paper claim.

A3 Pixel-level error-overlap breakdown

Table 8: Pixel-level error overlap metrics with FP/FN/boundary breakdowns. For error-Dice, no bootstrap replicate out of $B=500$ produced a non-positive gap on any task cell, giving the finite- B bound $\Pr[\text{gap} \leq 0] < 1/501$. FP/FN/boundary rows are point-estimate breakdowns.

Metric	RIGA Cup	RIGA Disc	ACDC LV
Error-Dice gap	0.485	0.481	0.485
Error-Jaccard gap	0.515	0.481	0.521
FP-only error-Dice gap	0.453	0.494	0.566
FN-only error-Dice gap	0.527	0.464	0.449
Boundary-band (5px)	0.430	0.439	0.433

FP = (pred foreground, GT background); FN = (pred background, GT foreground); boundary-band restricts error masks to pixels within 5 pixels of the GT edge (trimap). All metrics are pair-wise Dice between the two models’ error masks (or empty-mask-excluded Jaccard). On RIGA Cup, RIGA Disc, and ACDC LV binary, the similar-magnitude gaps across FP, FN, and boundary variants indicate the spatial co-failure is not driven by a single error type on those three tasks (Kvasir and BraTS not tabulated in this split).

A4 Bootstrap sensitivity of the gap

Table 1 reports 95% CIs from a three-level nested hierarchical bootstrap that resamples family labels, seeds within each family, and cases ($B=2000$). We additionally compute a *case-only* bootstrap that fixes the family/seed structure and resamples the N test cases with replacement ($B=2000$). The two procedures target different sampling dimensions: the hierarchical CI jointly perturbs the finite FM/seed/case surface, whereas the case-only CI asks “what if we had a different test set” conditional on the evaluated pool. We report both as sensitivity evidence for the same finite audit rather than as population-level guarantees.

Table 9: **Case-level bootstrap sensitivity.** Regenerated primary-row point estimates and gap CIs from the case-level bootstrap ($B=2,000$, fixed family/seed aggregation). BraTS uses slice-level resampling; subject-level clustering is reported separately in Appendix A11.

Task	Functional pool	Within	Cross	Gap [95% case CI]
Kvasir	11-bundle	0.958	0.697	0.261 [0.211, 0.320]
ACDC LV	10-bundle	0.959	0.639	0.321 [0.289, 0.355]
BraTS foreground [†]	10-bundle	0.938	0.623	0.315 [0.302, 0.329]
RIGA Cup	9-bundle	0.834	0.361	0.473 [0.395, 0.556]
RIGA Disc	11-bundle	0.870	0.232	0.638 [0.569, 0.687]

[†] Completed 11-bundle source pool with 10 post-floor bundles; see Appendix A11 for slice-vs-subject resampling caveats.

BraTS foreground pool reconciliation. The headline BraTS foreground result is the 10-bundle Table 1 value of $\Delta=0.315$ [0.302, 0.329] (case-level CI), from the completed 11-bundle source pool with 10 post-floor bundles (BiomedCLIP, CLIP-B/16, CLIP-L/14, ConvNeXt-T, DeiT-B/16, DINOv2-B/S, EfficientNet-B0, MAE-B/16, ResNet-50; ResNet-18 is below the Dice floor). Source: `results/_merged/paper_table.json`. Subset diagnostics whose exact aggregate files are not included in the submission artifact are omitted from the released numeric evidence. The headline reported throughout the main text and abstract is therefore the Table 1 value.

All gap CIs strictly exclude zero under the case-level bootstrap, matching the family/seed-level CIs in the main text.

A5 Estimator-robustness and negative-control checks

The primary row uses Pearson correlation and a Dice ≥ 0.30 per-FM functional floor. To make this choice auditable, the released bundle includes a CPU-only sensitivity artifact, `results/_merged/diagnostics/audit_sensitivity.json`, generated from the same per-case Dice JSONs by `code/scripts/compute_audit_sensitivity.py`. It recomputes the within-cross gap under a floor sweep, an ICC(A,1) alternative, leave-one-FM-out summaries, leave-one-task-out meta summaries, and a random family-label permutation control.

Table 10: **Primary-trace estimator sensitivity.** Pearson values are the Table 1 point estimates at the Dice ≥ 0.30 floor. ICC(A,1) is an alternate agreement statistic on the same post-floor members. The floor range reports the minimum/maximum Pearson gap over $\{0, 0.20, 0.25, 0.30, 0.35, 0.40\}$ where the pool is non-empty. The permutation note reports a random family-label rank fraction over completed members, using $(1 + \#\{\Delta_{\text{null}} \geq \Delta_{\text{obs}}\})/(1001)$.

Task	Members at 0.30	Pearson gap	ICC(A,1) gap	Floor-sweep gap range
Kvasir	44	0.261	0.333	0.261–0.261
ACDC LV	40	0.321	0.401	0.242–0.355
BraTS foreground	40	0.315	0.404	0.315–0.362
RIGA Cup	36	0.473	0.539	0.384–0.529
RIGA Disc	44	0.638	0.716	0.638–0.638

All five rows have observed-gap-exceeded rank fraction 0.001 at the 0.30 floor.

The five-task mean Pearson gap at the paper floor is 0.401 (median 0.321, range 0.261–0.638). Leave-one-task-out means stay in 0.342–0.437; the smallest value occurs when the largest-gap RIGA Disc row is held out. These checks do not add new cases or new model training. They instead reduce the risk that the headline sign is an artifact of the exact functional floor, the Pearson statistic, or arbitrary family labels.

Calibration-selected functional pools. The primary table defines a functional pool using all released evaluation traces. To check whether the sign depends on this convention, the scorer also makes a deterministic split of each task’s aligned units, selects bundles by mean Dice on the calibration subset, and evaluates the same/cross gap on disjoint cases. This is a leakage-sensitivity analysis rather than a prospective validation protocol, but it addresses the specific concern that the functional-floor sign is induced by evaluating the floor and the gap on exactly the same cases.

Table 11: **Calibration-selected functional-pool sensitivity.** Functional bundles are selected on the calibration subset and evaluated on disjoint cases under the Table 1 same/cross grouping. CIs are case-bootstrap intervals on the evaluation subset. Source: `results/_merged/calibration_split/*.json`, regenerated by `code/scripts/compute_all.py`.

Task	Cal./eval units	Selected bundles	Eval. Δ [95% CI]	Unit-agg. Δ
Kvasir	42/78 images	11	0.250 [0.196, 0.322]	–
ACDC LV	7/13 patients	11	0.350 [0.313, 0.392]	0.206
BraTS foreground	113/211 subjects	10	0.307 [0.292, 0.323]	0.205
RIGA Cup	33/62 images	9	0.478 [0.369, 0.599]	–
RIGA Disc	33/62 images	11	0.630 [0.534, 0.695]	–

A6 Quality-controlled pairwise-correlation diagnostic

A possible alternative explanation is that same-FM pairs have higher correlation merely because they have matched marginal quality. We therefore add a CPU-only diagnostic that uses the same released post-floor per-case Dice traces as Table 1 and fits a pair-level OLS model:

$$r_{ab} \sim I_{\text{same}} + \bar{D}_{ab} + |\Delta \bar{D}_{ab}| + \bar{V}_{ab} + |\Delta V_{ab}|.$$

Here $\bar{D}_{ab}$ and $\bar{V}_{ab}$ are the pair mean Dice and average per-case Dice variance, with absolute pair differences as additional controls. The same-group indicator is exactly the primary grouping rule: same FM plus the DINOv2-B/S pair. The coefficient is not a causal estimand, but it

552 checks that the main sign is not removed by simple pair-quality and variance controls. The diag-
553 nostic is generated by `code/scripts/compute_quality_controlled_pairs.py` and released
554 as `results/_merged/diagnostics/quality_controlled_pairs.json`.

Table 12: **Quality-controlled pairwise-correlation diagnostic.** Same-group coefficients remain positive after controlling for pair mean Dice, absolute mean-Dice difference, average per-case variance, and absolute variance difference. CIs are case-bootstrap intervals from $B=1000$ refits.

Task	Members	Same pairs	Cross pairs	β_{same} [95% CI]
Kvasir	44	82	864	0.169 [0.124, 0.220]
ACDC LV	40	76	704	0.208 [0.175, 0.242]
BraTS foreground	40	76	704	0.224 [0.210, 0.239]
RIGA Cup	36	70	560	0.372 [0.297, 0.456]
RIGA Disc	44	82	864	0.582 [0.495, 0.639]

555 A7 Failure threshold sensitivity

556 The primary estimand is continuous per-case Dice correlation. To make the co-failure interpretation
557 less dependent on continuous-score correlation, the released artifact also recomputes the Table 1
558 same/cross grouping on binary failure events. Table 13 reports the empirical P_{33} readout across all
559 five primary task rows; the released JSON also contains P_{25} and P_{50} thresholds. The RIGA Cup
560 all-fail table below is retained as a pool-level threshold diagnostic, and absolute-threshold lift with
561 support flags is in Appendix A9.

Table 13: **Binary failure-event same/cross diagnostic at the empirical P_{33} threshold.** Failure is defined as per-case Dice below the 33rd percentile of all model-case Dice values in the post-floor functional pool. ϕ is binary-event Pearson correlation averaged over same-group or Table-1 cross-group pairs. CIs are case-bootstrap intervals for the same-cross ϕ gap. Source: `results/_merged/diagnostics/failure_event_summary.json`, regenerated by `code/scripts/compute_failure_event_summary.py`.

Task	τ_{P33}	Same ϕ	Cross ϕ	Gap [95% CI]	Joint fail same/cross
Kvasir	0.659	0.869	0.540	0.329 [0.278, 0.388]	27.9/21.6%
ACDC LV	0.548	0.875	0.489	0.386 [0.354, 0.418]	27.7/19.9%
BraTS foreground	0.636	0.838	0.425	0.413 [0.401, 0.425]	27.3/18.7%
RIGA Cup	0.575	0.706	0.253	0.453 [0.399, 0.509]	24.4/14.4%
RIGA Disc	0.857	0.703	0.125	0.578 [0.511, 0.631]	24.8/12.0%

562 **Full τ -sweep of conditional recovery.** Table 14 tabulates grand-mean conditional recovery
563 rates $P[\text{pool-mean Dice} \geq \tau \mid \text{any member Dice} < \tau]$ for monoculture (4 seeds of one
564 FM, averaged over all FMs passing the per-task Dice ≥ 0.30 functional floor) and di-
565 verse pools (all $\binom{K}{4}$ functional-subset compositions and all released seed tuples) at $\tau \in$
566 $\{0.70, 0.75, 0.80, 0.85, 0.90\}$. Source: `results/_merged/diagnostics/tau_sweep.json`, re-
567 generated by `code/scripts/compute_tau_sweep.py`. Recovery is τ -fragile in both regimes; the
568 diverse – mono gap (column Δ) is non-monotone in τ .

569 A8 Kvasir secondary split diagnostic

570 Pixel-level diagnostic on this split: mono 0.80 vs. diverse 0.40, gap 0.394 [95% CI 0.38, 0.40].
571 Absolute-threshold lift at Dice<0.7: mono $252\times$ vs. diverse $120\times$ (ratio $2.1\times$); at Dice<0.8: $48.4\times$
572 vs. $14.6\times$ ($3.3\times$); at Dice<0.9: $4.6\times$ vs. $1.8\times$ ($2.5\times$). The numbers are slightly smaller than
573 RIGA/ACDC, likely reflecting the larger test set (300 cases vs. 224 / 90) and more variable polyp
574 presentations, but qualitatively identical: mono still shows substantially higher correlated-failure lift
575 than diverse.

Table 14: **Full τ -sweep of conditional recovery rate $P[\text{ens pass} \mid \text{any fail}]$.** Grand-mean across all functional 4-of- K diverse compositions and all functional monoculture baselines. These secondary tau-sweep values are summarized here rather than used as primary evidence.

Task	τ	Mono	Diverse	Δ
RIGA Cup (9-bundle)	0.70	0.061	0.214	+0.153
	0.75	0.073	0.120	+0.047
	0.80	0.055	0.050	−0.005
	0.85	0.010	0.007	−0.003
	0.90	0.001	0.000	−0.001
ACDC LV (10-bundle)	0.70	0.074	0.285	+0.211
	0.75	0.062	0.237	+0.175
	0.80	0.049	0.163	+0.114
	0.85	0.040	0.072	+0.032
	0.90	0.020	0.012	−0.007
Kvasir-SEG (11-bundle)	0.70	0.092	0.388	+0.295
	0.75	0.074	0.360	+0.285
	0.80	0.058	0.267	+0.209
	0.85	0.050	0.156	+0.106
	0.90	0.025	0.030	+0.005
BraTS foreground (2D FLAIR $\text{seg}>0$)	0.70	0.080	0.326	+0.245
	0.75	0.061	0.233	+0.172
	0.80	0.039	0.129	+0.090
	0.85	0.019	0.035	+0.016
	0.90	0.004	0.001	−0.003

Table 15: Failure correlation and ALL-fail rates at multiple empirical thresholds in the released primary RIGA Cup functional pool. Percentile thresholds are computed over all model-case Dice values after the $\text{Dice} \geq 0.30$ functional floor. Source: `results/_merged/diagnostics/threshold_table_riga_cup.json`.

Threshold	Mono fail $\bar{r}$	Diverse fail $\bar{r}$	Mono ALL-fail	Diverse ALL-fail
P_{25} ($\tau=0.513$)	0.833	0.209	19.9%	2.2%
P_{33} ($\tau=0.575$)	0.803	0.257	26.5%	5.3%
P_{50} ($\tau=0.686$)	0.840	0.268	44.8%	15.4%

Table 16: Kvasir-SEG (polyp endoscopy) secondary 8-bundle split diagnostic — *generic-8-bundle summary*. This row uses a 1,000-image, 700/300 train/test split and is not the 11-bundle, 880/120 Kvasir row in Table 1. 7 of 8 generic encoder bundles produce functional segmentation ($\text{Dice} > 0.40$); CLIP ViT-L/14 $\text{Dice}=0.17$ is non-functional and excluded from cross-family analysis. MedSAM-only traces are reported separately as an encoder-probe summary in Appendix A10; no RETFound Kvasir trace is included in the released artifact bundle.

FM Family	Dice	Within $\bar{r}$	Status
DINOv2 ViT-B/14	0.819	0.972	functional
DINOv2 ViT-S/14	0.807	0.951	functional
BiomedCLIP ViT-B/16	0.717	0.954	functional in this split
ConvNeXt-Tiny	0.651	0.949	functional
EfficientNet-B0	0.582	0.945	functional
ResNet-50	0.540	0.880	functional
ResNet-18	0.468	0.876	functional
CLIP ViT-L/14	0.174	0.839	non-functional

576 A9 Matched- M lift comparison

577 To ensure the mono vs. diverse comparison is not driven by pool size, we construct an *oracle*
 578 matched- $M=4$ diverse pool by selecting the top-4 bundles by *test-set seed-averaged* mean Dice
 579 (rank computed across all four decoder seeds, avoiding seed-42 selection circularity; selected
 580 bundles per task are recorded in the appendix diagnostic output—e.g. Kvasir top-4 = {DINOv2-
 581 B, DINOv2-S, BiomedCLIP, ConvNeXt-T}, Cup top-4 = {BiomedCLIP, DINOv2-B, DINOv2-S,
 582 EfficientNet-B0}). This is deliberately favourable to the diverse pool and useful for a matched-size
 583 structural comparison, but it is not a deployable pool-selection rule because the ranking uses the
 584 same test set on which lift is evaluated. For each task / threshold we report three lift variants:

- 585 • **Matched-4 seed-42 (single-seed diagnostic)**: the single-seed pool; retained for comparability
 586 with a single-seed estimator.
- 587 • **Matched-4 symmetric permuted (preferred)**: each of the 4 FMs takes a *distinct* decoder
 588 seed from {42, 43, 44, 45}; the lift is averaged on the log scale over all $4! = 24$ seed-to-FM per-
 589 mutations, then exponentiated. This is the matched-size structural analog of the main paper’s
 590 symmetric Pearson $\bar{r}$ estimator (mono = DINOv2-B with 4 distinct seeds matches matched-4 =
 591 4 FMs with 4 distinct seeds).
- 592 • **Matched-4 symmetric independent (sensitivity)**: each FM’s seed is drawn i.i.d. with replace-
 593 ment from {42..45}, $4^4 = 256$ tuples; reported as a robustness check (the permuted variant is
 594 preferred because mono uses four *distinct* seeds).

595 All bootstrap CIs are computed on *cases* only (patient-clustered on ACDC); the seed tuple space is
 596 finite and fully enumerated, not bootstrapped. We report 95% case-bootstrap CIs on the log-scale lift
 597 and exponentiate to the reported interval ($B=2000$). We also report the raw all-fail count (AF_{cnt})
 598 and the Poisson-binomial independence expectation $\mathbb{E}[AF_{ind}]$; when $N \cdot \mathbb{E}[AF_{ind}] < 5$ we flag
 599 the row as support-limited. Under the matched-size comparison, the adequately supported rows are
 600 RIGA Cup $\tau=0.9$, RIGA Disc $\tau=0.95$, and ACDC LV $\tau=0.9$; Kvasir $\tau=0.8$ is mono-sparse and
 601 retained as a descriptive diagnostic to match Fig. 1. Across these selected rows the mono/matched-4
 602 lift ratio is 1.5–5.8 $\times$ (Table 17); lower-support rows give larger ratios (5.9–17 $\times$) treated as descrip-
 603 tive order-of-magnitude only. We therefore treat τ values well below the task-specific Dice ceiling
 604 as support-limited and report them as descriptive rather than precise lift estimates.

Table 17: **Matched- M symmetric lift comparison.** Mono is DINOv2-B $\times$ 4 seeds; matched-4
 diverse uses the top-4 bundles by seed-averaged mean Dice. CIs are 95% case-bootstrap intervals
 on log-lift, exponentiated. Rows without adequate support ($\checkmark$ missing) are descriptive order-of-
 magnitude indicators.

Task	τ	Mono [CI]	Diverse-42	Diverse-perm [CI]	Ratio*	Supp.
RIGA Cup	0.8	4450 [1408, 24012]	174 [79, 467]	265 [130, 696]	17 $\times$	–
RIGA Cup	0.9	2.18 [1.81, 2.69]	1.33 [1.20, 1.47]	1.44 [1.31, 1.60]	1.5 $\times$	$\checkmark$
RIGA Disc	0.95	9.80 [6.82, 14.6]	2.87 [2.19, 3.82]	3.11 [2.44, 3.98]	3.2 $\times$	$\checkmark$
ACDC LV	0.5	2149 [462, 34449]	200 [51, 1026]	188 [62, 817]	11 $\times$	–
ACDC LV	0.7	189 [55.6, 1064]	31.2 [15.5, 77.7]	25.3 [13.1, 57.8]	7.5 $\times$	–
ACDC LV	0.8	33.8 [15.1, 107]	5.87 [3.60, 10.6]	5.71 [3.63, 9.93]	5.9 $\times$	–
ACDC LV	0.9	2.23 [1.63, 3.32]	1.39 [1.22, 1.60]	1.40 [1.25, 1.61]	1.6 $\times$	$\checkmark$
Kvasir	0.7	246 [128, 570]	33.9 [22.3, 54.5]	30.8 [20.4, 48.3]	8.0 $\times$	–
Kvasir	0.8	48.0 [28.9, 85.1]	8.72 [6.62, 12.0]	8.21 [6.35, 11.1]	5.8 $\times$	mono-sparse

* Ratio = Mono / Matched-4 PERM. **Bold** = preferred symmetric-permuted estimator (24 seed-to-FM tuples). Matched-4
 seed 42 is a single-seed diagnostic. Matched-4 INDEP ($4^4=256$ iid tuples) produces values within $\pm 1.5\%$ of PERM on
 every row and is omitted for compactness; the full table is a secondary diagnostic rather than the `results/_merged`
 primary scorer output. Supp. $\checkmark$ iff $N \cdot \mathbb{E}[AF_{ind}] \geq 5$ for the PERM estimator.

605 A10 Medical-domain encoder probe details

606 We implemented two medical-domain encoder paths, but only the MedSAM image-encoder probe
 607 is reported in the released evidence bundle:

608 **RETFound** is a ViT-L/16 MAE pretrained on 1.6M retinal images [Zhou et al., 2023]. The code path
 609 builds the RETFound-MAE CFP ViT-L/16 trunk and expects the gated HuggingFace checkpoint
 610 YukunZhou/RETFound_mae_natureCFP (or an equivalent local RETFOUND_CFP_CHECKPOINT);
 611 the official upstream project is rmaphoh/RETFound_MAE. RETFound MAE pretraining uses Ima-

geNet [0.485, 0.456, 0.406]/[0.229, 0.224, 0.225] normalisation; our frozen-probe adaptation would drop the CLS token and reshape the $196=14\times 14$ patch tokens into a spatial 14×14 feature map of dimension 1024. That patch-token extraction is not RETFound’s official global/CLS-style classification fine-tuning path. Because the gated checkpoint/authentication was unavailable for the reported artifact run, RETFound cells are not reported and no RETFound result supports the paper’s empirical claims.

MedSAM is a SAM ViT-B fine-tuned on 1.5M medical image-mask pairs [Ma et al., 2024]. The code path uses the official bowang-lab/MedSAM medsam_vit_b.pth checkpoint through `segment_anything.sam_model_registry['vit_b']` with MEDSAM_CHECKPOINT pointing to that checkpoint. In this audit, we use *only the image encoder* as a frozen feature extractor with our same 1×1 -conv decoder. We discard MedSAM’s prompt encoder and mask decoder entirely—i.e. MedSAM enters this boundary probe on the same frozen-feature terms as the other FMs (no prompt-based path, no bounding-box input, no iterative refinement). The wrapper strips the SAM image-encoder neck and uses the pre-neck ViT-B representation: the official 1024×1024 positional grid is bicubically interpolated to the 224×224 input grid, yielding a 14×14 spatial feature map of dimension 768. This pre-neck $768\times 14\times 14$ tensor is an internal representation, not the official SAM/MedSAM post-neck image embedding at 1024 input. This means MedSAM in our experiment is a “MedSAM image-encoder probe”, not its intended prompt-conditioned segmentation system or official inference preprocessing path. The comparison is therefore a controlled encoder-only probe for the *monoculture* question we ask, not a fair evaluation of MedSAM as released or of its best-case performance with its own prompt-based decoder.

All other aspects of the secondary sensitivity use the same four decoder seeds, 30-epoch Adam 10^{-3} training, and identical 1×1 conv decoder. The completed MedSAM-only probe covers ACDC LV, Kvasir, RIGA Cup, and RIGA Disc (four seeds each). Mean Dice across seeds is 0.517, 0.513, 0.734, and 0.908 respectively; the aggregate diagnostic is released as `results/_merged/diagnostics/medsam_only_summary.json`. We use this as scope context rather than primary evidence.

A11 BraTS/ACDC clustering sensitivity

The Table 1 BraTS foreground row is a BraTS 2024 GLI 2D FLAIR derivative with $n=3,228$ test slices from 324 held-out subjects: per held-out subject, the loader keeps up to the top-10 axial slices by foreground area, requires at least 64 positive pixels, and uses a subject-level 80/20 split with NumPy seed 42. The 3,228 count is traceable to the released `test_ids`: 321 held-out subjects contribute ten eligible slices, one contributes nine, one contributes six, and one contributes three, giving a 12-slice deficit from the 324×10 maximum. ACDC LV similarly evaluates 361 LV-positive slices nested under 20 held-out patients. The main table is therefore a case-level audit, but the released `subject_level/*.json` files include the patient/subject unit IDs linked to `test_ids` and two cluster-aware checks. **(i) Cluster bootstrap over the case-scale statistic.** Resampling patients/subjects with replacement and keeping all slices within each sampled unit gives ACDC $\Delta=0.321$ [0.269, 0.374] and BraTS foreground $\Delta=0.315$ [0.283, 0.351], so the case-scale gap remains positive under unit-level resampling. **(ii) Subject-aggregated statistic.** Averaging each member’s per-slice Dice within patient/subject before computing Pearson gives smaller but still positive gaps: ACDC $\Delta=0.262$ [0.151, 0.413] and BraTS foreground $\Delta=0.214$ [0.182, 0.251]. The conservative reading is that slice dependence inflates the magnitude relative to subject means, especially on BraTS, but it does not explain the within-vs-cross separation away.

A12 BraTS sparse-class filtering caveat and released diagnostics

Per-class BraTS sensitivity can be inflated by sparse labels: if a class is absent from a slice, empty-vs-empty Dice can return a perfect score and obscure the foreground-error structure that matters for a segmentation audit. For this reason, the primary BraTS row uses the foreground derivative `seg>0` rather than a sparse NCR-only target. NCR filtering diagnostics from a separate sparse-class sweep are not used as numeric evidence here because their exact aggregate files are not included in the submission artifact. The released BraTS foreground train-size and multimodal diagnostics are instead reported later as M19 and M20 (Appendices A33.5–A33.6) from `results/_merged/late_brats_summaries.json`.




Figure 2: **Same-backbone bootstrap-bagging data swaps decorrelate the pool less than swapping the encoder family in this diagnostic.** Rose: within-subset cross-seed $\bar{r}$; sand: cross-subset any-seed $\bar{r}$; indigo: cross-FM reference from Table 1. The ratio panel reports backbone-swap decorrelation divided by data-swap decorrelation; all shown configurations exceed parity.

A13 Mechanism-probe scope

This block keeps released diagnostics separate from exploratory mechanism probes whose source prediction caches are not included in the submission artifact. Only released traces or aggregate JSONs are used as numerical support. The shared reading is descriptive: same-backbone perturbations can reduce correlation, but the released diagnostics do not make same-backbone pools behave like the primary cross-family references.

ACDC architecture caveat. Secondary ACDC exploratory tables used a separate 8-bundle split and architecture-pair subsets that are not included as replayable aggregate artifacts. The released support for the architecture-caveat wording is instead the same-architecture-family ViT / different-pretraining diagnostic in Appendix A30 and the expanded same-family cross-checkpoint decomposition in Appendix A31.

Random-init probes. DINOv2-B random-initialisation probes used secondary splits and source prediction caches that are not included in the submission artifact. The released no-pretraining evidence in this bundle is limited to the CNN random-init diagnostic in Appendix A33.2, which should be read as a descriptive same-recipe sanity check rather than a causal isolation of architecture, optimiser, and data effects.

LoRA and cross-architecture probes. LoRA probes and ViT-B cross-architecture random-init sweeps require source prediction caches and aggregate summary files that are not included in the submission artifact. We do not tabulate their values or use them as bundle-supported evidence. The released same-backbone levers are the head-capacity sweep, UNet-skip replication, fold-reslicing diagnostic, CNN random-init diagnostic, and M22 MC dropout.

A14 Same-backbone, different-data bootstrap-bagging audit

For each (FM, dataset) pair, four bootstrap replicates of the training set (image-level on RIGA/Kvasir; subject-level on BraTS) are crossed with four decoder seeds. Bootstrap data-swap reduces within-pool $\bar{r}$ by only 0.05–0.12, while backbone swap reduces it by 0.33–0.50. This supports the scoped claim that, in the frozen-feature protocol, resampling the training cases does not provide the same decorrelation as changing encoder family.

A15 Distribution-shift audit

We evaluate one released cross-distribution robustness diagnostic for the within–cross correlation gap using an 8-bundle RIGA pool and the frozen-feature protocol from Section 4.

RIGA leave-MESSIDOR-out (cross-vendor fundus). RIGA+ [Almazroa, 2018] combines three imaging sources: *BinRushed* (Bin Rushed Ophthalmic Center, Riyadh, Saudi Arabia), *Magrabia* (Magrabi Eye Center, Riyadh, Saudi Arabia), and *MESSIDOR* (French public corpus, different

cameras and populations). We train the frozen-FM + 1×1 conv decoder on BinRushed+Magrabia (union) and evaluate on MESSIDOR (454 cases after local preprocessing/eligibility filtering from the 460-image MESSIDOR subset) as cross-vendor out-of-distribution. Per-case traces for this diagnostic are released under `results/_merged/per_case_dice/riga_cup_ood_messidor`. The released scorer uses a fixed 8-bundle $\times$ 4-seed pool without a functional-floor change between ID and OOD rows. Under that directly reproducible convention, the ID row has within/cross/gap 0.659/0.282/0.377 and the MESSIDOR OOD row has 0.743/0.390/0.353.

Pool-level summary. In this fixed-pool RIGA cross-vendor diagnostic, the within-cross gap changes by -0.024 and remains positive. Both within and cross correlations are higher on the MESSIDOR row, so this diagnostic should be read only as evidence that the gap survives one released cross-vendor stress test, not as evidence that distribution shift generally contracts dependence. Broader cross-scanner, cross-demographic, and longitudinal shifts remain out of scope.

Table 18: **RIGA cross-vendor distribution-shift pool-level $\bar{r}$ under the released fixed 8-bundle scorer.** Within = same-FM different-seed $\cup$ same-family different-FM (DINOv2-B/S and ResNet-18/50); cross = different-family different-FM. Brackets are case-level bootstrap 95% CIs ($B=500$). Source: `results/_merged/diagnostics/distshift_riga_messidor.json`.

Shift regime	Within- $\bar{r}$	Cross- $\bar{r}$	Gap
ID reference (fixed 8-bundle pool)	0.659 [0.603, 0.715]	0.282 [0.197, 0.369]	0.377 [0.301, 0.463]
OOD MESSIDOR (fixed 8-bundle pool)	0.743 [0.722, 0.765]	0.390 [0.345, 0.430]	0.353 [0.318, 0.388]

This diagnostic uses a paired fixed 8-bundle RIGA pool and should not be read as a case-identical comparison to the Magrabia-source held-out primary RIGA rows in Table 1.

A16 UNet-with-skip decoder replication

The paper’s primary decoder is a 1×1 conv with $C+1$ trainable parameters for C frozen channels (for example, 769 parameters for a 768-channel feature map), which is close to a linear probe on frozen features. Dense decoders with skip connections—nnU-Net, U-Net, SegFormer—are closer to deployed medical segmentation. We therefore ran a separate M4 diagnostic in which the frozen backbone is unchanged but the head is replaced by a multi-scale UNet-style decoder over four encoder feature levels. This is a *two-seed diagnostic grid* (seeds 42/43), not an inferential table with the same status as the four-seed primary pool; the reported correlation is the single seed-pair Pearson coefficient for each cell.

Architecture. For hierarchical CNNs (ConvNeXt-T, ResNet-50) the decoder uses the native feature pyramid as skip sources; for ViT-style encoders (DINOv2-B/S and BiomedCLIP ViT) it uses sampled intermediate token maps reassembled into a four-level pyramid. The released head first applies lateral 1×1 projections to a common decoder width, then runs a top-down bilinear-upsample path with three additive skip merges (deep $\rightarrow$ mid $\rightarrow$ shallow $\rightarrow$ shallowest), two conv-BN-ReLU blocks after each merge, a final conv-BN-ReLU after upsampling to the output grid, and a final 1×1 output conv. Training matches the frozen-feature recipe where possible: Adam lr 10^{-3} , batch 16, 30 epochs, BCE, seeds {42, 43}.

Finding. M4 improves mean Dice substantially over several weak 1×1 probe cells, but it does not remove same-backbone co-failure: across the 20 evaluated cells, the seed-pair correlation remains 0.774–0.979. On rows where the matched 1×1 seeds 42/43 are both available, UNet-skip usually lowers seed-pair r modestly, but the effect is heterogeneous: Kvasir ResNet-50 drops 0.973 $\rightarrow$ 0.774, while RIGA Cup DINOv2-B is flat (0.965 $\rightarrow$ 0.967) and RIGA Cup ResNet-50 increases (0.908 $\rightarrow$ 0.952). The BraTS cells are consistent with the same boundary: DINOv2-B/S UNet-skip reaches mean Dice 0.875–0.877, but seed-pair r remains 0.886–0.936. Thus the conclusion is narrower than a decoder-family claim: a stronger skip decoder can improve accuracy and sometimes reduce seed correlation, but under this frozen-backbone two-seed grid it does not make same-backbone pools resemble the cross-family references in Table 1. A full nnU-Net v2 5-fold audit on RIGA Cup is reported in Appendix A17.

Table 19: **M4 UNet-with-skip diagnostic grid.** Each cell uses the two completed seeds 42/43 from `results/_merged/per_case_dice_m4`; “seed-pair r ” is therefore one Pearson coefficient, not a bootstrapped four-seed within-family estimate.

Task	Backbone	Mean Dice	Seed-pair r
Kvasir	DINOv2-B	0.883	0.877
Kvasir	DINOv2-S	0.866	0.919
Kvasir	BiomedCLIP	0.806	0.926
Kvasir	ConvNeXt-T	0.860	0.775
Kvasir	ResNet-50	0.816	0.774
ACDC LV	DINOv2-B	0.864	0.979
ACDC LV	DINOv2-S	0.863	0.970
ACDC LV	BiomedCLIP	0.827	0.950
ACDC LV	ConvNeXt-T	0.885	0.917
ACDC LV	ResNet-50	0.882	0.945
RIGA Cup	DINOv2-B	0.902	0.967
RIGA Cup	DINOv2-S	0.894	0.893
RIGA Cup	BiomedCLIP	0.880	0.892
RIGA Cup	ConvNeXt-T	0.895	0.970
RIGA Cup	ResNet-50	0.881	0.952
BraTS foreground	DINOv2-B	0.877	0.936
BraTS foreground	DINOv2-S	0.875	0.886
BraTS foreground	BiomedCLIP	0.858	0.929
BraTS foreground	ConvNeXt-T	0.895	0.938
BraTS foreground	ResNet-50	0.889	0.934

738 A17 Fold-reslicing and nnU-Net complements

739 **Motivation.** nnU-Net [Isensee et al., 2021] is the de facto standard for medical image segmenta-
740 tion; its 5-fold cross-validation ensemble is a routinely-used uncertainty / robustness baseline. The
741 full nnU-Net pipeline is *self-configuring* (2D/3D selection, preprocessing, patch-based sampling, de-
742 fault geometric/intensity augmentation, cascaded network, post-processing). We keep two evidence
743 surfaces separate: (i) the released M15 frozen-feature fold-reslicing diagnostic, which tests split sen-
744 sitivity with one decoder seed per fold and therefore cannot estimate within-fold versus cross-fold
745 same-backbone pairs; and (ii) released task-level nnU-Net summary JSONs, which are full-pipeline
746 complements but not case-identical contrasts to the primary FM rows.

747 **Released M15 fold-reslicing design.** M15 re-slices the frozen-feature cache into 5 folds
748 and reruns the 8-bundle pool {DINOv2-B/S, BiomedCLIP, CLIP-L, ConvNeXt-T, EfficientNet-
749 B0, ResNet-50/18} on ACDC LV and RIGA Cup, single seed per fold cell. Be-
750 cause the released artifact contains one seed per fold, this diagnostic reports per-fold
751 cross-bundle $\bar{r}$ and pool-mean Dice, not a 20-decoder fixed-backbone within/cross esti-
752 mator. Source: `results/_merged/diagnostics/m15_cv_summary.json`, regenerated by
753 `code/scripts/compute_m15_cv_summary.py`.

Table 20: **M15 5-fold CV: per-fold cross-bundle $\bar{r}$.** 8 bundles $\times$ 5 folds; single seed per fold cell. Cross-bundle $\bar{r}$ is the Pearson correlation averaged over the 28 bundle pairs on the held-out fold.

Task	fold 0	fold 1	fold 2	fold 3	fold 4	mean $\pm$ std
ACDC LV	0.580	0.548	0.583	0.558	0.558	0.565 $\pm$ 0.015
RIGA Cup	0.359	0.317	0.325	0.315	0.261	0.315 $\pm$ 0.035

754 **Finding.** Train/test split sensitivity is small in this fold-reslicing check: cross-fold pool-mean Dice
755 is stable (ACDC 0.502 $\pm$ 0.022; Cup 0.610 $\pm$ 0.012), and per-fold cross-bundle $\bar{r}$ is stable across
756 folds (ACDC 0.565 $\pm$ 0.015; Cup 0.315 $\pm$ 0.035). This supports only the narrower statement that the
757 cross-bundle diagnostic is not dominated by one particular 5-fold re-slicing. It does *not* support a
758 claim about within-fold vs cross-fold same-backbone decorrelation, because those pair types require
759 multiple decoder seeds per fold that are not present in the released M15 traces.

760 **Full nnU-Net v2 complement.** To complement the frozen-feature proxy above and answer
761 whether nnU-Net’s full self-configuring pipeline decorrelates the pool more, we trained the full
762 nnU-Net v2 2D U-Net on RIGA Cup with a 520/224 split, using the generated v2 2D nnUNetPlans
763 configuration after `nnUNetv2_plan_and_preprocess` (self-configured preprocessing and architec-
764 ture, official default augmentation/loss schedule, 250 iterations per epoch). This split differs from the

Magrabia-source held-out FM row and is therefore a task-level complement. We train two seeds per fold via custom `nnUNetTrainerSeed13_100epochs` and `nnUNetTrainerSeed37_100epochs` subclasses that pre-seed random, np.random, torch, and cuda before any nnU-Net sampler touches the RNG; each is trained for 100 nnU-Net-epochs (25,000 iterations on two A100s in parallel), then all $5 \times 2 = 10$ models predict on the held-out 224-image test set via `nnUNetv2_predict` and per-case Dice is computed against the held-out labels staged in `labelsTr`. The local nnU-Net export used a split-index file to enumerate held-out IDs; the artifact bundle releases the summary arrays/JSONs for this complement but not that local split manifest.

Across all 10 models the mean Dice is 0.949 ± 0.001 (exceeding any of our frozen-feature FM+ 1×1 -conv combinations on RIGA Cup). The pairwise per-case Dice $\bar{r}$ is 0.972 within-fold over the 5 same-fold different-seed pairs and 0.925 cross-fold over the 40 different-fold any-seed pairs, with $\Delta_{\text{within-cross}} = 0.047$ and a case-level bootstrap 95% CI of [0.030, 0.069] that excludes zero. Full nnU-Net v2 is structurally a same-architecture, different-training-data ensemble on one 2D Plain-ConvUNet, so it is not comparable in kind to the primary frozen-FM cross-family reference, which swaps encoder bundles. The scoped conclusion is therefore: *within this task-level same-architecture nnU-Net complement, fold identity leaves a small but positive correlation gap*; we do not claim nnU-Net substitutes for or surpasses backbone-family diversification. The 100-epoch budget, single-task coverage, and 2D-only scope are closed in Appendix A18 (1000-epoch retrain on the same RIGA Cup setup; 3D fullres nnU-Net on ACDC; higher-resolution/Dice+BCE/aug frozen-FM replication).

Data-swap comparison caveat. Appendix A14 reports bootstrap-bagging within-pool decorrelation with four decoder seeds per bootstrap replicate. M15 instead reports single-seed cross-bundle correlations across held-out folds. These diagnostics both stress training data/split sensitivity, but they are not dose-matched estimators and are not combined into a causal ranking.

A18 Segmentation-recipe and nnU-Net scope audit

Our main protocol uses 224×224 2D inputs, frozen backbone, 1×1 -conv decoder, BCE loss, no augmentation, and 30-epoch head training. Routine medical-imaging pipelines often run with compound Dice+CE loss, geometric augmentation, 3D input, and full-length (1000-epoch) training, so we test whether the within/cross correlation structure we report is specific to the simplified protocol. We close these axes as task-level complements; the RIGA nnU-Net rows use a different 520/224 split rather than the Magrabia-source held-out FM split, so they should not be read as case-identical contrasts.

Dimensionality axis: 3D fullres nnU-Net at 1000 epochs on ACDC. Training nnU-Net v2 in its self-configuring `3d_fullres` setting on ACDC with a seed-42 105/45-patient split (210 training, 90 test volumes) using `nnUNetTrainerSeed{13,37}_1000epochs` (seed injected in `on_train_start` before sampler initialisation) gives 10 fully-trained 3D models across 5 folds. `nnUNetv2_predict` on all 90 held-out volumes yields per-case Dice on mean foreground and per class. Across 10 models the mean Dice is 0.920 ± 0.001 , within the range reported for strong nnU-Net baselines on ACDC. On per-case mean Dice the within-fold (5 pairs) vs cross-fold (40 pairs) gap is $\Delta = 0.043$ [0.021, 0.085] with within $\bar{r} = 0.978$ and cross $\bar{r} = 0.935$; per class, $\Delta_{\text{RV}} = 0.041$ [0.011, 0.110], $\Delta_{\text{Myo}} = 0.029$ [0.017, 0.046], $\Delta_{\text{LV}} = 0.067$ [0.033, 0.098]. The artifact bundle releases the aggregate summary for this 150-patient complement but not the exact split manifest; the released public ACDC converter targets patients 001–100 for the primary 2D row. We therefore treat this 3D row as a task-level scope summary rather than a from-bundle training recipe. In this ACDC nnU-Net complement, the 2D→3D fullres shift changes the small fold-gap regime by less than 0.01 relative to the 2D 100-epoch baseline ($\Delta = 0.047$ in Appendix A17).

Loss / augmentation axis: RIGA Cup frozen-FM at encoder-native 224×224 . We retrain the seven-bundle M14 RIGA Cup subset (DINOv2-B/S, CLIP ViT-L/14, ConvNeXt-T, EfficientNet-B0, ResNet-50/18) with four seeds each under a more standard segmentation loss/augmentation recipe: encoder-native input size 224×224 , composite soft Dice+BCE loss, paired random horizontal flip and rotation $\pm 15^\circ$, for 30 epochs. The frozen encoder wrappers use 224×224 inputs; we therefore treat this as a loss/augmentation test, not a resolution test. All other factors (frozen backbone, 1×1 -conv segmentation head, Adam 10^{-3}) match the main protocol. Per-FM mean Dice ranges from 0.526 (ResNet-18) to 0.758 (DINOv2-B). Under the same symmetric estimator used for the

main table, within-family $\bar{r}$ is 0.903 (58 within pairs, including same-FM seed pairs and DINOv2-B $\times$ S), cross-FM $\bar{r}$ is 0.556 (320 cross pairs), and the within/cross gap is $\Delta=0.347$ [0.275, 0.428]. The gap is smaller than the main RIGA Cup interval (0.473 [0.395, 0.556]) but remains positive; loss/augmentation changes do not remove the gap.

Training-length axis: full nnU-Net v2 at the 1000-epoch default on RIGA Cup. We retrain the 10-model RIGA Cup 2D nnU-Net pool from Appendix A17 at the nnU-Net v2 default 1000 epochs, using the same seeded trainer pattern (nnUNetTrainerSeed{13,37}_1000epochs) on the same 520/224 split with identical dataset/plan files. Mean Dice across the 10 models is 0.946 ± 0.0004 (vs 0.949 ± 0.001 at 100 epochs). Within-fold $\bar{r}=0.983$ (5 pairs), cross-fold $\bar{r}=0.949$ (40 pairs), and $\Delta=0.034$ [0.021, 0.058]. Relative to the 100-epoch run both estimators tighten slightly (within $0.972 \rightarrow 0.983$; cross $0.925 \rightarrow 0.949$) and the within/cross gap contracts from 0.047 to 0.034; measured predictions are more correlated under longer training, and the gap does not disappear. Ten-fold longer training therefore does not change the replicability structure of nnU-Net on this task-level complement; it is not a case-identical comparison to the Magrabia-source held-out FM row.

Synthesis. Across these scope analyses, the frozen-FM loss/augmentation row remains in the same large-gap regime as the RIGA Cup main row, while the same-architecture nnU-Net fold axis remains in the small-gap regime ($\Delta \approx 0.04$). Because the nnU-Net RIGA rows use a different split, the synthesis is a task-level scope check rather than a paired case-identical comparison. These rows do not support a simpler explanation that the main-table pattern is solely an artifact of BCE loss, no augmentation, short training, or the tested full-pipeline nnU-Net complements.

Table 21: **Scope audit.** Within/cross gap at simplified vs standard segmentation-recipe protocols on task-level complements. The frozen-reference row uses the same seven candidate FMs as the Dice+BCE+augmentation row before the primary functional-floor exclusion; the primary post-floor RIGA Cup row is Table 1. The RIGA nnU-Net rows use a different 520/224 split and are not case-identical comparisons to the Magrabia-source held-out FM row.

Axis	Protocol	Task	Within	Cross	Δ	95% CI
Frozen reference	7-bundle candidate set, 224, BCE, no aug, 30 epochs	RIGA Cup	0.798	0.266	0.532	[0.456, 0.614]
Loss / aug	7-bundle frozen, 224, Dice+BCE, aug, 30 epochs	RIGA Cup	0.903	0.556	0.347	[0.275, 0.428]
Baseline (App. A17)	nnU-Net 2D, 100 epochs	RIGA Cup	0.972	0.925	0.047	[0.030, 0.069]
Training length	nnU-Net 2D, 1000 epochs	RIGA Cup	0.983	0.949	0.034	[0.021, 0.058]
Dimensionality	nnU-Net 3D fullres, 1000ep	ACDC (mean)	0.978	0.935	0.043	[0.021, 0.085]
	per-class RV		0.967	0.927	0.041	[0.011, 0.110]
	per-class Myo		0.987	0.957	0.029	[0.017, 0.046]
	per-class LV		0.978	0.911	0.067	[0.033, 0.098]

A19 Held-out multi-task stress test

Protocol. In addition to the four primary benchmark datasets (five task rows: Cup, Disc, ACDC, Kvasir, BraTS foreground), we evaluate five *held-out* tasks covering additional imaging domains: MSD Hippocampus [Antonelli et al., 2022] (T1 MRI, native 3-class mask {background, anterior, posterior} binarized to {background, any-hippocampus}), ISIC-2018 Task 1 Lesion Boundary [Codella et al., 2019] (dermoscopy, 2D RGB, $N=80$ test cases held out from a 400-image random subset of the 2,594-image training release), MSD Heart (T2 MRI cardiac), MSD Prostate (T2 MRI prostate), and MSD Spleen (CT portal). This held-out stress test uses the same 11 generic encoder bundles and four decoder seeds as the primary source pool where possible, with the same 1×1 -conv decoder, no per-task hyperparameter tuning, and no model substitution. These are task-specific 2D binary derivatives, not official challenge-test evaluations. The Dice ≥ 0.30 functional floor is applied uniformly, and the released aggregate summary records 220/220 expected per-seed traces.

Interpretation. The 11-bundle held-out protocol yields three classes of outcome: (i) ISIC-2018 and Hippocampus both show direction-positive gaps under the same broad estimator family as the primary audit, but without case-bootstrap CIs and with smaller held-out support than the primary rows. (ii) Prostate and Spleen illustrate why held-out diagnostics are scope boundaries rather than headline evidence: Prostate is direction-positive but has only six test cases, while Spleen passes the floor yet has a near-zero gap and $M_{\text{eff}} \approx 1$ because the members fail together. (iii) Heart does

Table 22: **Held-out multi-task stress test** (11 generic encoder bundles $\times$ 4 decoder seeds; aggregate diagnostic only). Hippocampus and ISIC-2018 are direction-positive, Prostate is positive but has only six test cases, Spleen is essentially near-zero despite passing the Dice floor, and Heart has no functional bundle at the Dice $\geq$ 0.30 floor. These rows are secondary scope evidence, not primary release evidence.

Task	N	Func./avail.	Within	Cross	Δ	M_{eff}
<i>Direction-positive point estimates with enough support to read qualitatively</i>						
MSD Hippocampus	52	11/11	0.909	0.419	0.491 [§]	3.94
ISIC-2018 [‡]	80	11/11	0.918	0.511	0.406 [§]	2.72
<i>Scope-limit rows: tiny held-out support or non-functional pool</i>						
MSD Prostate	6	5/11	0.788	0.642	0.146 [§]	1.23
MSD Spleen	8	11/11	0.983	0.979	0.003 [§]	1.02
MSD Heart	4	0/11	—*	—*	—*	—

*No bundle passes the Dice $\geq$ 0.30 floor on the four-case Heart derivative, so no within/cross gap is interpretable. [§]Values are point estimates from `results/_merged/diagnostics/heldout_11fm_summary.json`; case-level CIs are not computed for this aggregate diagnostic, and the MSD Prostate/Spleen rows have too few held-out cases to support a primary claim. [‡]ISIC-2018 uses a single random $N=400$ -image subset of the 2,594-image ISIC-2018 training release with $N=80$ test cases under the seed-42 80/20 split; we do not report an independent redraw and we do not use the official ISIC test server in this study.

not yield a functional pool under the frozen-feature 2D derivative. Thus the held-out stress test is compatible with transfer to some external rows, but it also records where the protocol is too weak or too saturated to audit monoculture.

Operational 8-bundle held-out probes. The following recovery, pool-size, abstention, and pool-selection diagnostics use an 8-bundle held-out candidate set stored under `results/_merged/diagnostics/held_out/`. They are retained as operational-method sensitivities, not as the provenance for Table 22 above.

CLIP/BiomedCLIP normalization robustness. Uniform ImageNet normalization can be a potential artifact for BiomedCLIP and CLIP-L because both have their own normalization statistics. We therefore evaluated those two FMs on Hippocampus and ISIC-2018 with the correct CLIP normalization. The released `summary_clipfix.json` reports BiomedCLIP mean Dice 0.840 (Hippocampus) / 0.908 (ISIC-2018), while CLIP-L remains below the functional floor on Hippocampus (mean Dice $<10^{-10}$) and reaches 0.106 on ISIC-2018, still below the 0.30 floor. Relative to `summary_5tasks.json`, the aggregate Δ changed by 0.002 (Hippocampus 0.731 $\rightarrow$ 0.729) and 0.003 (ISIC-2018 0.552 $\rightarrow$ 0.555). In this held-out diagnostic, correcting CLIP/BiomedCLIP normalization does not materially change the Hippocampus/ISIC point estimates.

Held-out recovery-rate sensitivity. To complement the primary-task $\Delta_{\text{rec}} = +0.29$ (Cup) / $+0.26$ (ACDC) grand means, we computed an exploratory held-out recovery gap at $\tau=0.85$ on the single ISIC-2018 subset: $\Delta_{\text{rec}} = +0.269$ [$+0.212, +0.331$] under a case-resampling calculation saved as `recovery_CI.json`. This supports the direction of the operational gap on one dermoscopy subset, but it is not a deployment-calibrated guarantee and does not resolve the held-out estimator-dependence above. Hippocampus at $\tau=0.85$ is recovery-saturated and gives $\Delta_{\text{rec}} = -0.01$ [$-0.03, +0.00$] (diverse and mono both near-zero recovery because only 2/7 functional bundles individually pass Dice $\geq$ 0.85). Lower thresholds are not evaluated here.

Pool-size sweep on held-out tasks. How does diverse-vs-mono conditional recovery scale with M ? For $M \in \{2, 3, 4, 5, 6\}$ on the two functional held-out tasks (Table 23), diverse pools beat monoculture at every $M \geq 3$ on the single ISIC-2018 subset ($\Delta_{\text{rec}}=0.27\text{--}0.46$), while Hippocampus at the paper’s $\tau=0.85$ threshold is recovery-rate-saturated (all pools <0.11) because only 2 of 7 functional bundles reach Dice $\geq$ 0.85 individually. These held-out operational analyses are exploratory: ISIC provides one supportive dermoscopy subset, Hippocampus is threshold-saturated at $\tau=0.85$, and the $M=5, 6$ monoculture rows are proxy constructions using the available four seeds rather than true larger monoculture pools. Released aggregate data: `results/_merged/diagnostics/held_out/pool_size_sweep.json`.

Calibration-tuned disagreement-based abstention (an exploratory alternative lever, not a conformal guarantee). A natural operational question is whether standard abstention-based uncertainty gating closes the monoculture gap with modest additional computation. We calibrate a per-task disagreement-score threshold on 20 cases (inter-member Dice standard deviation as score; grid-search the quantile q to meet a target case-risk $\alpha=0.1$ on calibration; apply the se-

Table 23: **Pool-size sweep.** Conditional recovery rate $P[\text{ens pass} \mid \text{any fail}]$ at $\tau=0.85$ vs pool size M . On ISIC-2018 diverse pools maintain a large gap across M ; on Hippocampus $\tau=0.85$ is above most functional bundles’ marginal Dice so the recovery-rate estimator saturates.

M	Hippocampus			ISIC-2018		
	mono	diverse	Δ	mono	diverse	Δ
$2^\dagger$	0.000	0.000	+0.00	0.000	0.000	+0.00
3	0.069	0.107	+0.04	0.111	0.498	+0.39
4	0.052	0.039	−0.01	0.091	0.359	+0.27
$5^\ddagger$	0.052	0.078	+0.03	0.091	0.553	+0.46
$6^\ddagger$	0.052	0.027	−0.03	0.091	0.454	+0.36

[†] $M=2$ is a definition artifact under the strict-majority ($>M/2$) rule conditioned on any-fail: recovery is zero by construction whenever $M=2$, regardless of pool composition. We include the row for completeness. [‡] At $M=5, 6$ the monoculture entries are computed on the $\min(M, S)=4$ available seeds per FM, so they do not represent true $M=5, 6$ mono pools; the diverse entries do use legitimate $M=5, 6$ distinct-family compositions.

lected threshold to the remaining test cases). *This is a tuned heuristic, not split-conformal prediction* [Angelopoulos and Bates, 2021]: we use a grid search over q that terminates on the first calibration risk meeting α and use inter-member variance rather than a proper conformity score, so no finite-sample risk guarantee holds. The pool is also chosen per-task on the whole task mean Dice (family-representative highest-Dice FMs), so the diverse-pool selection step leaks task labels—this biases the comparison toward diverse pools. We report it as a sensitivity. On ISIC-2018 at matched abstention ≈ 0.74 , diverse pools reach test-risk 0.129 ± 0.027 vs. monoculture grand mean 0.174 ± 0.209 (conformal_vs_meff.json). The *point estimate* is $\sim 25\%$ lower for diverse, but the monoculture grand-mean SD (0.21) is $\sim 8\times$ the diverse SD (0.027); the per-FM monoculture risk means span 0.000–0.616 across the seven functional FMs, so the grand-mean comparison is highly composition-sensitive and we do not claim a tight ranking. Hippocampus cannot hit $\alpha=0.1$ for either pool, giving residual risk 0.694 (diverse) vs. 0.665 (mono grand mean) ($\Delta = +0.03$, with mono SD 0.28, indistinguishable). A label-free conformal procedure with a proper conformity score is not evaluated here. Released aggregate data: results/_merged/diagnostics/held_out/conformal_vs_meff.json.

Calibration-set pool-selection probe on held-out tasks. We operationalise a small calibration exercise: given only a 20-case calibration set on a held-out task, can a rule pick a 4-of-8 FM pool from the same secondary held-out candidate set that beats random? We evaluate five rules—(R1) random, (R2) top-4 by mean Dice on calibration, (R3) max- M_{eff} , (R4) quality-gated max- M_{eff} (Dice ≥ 0.30 floor then max- M_{eff} among qualifying), (R5) diverse-by-architecture (2 best-Dice ViT + 2 best-Dice CNN; reported as an exploratory hand-crafted baseline, not a general proposal)—and report conditional recovery rate $P[\text{ens pass} \mid \text{any fail}]$ at $\tau=0.85$ on the remaining held-out cases, averaged over 100 calibration-set bootstraps (Table 24). All rule selectors and the recovery estimator use seed-42 per-case Dice (not seed-averaged), matching a one-model-per-FM pool construction. Hippocampus has lower effective recovery support than ISIC at $\tau=0.85$, so its recovery numbers are exploratory. In this diagnostic, simple top-4-by-Dice has the highest mean on both tasks (Hippocampus 0.161 vs. 0.043 random; ISIC-2018 0.581 vs. 0.378 random), while max- M_{eff} alone is worse than random (0.011 / 0.327). This supports the narrower point that M_{eff} is a dependence readout rather than a standalone operational objective; it is not a validated clinical pool-selection rule. Released aggregate data: results/_merged/diagnostics/held_out/pool_selection_challenge.json.

Table 24: **Pool-selection probe on held-out tasks.** Conditional recovery rate at $\tau=0.85$ (mean over 100 calibration bootstraps, 20 cases each). Top-4-by-Dice has the highest mean on both tasks; naive max- M_{eff} underperforms random.

Rule	Hippocampus	ISIC-2018
(R1) Random 4-of-8	0.043 ± 0.050	0.378 ± 0.125
(R2) Top-4 by calibration Dice	0.161 ± 0.017	0.581 ± 0.030
(R3) Max M_{eff}	0.011 ± 0.011	0.327 ± 0.100
(R4) Quality-gated max M_{eff}	0.016 ± 0.014	0.381 ± 0.112
(R5) Diverse-by-arch (2 ViT + 2 CNN)	0.096 ± 0.017	0.445 ± 0.036

Held-out estimator note. Split-stability and alternate-estimator numerical summaries whose exact aggregate files are not included in the submission artifact are omitted from the numeric evidence. This section therefore relies on the released 11-bundle held-out aggregate and the released operational JSONs listed above.

ISIC-2018 single-subset note. The ISIC-2018 row above is a single random $N=400$ -image subset of the ISIC-2018 Task 1 training release with $N=80$ held-out test cases under a seed-42 80/20 split. We do not report an independent replication of this subset; a fully random redraw of the subset or use of the official ISIC evaluation server would be a natural next step but is out of scope for this paper.

A20 MC dropout diagnostic scope

An unbundled MC-dropout probe suggested lower any-draw correlations, but its traces are not included in the submission artifact and its scorer differs from the released M22 averaged-readout diagnostic. We therefore do not tabulate it or use it as evidence. The released MC-dropout diagnostic is Appendix A33.7: under the deterministic-vs-MC-averaged readout, $p=0.1$ – 0.4 reduces within-FM $\bar{r}$ only from roughly 0.97 – 0.99 to 0.95 – 0.89 , so MC dropout is a weak within-family lever in this diagnostic and does not approach the tested backbone-family swap.

A21 Full fine-tuning scope

RIGA Cup full fine-tuning epoch-sweep traces and aggregate summaries are not included in the submission artifact. We therefore do not tabulate those values or use them as numerical evidence. The released adaptation evidence is limited to released traces and summaries named elsewhere in the appendix, including a small set of ACDC DINOv2-B full-fine-tuning provenance traces but not the RIGA epoch-sweep matrix; full-fine-tuning numbers not backed by released paths should be read only as exploratory mechanism notes and not as support for the paper’s claims.

A22 Same-backbone and scope diagnostic index

Table 25: **Same-backbone and scope diagnostic index.** Qualitative RIGA-oriented map anchored on the DINOv2-B 4-seed baseline ($\bar{r} \approx 0.96$) unless marked. Rows pool from released or summary-supported diagnostics; this is not an ordered causal hierarchy, CI-comparable ranking, or Table 1 provenance. Unbundled LoRA, full-FT epoch-sweep, and ViT random-init probes are omitted.

Diagnostic	Readout	Effect	Evidence
Mono baseline (same FM, 4 seeds)	0.96	0.00	RIGA diagnostic baseline
UNet-skip decoder (5 backbones, 4 tasks)	0.774–0.979	heterogeneous	Appx. A16, Tab. 19
M15 fold-reslicing cross-bundle check	0.315–0.565	split-stable	Appx. A17; <code>m15_cv_summary.json</code>
MC dropout $p=0.1$ – 0.4	0.89–0.95	–0.02 to –0.10	Appx. A33.7
Same-backbone bootstrap-bagging	0.85–0.92	–0.05 to –0.12	Fig. 2
RIGA cross-vendor shift	OOD w/c 0.743/0.390	gap 0.353; $\Delta_{\text{gap}} - 0.024$	Appx. A15
Decoder capacity sweep (M3 average)	~ 0.69	~ -0.27	Appx. A33.1
Cross-family pretrained 8-bundle pool	0.38 [0.29, 0.46]	–0.58	RIGA mechanism diagnostic

A23 Leave-one-family-out scope

Leave-one-family-out summaries were used as exploratory robustness checks, but the exact aggregate source is not included in the submission artifact. We therefore omit the numeric table from the released evidence record. The primary robustness check for the primary rows is the released `diagnostics/audit_sensitivity.json`, which includes leave-one-task and leave-one-FM summaries under the same scorer path.

A24 Joint-failure lift at support limits

Joint-failure lift $L = P_{\text{empirical joint}} / \prod_m P_{\text{marginal}}$ is unstable when marginal fail rates are near 0 or near 1. When $P_{\text{marginal}} \approx 0.3$, independence-implied joint ≈ 0.01 ; small sampling variation in the empirical joint produces large lift ratios. Representative near-ceiling cells: Only the ACDC 2.8×

Table 26: **Representative joint-failure lift values at support limits.** Lifts $\geq 50\times$ are descriptive (support-limited, wide noise envelope) and are *not* paper-headline statistics.

Task / Pool	τ	Indep. joint	Empirical joint	Lift
RIGA Cup mono BiomedCLIP	0.85	0.0010	0.125	123.2×
RIGA Cup mono DINOv2-B	0.85	0.0047	0.183	38.8×
RIGA Cup diverse 4-bundle	0.85	0.0160	0.098	6.1×
ACDC LV mono DINOv2-B	0.70	0.0002	0.038	189×
ACDC LV diverse 4-bundle	0.85	0.135	0.378	2.8×

row is adequately supported by the $N \cdot \mathbb{E}[\text{AF}_{\text{ind}}] \geq 5$ rule. The RIGA Cup 6.1× and 38.8× rows are support-limited, and the 123.2× and 189× rows are extreme sparse-tail diagnostics.

A25 3D segmentation-architecture pilot

This section provides scope context only. Its per-case outputs and scorer version are not included in the submission artifact, so the pilot is not used as primary evidence for Table 1, the abstract, or the released artifact claims. It motivates why we avoid generalizing the frozen 2D FM-pool result to 3D segmentation-architecture pools.

We assembled a 3D segmentation-architecture pool for BraTS 2024 3D volumes (271 subjects, 216/55 train/test, seed 42) using three architectures in the implementation path: SegResNet (optionally initialised from the MONAI BraTS bundle), DynUNet (random-init), and UNETR (random-init). SwinUNETR random-init pilot cells were included initially but dropped after degenerate predictions across all seeds. For each retained (architecture, seed $\in \{42, 43, 44, 45\}$) cell, the pilot trained the full MONAI segmentation model for 100 epochs with Adam 10^{-3} and BCE loss; UNETR-style models used sliding-window inference where required by memory. Reported at three filter modes:

Table 27: **3D segmentation-architecture pilot on BraTS foreground.** Three functional-filter modes applied uniformly in the pilot. Strict filter (Dice ≥ 0.50) retains only SegResNet pretrained and DynUNet random-init and yields $\Delta=0.045$, substantially below the 2D BraTS foreground $\Delta=0.315$; this row is appendix scope context only.

Filter	FMs retained	Within	Cross	Δ [95% CI]
Mixed (all cells)	3 (12 cells)	0.772	0.595	0.177 [0.117, 0.248]
Dice ≥ 0.30	3 (12 cells, drop UNETR seed 43 degenerate)	0.919	0.703	0.216 [0.134, 0.271]
Dice ≥ 0.50 (strict functional)	2 (SegResNet pretrained, DynUNet random)	0.916	0.871	0.045 [0.015, 0.112]
Random-init only (DynUNet + UNETR)	2 random-init	0.724	0.479	0.245 [0.155, 0.351]

Interpretation. Within the strict functional regime (only converged pretrained + converged random-init CNN-style 3D architectures), the 3D segmentation-architecture pilot shows $\Delta=0.045$ —substantially below the 2D BraTS foreground $\Delta=0.315$. We treat this only as a scope-warning: the 2D frozen-feature monoculture claim should not be extended to 3D segmentation-architecture pools without a released, case-identical 3D benchmark. Random-init 3D architectures cross-correlate *more* than pretrained 3D architectures (within random 0.724 vs cross-architecture random 0.479, $\Delta=0.245$), so architecture-family effects again appear large in this pilot. The degenerate SwinUNETR pilot cells limit architecture coverage in the 3D pilot; expanded 3D foundation-model evaluation (Models Genesis, Vista3D, Triad-MRI) is not evaluated here.

A26 Matched-recipe adaptation scope

Matched-recipe adaptation factorials require source prediction caches and aggregate summary JSONs that are not included in the submission artifact. We therefore do not tabulate those values and do not use them as bundle-supported evidence. The released adaptation evidence is limited

to explicit released traces under `per_case_dice_adapted` and the aggregate summaries named in the surrounding appendix sections.

A27 Item-difficulty residualisation and held-out nulls

The main-text Section 5 reports a cross-fitted residualised Δ diagnostic on all five primary rows and keeps the two-task held-out-partition diagnostic below as a leakage warning.

Naive (leaked) null. Define difficulty $d_i = 1 - \frac{1}{K} \sum_{k=1}^K D_k^{(42)}(i)$ where $D_k^{(42)}$ is bundle k 's per-case Dice at seed 42, then residualise each $D_k^{(s)}$ against d by OLS and then evaluate $\bar{r}$. This instantiates one item-difficulty null inspired by Jo et al. [2026], but has a structural flaw: each bundle's Dice is a component of its own regressor (since d averages over all 8 bundles including k), so residuals are mechanically constrained to sum to approximately zero across bundles at each case. For K bundles at seed 42, cross-bundle residual correlations are pushed toward $-1/(K-1) \approx -0.14$, and we observe cross-family residuals of -0.088 (Cup) and -0.104 (ACDC), which sharply widens Δ compared to the raw statistic. This is a mechanical leakage artifact, not a cleaner measurement.

Two-task held-out-partition null. To remove the leakage we split the 8 bundles into halves A and B , estimate d from A 's seed-42 Dice only, and residualise B 's Dice on this d . The $\bar{r}$ statistics and Δ are then computed on B . The released diagnostic JSON records a fixed set of 10 held-out partitions and reports the mean and range of the resulting family-grouped Δ values across those partitions.

Table 28: **Difficulty-residualised Δ under two-task nulls.** Naive null leaks; the held-out-partition null avoids leakage on Cup/ACDC only; raw is reported for reference. Source: `r7_item_difficulty_holdout.json`.

Null	RIGA Cup Δ	ACDC Δ	Notes
Raw	0.499	0.600	Family grouping
Naive difficulty (all-bundle leaked)	0.886	0.772	Cross-family residuals ≈ -0.10 (leakage signature)
Held-out difficulty (released held-out partitions)	0.639	0.622	Family grouping; $n=10$ per task
Held-out range	[0.482, 0.823]	[0.472, 0.941]	Min / max across partitions

The two-task held-out estimate is above raw Δ on Cup (+0.14) and remains slightly above raw on ACDC (+0.02). Under this *one-directional* null, item difficulty does not explain the gap away. The naive inflation was a mechanical leakage artifact regardless.

All-row cross-fitted diagnostic. Because the held-out-partition diagnostic above covers only Cup and ACDC, the released `item_difficulty_robustness.json` adds an all-row cross-fitted diagnostic. For each trace, item difficulty is estimated from traces outside that trace's broad upstream family, that trace is residualised by OLS, and the Table 1 same/cross rule is recomputed on residuals. The residualised gap is positive on all five primary rows (Table 29). Residualisation changes the scale of correlations and can increase gaps, so this is a sign and attenuation/scale-sensitivity check, not a causal decomposition.

Table 29: **All-row cross-fitted item-difficulty residualisation.** Each row uses the paper's $\text{Dice} \geq 0.30$ functional pool and $B=2,000$ case bootstrap intervals from `item_difficulty_robustness.json`.

Task	Raw Δ	Residualised Δ	Residual CI	Mean fit R^2
RIGA Disc	0.638	0.874	[0.805, 0.920]	0.138
RIGA Cup	0.473	0.830	[0.782, 0.882]	0.318
Kvasir polyp	0.261	0.960	[0.895, 1.008]	0.663
ACDC LV	0.321	0.976	[0.950, 1.000]	0.607
BraTS foreground	0.315	0.921	[0.910, 0.932]	0.574

Residualisation caveat. The released held-out partitions are one-directional diagnostics rather than a symmetric enumeration over every possible difficulty/evaluation split. They should be read as evidence that this particular leakage-free residualisation does not absorb the gap, not as a two-sided

causal decomposition of difficulty, architecture, and pretraining. The all-row cross-fitted diagnostic broadens coverage to every primary row, but it still estimates difficulty from the same benchmark panel and therefore remains a robustness check rather than a pre-registered causal adjustment.

A28 Correlation-equivalent effective pool size M_{eff}

Definition. For a pool of M members whose per-case failure indicators $F_1, \dots, F_M$ at threshold τ are exchangeable with mean pairwise Pearson correlation ρ_{fail} , the variance of the per-case mean indicator is $\sigma_F^2[1 + (M - 1)\rho_{\text{fail}}]/M$ where σ_F^2 is the marginal variance. Equating this with the variance of a hypothetical independent pool of size M_{eff} , i.e. $\sigma_F^2/M_{\text{eff}}$, gives

$$M_{\text{eff}} = \frac{M}{1 + (M - 1)\rho_{\text{fail}}}.$$

This is the equicorrelation approximation also used in climate-model ensembles (“effective number of models”). We report it as a second-moment readout of the audit; it is *not* a literal independent-member count.

Table 30: M_{eff} at $\tau=0.85$ across monoculture and diverse pools. Monoculture pools are each FM’s completed decoder seeds. Diverse rows enumerate every 4-bundle composition of the post-floor functional pool and every available one-seed-per-FM tuple. Undefined zero-variance failure-correlation tuples are excluded rather than imputed. Source: `results/_merged/m_eff/*.json`.

Task	Func. bundles	Mono mean [†]	Diverse ρ_{fail}	Diverse M_{eff}	Valid tuples
Kvasir	11	1.12	0.333 ± 0.104	2.05 ± 0.32	84,480 / 84,480
ACDC LV	10	1.14	0.298 ± 0.075	2.14 ± 0.24	53,760 / 53,760
BraTS foreground	10	1.13	0.273 ± 0.089	2.24 ± 0.32	53,760 / 53,760
RIGA Cup	9	1.16	0.246 ± 0.121	2.43 ± 0.63	31,488 / 32,256
RIGA Disc	11	1.17	0.126 ± 0.067	2.97 ± 0.44	84,480 / 84,480

[†]Mean across per-FM seed pools after excluding undefined zero-variance failure indicators. The RIGA Cup row has 7/9 defined mono pools and the RIGA Disc row has 10/11; the other rows have all mono pools defined.

Reading the table. The scorer uses the same estimand as the text: a diverse pool is four distinct FMs with one decoder-seed member from each FM. Under that definition, the nominal $M=4$ diverse pools yield $M_{\text{eff}} \approx 2\text{--}3$ under this equicorrelation readout, not four; same-FM seed pools are near one under the same calculation. The exact level varies by task and failure threshold, so we use M_{eff} as a dependence readout, not as a literal deployment objective.

Pairwise ρ_{fail} heterogeneity disclosure (equicorrelation caveat). The M_{eff} formula assumes constant pairwise correlation across member pairs (equicorrelated pool). In practice the per-pair ρ_{fail} varies across the bundles in each task. From the `m_eff/*.json` outputs, mono-pool ρ_{fail} across the per-task functional bundles has mean / SD / range: Kvasir mean 0.865, SD 0.103, range [0.601, 0.948] ($n=11$); ACDC LV mean 0.847, SD 0.122, range [0.518, 0.931] ($n=10$); BraTS foreground mean 0.860, SD 0.117, range [0.553, 0.958] ($n=10$); RIGA Cup mean 0.833 over 7 defined mono pools, and RIGA Disc mean 0.812 over 10 defined mono pools. Undefined rows occur when an FM’s $\tau=0.85$ failure indicator is constant, giving zero-variance Pearson. Diverse 4-of- K compositions span a wider ρ_{fail} distribution, which is why Table 30 reports means and standard deviations across all valid seed tuples rather than selecting a representative pool.

A29 Shared-backbone random-effects representation

Model. For each FM m in a frozen-feature pool and test case i , let the per-case score (or its logit) decompose as

$$s_m(i) = u(i) + a_{\text{arch}(m)}(i) + p_{\text{family}(m)}(i) + \varepsilon_m(i),$$

where u is task-inherent item difficulty, a is a zero-mean architecture-family random effect, p is a zero-mean pretraining-family random effect nested within architecture, and ε_m is an FM-specific adaptation noise. Assume u, a, p, ε are mutually uncorrelated with variances $\sigma_u^2, \sigma_a^2, \sigma_p^2, \sigma_\varepsilon^2$. Then

the pairwise covariance of $s_m, s_{m'}$ takes a nested form:

$$\text{Cov}(s_m, s_{m'}) = \begin{cases} \sigma_u^2 + \sigma_a^2 + \sigma_p^2, & (\text{same family}) \\ \sigma_u^2 + \sigma_a^2, & (\text{same architecture, different pretraining}) \\ \sigma_u^2, & (\text{different architecture}). \end{cases}$$

Diagnostic expectation. Under this second-moment description, and absent parameter cancellations, one expects the ordering

$$\rho_{\text{within-FM}} \geq \rho_{\text{same-arch-diff-pretraining}} \geq \rho_{\text{cross-arch}}$$

as a qualitative diagnostic rather than a theorem. Empirically on 5 same-architecture-family ViT variants (Table 3) we observe the ordering on Cup, and on ACDC after applying the functional filter: $0.988 \geq 0.807 \geq 0.639$; the all-5 ACDC row is a sensitivity to including nonfunctional CLIP-B. This is consistent with the predicted ordering and with trajectory/same-basin evidence from deep ensembles and model soups [Fort et al., 2019, Neyshabur et al., 2020, Wortsman et al., 2022]; Abe et al. [2024] is better read here as motivation that predictive diversity can remain constrained in high-capacity ensembles, not as a trajectory-bound theorem.

Implication for M_{eff} . When the binary failure thresholds preserve the same broad ordering as the continuous Dice scores, M_{eff} is monotone decreasing in the observed ρ_{fail} by definition. The empirical diagnostics therefore suggest the descriptive ordering $M_{\text{eff}}^{\text{within-FM}} \leq M_{\text{eff}}^{\text{same-arch}} \leq M_{\text{eff}}^{\text{cross-arch}}$ in the tested regime, but the paper estimates this quantity from released per-case traces rather than inferring it from labels alone. Monoculture pools are same-FM by construction and sit near the lower end of the measured M_{eff} range; diverse pools that cross both architecture and pretraining bundles sit closer to the upper end.

Caveats. This is a second-moment model at the per-case-score level, not a first-principles claim about SGD dynamics. The random-effects representation is descriptive; it gives a language for the observed ordering and a hook for bounding pool redundancy, not a proof that all shared-pipeline pools must have the same structure. The item-difficulty residualisation of Appendix A27 removes u and leaves $a + p + \varepsilon$; our held-out null is consistent with $a + p$ contributing after u is removed.

A30 Architecture-confound diagnostics

The main-text Table 3 reports the released 5-variant same-architecture-family ViT / different-pretraining diagnostic on Cup and ACDC. A secondary Kvasir-SEG replication used the same 5 ViT-B variants, but its aggregate JSON is not distributed with the submission artifact, so it is not tabulated as released evidence here.

Table 31: **Released same-architecture-family ViT / different-pretraining diagnostic.** Within-FM is the decoder-seed floor; same-architecture-family ViT cross-pretraining holds the broad architecture family fixed while varying pretraining. Source: released aggregate diagnostic results/_merged/diagnostics/r7_arch_confound_analysis.json; raw NPZ prediction caches are not included.

Task	ViT subset	Used	Seed floor	ViT cross-pretrain	Primary cross	Gap
RIGA Cup	all 5 ViTs	5/5	0.968	0.680	0.361	0.288
ACDC LV	all 5 ViTs	5/5	0.965	0.631	0.639	0.334
ACDC LV [†]	CLIP-B dropped	4/5	0.988	0.807	0.639	0.182

The same-architecture cross-pretraining separation is 0.18–0.29 across the released Cup/ACDC diagnostics under the paper’s functional filter ([†] ACDC row drops the non-functional CLIP-B at Dice=0.21<0.30, which otherwise compresses same-arch cross $\bar{r}$ and inflates the separation to 0.33). The direction is preserved, but the claim is mechanism-diagnostic rather than causal. Results are consistent with the random-effects representation in Appendix A29.

A31 Same-family cross-checkpoint decomposition

The main text explicitly cautions that the primary “same” column is dominated by same-bundle different-decoder-seed pairs. To test whether the signal collapses to that seed floor, we recompute

the released same-task per-case traces, including additional same-family checkpoint variants beyond the primary pool, under a three-stratum decomposition: (i) same-FM decoder-seed pairs, (ii) distinct checkpoints within a broad upstream family (DINOv2-B/L/S, OpenAI CLIP B/16-B/32-L/14, ConvNeXt-T/S/B, and ResNet-18/34/50/101 where functional), and (iii) cross-family pairs. This is a diagnostic broad-family grouping rather than a causal architecture/pretraining decomposition. The distinct-checkpoint same-family gap remains positive on all five primary rows, but it is substantially below the same-FM seed floor.

Table 32: **Released same-family cross-checkpoint decomposition.** Same-FM seed pairs form the upper dependence floor; distinct-checkpoint same-family pairs remain above cross-family pairs on all five primary rows. Source: `results/_merged/diagnostics/cross_checkpoint_family.json`.

Task	Same-FM	Xckpt same-family	Diagnostic cross	Xckpt-cross	95% CI
Kvasir	0.987	0.777	0.697	0.080	[0.043, 0.122]
ACDC LV	0.986	0.750	0.616	0.134	[0.113, 0.155]
BraTS foreground	0.977	0.682	0.589	0.093	[0.085, 0.100]
RIGA Cup	0.963	0.537	0.271	0.267	[0.204, 0.332]
RIGA Disc	0.971	0.421	0.194	0.227	[0.149, 0.293]

CI's are case-bootstrap intervals with $B=2,000$. The cross column is the diagnostic broad-family cross stratum from this expanded checkpoint grid, not the Table 1 analytic-pool cross-family column. For the slice-derived rows, subject-aggregated

Xckpt-cross gaps remain positive: ACDC 0.166 [0.091, 0.262] and BraTS 0.098 [0.078, 0.119]. BiomedCLIP is kept separate from OpenAI CLIP because it changes both pretraining corpus and objective. The OpenAI CLIP same-family row mixes CLIP-B/16 QuickGELU-style checkpoints with the released CLIP-L/14 `open_clip` standard-GELU dense wrapper, so that family-specific xckpt readout also absorbs this loader/activation difference.

A32 Release manifest and artifact metadata

The artifact bundle includes the primary-task result JSONs, source code, metadata, selected figure / appendix inputs, and selected aggregate diagnostic JSONs needed to audit the numerical claims from saved artifacts where redistribution is permitted. The primary rows are stored under `results/_merged`: the per-seed `per_case_dice` JSONs carry `test_ids`, `per_case_dice`, FM name, seed, and checksum fields; `within_cross/*.json`, `m_eff/*.json`, and `paper_table.json` store the primary summaries. The same directory also includes `subject_level/*.json` for the ACDC/BraTS cluster checks, `calibration_split/*.json` as split-half sensitivity (not the primary protocol), `provenance_summary.json` for source counts and case-unit metadata, and `diagnostics/*.json` for the pixel-error, matched-lift, architecture, quality-controlled pair-regression, item-difficulty, and held-out aggregate diagnostics used only as secondary evidence. The bundle does *not* include raw medical images, raw labels, full prediction NPZs, feature caches, or held-out raw prediction arrays; full split manifests beyond the case IDs embedded in the result JSONs are not included in this artifact release.

ID and source-label scope. Kvasir rows expose only image-level file identifiers in the released `test_ids`; no patient, video, or near-duplicate grouping identifier is available in the released trace bundle, so we do not claim patient-level duplicate control for that row. RIGA Cup/Disc rows retain Magrabi source-path labels such as `Magrabi/MagrabiFemale/image1.tif` and `Magrabi/MagrabiMale/image46.tif`; these labels are used only to make the $n=95$ Magrabi-source test split auditable and must not be used for individual profiling.

Encoder-loader boundary. The released traces should be read as frozen-feature probes through the bundled loader wrappers, not as official vendor inference runs. In particular, the `clip_vitl14` traces use the version-pinned `ViT-L-14/openai` path exposed by the installed `open_clip` package; `open_clip` also exposes a `ViT-L-14-quickgelu/openai` model definition and warns that OpenAI weights are associated with QuickGELU. We therefore treat CLIP-L/14 as the bundle's documented dense-feature wrapper and do not use it to make a claim about official OpenAI CLIP inference accuracy. The code keeps this loader unchanged so that the released per-case traces remain reproducible.

Functional-floor members are retained in the source files. The paper applies a $\text{Dice} \geq 0.30$ functional floor downstream in the estimator, not at result ingestion. Thus the merged `per_case_dice`

Table 33: **Primary per-case Dice inventory.** Merged members count all completed (bundle, decoder-seed) JSON entries under `results/_merged/per_case_dice`; Table 1 uses the original 44-trace primary source subset and applies the post-floor analytic pool shown in the final column.

Task	N	Merged bundles	Merged members	Table 1 source	Analytic pool
kvasir	120	17	68	44	11 bundles
acdc_lv	361	17	68	44	10 bundles
riga_cup	95	17	68	44	9 bundles
riga_disc	95	17	68	44	11 bundles
brats_wt	3228	17	68	44	10 bundles / 40 members

tree can list more completed JSONs than the analytic pool used by Table 1; the extra cross-checkpoint members support Appendix A31. For example, ACDC LV has 68 merged source JSONs but Table 1 retains the 44-trace primary source subset and 10 bundles after the floor. The authoritative row-level provenance for the table is `within_cross/*.json`, especially the `pool` and `fms_kept_after_floor` fields.

What else is in the bundle. In addition to the primary `per_case_dice` tree, the bundle includes merged ablation JSON trees (`per_case_dice_m3`, `per_case_dice_m4`, `per_case_dice_m14`, `per_case_dice_m15`, `per_case_dice_m19`, `per_case_dice_m22`, `per_case_dice/brats_wt_multimodal`, released random-init adapted-head traces, and a small set of ACDC DINOv2-B full-FT provenance traces), selected aggregate diagnostic JSONs under `results/_merged/diagnostics`, plus code sufficient to inspect and regenerate the included primary summaries and the M4/M19/M20 appendix summaries. Appendix-only diagnostics not listed here, raw prediction arrays, and full pixel-mask caches are not included in the supplement and are treated as reported diagnostics rather than fully bundled reproduction targets.

What is not in the bundle and why. Raw medical images, raw GT labels, feature caches, and full prediction NPZ caches are not included. For restricted datasets such as BraTS, the release is limited to derived per-case Dice and aggregate dependence readouts; users must obtain the raw data from the upstream provider. Larger derived caches and fuller split manifests are not included in the released artifact.

A33 Additional mechanism and scope diagnostics

This section reports seven matrix diagnostics (M3, M9, M14, M15, M19, M20, and M22) that close scope axes around the main result. Sources are individual per-case Dice JSONs under `results/_merged/per_case_dice_m{3,14,15,19,22}/`, `results/_merged/per_case_dice/brats_wt_multimodal/`, and `results/_merged/per_case_dice_adapted/rand_init/`. The BraTS summaries are generated by `code/scripts/compute_late_brats_summaries.py`. All numbers are computed directly from the per-case traces; per-FM means use seed averaging where multiple seeds are present, and Pearson $\bar{r}$ is symmetric multi-seed.

A33.1 Decoder-capacity sweep (M3)

We rerun the Section 6 (C1) decoder sweep at four explicit head designs spanning 769 to 3.94M trainable parameters: `linear_1x1`, `mlp_2layer`, `unet_lite`, and `transformer_decoder`. Two FMs (DINOv2-B, BiomedCLIP) $\times$ two seeds {42, 43} on RIGA Cup. Across the $\sim 5,000\times$ capacity range, pool-mean Dice ranges over 0.725–0.848 but cross-decoder within-FM Pearson $\bar{r}$ remains 0.69 on average ($n=48$ pairs across the four head designs and the two FMs)—above the RIGA Cup functional-pool cross-family reference of 0.361 (Table 1). Head capacity changes accuracy and some correlations, but does not eliminate same-backbone co-failure in this diagnostic.

A33.2 Random-init CNN decorrelation (M9)

To check whether within-vs-cross structure persists in the absence of pretraining, we rerun the released random-init diagnostic on two CNN backbones (ConvNeXt-Tiny, ResNet-50) trained from scratch, 8 seeds {42–49}, on ACDC LV and RIGA Cup. Even with no shared pretraining, within-FM (cross-seed) $\bar{r}$ stays clearly above cross-FM (different-architecture) $\bar{r}$: 0.916 vs. 0.830 on ACDC

Table 34: **M3 head-capacity sweep on RIGA Cup** (2 FMs $\times$ 2 seeds, 30-epoch frozen probe). Pool-mean Dice and trainable-parameter count per head design.

Head design	Pool-mean Dice	Trainable params
linear_1x1	0.725	769
mlp_2layer	0.810	197,121
unet_lite	0.848	3,319,105
transformer_decoder	0.805	3,940,865

1161 LV ($\Delta=0.086$), 0.841 vs. 0.720 on RIGA Cup ($\Delta=0.121$). The gap is much smaller than the
 1162 frozen-FM Table 1 numbers (0.834/0.361 on Cup, $\Delta=0.473$), but it is still positive, consistent with
 1163 architecture and shared recipe contributing to within-vs-cross structure in this diagnostic.

Table 35: **M9 random-init within/cross $\bar{r}$** . 2 CNN backbones $\times$ 8 seeds; 30 epochs from scratch; same 1×1 probe protocol.

Task	within-FM $\bar{r}$	cross-FM $\bar{r}$	Δ
ACDC LV	0.916	0.830	0.086
RIGA Cup	0.841	0.720	0.121

1164 A33.3 Clinical-recipe sensitivity (M14)

1165 The Section 4 protocol uses BCE loss, 224 inputs, no augmentation. To check whether the gap
 1166 survives a more clinically conventional recipe, we evaluate the seven-bundle subset {DINOv2-B/S,
 1167 CLIP-L, ConvNeXt-T, EfficientNet-B0, ResNet-50/18} $\times$ 4 seeds with combined soft Dice+BCE
 1168 loss and paired {hflip, $\text{rot} \pm 15^\circ$ } augmentation. All runs keep the encoder frozen, use the same 1×1 -
 1169 conv decoder, Adam 10^{-3} , and encoder-native 224×224 input. The original RIGA Cup M14 row
 1170 remains the primary loss/augmentation diagnostic in Appendix A18; we also report three task-level
 1171 diagnostic extensions (RIGA Disc, ACDC LV, and an ISIC-2018 Task 1 fixed 400-image held-out
 1172 subset) to check whether the positive gap is an isolated Cup result. ISIC here is a supplementary per-
 1173 case Dice diagnostic on a fixed 400-image subset and is not the official ISIC challenge test/server
 1174 protocol or threshold-Jaccard leaderboard score.

Table 36: **M14 standard segmentation-recipe (Dice+BCE+aug) on RIGA Cup, 7 FMs $\times$ 4 seeds**. Per-FM seed-averaged Dice; pool $\bar{r}$ at the bottom. CLIP-L uses the same caveated OpenCLIP ViT-L-14/openai dense-token wrapper as the primary bundle.

FM	mean Dice	σ (seed)
DINOv2-B	0.758	0.005
DINOv2-S	0.733	0.003
CLIP-L	0.720	0.006
EfficientNet-B0	0.626	0.002
ResNet-50	0.618	0.010
ConvNeXt-T	0.583	0.004
ResNet-18	0.526	0.027
Pool: within $\bar{r}=0.903$, cross $\bar{r}=0.556$, $\Delta=0.347$ [0.275, 0.428].		

Table 37: **M14 task-level extensions**. Same seven-bundle, four-seed Dice+BCE+augmentation recipe as Table 36. Dice columns summarise per-bundle seed-averaged Dice across the seven bundles. The ISIC row is a fixed 400-image-subset held-out diagnostic, not the official ISIC test/server evaluation.

Task	N_{test}	FM-mean Dice	FM min	FM max	Δ	95% CI
RIGA Cup	95	0.652	0.526	0.758	0.347	[0.275, 0.428]
RIGA Disc	95	0.853	0.756	0.922	0.617	[0.535, 0.693]
ACDC LV	361	0.565	0.389	0.740	0.326	[0.291, 0.366]
ISIC-2018 staged	80	0.864	0.801	0.918	0.473	[0.375, 0.562]

1175 Across all four M14 rows, the Dice+BCE+augmentation recipe leaves a positive within-vs-cross gap
 1176 ($\Delta = 0.326\text{--}0.617$). The magnitude varies by task: ACDC remains lower accuracy and lower-gap

than the RIGA/ISIC rows, while RIGA Disc shows the largest extension gap. These rows are treated as scope checks rather than primary evidence because they use a seven-FM subset and, for ISIC, a fixed subset/split; they indicate that the simplified BCE/no-augmentation protocol is not the sole source of the monoculture signal.

A33.4 Five-fold CV robustness (M15)

M15 is reported with the nnU-Net complements in Appendix A17 and Table 20. The key scope point is that the released M15 artifact has one decoder seed per fold; it supports split-stability of cross-bundle $\bar{r}$, not a within-fold vs cross-fold same-backbone estimator.

A33.5 BraTS foreground train-size sweep (M19)

M19 is a train-size diagnostic over the five BraTS foreground bundles with completed feature caches (DINOv2-B, BiomedCLIP, CLIP-L/14, EfficientNet-B0, ResNet-50). It uses one decoder seed (42) and fractions $\{0.10, 0.25, 0.50, 1.00\}$ of the same cached training set, so it is an *accuracy/saturation* check rather than a within-vs-cross dependence estimator. The result is non-monotone: the 5-bundle mean Dice rises from 0.579 at 10% to 0.621–0.622 at 25–50%, then stays flat/slightly lower at full data (0.617). Thus, in this cached five-bundle subset, adding more 2D FLAIR slices improves the weakest low-data setting but does not by itself create a new high-performing, decorrelated pool; the correlation claim still has to be audited directly rather than inferred from sample size.

Table 38: **M19 BraTS foreground train-size sweep** (5 cached bundles, seed 42). Mean/min/max Dice are across bundles at each train fraction; this table is not a within-vs-cross estimator because each bundle has one seed.

Train fraction	Cells	Mean Dice	Min Dice	Max Dice
0.10	5	0.579	0.418	0.721
0.25	5	0.621	0.455	0.750
0.50	5	0.622	0.436	0.751
1.00	5	0.617	0.458	0.737

A33.6 BraTS multimodal RGB derivative (M20)

M20 checks whether the BraTS foreground monoculture signal is an artifact of using a unimodal FLAIR derivative. We build a deliberately simple multimodal 2D derivative with the same held-out subjects and top-10 foreground-area slice rule as Table 1, but pack three MRI modalities into RGB channels: R=T1c, G=T2w, B=T2-FLAIR. This is *not* the official 3D BraTS challenge protocol; it is a controlled frozen-feature stress test of the input-channel choice.

The evaluated grid contains 20 cells (5 FMs $\times$ 4 seeds). Multimodal input narrows the dependence gap relative to the unimodal Table 1 BraTS row, but does not remove it: case-level within $\bar{r}=0.904$, cross $\bar{r}=0.654$, $\Delta=0.249$ [0.238, 0.262]. Aggregating each subject’s slices before correlation gives within $\bar{r}=0.946$, cross $\bar{r}=0.772$, $\Delta=0.174$. The conservative interpretation is therefore narrower than “BraTS input modality does not matter”: adding T1c/T2w raises cross-family agreement and shrinks the gap, yet same-backbone/family dependence remains positive under the same frozen-head audit.

Table 39: **M20 BraTS multimodal derivative** (R=T1c, G=T2w, B=T2-FLAIR; 5 FMs $\times$ 4 seeds). Dice is seed-averaged by FM; dependence uses the same Pearson estimator as Table 1.

FM	Mean Dice	Min Dice	Max Dice	Seeds
BiomedCLIP	0.713	0.669	0.739	4
DINOv2-B	0.703	0.675	0.721	4
DINOv2-S	0.702	0.690	0.714	4
CLIP-L/14	0.686	0.682	0.692	4
ResNet-50	0.461	0.439	0.473	4
Case-level pool: within $\bar{r}=0.904$, cross $\bar{r}=0.654$, $\Delta=0.249$ [0.238, 0.262].				
Subject-aggregated pool: within $\bar{r}=0.946$, cross $\bar{r}=0.772$, $\Delta=0.174$.				

1207 A33.7 MC dropout as a within-family decorrelator (M22)

1208 A natural pool-design question is whether MC dropout [Gal and Ghahramani, 2016] can act as
 1209 a lightweight within-family decorrelation lever. We rerun DINOv2-B on RIGA Cup with three
 1210 dropout rates $p \in \{0.1, 0.2, 0.4\}$, 4 seeds, 10 MC samples per case. The deterministic within-FM
 1211 cross-seed $\bar{r} \approx 0.97\text{--}0.99$; under MC averaging, $\bar{r}$ drops to 0.95 ($p=0.1$), 0.93 ($p=0.2$), 0.89 ($p=0.4$).
 1212 The largest tested drop is ≈ 0.10 and remains materially smaller than the same-architecture-family
 1213 ViT-vs-primary-cross separation seen on Cup (Table 3). MC dropout is therefore a weak within-
 1214 family lever in this diagnostic; consistent with Appendix A13, it does not approach the architecture-
 1215 family separation.

Table 40: **M22 MC dropout, DINOv2-B / RIGA Cup, 4 seeds, 10 MC samples per case.** Within-FM cross-seed $\bar{r}$ on deterministic vs. MC-averaged per-case Dice; mean MC predictive entropy.

Dropout p	MC mean Dice	Det. within $\bar{r}$	MC within $\bar{r}$	H_{pred}
0.1	0.666	0.972	0.947	0.062
0.2	0.657	0.972	0.925	0.067
0.4	0.633	0.987	0.893	0.076

1216 NeurIPS Paper Checklist

1217 1. Claims

1218 Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contribu-
 1219 tions and scope?

1220 Answer: [Yes]

1221 Justification: The abstract and introduction claim that *in the frozen and light-adaptation regime of medical*
 1222 *segmentation pipelines*, same-FM/same-recipe pools show higher per-case Dice-score correlation, used as
 1223 a failure-dependence proxy, than the Table-1 cross-family reference pools. The primary table uses task-
 1224 specific post-floor analytic pools: Kvasir 11 FMs, ACDC LV 10, BraTS foreground 10, RIGA Cup 9, and
 1225 RIGA Disc 11. The scope is explicitly restricted to retrospective public-benchmark dependence audits;
 1226 secondary pixel/lift readouts, released nnU-Net complements, and non-primary mechanism diagnostics are
 1227 diagnostic or scope context rather than causal clinical-outcome evidence.

1228 2. Limitations

1229 Question: Does the paper discuss the limitations of the work performed by the authors?

1230 Answer: [Yes]

1231 Justification: The main paper gives high-level limitations in Section 8, and the appendix expands them with
 1232 scoped protocol notes: frozen/light-adaptation only; unbundled LoRA and RIGA full-fine-tuning mecha-
 1233 nism probes retained only as transparency notes; released CNN random-init and small ACDC full-fine-
 1234 tuning provenance traces are diagnostic, not primary evidence; UNet-with-skip and task-level nnU-Net v2
 1235 complements reported but not fully case-identical dense 3D FM/full-pipeline baselines for every primary
 1236 row; the unbundled 3D architecture pilot is appendix scope context only; 2D axial slices at 224×224 for
 1237 the primary FM audit; MedSAM used only as a frozen vision-encoder probe; broader cross-scanner / cross-
 1238 demographic shifts remain untested; clinical Dice thresholds are operational conventions and many lift
 1239 thresholds are support-limited.

1240 3. Theory assumptions and proofs

1241 Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and
 1242 correct) proof?

1243 Answer: [N/A]

1244 Justification: The paper is empirical and benchmark-oriented; it does not include theoretical results or proofs.
 1245 Section 3 introduces estimand notation but the $\bar{r}$, Δ , and lift statistics are descriptive.

1246 4. Experimental result reproducibility

1247 Question: Does the paper fully disclose all the information needed to reproduce the main experimental
 1248 results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless
 1249 of whether the code and data are provided or not)?

1250 Answer: [Yes]

1251 Justification: Section 4 and Appendix A13 specify the public FM checkpoints, decoder architecture (1×1 -
 1252 conv default plus stated sweeps), training protocol (Adam 10^{-3} , 30 epochs for the frozen-head primary runs,
 1253 batch 16, BCE for the binary primary rows), seed set ($\{42, 43, 44, 45\}$ with σ -seeded determinism), and
 1254 the actual primary split protocols: Kvasir 880/120, RIGA train = BinRushed+MESSIDOR with Magrabi
 1255 (Magrabi Eye Center) source held out ($n=95$), ACDC public patients 001–100 with patient-level fold-0

1256 LV-positive slices ($n=361$), and the BraTS 2024 2D FLAIR $\text{seg}>0$ foreground derivative: per held-out
1257 subject, up to the top-10 axial slices by foreground area with at least 64 positive pixels, subject-level 80/20
1258 split using NumPy seed 42 ($n=3228$). This BraTS target includes the GLI resection-cavity label and is not
1259 the official 3D BraTS WT challenge target. The supplementary artifact releases primary-task per-case Dice
1260 JSONs, merged summaries, selected aggregate diagnostic JSONs, and code; full raw data, full prediction
1261 caches, raw held-out prediction arrays, and fuller split manifests are outside the supplementary artifact.

1262 **5. Open access to data and code**

1263 Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully
1264 reproduce the main experimental results, as described in supplemental material?

1265 Answer: [Yes]

1266 Justification: All source datasets are public or access-controlled research resources from their upstream
1267 providers, subject to their licences, research-use terms, or DUAs (RIGA+ via Deep Blue; ACDC via the
1268 ACDC challenge; Kvasir-SEG via Simula; BraTS 2024 adult glioma via Synapse; CVC-ClinicDB via an
1269 operational HuggingFace mirror for a no-raw-redistribution OOD diagnostic). All backbones are publicly
1270 available or documented as gated research checkpoints. The supplementary artifact releases benchmark
1271 code, estimator implementations, primary-task per-case Dice traces, merged summaries, and selected ap-
1272 pendix result JSONs; authored code and metadata use the bundle licence, while derived scalar traces remain
1273 subject to upstream dataset/model terms. Raw medical images, full prediction caches, some held-out-task
1274 artifacts, and complete split manifests are not redistributed in the supplementary artifact.

1275 **6. Experimental setting/details**

1276 Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how
1277 they were chosen, type of optimizer) necessary to understand the results?

1278 Answer: [Yes]

1279 Justification: Section 3 gives the formal estimands (within-/cross-family $\bar{r}$, spatial Δ_{pix} , threshold lift L)
1280 and the symmetric 4×4 seed-combo averaging rule; Section 4 gives the datasets (with patient-level splits),
1281 foundation-model list, and protocol (image size, encoder-specific upstream normalisation, decoder, opti-
1282 miser, loss, epochs, batch, seed control).

1283 **7. Experiment statistical significance**

1284 Question: Does the paper report error bars suitably and correctly defined or other appropriate information
1285 about the statistical significance of the experiments?

1286 Answer: [Yes]

1287 Justification: Table 1 reports paired case-level and hierarchical family/seed bootstrap 95% CIs ($B=2000$).
1288 The BraTS foreground derivative is analysed at the slice level with an explicit within-subject dependence
1289 caveat, Appendix A11 reports subject-level sensitivity checks, and Appendix A12 explains why sparse-class
1290 NCR filtering diagnostics are not used as numeric evidence.

1291 **8. Experiments compute resources**

1292 Question: For each experiment, does the paper provide sufficient information on the computer resources
1293 (type of compute workers, memory, time of execution) needed to reproduce the experiments?

1294 Answer: [Yes]

1295 Justification: Experiments were run on NVIDIA A100-80GB, A6000, and L40S GPUs across a shared
1296 compute cluster. Feature extraction per FM per domain ranges from minutes for CNN backbones to longer
1297 ViT/BraTS jobs; cached 1×1 -head training per seed is lightweight. Exact wall-clock varies by completed
1298 matrix; the release reports per-job elapsed times in result JSONs where recorded rather than claiming a
1299 single exact compute total.

1300 **9. Code of ethics**

1301 Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of
1302 Ethics <https://neurips.cc/public/EthicsGuidelines>?

1303 Answer: [Yes]

1304 Justification: The research uses only de-identified public or access-controlled research medical imaging
1305 datasets (RIGA, ACDC, Kvasir-SEG, BraTS 2024) under their original research licences, terms, or DUAs.
1306 No additional patient data are collected. The paper audits retrospective failure-correlation properties of
1307 benchmark segmentation pipelines and does not propose new clinical tools.

1308 **10. Broader impacts**

1309 Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of
1310 the work performed?

1311 Answer: [Yes]

1312 Justification: Section 8 discusses deployment implications: shared-pipeline model pools (same FM, same ar-
1313 chitecture family, same recipe) in frozen/light-adaptation medical AI deployments may be over-estimated as
1314 providing independent safety margins. The paper quantifies a retrospective evaluation assumption (independ-
1315 ence among FM-pool members) and does not measure clinician decisions, diagnoses, or patient outcomes.

Backbone-family diversification is observed to have lower dependence in the tested comparisons, not to be a complete clinical mitigation.

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [\[Yes\]](#)

Justification: The supplementary artifact does not release raw medical images, voxel labels, full prediction caches, upstream checkpoints, trained clinical services, or patient-facing models. Released artifacts are derived per-case scalar traces, aggregate summaries, metadata, hashes, and scorer code under upstream dataset/model constraints. The paper states that the audit is not a clinical safety guarantee, does not support re-identification or patient profiling, and should not be used as a standalone model-selection or deployment rule.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: RIGA+ (CC BY-NC 4.0 via Deep Blue; non-commercial use only; raw images and masks not redistributed), ACDC (research licence via challenge website), Kvasir-SEG (research/educational use via Simula; raw images not redistributed), BraTS 2024 (research licence/DUA via Synapse; no raw images or voxel-level labels redistributed), and CVC-ClinicDB used only through a no-raw-redistribution OOD diagnostic; we do not redistribute raw CVC images or masks, and any future raw redistribution would require checking the original CVC terms. Backbones/checkpoints: DINOv2 (Apache-2.0), MAE ViT-B/16 (CC BY-NC 4.0), DeiT-B/16 (Apache-2.0), OpenAI CLIP ViT-B/16 and ViT-L/14 weights (OpenAI CLIP license; open_clip code MIT), BiomedCLIP checkpoint under its model-card terms rather than the open_clip library license, ConvNeXt-T/EfficientNet-B0/ResNet-50/18 loaded via torchvision weights/code under PyTorch/torchvision terms, RETFound checkpoint (gated; CC BY-NC 4.0/research-use), MedSAM/SAM (Apache-2.0). No upstream checkpoints are redistributed. **BraTS redistribution compliance:** in the supplementary artifact we release only derived scalar traces and aggregate statistics needed to reproduce the estimators; we do *not* redistribute raw BraTS images, raw voxel labels, or full prediction caches. All assets are cited.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The paper releases (i) a benchmark protocol (estimands + symmetric multi-seed + bootstrap estimators); (ii) primary-task per-case Dice JSONs and merged summary JSONs for the released rows; (iii) selected appendix/ablation result JSONs and reproduction scripts. Raw prediction masks/caches and full split manifests are not part of the supplementary artifact because upstream redistribution rights and artifact size constraints vary across datasets.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: No crowdsourcing; no new human-subjects research. We use existing public or access-controlled research datasets whose upstream governance, approvals, or access terms are handled by the dataset providers.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: No new human-subjects data were collected. Upstream datasets carry their own governance, approvals, or access terms (RIGA, ACDC, Kvasir-SEG, BraTS 2024).

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

1376 Answer: [\[Yes\]](#)
1377 Justification: LLMs were used as editorial and programming assistants during manuscript and artifact prepa-
1378 ration. They are not a component of the benchmark methodology, estimator, training pipeline, or evaluation
1379 protocol: all reported numbers come from the released scripts, logs, and JSON artifacts, and all manuscript
1380 text, code, citations, and claims were reviewed and approved by the authors.