
Auditing Correlated Failures in Frozen-Feature Pretrained-Encoder Pools for Medical Segmentation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Medical segmentation ensembles are credited with complementary failure cov-
2 erage when independently trained members make different mistakes. We au-
3 dit this assumption for frozen and lightly-adapted pretrained-encoder pipelines,
4 where independently trained decoders can share the same backbone, pretrain-
5 ing bundle, and recipe. Across five primary task rows (Kvasir polyp, ACDC
6 LV, BraTS foreground, RIGA Cup, RIGA Disc), our 11-bundle $\times$ 4-decoder-
7 seed pool yields per-case Dice correlation gaps of $\Delta = 0.261\text{--}0.638$ between
8 same-bundle and cross-bundle pairs after a Dice ≥ 0.30 functional floor; all
9 paired case and hierarchical bootstrap intervals exclude zero. The result is sta-
10 ble under floor sweeps, ICC(A,1), leave-one-out, subject aggregation, quality-
11 controlled pair regression, and cross-fitted item-difficulty residualisation. Diag-
12 nostic probes (pixel error, joint-failure lift, same-architecture-family ViT, same-
13 family cross-checkpoint, same-backbone) locate the dependence structure but do
14 not causally separate architecture, pretraining, and recipe. The claim is scoped to
15 2D frozen/light-adaptation encoder pools and retrospective benchmark audits. We
16 release per-case Dice traces, summary JSONs, scorer code, hashes, and metadata
17 as the audit artifact.

18 1 Introduction

19 Ensembles are often used in safety-relevant medical segmentation because different members may
20 fail on different cases. Aggregate Dice can tell us whether members are accurate, but it does not by
21 itself show whether additional members provide additional failure evidence on the same cases. This
22 paper audits that missing quantity for frozen or lightly adapted pretrained-encoder segmentation
23 pools: the aligned case-level correlation of segmentation quality across pool members.

24 The medical-imaging supply chain now reuses a small set of foundation-model encoders and con-
25 ventional pretrained backbones: DINOv2, CLIP/BiomedCLIP, SAM/MedSAM, ConvNeXt, Effi-
26 cientNet, ResNet, and domain-specific pretrainings such as RETFound. A pool of “different” seg-
27 menters can therefore share backbone family, pretraining bundle, input statistics, and training recipe.
28 This creates an evaluation gap: aggregate Dice and common uncertainty summaries usually do not
29 directly report aligned case-level co-failure, which is the evidence needed for independent failure-
30 coverage claims. The intuition that near-duplicate pipelines should share errors is not new. The
31 missing piece is a released, case-aligned audit that quantifies how large the dependence is, tests
32 whether it survives simple alternative explanations, and states which evaluation claims the evidence
33 can and cannot support.

34 **Primary contribution.** We contribute a reusable E&D audit protocol for fixed segmentation pools:
35 compute per-case Dice traces, define a prespecified primary-table same/release-recipe group and
36 cross-bundle reference, average symmetrically over decoder seeds, and report paired bootstrap in-
37 tervals for the same/cross correlation gap. The protocol asks a concrete reporting question: before

a pool is credited with complementary failure coverage, have same/release-recipe and cross-bundle co-failure correlations been measured on aligned cases? Out of scope: prompt-conditioned MedSAM, full nnU-Net pipelines, clinical outcomes, demographic fairness, and SOTA segmentation performance.

Artifact and diagnostic contribution. We release the trace-level files and scorer code needed to recompute the primary audit statistics. Additional diagnostics test estimator sensitivity, subject-level clustering, item difficulty, pixel error overlap, same-family cross-checkpoint structure, same-backbone perturbations, adaptation and full-pipeline boundaries, and selected scope checks. The reusable object is therefore an audit protocol plus a trace-level evidence bundle for measuring whether a fixed pool of segmentation models supplies independent failure evidence on aligned cases. We position the audit against monoculture, error-correlation, representation-similarity, and medical-segmentation literatures next.

2 Related work

Relation to prior work and scope of novelty. Prior work has established that correlated errors and monoculture matter. Our contribution is to make this issue auditable for medical segmentation by releasing aligned per-case traces, same/cross estimators, and a primary/diagnostic evidence hierarchy. Our risk framing follows algorithmic-monoculture and foundation-model ecosystem work [Kleinberg and Raghavan, 2021, Bommasani et al., 2021, 2022, Toups et al., 2023]. The correlation gap we measure has modelling antecedents: trial-by-trial error consistency asks whether systems fail on the same inputs [Geirhos et al., 2020], error correlations arise when deployed models share algorithms, datasets, or foundation models [Li et al., 2025], shared training trajectories and weight-space basins can constrain checkpoint diversity [Fort et al., 2019, Wortsman et al., 2022], transfer/representation work studies what is reused across pretrained models [Neyshabur et al., 2020, Gontijo-Lopes et al., 2022], pretrained-model ensemble work treats pretraining as a possible diversity source [Mustafa et al., 2020], same-architecture models can align strongly across initialisations and recent surveys consolidate functional/representational similarity measures [Kornblith et al., 2019, Klabunde et al., 2024], and predictive-diversity pathologies can limit deep-ensemble gains [Abe et al., 2024]. Recent LLM monoculture and null-model papers are relevant because they show that agreement/correlation claims depend on the choice of baseline and unit of analysis; our paper applies that lesson to case-level medical segmentation traces [Jiang et al., 2025, Kim et al., 2025a, Gorecki and Hardt, 2025, Goel et al., 2025, Jo et al., 2026]. Pathology-FM and precision-oncology studies provide adjacent evidence about center robustness, benchmark behavior, and complementary domain-specialised features [de Jong et al., 2025, Campanella et al., 2025, Luo et al., 2025, Neidlinger et al., 2025]. The distinction is threefold: our unit is an aligned segmentation Dice trace rather than a classification decision or deployment outcome, our artifact releases the trace-level evidence needed to recompute the audit rather than only a summary statistic, and the paper separates primary same-checkpoint/release-recipe evidence from mechanism-adjacent diagnostics. Our contribution is not the general claim that correlated errors exist. It is a released medical-segmentation audit of frozen-feature pretrained-encoder pools that distinguishes primary evidence from diagnostics and reports where same-checkpoint/recipe, same-architecture-family, and broader cross-family correlations appear under a shared protocol.

Medical segmentation, pretrained/foundation-model benchmarks, and ensemble baselines. nnU-Net, deep ensembles, MC dropout, adjacent uncertainty or parameter-efficient ensembles, and hyperparameter ensembles are usually evaluated through Dice, calibration, or uncertainty quality [Isensee et al., 2021, 2024, Lakshminarayanan et al., 2017, Gal and Ghahramani, 2016, Mehrtaash et al., 2020, Jungo and Reyes, 2019, Mühlematter et al., 2026, Gu et al., 2024]. Classical ensemble-diversity metrics motivate this audit, but our reporting target is the case-aligned correlation of task-performance traces because that statistic directly tests redundant failure evidence on the evaluated benchmark [Kuncheva and Whitaker, 2003, Ganaie et al., 2022, Wood et al., 2023]. Medical-segmentation diversity and failure-detection benchmarks ask whether a method can detect segmentation failures; our audit asks whether multiple pool members fail on the same cases, making the questions complementary rather than substitutable [Georgescu et al., 2023, Zenk et al., 2025]. The choice of per-case Dice as the unit of analysis follows recent recommendations for image-analysis validation that case-level distributions, not aggregate scalars, drive evaluation reliability [Maier-Hein et al., 2024]. Recent medical-segmentation benchmarks and robustness-specification work evalu-

ate or discuss generic-vision probes, SAM medical-image evaluations, standardized segmentation datasets, promptable SAM adaptations, universal 3D/medical segmenters, task-specific robustness requirements, transferability estimation, and VFM segmentation benchmarking [Baharoon et al., 2023, Huang et al., 2024, Kuş and Aydin, 2024, Kirillov et al., 2023, Ma et al., 2024, Cheng et al., 2023, Archit et al., 2025, Zhu et al., 2024, Ma et al., 2025, Gu et al., 2025, He et al., 2025, Kim et al., 2025b, Xian et al., 2025, Yang et al., 2023, Keressies et al., 2024]. These benchmarks are complementary to our audit: they ask which model or foundation model performs well, whereas we ask whether multiple members of a pool provide non-redundant failure evidence on the same cases. Promptable and 3D medical foundation models are deployment alternatives rather than baselines for this co-failure audit.

3 Formal setup and benchmark estimands

Pool and estimands. Let $\mathcal{F}=\{F_1, \dots, F_K\}$ be the operational pretrained-encoder bundles used by the Table 1 scorer. In the primary 11-bundle source pool, a bundle is the concrete encoder / pretraining / input-statistic wrapper used for frozen-feature extraction. Same-bundle decoder seeds form most of the main same group, and DINOv2-B×DINOv2-S is the only recurring cross-checkpoint member of that group because both variants share the same upstream DINOv2 release and pretraining recipe. Other broad-family groupings (for example OpenAI CLIP sizes, ResNet sizes, and expanded DINOv2/ConvNeXt variants) are not folded into the Table 1 primary estimator; they are separated in the same-family cross-checkpoint diagnostic in Appendix A30. BiomedCLIP (ViT-B/16, biomedical image-text pretraining) is treated as a different bundle from DINOv2 despite being a ViT. A model trace is a tuple (F_k, σ) (bundle, decoder seed) producing per-case Dice $D_{(k,\sigma)}(i, t)$. Appendix A29 and Appendix A30 report diagnostics that separate primary product bundles, same-architecture-family ViT / different-pretraining, and broader same-family cross-checkpoint readouts; these clarify that a “cross” column can still contain broad architectural structure. For a pair of models $a \neq b$ we define $r_{a,b}(t) = \text{Pearson}_i[D_a(i, t), D_b(i, t)]$ (and a McGraw–Wong ICC(A,1) variant $\iota_{a,b}$ [McGraw and Wong, 1996], chosen to penalise per-bundle Dice offsets that Pearson absorbs). The primary-table same/cross means at task t are the expectations of $r_{a,b}$ over pairs in the same/release-recipe group (same bundle, plus DINOv2-B/S) and the remaining cross-bundle reference under this scorer. We define $\Delta(t) = \bar{r}_{\text{same}/\text{release}}(t) - \bar{r}_{\text{cross}}(t)$ as the co-failure gap and refer to it as the *primary same/cross gap*, Δ , throughout. Secondary estimands are pixel-level error-mask Dice, computed the same way over binarised error masks, and joint-failure *lift* over independence $L_{\text{mono}}(\tau) = \Pr[\text{all fail} \mid \text{mono pool}] / \prod_m \Pr[\text{fail}_m \mid \tau]$.

Terminology used by the audit. The primary claim uses only the primary-table same/release-recipe group and the cross-bundle reference; broader family rows are diagnostics.

Term	Meaning in this paper	Primary claim?
Bundle	Concrete encoder checkpoint plus wrapper, pretraining, input normalisation, and feature-extraction path	Yes
Model trace	Bundle plus decoder seed and aligned per-case Dice vector	Yes
127 Primary-table same/release-recipe group	Same-bundle seed pairs plus the prespecified DINOv2-B/S same-release cross-checkpoint pair	Yes
Primary-table cross reference	Remaining primary-pool pairs under the primary grouping	Yes
Same-family cross-checkpoint	Broader DINOv2, CLIP, ConvNeXt, and ResNet diagnostic grid	Diagnostic only
Same-architecture ViT cross-pretraining	ViT-B diagnostic varying pretraining source	Diagnostic only

Symmetric multi-seed averaging. For each bundle pair we enumerate all $|\Sigma|^2=16$ seed combinations (4 per model) rather than seed-matched pairs, removing the one-seed-per-bundle noise that inflates ad-hoc gap estimates. Same-bundle averaging uses $\binom{|\Sigma|}{2}=6$ same-bundle seed pairs per F_k plus the full $|\Sigma|^2$ on DINOv2-B×DINOv2-S; the cross group uses all remaining primary-table bundle pairs.

CI and scope. A 3-level nested hierarchical bootstrap (primary groups $\rightarrow$ seeds $\rightarrow$ cases, $B=2,000$, deterministic seed 0) yields 95% CIs as the 2.5th/97.5th percentiles of bootstrap-sampled $\bar{r}_{\text{same}}, \bar{r}_{\text{cross}}, \Delta$. Because the number of encoder-bundle clusters is small, neither interval should be read as a population-level guarantee over all possible pretrained encoders. We report both intervals as finite-audit uncertainty summaries: the case bootstrap asks about new cases conditional on the evaluated pool, while the hierarchical bootstrap also perturbs the finite group/seed surface. The estimands assume a shared-decoder, shared-task pool and measure *observed* correlation within

such pools; a CNN random-init diagnostic (Appendix A31.2) shows positive same-vs-cross structure without pretrained weights, so we report the pretrained-bundle gap descriptively rather than as a causal isolation of pretrained-weight effects.

Primary-group structural caveat. The primary rows retain task-specific functional pools of 11/10/10/9/11 bundles after the Dice floor. In these pools, the Table 1 “same” column is dominated by same-bundle, different-seed decoder-head pairs, with DINOv2-B $\times$ DINOv2-S providing the only recurring same-family cross-checkpoint pair when both variants are functional. Thus the high same-group $\bar{r}$ values in Table 1 (0.834–0.959) are best read as same-frozen-checkpoint/recipe-dominated dependence readouts, not a broad sample of many same-family cross-checkpoint comparisons. Section 6.1 and Appendix A29 report a 5-variant same-architecture-family ViT / different-pretraining diagnostic (DINOv2-B, BiomedCLIP, MAE-B, DeiT-B, CLIP-B) that measures cross-pretraining correlation under a shared broad architecture family and is the closest “within-architecture-family, cross-pretraining” control in this benchmark.

4 Experimental design

Datasets & targets. The primary rows are summarized below. The splits are recorded in the released JSON files’ `test_ids` fields; subject-level sensitivity for ACDC and BraTS is in Appendix A11.

Primary task rows and evaluation units. The audit uses benchmark-derived targets and aligned per-case Dice traces, not clinical deployment cohorts.

Task row	Raw source	Target / split	Unit	Main caveat
Kvasir polyp	Kvasir-SEG [Jha et al., 2020]	Loader-defined 880/120 image split	image, $N=120$	No patient/video IDs available; no patient-level duplicate control claimed.
ACDC LV	ACDC [Bernard et al., 2018]	LV-positive ED/ES slices from public patients 001–100, fold 0	slice, $N=361$; 20 patients	Slice row; patient-level sensitivity reported in Appendix A11.
BraTS foreground	BraTS 2024 GLI [BraTS 2024 Organizers, 2024, de Verdier et al., 2024]	2D FLAIR <code>seg>0</code> ; top-10 eligible slices per held-out subject, NumPy seed 42	slice, $N=3,228$; 324 subjects	Includes resection-cavity label; not the official 3D BraTS WT challenge target.
RIGA Cup/Disc	RIGA+ [Almazroa et al., 2018, Almazroa, 2018]	BinRushed+MESSIDOR Magrabia source test	train, fundus, $N=95$	Deterministic hashed alignment IDs retained only for split alignment.

Encoder bundles. *Generic pretrained-encoder bundles* (11, mixing foundation-model encoders and conventional pretrained backbones): DINOv2-B/S, BiomedCLIP, CLIP ViT-L/14, CLIP ViT-B/16, MAE ViT-B/16, DeiT ViT-B/16, ConvNeXt-T, EfficientNet-B0, ResNet-50/18. Encoder names refer to the exact released wrapper/checkpoint paths in the artifact; in particular, CLIP-L/14 is the documented OpenCLIP dense-token wrapper rather than a claim of official OpenAI inference parity. *Medical-domain encoder probe:* MedSAM vision-encoder-only (SAM ViT-B; prompt encoder and mask decoder discarded) is reported as a secondary boundary probe on RIGA/Kvasir/ACDC. RETFound support is documented as a gated-checkpoint path, but RETFound cells are not reported because the gated checkpoint was unavailable in the release environment.

Protocol. The primary protocol resizes images and masks to 224×224 with nearest-neighbor mask resizing, applies encoder-specific upstream normalisation (ImageNet-style statistics for DINOv2/MAE/DeiT/torchvision CNNs and OpenAI CLIP statistics for CLIP/BiomedCLIP), freezes the encoder, and trains a 1×1 -conv decoder with Adam 10^{-3} for 30 epochs, batch 16, BCE loss on the binary primary rows, no early stopping, no validation-based model selection, and fixed sigmoid threshold 0.5 at inference. Table 1 starts from the 11-bundle generic source pool where traces are available and applies the estimator’s Dice ≥ 0.30 retrospective functional floor to remove non-segmenting, near-constant traces: Kvasir retains 11 bundles, ACDC 10, BraTS foreground 10, RIGA Cup 9, and RIGA Disc 11. This test-trace-based analytic convention defines the reported functional pool; it is not a training, deployment, or pre-registered hypothesis threshold. The floor systematically reduces cross-bundle correlation (near-constant outputs would be near-zero correlation with everything), so the post-floor pool biases the gap upward by construction; the released floor sweep (Appendix A5) shows that primary-row signs do not flip across $\{0, 0.20, 0.25, 0.30, 0.35, 0.40\}$, but the magnitudes are floor-conditional. The MedSAM sensitivity is a secondary encoder-only boundary probe, not part of the primary pool or a prompt-conditioned MedSAM evaluation.

5 Main result: same-frozen-checkpoint pools show higher co-failure correlation

Table 1 reports the primary finite-pool audit on five medical segmentation task rows. Same-bundle/release-recipe groups have substantially higher per-case Dice correlation than the cross-bundle reference: $\rho_{\text{same/release}} = 0.834\text{--}0.959$ vs. $\rho_{\text{cross}} = 0.232\text{--}0.697$, with $\Delta_{0.30} = 0.261\text{--}0.638$. Every case-level and hierarchical bootstrap interval excludes zero (smallest hierarchical lower bound 0.174, Kvasir).

Table 1: Primary finite-pool co-failure audit. For each filtered functional pool we report average pairwise Pearson correlation of per-case Dice traces for the primary same/release-recipe group and the cross-bundle reference. K is the retained bundle count after the $\text{Dice} \geq 0.30$ functional floor; with four decoder seeds per bundle the trace count is $4K$ and same/cross pair counts follow deterministically from K . $\Delta_{0.30}$ is conditional on the retained functional pool; $\Delta_{\text{no floor}}$ recomputes the same scorer without the floor. CIs are 95% case-level paired bootstrap ($B=2,000$), and the conservative finite-pool interval Δ hier CI[†] is a hierarchical bootstrap over primary grouping and seed clusters; all displayed hierarchical CIs exclude zero. ACDC and BraTS are slice-level rows nested within 20 ACDC patients and 324 BraTS subjects; subject-aggregated gaps are ACDC $\Delta = 0.262$ [0.151, 0.413] and BraTS $\Delta = 0.214$ [0.182, 0.251] (Appendix A11). The BraTS row uses the 11-bundle source pool for the 2D FLAIR $\text{seg}>0$ derivative; ResNet-18 is below the Dice floor, leaving 10 retained bundles. The primary ACDC pool retains CLIP-L/14 because it passes the Dice floor; a secondary 8-FM ACDC diagnostic with a different split (Appendix A13) reports CLIP-L/14 collapsed at Dice 0.007 on that split.

Task row	Unit N	K	$\rho_{\text{same release}}$	ρ_{cross}	$\Delta_{0.30}$	$\Delta_{\text{no floor}}$	$\Delta_{\text{case CI}}$	$\Delta_{\text{hier CI}^\dagger}$
Kvasir polyp	image 120	11	0.958	0.697	0.261	0.261	[0.211,0.320]	[0.174,0.378]
ACDC LV	slice 361 (20 pts)	10	0.959	0.639	0.321	0.355	[0.289,0.355]	[0.211,0.435]
BraTS foreground	slice 3,228 (324 subj)	10	0.938	0.623	0.315	0.362	[0.302,0.329]	[0.224,0.425]
RIGA Cup	fundus 95	9	0.834	0.361	0.473	0.529	[0.395,0.556]	[0.280,0.780]
RIGA Disc	fundus 95	11	0.870	0.232	0.638	0.638	[0.569,0.687]	[0.475,0.831]

Within these filtered 36–44-trace pools the gap is sign-stable across robustness checks: no-floor gaps are also positive (Kvasir 0.261, ACDC 0.355, BraTS 0.362, RIGA Cup 0.529, RIGA Disc 0.638), and subject aggregation on the slice-level rows preserves the sign (ACDC 0.262, BraTS 0.214; Kvasir remains image-level evidence because patient/video IDs are unavailable). The same column is dominated by same-bundle decoder-seed pairs because each $K=9\text{--}11$ functional pool admits only DINOv2-B $\times$ DINOv2-S as a recurring same-family cross-checkpoint pair; the broader same-family cross-checkpoint stratum is decomposed separately in Appendix A30 and the distinct-checkpoint minus cross-family gap remains positive on all five rows (0.080–0.267). Table 1 therefore audits the same-frozen-checkpoint/release-recipe failure-evidence assumption rather than claiming broad same-family checkpoints are equally redundant. Table 2 summarizes four sign-stability diagnostics: a calibration-selected functional pool evaluated on disjoint cases, a binary failure-event readout at the empirical P_{33} threshold, quality-controlled pair regression, and distinct-checkpoint same-family decomposition.

Estimator robustness. A sensitivity summary (Appendix A5) recomputes the primary Pearson gap under a functional-floor sweep, an ICC(A,1) alternative, leave-one-bundle/task-out summaries, and a random family-label permutation stress test. Analytic-floor checks show the gap remains positive at no-floor and stricter-floor settings where the pool is non-empty. Event and ICC checks show the same sign after binarising failures at the empirical P_{33} threshold and after replacing Pearson with ICC(A,1). Quality and residualisation checks show that simple marginal-quality controls and held-out item-difficulty residualisation do not absorb the same/cross separation. The random-label permutation is retained only as an appendix stress test because family labels are not exchangeable sampling units; it is not used as a main-text significance claim or population-level p -value.

Spatial and lift readouts. Figure 1 corroborates the Table 1 gap at the pixel and event level. Same-bundle seeds produce spatially overlapping false positives and false negatives (panel a); aggregate pixel-level error-mask Dice (panel b) is 0.80–0.87 same-family versus 0.32–0.40 cross-family on RIGA Cup/Disc ($N=224$), ACDC LV ($N=90$), and Kvasir ($N=300$) (Appendices A3, A8); matched-size joint-failure lift ratios (panel c) follow the same ordering. These aggregate diagnos-




Figure 1: **Spatial and lift evidence corroborating the per-case Dice gap in Table 1.** Panel (a) shows single-case visualisations on RIGA Cup and BraTS foreground (red=FP, blue=FN, green=GT boundary): four DINOv2-B decoder seeds versus four distinct-bundle members produce visibly different error-overlap structure even when their mean Dice matches. Panel (b) shows pixel-level error-mask Dice on aggregate splits (RIGA Cup/Disc $N=224$, ACDC LV $N=90$, Kvasir $N=300$): same-family 0.80–0.87 versus cross-family 0.32–0.40. Panel (c) shows matched-size joint-failure lift ratios. Panels (b–c) use non-case-identical aggregate splits and complement the per-case-Dice rows of Table 1 at the pixel and event levels; the artifact distributes summary JSONs rather than raw masks.

tics use non-case-identical splits and therefore complement rather than replace the per-case Pearson result in Table 1.

Item-difficulty diagnostic. Following the warning that co-failure claims depend on the null model [Jo et al., 2026], we residualise per-case Dice against held-out item-difficulty estimates. A naive all-bundle difficulty regressor uses information from the target trace and inflates Δ , so we use a cross-fitted estimator that fits each trace’s item-difficulty regressor on traces outside the trace’s own broad upstream family and then recomputes the primary same/cross rule on residuals. The residualised gap is positive on every primary row (Appendix A26); residualisation removes a shared per-case difficulty signal common to both pair members, so the residual statistic is no longer a canonical Pearson correlation in the $[-1, 1]$ sense and the magnitudes are read ordinarily rather than on the raw correlation scale.

Because the primary same column is dominated by same-bundle decoder-seed pairs, we also report a same-family cross-checkpoint decomposition over the five primary rows (Appendix A30). It separates the same-bundle seed floor, distinct checkpoints within broad upstream families (DINOv2-S/B, OpenAI CLIP-B/L, ResNet-18/50 when functional), and cross-family pairs. The distinct-checkpoint

Table 2: **Sign-stability diagnostics for the primary rows.** Values are point estimates from distinct diagnostics and are not directly comparable in scale. Cal.-split Δ selects functional bundles on a deterministic calibration subset and evaluates the same/cross gap on disjoint cases. P_{33} fail ϕ gap is the same/cross gap in binary failure-event correlation at the empirical 33rd percentile of model-case Dice after functional-floor filtering. Quality-control β_{same} is the same-group coefficient in a pair-level OLS diagnostic with case-bootstrap refits. The expanded-grid same-family checkpoint gap is the same-family cross-checkpoint diagnostic from Appendix A30, not the Table 1 analytic-pool cross-family stratum. Full definitions and CIs for the Table 2 columns are in Appendices A5–A7 and Appendix A30; the separate item-difficulty diagnostic is in Appendix A26.

Task	Cal.-split Δ	P_{33} fail ϕ gap	QC β_{same}	Checkpoint- family gap
Kvasir polyp	0.250	0.329	0.169	0.080
ACDC LV	0.350	0.386	0.208	0.134
BraTS foreground	0.307	0.413	0.224	0.093
RIGA Cup	0.478	0.453	0.372	0.267
RIGA Disc	0.630	0.578	0.582	0.227

same-family minus cross-family gap remains positive on all five rows (0.080–0.267), while staying well below the same-bundle seed floor. Thus the primary effect is strongest for shared checkpoints and recipes, but the per-case Dice traces do not collapse to a same-decoder-seed artifact. A limited ViT same-architecture/different-pretraining diagnostic is taken up in Section 6 as a control on architecture-vs-pretraining attribution.

6 Same-backbone controls and scope boundaries

The following diagnostics do not decompose the Table 1 gap into causal shares for architecture, pretraining corpus, checkpoint lineage, wrapper normalisation, and decoder recipe. They answer a narrower evaluation question: if a pool shares one of these upstream factors, does the trace audit still show aligned co-failure that should be reported before claiming independent failure coverage? The practical reporting recommendation does not require identifying which shared factor dominates.

6.1 ViT diagnostic: cross-pretraining correlations sit close to the same-bundle floor

A representation-similarity concern, following Kornblith et al. [2019], is that two ViTs produce similar representations *because they share architecture family*, not because they share pretraining. We hold broad architecture family fixed in a 5-variant ViT pool—DINOv2-B (SSL on LVD-142M), BiomedCLIP (CLIP on biomed image-text), MAE-B (ImageNet MAE), DeiT-B (ImageNet supervised), OpenAI CLIP-B—trained with the identical frozen 1×1 -conv probe on RIGA Cup and ACDC LV (Table 3). ViT cross-pretraining correlations sit between the same-bundle seed floor and the primary cross-family reference; the ordering does not identify which factor is causal but does show that holding architecture family fixed leaves a substantial gap to the cross-family reference.

Table 3: **Limited ViT same-architecture/different-pretraining diagnostic.** The table compares same-bundle seed-pair correlations with cross-pretraining correlations among ViT-style encoders. It is a two-task diagnostic, not a causal decomposition of architecture and pretraining. Seed floor here is the mean within-FM cross-seed Pearson restricted to the functional 5-ViT subset; this differs from Table 1’s ρ_{same} , which averages over the full 9/10-bundle post-floor pool including CNN bundles whose lower within-FM correlation pulls the mean down.

Dataset	ViT subset	Func. ViTs	Seed floor	ViT cross- pretrain	Cross ref.	Floor – cross-pretrain
RIGA Cup	all 5 ViTs	5/5	0.968	0.680	0.361	0.288
ACDC LV	CLIP-B dropped [†]	4/5	0.988	0.807	0.639	0.182

Values are from the architecture-confound diagnostic and primary table summary. [†]The all-5 ACDC same-architecture-family ViT cross-pretraining value is 0.631 because CLIP-B is nonfunctional on ACDC (mean Dice 0.205); the functional-filtered readout is 0.807. ViT $\times$ CNN comparisons are reported only as appendix diagnostics and are not used as main-table numeric evidence.

Table 4: **Same-backbone perturbation diagnostics.** These comparisons vary decoder capacity, skip structure, or MC dropout under limited grids. They test whether simple same-backbone perturbations remove high correlation; they are not dose-matched causal controls.

Diagnostic	Released source	Main readout
Head-capacity sweep, RIGA Cup	Released App. A31.1	traces, 4 heads, 769–3.94M params; Dice 0.725–0.848; cross-head same-bundle $\bar{r}=0.690$ (range 0.510–0.896)
UNet-skip decoder replication	Released App. A16	summary, 20 cells; Dice 0.806–0.902; seed-pair $r=0.774$ –0.979
MC dropout, DINOv2-B/RIGA Cup	Released App. A31.7	traces, MC $\bar{r}=0.947, 0.925, 0.893$ for $p=0.1, 0.2, 0.4$

Readings: head-capacity and UNet-skip stay above the primary RIGA cross-family $\bar{r}=0.361$; MC dropout is a weak decorrelation lever (largest tested drop ≈ 0.10).

This diagnostic alone does not isolate pretraining-bundle reuse from broad architecture-family imprint; in particular, ViT cross-pretraining correlations sit substantially closer to the same-bundle floor (gap 0.288 on RIGA Cup, 0.182 on ACDC LV) than to the primary cross-family reference, consistent with shared architecture and shared 1×1 -conv readout contributing materially to the measured dependence.

The released CNN random-init diagnostic (Appendix A31.2) shows positive same-vs-cross structure even without shared pretrained weights ($\Delta = 0.086$ on ACDC LV, $\Delta = 0.121$ on RIGA Cup); a more substantive 3-architecture from-scratch replication (Appendix A13) reaches the same conclusion. The pretrained-bundle effect therefore operates on top of, not in isolation from, an architecture-and-recipe baseline.

6.2 Decoder-capacity, stochastic, and medical-domain probes

(C1) Head architecture/capacity. The head-capacity and UNet-skip decoder perturbation diagnostics change accuracy and reduce some pairwise correlations, but their same-backbone correlations remain above the primary RIGA cross-family reference. Thus the tested decoder-capacity changes did not remove same-backbone co-failure in these diagnostics.

(C2) Stochastic/adaptation levers. In the RIGA Cup diagnostic, MC averaging modestly lowers within-bundle correlation but remains far above the cross-family reference. Bootstrap-bagging and fold-reslicing are appendix diagnostics with different sampling surfaces; the fold-reslicing row is a single-seed split-stability check, not a within-fold versus cross-fold same-backbone estimator. Full-fine-tuning and LoRA probes without trace artifacts are described only to delimit scope and are not used to support the main claims.

(C3) Medical-domain encoder-only probe. A MedSAM image-encoder-only sensitivity is described in Appendix A10 as a summary diagnostic. The probe strips MedSAM’s prompt encoder and mask decoder, so it does not evaluate prompt-conditioned MedSAM. RETFound is documented only as a gated-checkpoint path; without access to the gated checkpoint, we do not report RETFound rows.

(C4) Full-pipeline scope boundary. nnU-Net v2 2D/3D complements give task-level different-fold gaps of $\Delta \approx 0.034$ –0.047 on RIGA Cup and $\Delta \approx 0.043$ on ACDC 3D; a case-identical RIGA Cup nnU-Net check gives $\Delta=0.0079$ [0.002, 0.039] at 100 epochs (correlations 0.984/0.976) and a matched 1000-epoch run leaves the gap essentially unchanged ($\Delta=0.0075$ [0.003, 0.034], 0.987/0.980). The full-pipeline gap is therefore much smaller than the primary frozen-feature gap; both diagnostics are descriptive scope checks rather than substitutes for the primary audit. Appendix A21 consolidates all same-backbone, adaptation, and scope levers in one map; Appendix A24 reports a 3D segmentation-architecture pilot. Appendix Table 38 encodes the evidence-tier structure: Table 1 is the primary frozen-encoder audit, and estimator, event, architecture/pretraining, same-backbone, full-pipeline, and held-out rows are diagnostic or scope evidence.

7 Discussion

Practical reading. For an operational pool of M members with mean failure correlation ρ_{fail} , report the equicorrelation effective pool size $M_{\text{eff}} = M/[1 + (M - 1)\rho_{\text{fail}}]$ (Appendix A27; here M

denotes operational pool size, distinct from the trace count $4K$ in Table 1) as a threshold-specific dependence readout before treating same-bundle seeds as independent evidence, not as a standalone pool-selection objective. In our released rows, diverse $M=4$ pools yield mean M_{eff} from 2.05 (Kvasir) to 2.97 (RIGA Disc) under the paper’s equicorrelation readout at $\tau = 0.85$, rather than 4; using the maximum observed diverse-pool ρ_{fail} gives conservative lower readouts of 1.30–1.69 for the most redundant valid 4-bundle tuples. Same-bundle seed pools yield M_{eff} near 1 under the same threshold. The readout should always be paired with marginal Dice and matched-quality comparisons before any pool-selection decision (Appendix A27). For procurement-style decisions where a clinical-imaging team is choosing between an N -bundle FM pool and a smaller curated pool, an in-domain co-failure audit on a small validation set should precede pool-size-based diversity claims.

Implications for ensemble methodology. Standard nnU-Net 5-fold validation under shared encoders is a within-bundle estimator and inherits the dependence structure measured here; cross-encoder ensembles should be preferred where co-failure is decision-critical. A structured co-failure audit reporting Δ , joint-failure lift, and M_{eff} together provides one such readout, and Appendix A28 gives a random-effects representation that motivates why same-backbone fold/seed surfaces concentrate above cross-encoder references.

Limitations. The evidence is retrospective and benchmark-based: this study audits a reporting assumption about failure evidence, not clinical deployment, and we do not claim population-level inference over future encoders, datasets, or recipes. The BraTS foreground row is a 2D derivative of the 2024 GLI release, not the official 3D challenge target. Table 1 is the confirmatory evidence surface; threshold, lift, and held-out rows are exploratory diagnostics. Per-case Dice Pearson is a continuous-monotone alignment statistic, so two members can have high Pearson Dice yet still disagree on which low-quality cases fail; pixel error-mask Dice and joint-failure lift in Section 5 supplement it at the failure-event level, and Geirhos-style binarised error consistency would be a complementary discrete estimand. The cross-bundle stratum draws from at most 8 broad family clusters across 11 bundles, so few-cluster bias is unavoidable; we report case-level and hierarchical bootstrap CIs jointly. The frozen 1×1 -conv decoder is deliberately minimal so the encoder dominates the readout; absolute correlation magnitudes are upper bounds for richer decoders, and the M3 head-capacity and M4 UNet-skip diagnostics show more capable heads reduce some pairwise correlations but do not eliminate the same-vs-cross gap. All ten sign tests (five rows $\times$ two CI methods) independently exclude zero, with the smallest hierarchical lower bound at 0.174 (Kvasir); we report directional gaps per row rather than a joint significance statement. Appendices A18, A19, and A15 extend these caveats.

What the evaluation supports. The E&D takeaway is a reporting requirement: when a medical segmentation pool reuses pretrained encoder checkpoints, input statistics, and decoder recipes, aggregate Dice should be accompanied by aligned case-level co-failure correlations before the pool is credited with independent failure coverage. The result is a finite audit of evaluation evidence, not a clinical safety claim or a causal decomposition of pretraining and architecture.

8 Artifact and reproducibility scope

The anonymized artifact is available through the supplementary submission and anonymous code URL https://anonymous.4open.science/r/NIPS_A_ACF-86CD: primary per-case Dice JSONs, summary tables, code, metadata/manifests, hashes, environment files, a data card, selected appendix summary JSONs, and Croissant metadata with the required Responsible AI fields. It does not introduce a new raw dataset; released JSONs are derived trace artifacts that retain upstream de-identified subject/patient/slice IDs and deterministic hashed RIGA alignment IDs solely for case alignment and subject-level resampling, and are not intended for individual-level analysis, subgroup profiling, re-identification, inference of protected attributes, clinical diagnoses, or patient outcomes. The bundle excludes raw medical images, raw masks, direct identifiers, full prediction caches, and held-out raw prediction arrays; raw data and checkpoints remain governed by upstream licenses and DUAs. Appendix Table 38 ties each main-text evidence class to its released JSON or recompute script.

References

- Taiga Abe, E. Kelly Buchanan, Geoff Pleiss, and John Patrick Cunningham. Pathologies of predictive diversity in deep ensembles. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=TQfQUksaC8>.
- Ahmed Almazroa. Retinal fundus images for glaucoma analysis: The riga dataset. University of Michigan - Deep Blue Data, 2018. URL <https://doi.org/10.7302/Z23R0R29>.
- Ahmed Almazroa, Sami Alodhayb, Essameldin Osman, Eslam Ramadan, Mohammed Hummadi, Mohammed Dlaim, Muhannad Alkatee, Kaamran Raahemifar, and Vasudevan Lakshminarayanan. Retinal fundus images for glaucoma analysis: The riga dataset. In *Medical Imaging 2018: Imaging Informatics for Healthcare, Research, and Applications*, volume 10579, page 105790B. SPIE, 2018. doi: 10.1117/12.2293584.
- Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *arXiv:2107.07511*, 2021.
- Michela Antonelli, Annika Reinke, Spyridon Bakas, et al. The medical segmentation decathlon. *Nature Communications*, 13:4128, 2022. doi: 10.1038/s41467-022-30695-9. URL <http://medicaldecathlon.com/>.
- Anwai Archit, Luca Freckmann, and Constantin Pape. Medicosam: Robust improvement of sam for medical imaging. *arXiv:2501.11734*, 2025.
- Mohammed Baharoon et al. Evaluating general purpose vision foundation models for medical image analysis: An experimental study of dinov2 on radiology benchmarks. *arXiv:2312.02366*, 2023.
- Olivier Bernard, Alain Lalande, Clement Zotti, et al. Deep learning techniques for automatic mri cardiac multi-structures segmentation and diagnosis: Is the problem solved? *IEEE Transactions on Medical Imaging*, 37(11):2514–2525, 2018. doi: 10.1109/TMI.2018.2837502.
- Rishi Bommasani, Drew A Hudson, Ehsan Adeli, et al. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021.
- Rishi Bommasani, Kathleen A. Creel, Ananya Kumar, Dan Jurafsky, and Percy Liang. Picking on the same person: Does algorithmic monoculture lead to outcome homogenization? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- BraTS 2024 Organizers. Brats 2024 glioma challenge data. Synapse: syn53708249, 2024. URL <https://www.synapse.org/Synapse:syn53708249/wiki/628815>.
- Gabriele Campanella, Shengjia Chen, Manbir Singh, et al. A clinical benchmark of public self-supervised pathology foundation models. *Nature Communications*, 16:3640, 2025. doi: 10.1038/s41467-025-58796-1.
- Junlong Cheng, Jin Ye, Zhongying Deng, Jianpin Chen, Tianbin Li, Haoyu Wang, Yanzhou Su, Ziyang Huang, Jilong Chen, Lei Jiang, Hui Sun, Junjun He, Shaoting Zhang, Min Zhu, and Yu Qiao. Sam-med2d. *arXiv:2308.16184*, 2023.
- Noel Codella, Veronica Rotemberg, Philipp Tschandl, M. Emre Celebi, Stephen Dusza, David Gutman, Brian Helba, Aadi Kallou, Konstantinos Liopyris, Michael Marchetti, Harald Kittler, and Allan Halpern. Skin lesion analysis toward melanoma detection 2018: A challenge hosted by the international skin imaging collaboration (isic). *arXiv:1902.03368*, 2019.
- Edwin de Jong, Eric Marcus, and Jonas Teuwen. Current pathology foundation models are unrobust to medical center differences. *arXiv:2501.18055*, 2025.
- Maria Correia de Verdier, Rachit Saluja, Louis Gagnon, Dominic LaBella, Ujjwall Baid, Nourel Hoda Tahon, Martha Foltyn-Dumitru, Jikai Zhang, Maram Alafif, Saif Baig, et al. The 2024 brain tumor segmentation (brats) challenge: Glioma segmentation on post-treatment mri. *arXiv preprint arXiv:2405.18368*, 2024. doi: 10.48550/arXiv.2405.18368.
- Stanislav Fort, Huiyi Hu, and Balaji Lakshminarayanan. Deep ensembles: A loss landscape perspective. *arXiv preprint arXiv:1912.02757*, 2019.
- Yarin Gal and Zoubin Ghahramani. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. In *ICML*, 2016.
- M A Ganaie, Minghui Hu, A K Malik, M Tanveer, and P N Suganthan. Ensemble deep learning: A review. *Engineering Applications of Artificial Intelligence*, 115, 2022.

Robert Geirhos, Kristof Meding, and Felix A. Wichmann. Beyond accuracy: Quantifying trial-by-trial behaviour of cnns and humans by measuring error consistency. In *Advances in Neural Information Processing Systems*, 2020.

Mariana-Iuliana Georgescu, Radu Tudor Ionescu, and Andreea-Iuliana Miron. Diversity-promoting ensemble for medical image segmentation. In *Proceedings of the 38th ACM/SIGAPP Symposium on Applied Computing*, pages 599–606, 2023. doi: 10.1145/3555776.3577822.

Shashwat Goel, Joschka Struber, Ilze Amanda Auzina, Karuna K. Chandra, Ponnurangam Kumaraguru, Douwe Kiela, Ameya Prabhu, Matthias Bethge, and Jonas Geiping. Great models think alike and this undermines AI oversight. arXiv:2502.04313, 2025.

Raphael Gontijo-Lopes, Yann Dauphin, and Ekin Dogus Cubuk. No one representation to rule them all: Overlapping features of training methods. In *International Conference on Learning Representations (ICLR)*, 2022.

Mila Gorecki and Moritz Hardt. Monoculture or multiplicity: Which is it? In *NeurIPS*, 2025. URL <https://openreview.net/forum?id=D05LtJc80w>.

Hanxue Gu, Haoyu Dong, Jichen Yang, and Maciej A. Mazurowski. How to build the best medical image segmentation algorithm using foundation models: A comprehensive empirical study with segment anything model. *MELBA*, 2025. arXiv:2404.09957.

Yi Gu, Yi Lin, Kwang-Ting Cheng, and Hao Chen. Revisiting deep ensemble uncertainty for enhanced medical anomaly detection. In *MICCAI*, 2024.

Yufan He et al. Vista3d: A unified segmentation foundation model for 3d medical imaging. In *CVPR*, 2025.

Yuhao Huang, Xin Yang, Lian Liu, Han Zhou, Ao Chang, Xinrui Zhou, Rusi Chen, Junxuan Yu, Jiongquan Chen, Chaoyu Chen, Sijing Liu, Haozhe Chi, Xindi Hu, Kejuan Yue, Lei Li, Vicente Grau, Deng-Ping Fan, Fajin Dong, and Dong Ni. Segment anything model for medical images? *Medical Image Analysis*, 92: 103061, 2024. doi: 10.1016/j.media.2023.103061.

Fabian Isensee, Paul F Jaeger, Simon A A Kohl, Jens Petersen, and Klaus H Maier-Hein. nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18:203–211, 2021.

Fabian Isensee et al. nnu-net revisited: A call for rigorous validation in 3d medical image segmentation. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, 2024.

Debesh Jha, Pia H. Smedsrud, Michael A. Riegler, Pål Halvorsen, Thomas de Lange, Dag Johansen, and Håvard D. Johansen. Kvasir-SEG: A segmented polyp dataset. In *International Conference on Multimedia Modeling (MMM)*, 2020.

Liwei Jiang et al. Artificial hivemind: The open-ended homogeneity of language models (and beyond). arXiv:2510.22954, 2025.

Nathanael Jo, Nikhil Garg, and Manish Raghavan. The subjectivity of monoculture. arXiv:2602.24086, 2026.

Alain Jungo and Mauricio Reyes. Assessing reliability and challenges of uncertainty estimations for medical image segmentation. In *MICCAI*, 2019.

Tommie Kerssies et al. How to benchmark vision foundation models for semantic segmentation? In *CVPR Workshop*, 2024.

Elliot Kim, Avi Garg, Kenny Peng, and Nikhil Garg. Correlated errors in large language models. In *International Conference on Machine Learning (ICML)*, 2025a. arXiv:2506.07962.

Yunhe Kim et al. Show and segment: Universal medical image segmentation via in-context learning. arXiv:2503.19359, 2025b.

Alexander Kirillov, Eric Mintun, Nikhila Ravi, et al. Segment anything. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3992–4003, 2023. doi: 10.1109/ICCV51070.2023.00371.

Max Klabunde, Tobias Schumacher, Markus Strohmaier, and Florian Lemmerich. Similarity of neural network models: A survey of functional and representational measures. *Transactions on Machine Learning Research*, 2024.

Jon Kleinberg and Manish Raghavan. Algorithmic monoculture and social welfare. *Proceedings of the National Academy of Sciences*, 118(22), 2021.

443 Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network
444 representations revisited. In *ICML*, 2019.

445 Ludmila I Kuncheva and Christopher J Whitaker. Measures of diversity in classifier ensembles and their rela-
446 tionship with the ensemble accuracy. *Machine Learning*, 51(2):181–207, 2003.

447 Zeki Kuş and Musa Aydin. Medsegbench: A comprehensive benchmark for medical image segmentation in
448 diverse data modalities. *Scientific Data*, 11(1):1283, 2024. doi: 10.1038/s41597-024-04159-2.

449 Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty
450 estimation using deep ensembles. In *NeurIPS*, 2017.

451 Yuanyuan Li, Neeraj Sarna, and Yang Lin. Quantifying correlations of machine learning models. *arXiv preprint*
452 *arXiv:2502.03937*, 2025.

453 Xiangde Luo, Xiyue Wang, Feyisope Eweje, Xiaoming Zhang, et al. Ensemble learning of foundation models
454 for precision oncology. *arXiv:2508.16085*, 2025. URL <https://arxiv.org/abs/2508.16085>.

455 Jun Ma, Yuting He, Feifei Li, Lin Han, Chenyu You, and Bo Wang. Segment anything in medical images.
456 *Nature Communications*, 15, 2024.

457 Jun Ma, Zongxin Yang, Sumin Kim, Bihui Chen, Mohammed Baharoon, Adibvafa Fallahpour, Reza Asakereh,
458 Hongwei Lyu, and Bo Wang. Medsam2: Segment anything in 3d medical images and videos. *arXiv preprint*
459 *arXiv:2504.03600*, 2025.

460 Lena Maier-Hein et al. Metrics reloaded: Recommendations for image analysis validation. *Nature Methods*,
461 21:195–212, 2024. doi: 10.1038/s41592-023-02151-z.

462 Kenneth O. McGraw and S. P. Wong. Forming inferences about some intraclass correlation coefficients. *Psy-*
463 *chological Methods*, 1(1):30–46, 1996.

464 Alireza Mehrtash, William M Wells, Clare M Tempny, Purang Abolmaesumi, and Tina Kapur. Confidence
465 calibration and predictive uncertainty estimation for deep medical image segmentation. *IEEE Transactions*
466 *on Medical Imaging*, 39(12):3868–3878, 2020.

467 Dominik J. Mühlematter, Michelle Halbheer, Alexander Becker, Dominik Narnhofer, Helge Aasen, Konrad
468 Schindler, and Mehmet Ozgur Turkoglu. Lora-ensemble: Efficient uncertainty modelling for self-attention
469 networks. *Transactions on Machine Learning Research*, 2026. URL [https://openreview.net/forum?](https://openreview.net/forum?id=yhXXm0mpSQ)
470 [id=yhXXm0mpSQ](https://openreview.net/forum?id=yhXXm0mpSQ).

471 Basil Mustafa, Carlos Riquelme, Joan Puigcerver, André Susano Pinto, Daniel Keysers, and Neil Houlsby.
472 Deep ensembles for low-data transfer learning. *arXiv preprint arXiv:2010.06866*, 2020.

473 Peter Neidlinger, Omar S. M. El Nahhas, Hannah Sophie Muti, et al. Benchmarking foundation models as
474 feature extractors for weakly-supervised computational pathology. *Nature Biomedical Engineering*, 2025.
475 doi: 10.1038/s41551-025-01516-3.

476 Behnam Neyshabur, Hanie Sedghi, and Chiyuan Zhang. What is being transferred in transfer learning? In
477 *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

478 Connor Toups, Rishi Bommasani, Kathleen Creel, Sarah H. Bana, Dan Jurafsky, and Percy Liang. Ecosystem-
479 level analysis of deployed machine learning reveals homogeneous outcomes. In *Advances in Neural Infor-*
480 *mation Processing Systems (NeurIPS)*, 2023.

481 Danny Wood, Tingting Mu, Andrew M. Webb, Henry Reeve, Mikel Lujan, and Gavin Brown. A unified theory
482 of diversity in ensemble learning. *arXiv:2301.03962*, 2023.

483 Mitchell Wortsman, Gabriel Ilharco, Samir Ya Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes, Ari S. Morcos,
484 Hongseok Namkoong, Ali Farhadi, Yair Carmon, Simon Kornblith, and Ludwig Schmidt. Model soups:
485 Averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In
486 *International Conference on Machine Learning (ICML)*, 2022.

487 R. Patrick Xian, Noah R. Baker, Tom David, Qiming Cui, A. Jay Holmgren, Stefan Bauer, Madhumita Sushil,
488 et al. Robustness tests for biomedical foundation models should tailor to specifications. *npj Digital Medicine*,
489 8:557, 2025. doi: 10.1038/s41746-025-01926-2.

490 Yuncheng Yang, Meng Wei, Junjun He, Jie Yang, Jin Ye, and Yun Gu. Pick the best pre-trained model: Towards
491 transferability estimation for medical image segmentation. In *Medical Image Computing and Computer*
492 *Assisted Intervention (MICCAI)*, 2023.

- 493 Maximilian Zenk, Klaus H. Maier-Hein, Jens Petersen, et al. Comparative benchmarking of failure detection
 494 methods in medical image segmentation: Unveiling the role of confidence aggregation. *Medical Image*
 495 *Analysis*, 101:103392, 2025. doi: 10.1016/j.media.2024.103392.
- 496 Yukun Zhou, Mark A. Chia, Siegfried K. Wagner, et al. A foundation model for generalizable disease detection
 497 from retinal images. *Nature*, 622:156–163, 2023.
- 498 Jiayuan Zhu et al. Medical sam 2: Segment medical images as video via segment anything model 2.
 499 arXiv:2408.00874, 2024.

500 A1 Appendix roadmap and evidence inventory

501 The appendix is organised by evidentiary role rather than by experiment chronology. It starts with
 502 diagnostic readouts, then documents primary estimators, uncertainty, task-level sources, and estima-
 503 tor sensitivity. Middle sections collect same-backbone perturbations, segmentation-pipeline scope
 504 checks, distribution-shift and held-out stress tests, operational metrics, and M_{eff} . Later sections
 505 collect mechanism diagnostics and additional scope tables, and the appendix ends with the release
 506 manifest and artifact metadata. Secondary rows are retained only when they qualify or explain a
 507 paper claim, and their captions state when they use non-primary pools or splits. Unless explicitly
 508 marked primary or sourced from the primary results/_merged summaries, appendix values are
 509 diagnostic rather than primary support for Table 1.

510 **CLIP-L notation.** Throughout the appendix, CLIP-L/14 refers to the preserved OpenCLIP
 511 ViT-L-14/openai dense-token wrapper used to generate the released traces. It is not a claim of
 512 OpenAI clip.load parity; a strict QuickGELU-parity CLIP-L loader would be a separate sensitiv-
 513 ity with a distinct FM identifier.

Table 5: **Appendix guide.** Blocks describe the actual order of sections below; cross-references point to the detailed tables rather than repeating them in the main paper.

Block	Sections (in order)	Evidence role
Estimators and provenance	App. A2, A3, A4, A5, A6	Diagnostic readouts, pixel overlap, case/hierarchical bootstrap, functional-floor/ICC/permutation sensitivity, pair-quality OLS.
Operational metrics	App. A7, A8, A9, A10	Failure-event τ -sweep, conditional recovery, Kvasir secondary split, matched- M lift, MedSAM probe details.
Same-backbone controls	App. A11, A12, A13, A14	Subject clustering, BraTS sparse caveat, ACDC architecture caveat / random-init / LoRA scope, bootstrap-bagging.
Adaptation & distribution shift	App. A15, A16, A17, A18	RIGA cross-vendor, UNet-skip decoder, nnU-Net fold reslicing + case-identical, segmentation-recipe scope (loss/aug, length, dimensionality).
Held-out and meta-analyses	App. A19, A21, A22, A23, A24, A25	Held-out 5-task stress test, lever index, leave-one-family-out, lift-ceiling support, 3D pilot, matched-recipe adaptation scope.
Mechanism diagnostics	App. A26, A27, A28, A29, A30	Item-difficulty residualisation (cross-fitted + holdout), M_{eff} , random-effects representation, ViT cross-pretraining, cross-checkpoint same-family.
Additional scope diagnostics	App. A31, A31.1, A31.2, A31.3, A31.4, A31.5, A31.6, A31.7	Decoder capacity, random-init CNN, loss/aug sensitivity, fold CV, BraTS train-size, BraTS multimodal, MC dropout.
Artifact record	App. A32	Bundle manifest, claim-to-artifact map, compute scope, trace inventory.

514 A2 Diagnostic readouts and scope

515 This section records diagnostic readouts that clarify scope without adding new primary claims.

RIGA annotation source. Table 1 uses the six-rater majority consensus as ground truth for both RIGA Cup and RIGA Disc, matching the released RIGA per-case Dice traces. RIGA also distributes individual ophthalmologist masks, but individual-rater replay artifacts are not included in the supplementary artifact. The consensus Cup and Disc rows are therefore the only RIGA annotation-source evidence used by the manuscript.

Probability-map diagnostics. The artifact releases per-case Dice traces and selected aggregate pixel/lift diagnostics, but not raw probability maps. We therefore do not use mean-confidence, probability-calibration, or failure-flagging AUROC values as evidence for any paper claim.

A3 Pixel-level error-overlap breakdown

Table 6: **Pixel-level error-overlap diagnostic on selected secondary rows.** Values are same/cross gaps in error-mask overlap on aggregate diagnostic splits, not primary case-level evidence. For error-Dice, no bootstrap replicate out of $B=500$ produced a non-positive gap on any task cell, giving the finite- B bound $\Pr[\text{gap} \leq 0] < 1/501$; FP/FN/boundary rows are point-estimate breakdowns.

Metric	RIGA Cup	RIGA Disc	ACDC LV
Error-Dice gap	0.485	0.481	0.485
Error-Jaccard gap	0.515	0.481	0.521
FP-only error-Dice gap	0.453	0.494	0.566
FN-only error-Dice gap	0.527	0.464	0.449
Boundary-band (5px)	0.430	0.439	0.433

FP = (pred foreground, GT background); FN = (pred background, GT foreground); boundary-band restricts error masks to pixels within 5 pixels of the GT edge (trimap). All metrics are pair-wise Dice between the two models’ error masks (or empty-mask-excluded Jaccard). On RIGA Cup, RIGA Disc, and ACDC LV binary, the similar-magnitude gaps across FP, FN, and boundary variants indicate the spatial co-failure is not driven by a single error type on those three tasks (Kvasir and BraTS not tabulated in this split).

A4 Bootstrap sensitivity of the gap

Table 1 reports 95% CIs from a three-level nested hierarchical bootstrap that resamples family labels, seeds within each family, and cases ($B=2000$). We additionally compute a *case-only* bootstrap that fixes the family/seed structure and resamples the N test cases with replacement ($B=2000$). The two procedures target different sampling dimensions: the hierarchical CI jointly perturbs the finite FM/seed/case surface, whereas the case-only CI asks “what if we had a different test set” conditional on the evaluated pool. We report both as sensitivity evidence for the same finite audit rather than as population-level guarantees.

Case-only bootstrap intervals match the Case CI column of Table 1 exactly (the case-only resampling uses the same fixed family/seed structure as the per-row case CI), so we do not retabulate them here. The remainder of this section reports BraTS pool reconciliation and the slice-vs-subject sampling caveats.

BraTS foreground pool reconciliation. The 5-FM, 8-FM, and 8-FM-alt-estimator BraTS values cited elsewhere differ because they use disjoint analytic pools (5-FM: a subset of the 11-bundle source pool; 8-FM: a completed-only subset; alt-estimator: scoring with a different family-grouping convention). The principled headline is the post-floor 10-FM value because it uses the $\text{Dice} \geq 0.30$ functional floor and the Table 1 same/cross definition consistent with every other primary row. The primary BraTS foreground result is the 10-bundle Table 1 value of $\Delta=0.315$ [0.302, 0.329] (case-level CI), from the 11-bundle source pool with 10 retained bundles after functional-floor filtering (BiomedCLIP, CLIP-B/16, CLIP-L/14, ConvNeXt-T, DeiT-B/16, DINOv2-B/S, EfficientNet-B0, MAE-B/16, ResNet-50; ResNet-18 is below the Dice floor). Source: `results/_merged/paper_table.json`. Subset diagnostics without released aggregate files are omitted from the released numeric evidence. The value reported throughout the main text and abstract is therefore the Table 1 value.

555 All gap CIs strictly exclude zero under the case-level bootstrap, matching the family/seed-level CIs
556 in the main text.

557 **Per-rater RIGA Cup sensitivity (legacy non-replayable pool).** A non-replayable legacy sensi-
558 tivity pool (containing RETFound traces not in the released artifact) repeated the symmetric multi-
559 seed analysis using each of RIGA’s six individual rater masks; per-rater Pearson within–cross gaps
560 were 0.46–0.53 across all six raters, with the range across raters an order of magnitude smaller than
561 the within-vs-cross gap itself. The released Table 1 RIGA Cup row remains the bundle-replayable
562 evidence.

563 A5 Estimator-robustness and negative-control checks

564 The primary row uses Pearson correlation and a $\text{Dice} \geq 0.30$ per-FM functional floor.
565 To make this choice auditable, the released bundle includes a sensitivity artifact,
566 `results/_merged/diagnostics/audit_sensitivity.json`, generated from the same
567 per-case Dice JSONs by `code/scripts/compute_audit_sensitivity.py`. It recomputes
568 the within–cross gap under a floor sweep, an ICC(A,1) alternative, leave-one-FM-out summaries,
569 leave-one-task-out meta summaries, and a random family-label permutation control.

Table 7: **Estimator sensitivity on primary traces.** Pearson values are the Table 1 point estimates at the $\text{Dice} \geq 0.30$ floor. ICC(A,1) is an alternate agreement statistic on the same retained members. The floor range reports the minimum/maximum Pearson gap over $\{0, 0.20, 0.25, 0.30, 0.35, 0.40\}$ where the pool is non-empty. The permutation note reports a non-exchangeable random family-label stress test over completed members, using $(1 + \#\{\Delta_{\text{null}} \geq \Delta_{\text{obs}}\})/(1001)$; this is not interpreted as a p -value. Retained counts at every floor are in `audit_sensitivity.json`.

Task	Members at 0.30	Pearson gap	ICC(A,1) gap	Floor-sweep gap range
Kvasir	44	0.261	0.333	0.261 ^b
ACDC LV	40	0.321	0.401	0.242–0.355
BraTS foreground	40	0.315	0.404	0.315–0.362
RIGA Cup	36	0.473	0.539	0.384–0.529
RIGA Disc	44	0.638	0.716	0.638 ^b

All five rows have observed-gap-exceeded rank fraction 0.001 at the 0.30 floor. ^b Value is constant across all valid floors in $\{0, 0.20, 0.25, 0.30, 0.35, 0.40\}$.

570 The five-task mean Pearson gap at the paper floor is 0.401 (median 0.321, range 0.261–0.638).
571 Leave-one-task-out means stay in 0.342–0.437; the smallest value occurs when the largest-gap
572 RIGA Disc row is held out. These checks do not add new cases or new model training. They
573 instead reduce the risk that the primary sign is an artifact of the exact functional floor, the Pearson
574 statistic, or arbitrary family labels.

575 **Calibration-selected functional pools.** The primary table defines a functional pool using all re-
576 leased evaluation traces. To check whether the sign depends on this convention, the scorer also
577 makes a deterministic split of each task’s aligned units, selects bundles by mean Dice on the cal-
578 ibration subset, and evaluates the same/cross gap on disjoint cases. This is a leakage-sensitivity
579 analysis rather than a prospective validation protocol, but it addresses the specific concern that the
580 functional-floor sign is induced by evaluating the floor and the gap on exactly the same cases.

581 A6 Quality-controlled pairwise-correlation diagnostic

582 A possible alternative explanation is that same-FM pairs have higher correlation merely because
583 they have matched marginal quality. We therefore add a CPU-only diagnostic that uses the same
584 released per-case Dice traces retained by Table 1 and fits a pair-level OLS model:

$$r_{ab} \sim I_{\text{same}} + \bar{D}_{ab} + |\Delta \bar{D}_{ab}| + \bar{V}_{ab} + |\Delta V_{ab}|.$$

585 Here $\bar{D}_{ab}$ and $\bar{V}_{ab}$ are the pair mean Dice and average per-case Dice variance, with absolute
586 pair differences as additional controls. The same-group indicator is exactly the primary group-
587 ing rule: same FM plus the DINOv2-B/S pair. The coefficient is not a causal estimand, but it

Table 8: **Calibration-selected functional-pool sensitivity.** Functional bundles are selected on the calibration subset and evaluated on disjoint cases under the Table 1 same/cross grouping. Unit-agg. Δ is reported only for nested slice datasets after patient/subject aggregation. ACDC can select 11 bundles here because selection is calibration-only, whereas Table 1 applies the retrospective all-evaluation functional floor. CIs are case-bootstrap intervals on the evaluation subset. Source: `results/_merged/calibration_split/*.json`, regenerated by `code/scripts/compute_all.py`.

Task	Cal./eval units	Selected bundles	Eval. Δ [95% CI]	Unit-agg. Δ
Kvasir	42/78 images	11	0.250 [0.196, 0.322]	–
ACDC LV	7/13 patients	11	0.350 [0.313, 0.392]	0.206
BraTS foreground	113/211 subjects	10	0.307 [0.292, 0.323]	0.205
RIGA Cup	33/62 images	9	0.478 [0.369, 0.599]	–
RIGA Disc	33/62 images	11	0.630 [0.534, 0.695]	–

588 checks that the main sign is not removed by simple pair-quality and variance controls. The diag-
589 nostic is generated by `code/scripts/compute_quality_controlled_pairs.py` and released
590 as `results/_merged/diagnostics/quality_controlled_pairs.json`.

Table 9: **Pair-quality controlled diagnostic.** β_{same} is the coefficient on the primary same-group indicator in a pair-level OLS check. Same-group coefficients remain positive after controlling for pair mean Dice, absolute mean-Dice difference, average per-case variance, and absolute variance difference. CIs are case-bootstrap intervals from $B=1000$ refits and should be read as diagnostic, not causal, uncertainty.

Task	Members	Same pairs	Cross pairs	β_{same} [95% CI]
Kvasir	44	82	864	0.169 [0.124, 0.220]
ACDC LV	40	76	704	0.208 [0.175, 0.242]
BraTS foreground	40	76	704	0.224 [0.210, 0.239]
RIGA Cup	36	70	560	0.372 [0.297, 0.456]
RIGA Disc	44	82	864	0.582 [0.495, 0.639]

591 A7 Failure threshold sensitivity

592 The primary estimand is continuous per-case Dice correlation. To make the co-failure interpretation
593 less dependent on continuous-score correlation, the released artifact also recomputes the Table 1
594 same/cross grouping on binary failure events. Table 10 reports the empirical P_{33} readout across all
595 five primary task rows; the released JSON also contains P_{25} and P_{50} thresholds. The RIGA Cup
596 all-fail table below is retained as a pool-level threshold diagnostic, and absolute-threshold lift with
597 support flags is in Appendix A9.

Table 10: **Binary failure-event diagnostic at the empirical P_{33} threshold.** Failure is defined as per-case Dice below the 33rd percentile of all model-case Dice values after functional-floor filtering. ϕ is binary-event Pearson correlation averaged over same-group or primary cross-group pairs. Joint fail same/cross is the pairwise joint-failure probability averaged over the corresponding model-trace pairs. CIs are case-bootstrap intervals for the same-cross ϕ gap. This is not a clinically chosen threshold. Source: `results/_merged/diagnostics/failure_event_summary.json`, regenerated by `code/scripts/compute_failure_event_summary.py`.

Task	τ_{P33}	Same ϕ	Cross ϕ	Gap [95% CI]	Joint fail same/cross
Kvasir	0.659	0.869	0.540	0.329 [0.278, 0.388]	27.9/21.6%
ACDC LV	0.548	0.875	0.489	0.386 [0.354, 0.418]	27.7/19.9%
BraTS foreground	0.636	0.838	0.425	0.413 [0.401, 0.425]	27.3/18.7%
RIGA Cup	0.575	0.706	0.253	0.453 [0.399, 0.509]	24.4/14.4%
RIGA Disc	0.857	0.703	0.125	0.578 [0.511, 0.631]	24.8/12.0%

598 **Full τ -sweep of conditional recovery.** Table 11 tabulates grand-mean conditional recovery
599 rates $P[\text{pool-mean Dice} \geq \tau \mid \text{any member Dice} < \tau]$ for same-bundle pools (4 seeds
600 of one FM, averaged over all FMs passing the per-task Dice ≥ 0.30 functional floor) and
601 diverse pools (all $\binom{K}{4}$ functional-subset compositions and all released seed tuples) at $\tau \in$

602 {0.70, 0.75, 0.80, 0.85, 0.90}. Source: results/_merged/diagnostics/tau_sweep.json, re-
603 generated by code/scripts/compute_tau_sweep.py. Recovery is τ -fragile in both regimes; the
604 diverse – same-bundle recovery-rate gap $\Delta_{\text{rec}}(\tau)$ (column header below) uses the opposite sign
605 convention to the primary Δ in Table 1 (recovery favours diverse so $\Delta_{\text{rec}} > 0$ is the operationally
606 desired direction) and is non-monotone in τ .

Table 11: **Secondary τ -sweep of conditional recovery rate $P[\text{ens pass} \mid \text{any fail}]$.** Grand-mean across all functional 4-of- K diverse compositions and all functional same-bundle baselines. Recovery gaps vary with threshold and can become near-zero or negative near ceiling; these values are operational diagnostics, not primary support for the same/cross correlation claim.

Task	τ	Mono	Diverse	Δ_{rec}
RIGA Cup (9-bundle)	0.70	0.061	0.214	+0.153
	0.75	0.073	0.120	+0.047
	0.80	0.055	0.050	−0.005
	0.85	0.010	0.007	−0.003
	0.90	0.001	0.000	−0.001
ACDC LV (10-bundle)	0.70	0.074	0.285	+0.211
	0.75	0.062	0.237	+0.175
	0.80	0.049	0.163	+0.114
	0.85	0.040	0.072	+0.032
	0.90	0.020	0.012	−0.007
Kvasir-SEG (11-bundle)	0.70	0.092	0.388	+0.295
	0.75	0.074	0.360	+0.285
	0.80	0.058	0.267	+0.209
	0.85	0.050	0.156	+0.106
	0.90	0.025	0.030	+0.005
BraTS foreground (2D FLAIR seg>0)	0.70	0.080	0.326	+0.245
	0.75	0.061	0.233	+0.172
	0.80	0.039	0.129	+0.090
	0.85	0.019	0.035	+0.016
	0.90	0.004	0.001	−0.003

607 At the strictest tested threshold ($\tau=0.90$), the diverse advantage in conditional recovery vanishes
608 and reverses on RIGA Cup ($\Delta = -0.001$) and ACDC LV ($\Delta = -0.007$), while remaining positive
609 on Kvasir (+0.005) and slightly negative on BraTS (−0.003). The diverse-pool benefit is concen-
610 trated in the moderate-failure regime ($\tau \in [0.70, 0.80]$), consistent with the primary correlation gap
611 predicting reduced *but not zero* coverage benefit at extreme thresholds.

Table 12: **RIGA Cup threshold diagnostic.** Values are for the released primary RIGA Cup functional pool only. Mono pools are same-bundle four-seed pools averaged over functional FMs; diverse pools are four distinct functional bundles. Percentile thresholds are computed over all model-case Dice values after the Dice ≥ 0.30 functional floor. Source: results/_merged/diagnostics/threshold_table_riga_cup.json.

Threshold	Mono fail $\bar{r}$	Diverse fail $\bar{r}$	Mono ALL-fail	Diverse ALL-fail
$P_{25} (\tau=0.513)$	0.833	0.209	19.9%	2.2%
$P_{33} (\tau=0.575)$	0.803	0.257	26.5%	5.3%
$P_{50} (\tau=0.686)$	0.840	0.268	44.8%	15.4%

612 A8 Kvasir secondary split diagnostic

613 Pixel-level diagnostic on this split: mono 0.80 vs. diverse 0.40, gap 0.394 [95% CI 0.38, 0.40].
614 Absolute-threshold lift at Dice<0.7: mono 252 $\times$ vs. diverse 120 $\times$ (ratio 2.1 $\times$); at Dice<0.8: 48.4 $\times$
615 vs. 14.6 $\times$ (3.3 $\times$); at Dice<0.9: 4.6 $\times$ vs. 1.8 $\times$ (2.5 $\times$). The numbers are slightly smaller than
616 RIGA/ACDC, likely reflecting the larger test set (300 cases vs. 224 / 90) and more variable polyp
617 presentations, but qualitatively identical: mono still shows substantially higher correlated-failure lift
618 than diverse.

Table 13: **Kvasir-SEG secondary split diagnostic.** This generic 8-bundle summary uses a 1,000-image, 700/300 train/test split and is not the 11-bundle, 880/120 Kvasir row in Table 1. Seven of eight generic encoder bundles produce functional segmentation (Dice>0.40); CLIP ViT-L/14 is non-functional and excluded from cross-family analysis. MedSAM-only traces are reported separately as an encoder-probe summary in Appendix A10; no RETFound Kvasir trace is included in the released artifact bundle.

FM Family	Dice	Within $\bar{r}$	Status
DINOv2 ViT-B/14	0.819	0.972	functional
DINOv2 ViT-S/14	0.807	0.951	functional
BiomedCLIP ViT-B/16	0.717	0.954	functional in this split
ConvNeXt-Tiny	0.651	0.949	functional
EfficientNet-B0	0.582	0.945	functional
ResNet-50	0.540	0.880	functional
ResNet-18	0.468	0.876	functional
CLIP ViT-L/14	0.174	0.839	non-functional

A9 Matched- M lift comparison

To ensure the mono vs. diverse comparison is not driven by pool size, we construct an *oracle* matched- $M=4$ diverse pool by selecting the top-4 bundles by *test-set seed-averaged* mean Dice (rank computed across all four decoder seeds, avoiding seed-42 selection circularity; selected bundles per task are recorded in the appendix diagnostic output—e.g. Kvasir top-4 = {DINOv2-B, DINOv2-S, BiomedCLIP, ConvNeXt-T}, Cup top-4 = {BiomedCLIP, DINOv2-B, DINOv2-S, EfficientNet-B0}). This is deliberately favourable to the diverse pool and useful for a matched-size structural comparison, but it is not a deployable pool-selection rule because the ranking uses the same test set on which lift is evaluated. For each task / threshold we report three lift variants:

- **Matched-4 seed-42 (single-seed diagnostic):** the single-seed pool; retained for comparability with a single-seed estimator.
- **Matched-4 symmetric permuted (preferred):** each of the 4 FMs takes a *distinct* decoder seed from {42, 43, 44, 45}; the lift is averaged on the log scale over all $4! = 24$ seed-to-FM permutations, then exponentiated. This is the matched-size structural analog of the main paper’s symmetric Pearson $\bar{r}$ estimator (mono = DINOv2-B with 4 distinct seeds matches matched-4 = 4 FMs with 4 distinct seeds).
- **Matched-4 symmetric independent (sensitivity):** each FM’s seed is drawn i.i.d. with replacement from {42..45}, $4^4 = 256$ tuples; reported as a robustness check (the permuted variant is preferred because mono uses four *distinct* seeds).

All bootstrap CIs are computed on *cases* only (patient-clustered on ACDC); the seed tuple space is finite and fully enumerated, not bootstrapped. We report 95% case-bootstrap CIs on the log-scale lift and exponentiate to the reported interval ($B=2000$). We also report the raw all-fail count (AF_{cnt}) and the Poisson-binomial independence expectation $\mathbb{E}[AF_{ind}]$; when $N \cdot \mathbb{E}[AF_{ind}] < 5$ we flag the row as support-limited. Under the matched-size comparison, the adequately supported rows are RIGA Cup $\tau=0.9$, RIGA Disc $\tau=0.95$, and ACDC LV $\tau=0.9$; Kvasir $\tau=0.8$ is mono-sparse and retained as a descriptive diagnostic to match Fig. 1. Across these selected rows the mono/matched-4 lift ratio is $1.5\text{--}5.8\times$ (Table 14); lower-support rows give larger ratios ($5.9\text{--}17\times$) treated as descriptive order-of-magnitude only. We therefore treat τ values well below the task-specific Dice ceiling as support-limited and report them as descriptive rather than precise lift estimates.

A10 Medical-domain encoder probe details

We implemented two medical-domain encoder paths, but only the MedSAM image-encoder probe is reported in the released evidence bundle:

RETFound is a ViT-L/16 MAE pretrained on 1.6M retinal images [Zhou et al., 2023]. The code path builds the RETFound-MAE CFP ViT-L/16 trunk and expects the gated HuggingFace checkpoint YukunZhou/RETFound_mae_natureCFP or a user-specified checkpoint path; the official upstream project is rmaphoh/RETFound_MAE. RETFound MAE pretraining uses ImageNet [0.485, 0.456, 0.406]/[0.229, 0.224, 0.225] normalisation; our frozen-probe adaptation would drop

Table 14: **Matched- $M=4$ lift diagnostic.** Mono is DINOv2-B $\times$ 4 seeds; matched-4 diverse uses the top-4 bundles by test-set seed-averaged mean Dice, so this is an oracle structural comparison, not a deployable selection rule. CIs are 95% case-bootstrap intervals on log-lift, exponentiated. Rows without adequate support ($\checkmark$ missing) are descriptive order-of-magnitude indicators.

Task	τ	Mono [CI]	Diverse-42	Diverse-perm [CI]	Ratio*	Supp.
RIGA Cup	0.8	4450 [1408, 24012]	174 [79, 467]	265 [130, 696]	$17\times$	–
RIGA Cup	0.9	2.18 [1.81, 2.69]	1.33 [1.20, 1.47]	1.44 [1.31, 1.60]	$1.5\times$	$\checkmark$
RIGA Disc	0.95	9.80 [6.82, 14.6]	2.87 [2.19, 3.82]	3.11 [2.44, 3.98]	$3.2\times$	$\checkmark$
ACDC LV	0.5	2149 [462, 34449]	200 [51, 1026]	188 [62, 817]	$11\times$	–
ACDC LV	0.7	189 [55.6, 1064]	31.2 [15.5, 77.7]	25.3 [13.1, 57.8]	$7.5\times$	–
ACDC LV	0.8	33.8 [15.1, 107]	5.87 [3.60, 10.6]	5.71 [3.63, 9.93]	$5.9\times$	–
ACDC LV	0.9	2.23 [1.63, 3.32]	1.39 [1.22, 1.60]	1.40 [1.25, 1.61]	$1.6\times$	$\checkmark$
Kvasir	0.7	246 [128, 570]	33.9 [22.3, 54.5]	30.8 [20.4, 48.3]	$8.0\times$	–
Kvasir	0.8	48.0 [28.9, 85.1]	8.72 [6.62, 12.0]	8.21 [6.35, 11.1]	$5.8\times$	mono-sparse

* Ratio = Mono / Matched-4 PERM. **Bold** = preferred symmetric-permuted estimator (24 seed-to-FM tuples). Matched-4 seed 42 is a single-seed diagnostic. Matched-4 INDEP ($4^4=256$ iid tuples) produces values within $\pm 1.5\%$ of PERM on every row and is omitted for compactness; the full table is a secondary diagnostic rather than the `results/_merged` primary scorer output. Supp. $\checkmark$ iff $N \cdot \mathbb{E}[\text{AF}_{\text{ind}}] \geq 5$ for the PERM estimator. Mono lift values span four orders of magnitude ($1.4\text{--}4,450\times$); precision is reported to the data’s underlying decimals and is not uniform across the lift column.

the CLS token and reshape the $196=14\times 14$ patch tokens into a spatial 14×14 feature map of dimension 1024. That patch-token extraction is not RETFound’s official global/CLS-style classification fine-tuning path. Because the gated checkpoint was unavailable in the release environment, RETFound cells are not reported and no RETFound result supports the paper’s empirical claims.

MedSAM is a SAM ViT-B fine-tuned on 1.5M medical image-mask pairs [Ma et al., 2024]. The code path uses the official `bowang-lab/MedSAM medsam_vit_b.pth` checkpoint through `segment_anything.sam_model_registry['vit_b']` with `MEDSAM_CHECKPOINT` pointing to that checkpoint. In this audit, we use *only the image encoder* as a frozen feature extractor with our same 1×1 -conv decoder. We discard MedSAM’s prompt encoder and mask decoder entirely—i.e. MedSAM enters this boundary probe on the same frozen-feature terms as the other FMs (no prompt-based path, no bounding-box input, no iterative refinement). The wrapper strips the SAM image-encoder neck and uses the pre-neck ViT-B representation: the official 1024×1024 positional grid is bicubically interpolated to the 224×224 input grid, yielding a 14×14 spatial feature map of dimension 768. This pre-neck $768\times 14\times 14$ tensor is an internal representation, not the official SAM/MedSAM post-neck image embedding at 1024 input. This means MedSAM in our experiment is a “MedSAM image-encoder probe”, not its intended prompt-conditioned segmentation system or official inference preprocessing path. The comparison is therefore an encoder-only boundary probe for the same-bundle dependence question we ask, and it does not evaluate MedSAM as released or its best-case performance with its own prompt-based decoder.

All other aspects of the secondary sensitivity use the same four decoder seeds, 30-epoch Adam 10^{-3} training, and identical 1×1 conv decoder. The completed MedSAM-only probe covers ACDC LV, Kvasir, RIGA Cup, and RIGA Disc (four seeds each). Mean Dice across seeds is 0.517, 0.513, 0.734, and 0.908 respectively; the aggregate diagnostic is released as `results/_merged/diagnostics/medsam_only_summary.json`. We use this as scope context rather than primary evidence.

A11 BraTS/ACDC clustering sensitivity

The Table 1 BraTS foreground row is a BraTS 2024 GLI 2D FLAIR derivative with $n=3,228$ test slices from 324 held-out subjects: per held-out subject, the loader keeps up to the top-10 axial slices by foreground area, requires at least 64 positive pixels, and uses a subject-level 80/20 split with NumPy seed 42. The 3,228 count is traceable to the released `test_ids`: 321 held-out subjects contribute ten eligible slices, one contributes nine, one contributes six, and one contributes three, giving a 12-slice deficit from the 324×10 maximum. ACDC LV similarly evaluates 361 LV-positive slices nested under 20 held-out patients. The main table is therefore a case-level audit, but the released `subject_level/*.json` files include the patient/subject unit IDs linked to `test_ids` and two cluster-aware checks. (i) **Cluster bootstrap over the case-scale statistic.** Resampling patients/subjects with replacement and keeping all slices within each sampled unit gives ACDC $\Delta=0.321$ [0.269, 0.374] and BraTS foreground $\Delta=0.315$ [0.283, 0.351], so the case-scale gap remains positive under unit-level resampling. (ii) **Subject-aggregated statistic.** Averaging each

694 member’s per-slice Dice within patient/subject before computing Pearson gives smaller but still
695 positive gaps: ACDC $\Delta=0.262$ [0.151, 0.413] and BraTS foreground $\Delta=0.214$ [0.182, 0.251].
696 The conservative reading is that slice dependence inflates the magnitude relative to subject means,
697 especially on BraTS, but it does not explain the within-vs-cross separation away.

698 **A12 BraTS sparse-class filtering caveat and released diagnostics**

699 Per-class BraTS sensitivity can be inflated by sparse labels: if a class is absent from a slice,
700 empty-vs-empty Dice can return a perfect score and obscure the foreground-error structure that
701 matters for a segmentation audit. For this reason, the primary BraTS row uses the foreground
702 derivative $\text{seg}>0$ rather than a sparse NCR-only target. NCR filtering diagnostics from a sep-
703 arate sparse-class sweep are not used as numeric evidence here because their exact aggregate
704 files are not included in the supplementary artifact. The released BraTS foreground train-
705 size and multimodal diagnostics are instead reported later in Appendices A31.5–A31.6 from
706 `results/_merged/late_brats_summaries.json`.

707 **A13 Scope-only exploratory probes**

708 *Note on CLIP-L on ACDC.* The secondary 8-FM ACDC diagnostic discussed in this section uses a
709 different split where ResNet-50, ResNet-18, and CLIP ViT-L/14 collapse below the $\text{Dice} \geq 0.30$ floor
710 (Dice 0.099, 0.230, and 0.007 respectively); the *primary* ACDC LV pool of Table 1 (post-floor 10-
711 bundle subset of the 11-bundle source pool, with only ResNet-18 dropped) retains CLIP ViT-L/14
712 because it passes the Dice floor on the primary split. Different splits, different functional outcomes
713 for the borderline bundles.

714 This block separates released diagnostics from exploratory mechanism probes without released trace
715 artifacts or aggregate JSONs. Only released traces or released aggregate JSONs are used as numeri-
716 cal support for the paper’s claims. The shared reading is descriptive: same-backbone perturbations
717 can reduce correlation, but the released diagnostics do not make same-backbone pools behave like
718 the primary cross-family references.

719 **ACDC architecture caveat.** Secondary ACDC exploratory tables used a separate 8-bundle split
720 and architecture-pair subsets that are not included as replayable aggregate artifacts. The released
721 support for the architecture-caveat wording is instead the same-architecture-family ViT / different-
722 pretraining diagnostic in Appendix A29 and the expanded same-family cross-checkpoint decompo-
723 sition in Appendix A30.

724 **Random-init probes.** DINOv2-B random-initialisation probes used secondary splits and source
725 prediction caches outside the released artifact. The released no-pretraining evidence is limited to the
726 CNN random-init diagnostic in Appendix A31.2, which should be read as a descriptive same-recipe
727 sanity check rather than a causal isolation of architecture, optimiser, and data effects.

728 **LoRA and cross-architecture probes.** LoRA probes and ViT-B cross-architecture random-init
729 sweeps require source prediction caches and aggregate summary files outside the released artifact.
730 We do not tabulate their values or use them as bundle-supported evidence. The released same-
731 backbone levers are the head-capacity sweep, UNet-skip replication, fold-reslicing diagnostic, CNN
732 random-init diagnostic, and MC dropout.

733 **A14 Same-backbone, different-data bootstrap-bagging audit**

734 For each (FM, dataset) pair, four bootstrap replicates of the training set (image-level on
735 RIGA/Kvasir; subject-level on BraTS) are crossed with four decoder seeds. Bootstrap data-swap
736 reduces within-pool $\bar{r}$ by only 0.05–0.12, while backbone swap reduces it by 0.33–0.50. This sup-
737 ports the scoped claim that, in the frozen-feature protocol, resampling the training cases does not
738 provide the same decorrelation as changing encoder family.




Figure 2: **Same-backbone bootstrap-bagging diagnostic.** For each shown dataset/backbone pair, we compare (i) same-backbone/same-data cross-seed correlation, (ii) same-backbone/bootstrap-resampled-data correlation, and (iii) a cross-backbone reference from the corresponding diagnostic pool. The ratio panel reports $(r_{\text{same data}} - r_{\text{cross backbone}})/(r_{\text{same data}} - r_{\text{bootstrap data}})$, so values above one mean that changing backbone reduces correlation more than bootstrap-resampling the training data in this diagnostic. The BraTS label refers to the 2D foreground derivative (seg>0), not the official 3D BraTS WT challenge target. This figure is descriptive and does not identify a causal mechanism.

739 A15 Distribution-shift audit

740 We evaluate one released cross-distribution robustness diagnostic for the within-cross correlation
741 gap using an 8-bundle RIGA pool and the frozen-feature protocol from Section 4.

742 **RIGA leave-MESSIDOR-out (cross-vendor fundus).** RIGA+ [Almazroa, 2018] combines three
743 imaging sources: *BinRushed* (Bin Rushed Ophthalmic Center, Riyadh, Saudi Arabia), *Magrabia*
744 (Magrabi Eye Center, Riyadh, Saudi Arabia), and *MESSIDOR* (French public corpus, different
745 cameras and populations). We train the frozen-FM + 1×1 conv decoder on BinRushed+Magrabia
746 (union) and evaluate on MESSIDOR (454 cases after the preprocessing and eligibility filters im-
747 plemented by the released scorer) as cross-vendor out-of-distribution. Per-case traces for this diag-
748 nostic are released under `results/_merged/per_case_dice/riga_cup_ood_messidor`. The
749 released scorer uses a fixed 8-bundle $\times$ 4-seed pool without a functional-floor change between
750 ID and OOD rows. Under that directly reproducible convention, the ID row has within/cross/gap
751 0.659/0.282/0.377 and the MESSIDOR OOD row has 0.743/0.390/0.353.

752 **Pool-level summary.** In this fixed-pool RIGA cross-vendor diagnostic, the within-cross gap
753 changes by -0.024 and remains positive. Both within and cross correlations are higher on the
754 MESSIDOR row, so this diagnostic should be read only as evidence that the gap survives one re-
755 leased cross-vendor stress test, not as evidence that distribution shift generally contracts dependence.
756 Broader cross-scanner, cross-demographic, and longitudinal shifts remain out of scope.

Table 15: **RIGA cross-vendor diagnostic under a fixed 8-bundle scorer.** This is not case-identical to Table 1. Within = same-FM different-seed $\cup$ same-family different-FM (DINOv2-B/S and ResNet-18/50); cross = different-family different-FM. The gap remains positive, but both within and cross correlations increase on MESSIDOR; no general distribution-shift conclusion is supported. Brackets are case-level bootstrap 95% CIs ($B=500$). Source: `results/_merged/diagnostics/distshift_riga_messidor.json`.

Shift regime	Within- $\bar{r}$	Cross- $\bar{r}$	Gap
ID reference (fixed 8-bundle pool)	0.659 [0.603, 0.715]	0.282 [0.197, 0.369]	0.377 [0.301, 0.463]
OOD MESSIDOR (fixed 8-bundle pool)	0.743 [0.722, 0.765]	0.390 [0.345, 0.430]	0.353 [0.318, 0.388]

This diagnostic uses a paired fixed 8-bundle RIGA pool and should not be read as a case-identical comparison to the Magrabia-source held-out primary RIGA rows in Table 1.

757 A16 UNet-with-skip decoder replication

758 The paper’s primary decoder is a 1×1 conv with $C+1$ trainable parameters for C frozen channels
759 (for example, 769 parameters for a 768-channel feature map), which is close to a linear probe on

frozen features. Dense decoders with skip connections—nnU-Net, U-Net, SegFormer—are closer to deployed medical segmentation. We therefore ran a separate diagnostic in which the frozen backbone is unchanged but the head is replaced by a multi-scale UNet-style decoder over four encoder feature levels. The diagnostic uses two released seeds (42/43) per cell; “seed-pair r ” is therefore a single Pearson coefficient, not a bootstrapped four-seed within-family estimate.

Architecture. For hierarchical CNNs (ConvNeXt-T, ResNet-50, ResNet-18, EfficientNet-B0) the decoder uses the native feature pyramid as skip sources; for ViT-style encoders (DINOv2-B/S, BiomedCLIP, MAE-B, DeiT-B) it uses sampled intermediate token maps reassembled into a four-level pyramid. The released head first applies lateral 1×1 projections to a common decoder width, then runs a top-down bilinear-upsample path with three additive skip merges (deep $\rightarrow$ mid $\rightarrow$ shallow $\rightarrow$ shallowest), two conv-BN-ReLU blocks after each merge, a final conv-BN-ReLU after up-sampling to the output grid, and a final 1×1 output conv. Training matches the frozen-feature recipe where possible: Adam $\text{lr } 10^{-3}$, batch 16, 30 epochs, BCE, seeds $\{42, 43\}$.

Coverage. The 9-backbone $\times$ 5-task expansion adds (i) ResNet-18, EfficientNet-B0, MAE-B and DeiT-B as four new backbones for Kvasir, ACDC LV, RIGA Cup and RIGA Disc, and (ii) RIGA Disc as a fifth task row for the original 5 backbones (DINOv2-B/S, BiomedCLIP, ConvNeXt-T, ResNet-50). The BraTS column for the four new backbones requires a different data-root layout on the host that ran the expansion (BraTS-2024-GLI subjects live under `training/training_data1_v2/` rather than at the data-root top level expected by `fmppool.datasets.brats`); those four cells are therefore marked -- and the BraTS row for the original 5 backbones is unchanged from the prior diagnostic. Five cells (kvasir/EfficientNet-B0, acdc_lv/ResNet-18, riga_cup/MAE-B, riga_disc/DINOv2-S, riga_disc/MAE-B) are marked TBD because one of the two seeds was still finalising at submission time; those cells are produced from the same protocol and will be filled in a follow-up commit.

Table 16: **UNet-skip same-backbone diagnostic, 9 backbones $\times$ 5 tasks (two seeds 42/43 per cell).** Each cell reports “mean Dice (averaged over the two seeds) / seed-pair Pearson r ”. Single-pair r is a high-variance estimator; the reading is that within-backbone correlations stay clearly positive on every reported cell while pool-mean Dice rises substantially over the matching 1×1 probe rows. Primary Table 1 cross-bundle references for context: Kvasir 0.697, ACDC LV 0.639, RIGA Cup 0.361, BraTS fg 0.623, RIGA Disc 0.232. Cells marked -- are not in scope (BraTS expansion for the four new backbones, see Coverage paragraph). Cells marked TBD are pending the second seed.

Backbone	Kvasir	ACDC LV	RIGA Cup	BraTS fg	RIGA Disc
DINOv2-B	0.883/0.877	0.864/0.979	0.902/0.967	0.877/0.936	0.963/0.426
DINOv2-S	0.866/0.919	0.863/0.970	0.894/0.893	0.875/0.886	TBD
BiomedCLIP	0.806/0.926	0.827/0.950	0.880/0.892	0.858/0.929	0.960/0.874
ConvNeXt-T	0.860/0.775	0.885/0.917	0.895/0.970	0.895/0.938	0.966/0.809
ResNet-50	0.816/0.774	0.882/0.945	0.881/0.952	0.889/0.934	0.965/0.873
ResNet-18	0.774/0.828	TBD	0.883/0.814	—	0.961/0.651
EfficientNet-B0	TBD	0.881/0.938	0.899/0.888	—	0.965/0.787
MAE-B	0.795/0.827	0.846/0.955	TBD	—	TBD
DeiT-B	0.837/0.920	0.849/0.945	0.893/0.941	—	0.963/0.880

Finding. The UNet-skip head improves mean Dice substantially over several weak 1×1 probe cells, but it does not remove same-backbone co-failure. Across the 36 evaluated cells, single seed-pair r values stay positive (range 0.426–0.979) while pool-mean Dice rises into the 0.77–0.97 range; the within-backbone seed correlation is therefore concentrated well above the primary cross-bundle references in Table 1 on every reported cell. The single-pair r is high-variance (one Pearson coefficient per cell), so individual outliers like the DINOv2-B RIGA Disc cell ($r = 0.426$) should not be over-read; what survives the noise is that no cell yields a near-zero seed-pair correlation and no cell approaches the cross-bundle reference for its task. The conclusion is narrower than a decoder-family claim: a stronger skip decoder can improve accuracy and sometimes reduce seed correlation, but under this frozen-backbone two-seed grid it does not make same-backbone pools resemble the cross-family references in Table 1. A full nnU-Net v2 5-fold audit on RIGA Cup is reported in Appendix A17.

A17 Fold-reslicing and nnU-Net complements

Motivation. nnU-Net [Isensee et al., 2021] is the de facto standard for medical image segmentation; its 5-fold cross-validation ensemble is a routinely-used uncertainty / robustness baseline. The full nnU-Net pipeline is *self-configuring* (2D/3D selection, preprocessing, patch-based sampling, default geometric/intensity augmentation, cascaded network, post-processing). We keep three evidence surfaces separate: (i) the released frozen-feature fold-reslicing diagnostic, which tests split sensitivity with one decoder seed per fold and therefore cannot estimate within-fold versus cross-fold same-backbone pairs; (ii) released task-level nnU-Net summary JSONs, which are full-pipeline complements but not case-identical contrasts to the primary FM rows; and (iii) a single case-identical RIGA Cup full-pipeline scope check on the primary Magrabia held-out cases.

Released fold-reslicing design. This diagnostic re-slices the frozen-feature cache into 5 folds and reruns the 8-bundle pool {DINOv2-B/S, BiomedCLIP, CLIP-L, ConvNeXt-T, EfficientNet-B0, ResNet-50/18} on ACDC LV and RIGA Cup, single seed per fold cell. Because the released artifact contains one seed per fold, this diagnostic reports per-fold cross-bundle $\bar{r}$ and pool-mean Dice, not a 20-decoder fixed-backbone within/cross estimator. The released artifact provides the corresponding fold-reslicing summary and scorer.

Table 17: **Fold-reslicing cross-bundle stability, one seed/fold.** 8 bundles $\times$ 5 folds; single seed per fold cell. Cross-bundle $\bar{r}$ is the Pearson correlation averaged over the 28 bundle pairs on the held-out fold. This table cannot support within-fold versus cross-fold same-backbone claims.

Task	fold 0	fold 1	fold 2	fold 3	fold 4	mean $\pm$ std
ACDC LV	0.580	0.548	0.583	0.558	0.558	0.565 ± 0.015
RIGA Cup	0.359	0.317	0.325	0.315	0.261	0.315 ± 0.035

Finding. Train/test split sensitivity is small in this fold-reslicing check: cross-fold pool-mean Dice is stable (ACDC 0.502 ± 0.022 ; Cup 0.610 ± 0.012), and per-fold cross-bundle $\bar{r}$ is stable across folds (ACDC 0.565 ± 0.015 ; Cup 0.315 ± 0.035). This supports only the narrower statement that the cross-bundle diagnostic is not dominated by one particular 5-fold re-slicing. It does *not* support a claim about within-fold vs cross-fold same-backbone decorrelation, because those pair types require multiple decoder seeds per fold that are not present in the released traces.

Full nnU-Net v2 complement. To complement the frozen-feature proxy above and answer whether nnU-Net’s full self-configuring pipeline decorrelates the pool more, we trained the full nnU-Net v2 2D U-Net on RIGA Cup with a 520/224 split, using the generated v2 2D nnUNetPlans configuration after nnUNetv2_plan_and_preprocess (self-configured preprocessing and architecture, official default augmentation/loss schedule, 250 iterations per epoch). This split differs from the Magrabia-source held-out FM row and is therefore a task-level complement. We train two seeds per fold via custom nnUNetTrainerSeed13_100epochs and nnUNetTrainerSeed37_100epochs subclasses that pre-seed random, np.random, torch, and cuda before any nnU-Net sampler touches the RNG; each is trained for 100 nnU-Net-epochs (25,000 iterations on two A100s in parallel), then all $5 \times 2 = 10$ models predict on the held-out 224-image test set via nnUNetv2_predict and per-case Dice is computed against the held-out labels staged in labelsTr. A split-index file enumerated held-out IDs; the artifact bundle releases the summary arrays/JSONs for this complement but not the split manifest.

Across all 10 models the mean Dice is 0.949 ± 0.001 (exceeding any of our frozen-feature FM+1 $\times$ 1-conv combinations on RIGA Cup). The pairwise per-case Dice $\bar{r}$ is 0.972 within-fold over the 5 same-fold different-seed pairs and 0.925 cross-fold over the 40 different-fold any-seed pairs, with $\Delta_{\text{within-cross}} = 0.047$ and a case-level bootstrap 95% CI of [0.030, 0.069] that excludes zero. Full nnU-Net v2 is structurally a same-architecture, different-training-data ensemble on one 2D Plain-ConvUNet, so it is not comparable in kind to the primary frozen-FM cross-family reference, which swaps encoder bundles. The scoped conclusion is therefore: *within this task-level same-architecture nnU-Net complement, fold identity leaves a small but positive correlation gap*; we do not claim nnU-Net substitutes for or surpasses backbone-family diversification. The 100-epoch budget, single-task coverage, and 2D-only scope are closed in Appendix A18 (1000-epoch retrain on the same RIGA Cup setup; 3D fullres nnU-Net on ACDC; higher-resolution/Dice+BCE/aug frozen-FM replication).

Case-identical RIGA Cup held-out scope check. We additionally train the same 5×2 seeded nnU-Net v2 2D grid on BinRushed+MESSIDOR folds and evaluate all models on the primary Magrabia-source held-out Cup cases ($n=95$). The exporter uses the official nnU-Net v2 natural-image channel convention, with RGB saved as _0000/_0001/_0002 channel files, imagesTs/labelsTs for held-out scoring, and the generated nnUNetPlans 2D configuration. Released case IDs are deterministic hashes rather than upstream path strings. Across the 10 models, mean Dice is 0.9162 ± 0.0019 . The same-fold, different-seed correlation is $\bar{r}=0.984$ over 5 pairs, the different-fold correlation is $\bar{r}=0.976$ over 40 pairs, and the gap is $\Delta=0.0079$ with case-bootstrap 95% CI [0.002, 0.039]. This confirms the boundary in a stricter paired setting: nnU-Net’s same-architecture fold/seed surface is highly correlated on the same held-out cases, and the fold gap is much smaller than both the primary RIGA Cup same/cross gap (0.473) and the separate-split task-level nnU-Net summary (0.047). It does not make nnU-Net a cross-family reference or a causal control for architecture versus pretraining.

Data-swap comparison caveat. Appendix A14 reports bootstrap-bagging within-pool decorrelation with four decoder seeds per bootstrap replicate. The fold-reslicing row instead reports single-seed cross-bundle correlations across held-out folds. These diagnostics both stress training data/split sensitivity, but they are not dose-matched estimators and are not combined into a causal ranking.

A18 Segmentation-recipe and nnU-Net scope audit

Our main protocol uses 224×224 2D inputs, frozen backbone, 1×1 -conv decoder, BCE loss, no augmentation, and 30-epoch head training. Routine medical-imaging pipelines often run with compound Dice+CE loss, geometric augmentation, 3D input, and full-length (1000-epoch) training, so we test whether the within/cross correlation structure we report is specific to the simplified protocol. Most nnU-Net rows below are task-level complements; the added case-identical RIGA held-out row uses the same Magrabia-source held-out cases as the primary frozen-FM row but remains a single-task, 2D, same-architecture nnU-Net scope check.

Dimensionality axis: 3D fullres nnU-Net at 1000 epochs on ACDC. Training nnU-Net v2 in its self-configuring 3d_fullres setting on ACDC with a seed-42 105/45-patient split (210 training, 90 test volumes) using nnUNetTrainerSeed{13,37}_1000epochs (seed injected in on_train_start before sampler initialisation) gives 10 fully-trained 3D models across 5 folds. nnUNetv2_predict on all 90 held-out volumes yields per-case Dice on mean foreground and per class. Across 10 models the mean Dice is 0.920 ± 0.001 , within the range reported for strong nnU-Net baselines on ACDC. On per-case mean Dice the within-fold (5 pairs) vs cross-fold (40 pairs) gap is $\Delta=0.043$ [0.021, 0.085] with within $\bar{r}=0.978$ and cross $\bar{r}=0.935$; per class, $\Delta_{RV}=0.041$ [0.011, 0.110], $\Delta_{Myo}=0.029$ [0.017, 0.046], $\Delta_{LV}=0.067$ [0.033, 0.098]. The artifact bundle releases the aggregate summary for this 150-patient complement but not the exact split manifest; the released public ACDC converter targets patients 001–100 for the primary 2D row. We therefore treat this 3D row as a task-level scope summary rather than a from-bundle training recipe. In this ACDC nnU-Net complement, the 2D→3D fullres shift changes the small fold-gap regime by less than 0.01 relative to the 2D 100-epoch baseline ($\Delta=0.047$ in Appendix A17).

Loss / augmentation axis: RIGA Cup frozen-FM at encoder-native 224×224 . We retrain the seven-bundle RIGA Cup subset (DINOv2-B/S, CLIP ViT-L/14, ConvNeXt-T, EfficientNet-B0, ResNet-50/18) with four seeds each under a more standard segmentation loss/augmentation recipe: encoder-native input size 224×224 , composite soft Dice+BCE loss, paired random horizontal flip and rotation $\pm 15^\circ$, for 30 epochs. The frozen encoder wrappers use 224×224 inputs; we therefore treat this as a loss/augmentation test, not a resolution test. All other factors (frozen backbone, 1×1 -conv segmentation head, Adam 10^{-3}) match the main protocol. Per-FM mean Dice ranges from 0.526 (ResNet-18) to 0.758 (DINOv2-B). Under the same symmetric estimator used for the main table, within-family $\bar{r}$ is 0.903 (58 within pairs, including same-FM seed pairs and DINOv2-B×S), cross-FM $\bar{r}$ is 0.556 (320 cross pairs), and the within/cross gap is $\Delta=0.347$ [0.275, 0.428]. The gap is smaller than the main RIGA Cup interval (0.473 [0.395, 0.556]) but remains positive; loss/augmentation changes do not remove the gap.

Training-length axis: full nnU-Net v2 at the 1000-epoch default on RIGA Cup. We retrain the 10-model RIGA Cup 2D nnU-Net pool from Appendix A17 at the nnU-Net v2 default 1000 epochs, using the same seeded trainer pattern (`nnUNetTrainerSeed{13,37}_1000epochs`) on the same 520/224 split with identical dataset/plan files. Mean Dice across the 10 models is 0.946 ± 0.0004 (vs 0.949 ± 0.001 at 100 epochs). Within-fold $\bar{r}=0.983$ (5 pairs), cross-fold $\bar{r}=0.949$ (40 pairs), and $\Delta=0.034$ [0.021, 0.058]. Relative to the 100-epoch run both estimators tighten slightly (within $0.972 \rightarrow 0.983$; cross $0.925 \rightarrow 0.949$) and the within/cross gap contracts from 0.047 to 0.034; measured predictions are more correlated under longer training, and the gap does not disappear. Ten-fold longer training therefore does not change the replicability structure of nnU-Net on this task-level complement; it is not a case-identical comparison to the Magrabia-source held-out FM row.

Synthesis. Across these scope analyses, the frozen-FM loss/augmentation row remains in the same large-gap regime as the RIGA Cup main row, while the same-architecture nnU-Net fold axis remains in a small-gap regime. The case-identical RIGA held-out row tightens this boundary: on the same Magrabia cases, both nnU-Net correlations are very high and the fold gap is only 0.0079 at 100 epochs. The matched 1000-epoch case-identical re-run gives mean Dice 0.9151 ± 0.0018 across the same ten models with within-fold $\bar{r}=0.987$, cross-fold $\bar{r}=0.980$, and $\Delta=0.0075$ [0.003, 0.034], so a tenfold-longer training budget on the Magrabia-source held-out split leaves the case-identical fold gap essentially unchanged ($0.0079 \rightarrow 0.0075$) while both correlations rise modestly. These rows do not support a simpler explanation that the main-table pattern is solely an artifact of BCE loss, no augmentation, short training, or the tested full-pipeline nnU-Net complements.

Table 18: **Scope audit across recipe and nnU-Net complements.** Rows are not all case-identical and should not be compared as a controlled ablation. The table asks whether selected recipe changes remove the qualitative gap. The frozen-reference row uses the same seven candidate FMs as the Dice+BCE+augmentation row before the primary functional-floor exclusion; the primary filtered RIGA Cup row is Table 1. The case-identical RIGA held-out row uses the Magrabia-source test cases; the separate-split RIGA nnU-Net rows use a 520/224 split.

Axis	Protocol	Task	Within	Cross	Δ	95% CI
Frozen reference	7-bundle candidate set, 224, BCE, no aug, 30 epochs	RIGA Cup	0.798	0.266	0.532	[0.456, 0.614]
Loss / aug	7-bundle frozen, 224, Dice+BCE, aug, 30 epochs	RIGA Cup	0.903	0.556	0.347	[0.275, 0.428]
Baseline (App. A17)	nnU-Net 2D, 100 epochs, separate split	RIGA Cup	0.972	0.925	0.047	[0.030, 0.069]
Case-identical held-out	nnU-Net 2D, 100 epochs, Magrabia test	RIGA Cup	0.984	0.976	0.0079	[0.002, 0.039]
Case-identical held-out	nnU-Net 2D, 1000 epochs, Magrabia test	RIGA Cup	0.987	0.980	0.0075	[0.003, 0.034]
Training length	nnU-Net 2D, 1000 epochs, separate split	RIGA Cup	0.983	0.949	0.034	[0.021, 0.058]
Dimensionality	nnU-Net 3D fullres, 1000ep	ACDC (mean)	0.978	0.935	0.043	[0.021, 0.085]
	per-class RV		0.967	0.927	0.041	[0.011, 0.110]
	per-class Myo		0.987	0.957	0.029	[0.017, 0.046]
	per-class LV		0.978	0.911	0.067	[0.033, 0.098]

A19 Held-out multi-task stress test

Summary. Across the five held-out 2D binary derivatives we tested, two rows give substantively positive gaps (Hippocampus $\Delta=0.491$, ISIC-2018 $\Delta=0.406$), one is tiny- N positive (Prostate, $n=6$, $\Delta=0.146$), one gives a near-zero gap (Spleen $\Delta=0.003$), and one fails the functional floor (Heart, 0/11 functional FMs). The held-out evidence *supports* rather than confirms the primary 2D rows and is consistent with task-specific structure modulating the gap magnitude.

Protocol. In addition to the four primary benchmark datasets (five task rows: Cup, Disc, ACDC, Kvasir, BraTS foreground), we evaluate five *held-out* tasks covering additional imaging domains: MSD Hippocampus [Antonelli et al., 2022] (T1 MRI, native 3-class mask {background, anterior, posterior} binarized to {background, any-hippocampus}), ISIC-2018 Task 1 Lesion Boundary [Codella et al., 2019] (dermoscopy, 2D RGB, $N=80$ test cases held out from a 400-image random subset of the 2,594-image training release), MSD Heart (T2 MRI cardiac), MSD Prostate (T2 MRI prostate), and MSD Spleen (CT portal). This held-out stress test uses the same 11 generic encoder bundles and four decoder seeds as the primary source pool where possible, with the same 1×1 -conv

928 decoder, no per-task hyperparameter tuning, and no model substitution. These are task-specific 2D
 929 binary derivatives, not official challenge-test evaluations. The $\text{Dice} \geq 0.30$ functional floor is applied
 930 uniformly, and the released aggregate summary records 220/220 expected per-seed traces.

Table 19: **Held-out aggregate stress test** (11 generic encoder bundles $\times$ 4 decoder seeds; point estimates only, no case-bootstrap CIs). Hippocampus and ISIC-2018 are direction-positive, Prostate is positive but has only six test cases, Spleen is essentially near-zero despite passing the Dice floor, and Heart has no functional bundle at the $\text{Dice} \geq 0.30$ floor. Tiny- N and floor-failure rows are scope boundaries and do not support the primary claim.

Task	N	Func./avail.	Within	Cross	Δ	M_{eff}
<i>Direction-positive point estimates with enough support to read qualitatively</i>						
MSD Hippocampus	52	11/11	0.909	0.419	0.491 [§]	3.94
ISIC-2018 [‡]	80	11/11	0.918	0.511	0.406 [§]	2.72
<i>Scope-limit rows: tiny held-out support or non-functional pool</i>						
MSD Prostate	6	5/11	0.788	0.642	0.146 [§]	1.23
MSD Spleen	8	11/11	0.983	0.979	0.003 [§]	1.02
MSD Heart	4	0/11	—*	—*	—*	—

*No bundle passes the $\text{Dice} \geq 0.30$ floor on the four-case Heart derivative, so no within/cross gap is interpretable. [§]Values are point estimates from `results/_merged/diagnostics/heldout_11fm_summary.json`; case-level CIs are not computed for this aggregate diagnostic, and the MSD Prostate/Spleen rows have too few held-out cases to support a primary claim. [‡]ISIC-2018 uses a single random $N=400$ -image subset of the 2,594-image ISIC-2018 training release with $N=80$ test cases under the seed-42 80/20 split; we do not report an independent redraw and we do not use the official ISIC test server in this study.

931 **Interpretation.** The 11-bundle held-out protocol yields three classes of outcome: (i) ISIC-2018
 932 and Hippocampus both show direction-positive gaps under the same broad estimator family as the
 933 primary audit, but without case-bootstrap CIs and with smaller held-out support than the primary
 934 rows. (ii) Prostate and Spleen illustrate why held-out diagnostics are scope boundaries rather than
 935 primary evidence: Prostate is direction-positive but has only six test cases, while Spleen passes the
 936 floor yet has a near-zero gap and $M_{\text{eff}} \approx 1$ because the members fail together. (iii) Heart does
 937 not yield a functional pool under the frozen-feature 2D derivative. Thus the held-out stress test is
 938 compatible with transfer to some external rows, but it also records where the protocol is too weak or
 939 too saturated to audit same-bundle dependence.

940 **Operational 8-bundle held-out probes.** The following recovery, pool-size, absten-
 941 tion, and pool-selection diagnostics use an 8-bundle held-out candidate set stored under
 942 `results/_merged/diagnostics/held_out/`. They are retained as operational-method sensitiv-
 943 ities, not as the provenance for Table 19 above.

944 **CLIP/BiomedCLIP normalization robustness.** Uniform ImageNet normalization can be a po-
 945 tential artifact for BiomedCLIP and CLIP-L because both have their own normalization statistics.
 946 We therefore evaluated those two FMs on Hippocampus and ISIC-2018 with the correct CLIP nor-
 947 malization. The released `summary_clipfix.json` reports BiomedCLIP mean Dice 0.840 (Hip-
 948 pocampus) / 0.908 (ISIC-2018), while CLIP-L remains below the functional floor on Hippocam-
 949 pus (mean Dice $< 10^{-10}$) and reaches 0.106 on ISIC-2018, still below the 0.30 floor. Relative
 950 to `summary_5tasks.json`, the aggregate Δ changed by 0.002 (Hippocampus 0.731 $\rightarrow$ 0.729) and
 951 0.003 (ISIC-2018 0.552 $\rightarrow$ 0.555). In this held-out diagnostic, correcting CLIP/BiomedCLIP nor-
 952 malization does not materially change the Hippocampus/ISIC point estimates.

953 **Held-out recovery-rate sensitivity.** To complement the primary-task $\Delta_{\text{rec}} = +0.29$ (Cup) / $+0.26$
 954 (ACDC) grand means, we computed an exploratory held-out recovery gap at $\tau=0.85$ on the sin-
 955 gle ISIC-2018 subset: $\Delta_{\text{rec}} = +0.269$ [$+0.212, +0.331$] under a case-resampling calculation saved
 956 as `recovery_CI.json`. This supports the direction of the operational gap on one dermoscopy
 957 subset, but it is not a deployment-calibrated guarantee and does not resolve the held-out estimator-
 958 dependence above. Hippocampus at $\tau=0.85$ is recovery-saturated and gives $\Delta_{\text{rec}} = -0.01$ [$-0.03,$
 959 $+0.00$] (diverse and mono both near-zero recovery because only 2/7 functional bundles individually
 960 pass $\text{Dice} \geq 0.85$). Lower thresholds are not evaluated here.

961 **Pool-size sweep on held-out tasks.** How does diverse-vs-same-bundle conditional recovery scale
 962 with M ? For $M \in \{2, 3, 4, 5, 6\}$ on the two functional held-out tasks (Table 20), diverse pools
 963 beat same-bundle seed pools at every $M \geq 3$ on the single ISIC-2018 subset ($\Delta_{\text{rec}} = 0.27\text{--}0.46$),
 964 while Hippocampus at the paper’s $\tau=0.85$ threshold is recovery-rate-saturated (all pools < 0.11)
 965 because only 2 of 7 functional bundles reach $\text{Dice} \geq 0.85$ individually. These held-out opera-

966 tional analyses are exploratory: ISIC provides one supportive dermoscopy subset, Hippocampus
 967 is threshold-saturated at $\tau=0.85$, and the $M=5, 6$ same-bundle rows are proxy constructions us-
 968 ing the available four seeds rather than true larger same-bundle pools. Released aggregate data:
 969 results/_merged/diagnostics/held_out/pool_size_sweep.json.

Table 20: **Exploratory held-out pool-size sweep.** Conditional recovery rate $P[\text{ens pass} \mid \text{any fail}]$ at $\tau=0.85$ vs pool size M . $M=2$ recovery is a definition artifact under the strict-majority rule, and $M=5, 6$ same-bundle values are four-seed proxies. This table is not a deployable pool-size recommendation.

M	Hippocampus			ISIC-2018		
	mono	diverse	Δ	mono	diverse	Δ
$2^\dagger$	0.000	0.000	+0.00	0.000	0.000	+0.00
3	0.069	0.107	+0.04	0.111	0.498	+0.39
4	0.052	0.039	-0.01	0.091	0.359	+0.27
$5^\ddagger$	0.052	0.078	+0.03	0.091	0.553	+0.46
$6^\ddagger$	0.052	0.027	-0.03	0.091	0.454	+0.36

$^\dagger M=2$ is a definition artifact under the strict-majority ($>M/2$) rule conditioned on any-fail: recovery is zero by construction whenever $M=2$, regardless of pool composition. We include the row for completeness. ‡ At $M=5, 6$ the same-bundle entries are computed on the $\min(M, S)=4$ available seeds per FM, so they do not represent true $M=5, 6$ same-bundle pools; the diverse entries do use legitimate $M=5, 6$ distinct-family compositions.

970 **Calibration-tuned disagreement-based abstention (an exploratory alternative lever, not a con-**
 971 **formal guarantee).** A natural operational question is whether standard abstention-based uncertainty
 972 gating closes the same-bundle gap with modest additional computation. We calibrate a per-task
 973 disagreement-score threshold on 20 cases (inter-member Dice standard deviation as score; grid-
 974 search the quantile q to meet a target case-risk $\alpha=0.1$ on calibration; apply the selected thresh-
 975 old to the remaining test cases). *This is a tuned heuristic, not split-conformal prediction* [An-
 976 gelopoulos and Bates, 2021]: we use a grid search over q that terminates on the first calibra-
 977 tion risk meeting α and use inter-member variance rather than a proper conformity score, so no
 978 finite-sample risk guarantee holds. The pool is also chosen per-task on the whole task mean Dice
 979 (family-representative highest-Dice FMs), so the diverse-pool selection step leaks task labels—
 980 this biases the comparison toward diverse pools. We report it as a sensitivity. On ISIC-2018
 981 at matched abstention ≈ 0.74 , diverse pools reach test-risk 0.129 ± 0.027 vs. same-bundle grand
 982 mean 0.174 ± 0.209 (conformal_vs_meff.json). The *point estimate* is $\sim 25\%$ lower for di-
 983 verse, but the same-bundle grand-mean SD (0.21) is $\sim 8\times$ the diverse SD (0.027); the per-FM
 984 same-bundle risk means span 0.000–0.616 across the seven functional FMs, so the grand-mean
 985 comparison is highly composition-sensitive and we do not claim a tight ranking. Hippocampus
 986 cannot hit $\alpha=0.1$ for either pool, giving residual risk 0.694 (diverse) vs. 0.665 (same-bundle
 987 grand mean) ($\Delta = +0.03$, with same-bundle SD 0.28, indistinguishable). A label-free conformal
 988 procedure with a proper conformity score is not evaluated here. Released aggregate data:
 989 results/_merged/diagnostics/held_out/conformal_vs_meff.json.

990 **Calibration-set pool-selection probe on held-out tasks.** We operationalise a small calibration ex-
 991 ercise: given only a 20-case calibration set on a held-out task, can a rule pick a 4-of-8 FM pool from
 992 the same secondary held-out candidate set that beats random? We evaluate five rules—(R1) random,
 993 (R2) top-4 by mean Dice on calibration, (R3) max- M_{eff} , (R4) quality-gated max- M_{eff} (Dice ≥ 0.30
 994 floor then max- M_{eff} among qualifying), (R5) diverse-by-architecture (2 best-Dice ViT + 2 best-Dice
 995 CNN; reported as an exploratory hand-crafted baseline, not a general proposal)—and report condi-
 996 tional recovery rate $P[\text{ens pass} \mid \text{any fail}]$ at $\tau=0.85$ on the remaining held-out cases, averaged over
 997 100 calibration-set bootstraps (Table 21). All rule selectors and the recovery estimator use seed-42
 998 per-case Dice (not seed-averaged), matching a one-model-per-FM pool construction. Hippocampus
 999 has lower effective recovery support than ISIC at $\tau=0.85$, so its recovery numbers are exploratory.
 1000 In this diagnostic, simple top-4-by-Dice has the highest mean on both tasks (Hippocampus 0.161
 1001 vs. 0.043 random; ISIC-2018 0.581 vs. 0.378 random), while max- M_{eff} alone is worse than random
 1002 (0.011 / 0.327). This supports the narrower point that M_{eff} is a dependence readout rather than a
 1003 standalone operational objective; it is not a validated clinical pool-selection rule. Released aggre-
 1004 gate data: results/_merged/diagnostics/held_out/pool_selection_challenge.json.

Table 21: **Exploratory pool-selection probe.** Conditional recovery rate at $\tau=0.85$ (mean $\pm$ SD over 100 calibration bootstraps, 20 cases each, using seed-42 traces). Selection uses calibration labels. This probes whether M_{eff} alone is a poor objective; it is not a validated clinical selection rule.

Rule	Hippocampus	ISIC-2018
(R1) Random 4-of-8	0.043 ± 0.050	0.378 ± 0.125
(R2) Top-4 by calibration Dice	0.161 ± 0.017	0.581 ± 0.030
(R3) Max M_{eff}	0.011 ± 0.011	0.327 ± 0.100
(R4) Quality-gated max M_{eff}	0.016 ± 0.014	0.381 ± 0.112
(R5) Diverse-by-arch (2 ViT + 2 CNN)	0.096 ± 0.017	0.445 ± 0.036

Held-out estimator note. Split-stability and alternate-estimator numerical summaries without released aggregate files are omitted from the numeric evidence. This section therefore relies on the released 11-bundle held-out aggregate and the released operational JSONs listed above.

ISIC-2018 single-subset note. The ISIC-2018 row above is a single random $N=400$ -image subset of the ISIC-2018 Task 1 training release with $N=80$ held-out test cases under a seed-42 80/20 split. We do not report an independent replication of this subset; a fully random redraw of the subset or use of the official ISIC evaluation server would be a natural next step but is out of scope for this paper.

A20 Full fine-tuning scope

RIGA Cup full fine-tuning epoch-sweep traces and aggregate summaries are outside the released artifact. We therefore do not tabulate those values or use them as numerical evidence. The released adaptation evidence is limited to released traces and summaries named elsewhere in the appendix, including a small set of ACDC DINOv2-B full-fine-tuning provenance traces but not the RIGA epoch-sweep matrix; full-fine-tuning numbers not backed by released paths should be read only as exploratory mechanism notes and not as support for the paper’s claims.

A21 Same-backbone and scope diagnostic index

Table 22: **Qualitative index of same-backbone and scope diagnostics.** Rows are heterogeneous and not CI-comparable; the table is a navigation aid, not an ordered causal hierarchy. The RIGA-oriented map is anchored on the DINOv2-B 4-seed baseline ($\bar{r} \approx 0.96$) unless marked. LoRA, full-fine-tuning epoch-sweep, and ViT random-init probes without released trace artifacts are omitted.

Diagnostic	Readout	Approx. change from stated baseline	Evidence
Mono baseline (same FM, 4 seeds)	0.96	0.00	RIGA diagnostic baseline
UNet-skip decoder (9 backbones, 5 tasks expansion)	0.426–0.979	heterogeneous	Appendix A16, Tab. 16
Fold-reslicing cross-bundle check	0.315–0.565	split-stable	Appendix A17; released summary
MC dropout $p=0.1$ –0.4	0.89–0.95	−0.02 to −0.10	Appendix A31.7
Same-backbone bootstrap-bagging	0.85–0.92	−0.05 to −0.12	Fig. 2
RIGA cross-vendor shift	OOD w/c 0.743/0.390	gap 0.353;	Appendix A15
		$\Delta_{\text{gap}} - 0.024$	
Decoder capacity sweep	~ 0.69	~ -0.27	Appendix A31.1
Cross-family pretrained 8-bundle pool	0.38 [0.29, 0.46]	−0.58	RIGA mechanism diagnostic

A22 Leave-one-family-out scope

Leave-one-family-out summaries were used as exploratory robustness checks, but the exact aggregate source is outside the released artifact. We therefore omit the numeric table from the released evidence record. The primary robustness check for the primary rows is the released `diagnostics/audit_sensitivity.json`, which includes leave-one-task and leave-one-FM summaries under the same scorer path.

A23 Joint-failure lift at support limits

Joint-failure lift $L = P_{\text{empirical joint}} / \prod_m P_{\text{marginal}}$ is unstable when marginal fail rates are near 0 or near 1. When $P_{\text{marginal}} \approx 0.3$, independence-implied joint ≈ 0.01 ; small sampling variation in the empirical joint produces large lift ratios. Representative near-ceiling cells: Only the ACDC 2.8×

Table 23: **Support-limit examples for joint-failure lift.** Extreme lifts are shown to explain instability, not to support a primary claim. Lifts $\geq 50\times$ are descriptive (support-limited, wide noise envelope) and are *not* primary evidence.

Task / Pool	τ	Indep. joint	Empirical joint	Lift
RIGA Cup mono BiomedCLIP	0.85	0.0010	0.125	123.2×
RIGA Cup mono DINOv2-B	0.85	0.0047	0.183	38.8×
RIGA Cup diverse 4-bundle	0.85	0.0160	0.098	6.1×
ACDC LV mono DINOv2-B	0.70	0.0002	0.038	189×
ACDC LV diverse 4-bundle	0.85	0.135	0.378	2.8×

row is adequately supported by the $N \cdot \mathbb{E}[\text{AF}_{\text{ind}}] \geq 5$ rule. The RIGA Cup 6.1×

 and 38.8× rows are support-limited, and the 123.2× and 189× rows are extreme sparse-tail diagnostics.

A24 3D segmentation-architecture pilot

This section provides scope-boundary context. Its per-case outputs and scorer version are not included in the supplementary artifact, so the pilot is not used as primary evidence for Table 1, the abstract, or the released artifact claims. It motivates why we avoid generalizing the frozen 2D FM-pool result to 3D segmentation-architecture pools.

We assembled a 3D segmentation-architecture pool for BraTS 2024 3D volumes (271 subjects, 216/55 train/test, seed 42) using three retained architectures: SegResNet (optionally initialised from the MONAI BraTS bundle), DynUNet (random-init), and UNETR (random-init). SwinUNETR random-init cells were excluded because predictions were degenerate across all seeds. For each retained (architecture, seed $\in \{42, 43, 44, 45\}$) cell, the pilot trained the full MONAI segmentation model for 100 epochs with Adam 10^{-3} and BCE loss; UNETR-style models used sliding-window inference where required by memory. Reported at three filter modes:

Table 24: **3D segmentation-architecture pilot.** Per-case outputs and scorer version are not included in the supplementary artifact; this table is excluded from the abstract and main-result evidence. Three functional-filter modes are applied uniformly in the pilot. Strict filter (Dice ≥ 0.50) retains only SegResNet pretrained and DynUNet random-init and yields $\Delta=0.045$, substantially below the 2D BraTS foreground $\Delta=0.315$.

Filter	FMs retained	Within	Cross	Δ [95% CI]
Mixed (all cells)	3 (12 cells)	0.772	0.595	0.177 [0.117, 0.248]
Dice ≥ 0.30	3 (12 cells, drop UNETR seed 43 degenerate)	0.919	0.703	0.216 [0.134, 0.271]
Dice ≥ 0.50 (strict functional)	2 (SegResNet pretrained, DynUNet random)	0.916	0.871	0.045 [0.015, 0.112]
Random-init only (DynUNet + UNETR)	2 random-init	0.724	0.479	0.245 [0.155, 0.351]

Interpretation. Within the strict functional regime (only converged pretrained + converged random-init CNN-style 3D architectures), the 3D segmentation-architecture pilot shows $\Delta=0.045$ —substantially below the 2D BraTS foreground $\Delta=0.315$. We treat this only as a scope warning: the 2D frozen-feature same-bundle dependence claim should not be extended to 3D segmentation-architecture pools without a released, case-identical 3D benchmark. Random-init 3D architectures cross-correlate *more* than pretrained 3D architectures (within random 0.724 vs cross-architecture random 0.479, $\Delta=0.245$), so architecture-family effects again appear large in this pilot. The degenerate SwinUNETR pilot cells limit architecture coverage in the 3D pilot; expanded 3D foundation-model evaluation (Models Genesis, Vista3D, Triad-MRI) is not evaluated here.

A25 Matched-recipe adaptation scope

Matched-recipe adaptation factorials require source prediction caches and aggregate summary JSONs that are not included in the supplementary artifact. We therefore do not tabulate those values

and do not use them to support the main claims. The released adaptation evidence is limited to explicit released traces under `per_case_dice_adapted` and the aggregate summaries named in the surrounding appendix sections.

A26 Item-difficulty residualisation and held-out nulls

The main-text Section 5 reports a cross-fitted residualised Δ diagnostic on all five primary rows. Three distinct residualisation procedures appear in this appendix and we tabulate them in order of provenance: (i) the *naive* all-bundle null, included only as a leakage warning; (ii) the *random A-vs-B family-split holdout* on Cup and ACDC (the existing `r7_item_difficulty_holdout.json` diagnostic, Table 26 below); and (iii) the *cross-fitted broad-family-leaveout* on all five primary rows (Table 25, the headline 5-row diagnostic referenced in the main text). The cross-fitted procedure keeps the same/cross gap positive on every primary row (95% CIs in Table 25: Kvasir [0.895, 1.008], ACDC LV [0.950, 1.000], BraTS fg [0.910, 0.932], RIGA Cup [0.782, 0.882], RIGA Disc [0.805, 0.920]; the Kvasir upper bound exceeds 1 because the residualised statistic is no longer constrained to canonical Pearson scale once item-difficulty signal has been subtracted from both members of the pair). The held-out-partition table below is retained as the historical Cup/ACDC-only diagnostic from `r7_item_difficulty_holdout.json`.

Cross-fitted broad-family-leaveout (headline, all five rows). For each trace, item difficulty is estimated from traces outside that trace’s broad upstream family, the trace is residualised by OLS against this leave-out estimator, and the Table 1 same/cross rule is recomputed on residuals. The residualised gap is positive on all five primary rows. Residualised gaps *exceed* raw gaps because the cross-fitted regressor removes shared item-difficulty signal that both same- and cross-bundle pairs already share at the per-case level; same-bundle residual continues to capture bundle-specific systematic deviations not absorbed by case difficulty, while cross-bundle residual decorrelates faster. The attenuation ratio (residual gap / raw gap) is 1.37–3.68 across the five tasks.

Table 25: **Cross-fitted item-difficulty residualisation, all five primary rows.** Each trace’s item-difficulty regressor is estimated using only traces from broad upstream families *other than* the trace’s own family; the same-vs-cross Pearson gap is then recomputed on residuals. Residualised gaps exceed raw gaps because the cross-fitted regressor removes shared item-difficulty signal that both same- and cross-bundle pairs already share at the per-case level; same-bundle residual continues to capture bundle-specific systematic deviations not absorbed by case difficulty, while cross-bundle residual decorrelates faster. Source: `results/_merged/diagnostics/item_difficulty_robustness.json` (schema `fmpool_item_difficulty_robustness_v1`); $B=2,000$ case bootstrap.

Task	Raw Δ 95% CI	Residualised Δ 95% CI	Attenuation ratio	Functional pool
Kvasir	[0.211, 0.320]	[0.895, 1.008]	3.68	11-bundle
ACDC LV	[0.289, 0.355]	[0.950, 1.000]	3.05	10-bundle
BraTS foreground	[0.302, 0.329]	[0.910, 0.932]	2.92	10-bundle
RIGA Cup	[0.395, 0.556]	[0.782, 0.882]	1.75	9-bundle
RIGA Disc	[0.569, 0.687]	[0.805, 0.920]	1.37	11-bundle

Naive (leaked) null. Define difficulty $d_i = 1 - \frac{1}{K} \sum_{k=1}^K D_k^{(42)}(i)$ where $D_k^{(42)}$ is bundle k ’s per-case Dice at seed 42, then residualise each $D_k^{(s)}$ against d by OLS and then evaluate $\bar{r}$. This instantiates one item-difficulty null inspired by Jo et al. [2026], but has a structural flaw: each bundle’s Dice is a component of its own regressor (since d averages over all 8 bundles including k), so residuals are mechanically constrained to sum to approximately zero across bundles at each case. For K bundles at seed 42, cross-bundle residual correlations are pushed toward $-1/(K-1) \approx -0.14$, and we observe cross-family residuals of -0.088 (Cup) and -0.104 (ACDC), which sharply widens Δ compared to the raw statistic. This is a mechanical leakage artifact, not a cleaner measurement.

Random A-vs-B family-split holdout (two-task null). To remove the leakage we split the 8 bundles into halves A and B , estimate d from A ’s seed-42 Dice only, and residualise B ’s Dice on this d . The $\bar{r}$ statistics and Δ are then computed on B . The released diagnostic JSON records a fixed set of 10 held-out partitions and reports the mean and range of the resulting family-grouped Δ values across those partitions. This procedure is procedurally distinct from the cross-fitted broad-family-

1094 leaveout above: it does a random within-task split of the panel rather than a leave-one-broad-family-
 1095 out cross-fit, and it covers only Cup and ACDC.

Table 26: **Two-task item-difficulty residualisation diagnostics (random A-vs-B family-split holdout).** The naive null is included only as a leakage warning. Raw values use this diagnostic’s family grouping and are not Table 1 values; the held-out-partition null avoids leakage on Cup/ACDC only. This is procedurally distinct from the cross-fitted broad-family-leaveout in Table 25. Source: `r7_item_difficulty_holdout.json`.

Null	RIGA Cup Δ	ACDC Δ	Notes
Raw	0.499	0.600	Family grouping
Naive difficulty (all-bundle leaked)	0.886	0.772	Cross-family residuals ≈ -0.10 (leakage signature)
Held-out difficulty (released held-out partitions)	0.639	0.622	Family grouping: $n=10$ per task
Held-out range	[0.482, 0.823]	[0.472, 0.941]	Min / max across partitions

1096 The two-task held-out estimate is above raw Δ on Cup (+0.14) and remains slightly above raw on
 1097 ACDC (+0.02). Under this *one-directional* null, item difficulty does not explain the gap away. The
 1098 naive inflation was a mechanical leakage artifact regardless.




Figure 3: **Three item-difficulty residualisation procedures on aligned per-case Dice.** Each procedure operates on the same released traces but uses a different difficulty regressor. *Raw* (dark grey) is Pearson Δ without residualisation, recomputed per task at the audit’s family grouping. *Cross-fitted broad-family-leaveout* (blue, all five primary rows) estimates each trace’s difficulty regressor from traces outside the trace’s own broad upstream family; the residualised gap CIs are taken from `item_difficulty_robustness.json` (`cross_fitted_item_difficulty_floor_0.30.bootstrap`). *Held-out family-split* (orange, Cup and ACDC only) is the random A-vs-B split-holdout from `r7_item_difficulty_holdout.json`; mean across released partitions. *Naive (leaked, illustrative)* (red hatched) is the all-bundle null whose mechanical leakage is documented in the prose above. The cross-fitted bars are the procedure referenced in the main text Section 5; the held-out family-split bars are a procedurally distinct two-task companion. Error bars are 95% case-bootstrap intervals where reported. Source script: `code/scripts/build_diffleak_procedures_figure.py`.

1099 **Residualisation caveat.** The released held-out partitions are one-directional diagnostics rather
 1100 than a symmetric enumeration over every possible difficulty/evaluation split. They should be read as
 1101 evidence that this particular leakage-free residualisation does not absorb the gap, not as a two-sided
 1102 causal decomposition of difficulty, architecture, and pretraining. The all-row cross-fitted diagnostic
 1103 broadens coverage to every primary row, but it still estimates difficulty from the same benchmark
 1104 panel and therefore remains a robustness check rather than a pre-registered causal adjustment.

1105 A27 Correlation-equivalent effective pool size M_{eff}

1106 **Definition.** For a pool of M members whose per-case failure indicators $F_1, \dots, F_M$ at threshold
 1107 τ are exchangeable with mean pairwise Pearson correlation ρ_{fail} , the variance of the per-case mean
 1108 indicator is $\sigma_F^2[1 + (M - 1)\rho_{\text{fail}}]/M$ where σ_F^2 is the marginal variance. Equating this with the
 1109 variance of a hypothetical independent pool of size M_{eff} , i.e. $\sigma_F^2/M_{\text{eff}}$, gives

$$M_{\text{eff}} = \frac{M}{1 + (M - 1)\rho_{\text{fail}}}.$$

1110 This is the equicorrelation approximation also used in climate-model ensembles (“effective number
 1111 of models”). We report it as a second-moment readout of the audit; it is *not* a literal independent-
 1112 member count. The reported M_{eff} values follow algebraically from the primary correlation values in
 1113 Table 1; we report them as a compact summary of the primary same-vs-cross gap, not as additional
 1114 independent evidence.

Table 27: **Threshold-specific effective pool-size diagnostic at $\tau=0.85$.** M_{eff} uses an equicorrelation approximation for binary failure indicators and is not a literal independent-member count. Same-bundle pools are each FM’s completed decoder seeds. Diverse rows enumerate every 4-bundle composition of the filtered functional pool and every available one-seed-per-FM tuple. ρ_{max} is the largest valid diverse-pool ρ_{fail} over that released enumeration and gives a conservative lower M_{eff} readout for the most redundant valid tuple. Undefined zero-variance failure-correlation tuples are excluded rather than imputed. Source: `results/_merged/m_eff/*.json`.

Task	Func. bundles	Mono mean [†]	Diverse ρ_{fail}	ρ_{max}	Mean M_{eff}	$M_{\text{eff}}(\rho_{\text{max}})$	Valid tuples
Kvasir	11	1.12	0.333±0.104	0.625	2.05±0.32	1.39	84,480 / 84,480
ACDC LV	10	1.14	0.298±0.075	0.528	2.14±0.24	1.55	53,760 / 53,760
BraTS foreground	10	1.13	0.273±0.089	0.506	2.24±0.32	1.59	53,760 / 53,760
RIGA Cup	9	1.16	0.246±0.121	0.695	2.43±0.63	1.30	31,488 / 32,256
RIGA Disc	11	1.17	0.126±0.067	0.455	2.97±0.44	1.69	84,480 / 84,480

[†] Mean across per-FM seed pools after excluding undefined zero-variance failure indicators. The RIGA Cup row has 7/9 defined mono pools and the RIGA Disc row has 10/11; the other rows have all mono pools defined.

1115 **Reading the table.** The scorer uses the same estimand as the text: a diverse pool is four distinct
 1116 FMs with one decoder-seed member from each FM. Under that definition, the nominal $M=4$ diverse
 1117 pools yield $M_{\text{eff}} \approx 2\text{--}3$ under the mean equicorrelation readout, not four. The ρ_{max} column shows
 1118 that the most redundant valid diverse tuples have lower conservative readouts (1.30–1.69), while
 1119 same-FM seed pools are near one under the same calculation. The exact level varies by task, tuple,
 1120 and failure threshold, so we use M_{eff} as a dependence readout, not as a literal deployment objective.

1121 **Pairwise ρ_{fail} heterogeneity disclosure (equicorrelation caveat).** The M_{eff} formula assumes constant
 1122 pairwise correlation across member pairs (equicorrelated pool). In practice the per-pair ρ_{fail}
 1123 varies across the bundles in each task. From the `m_eff/*.json` outputs, mono-pool ρ_{fail} across
 1124 the per-task functional bundles has mean / SD / range: Kvasir mean 0.865, SD 0.103, range
 1125 [0.601, 0.948] ($n=11$); ACDC LV mean 0.847, SD 0.122, range [0.518, 0.931] ($n=10$); BraTS
 1126 foreground mean 0.860, SD 0.117, range [0.553, 0.958] ($n=10$); RIGA Cup mean 0.833 over 7
 1127 defined mono pools, and RIGA Disc mean 0.812 over 10 defined mono pools. Undefined rows
 1128 occur when an FM’s $\tau=0.85$ failure indicator is constant, giving zero-variance Pearson. Diverse
 1129 4-of- K compositions span a wider ρ_{fail} distribution, which is why Table 27 reports both means and
 1130 the released-enumeration ρ_{max} rather than selecting a single representative pool.

1131 A28 Shared-backbone random-effects representation

1132 **Model.** For each FM m in a frozen-feature pool and test case i , let the per-case score (or its logit)
 1133 decompose as

$$s_m(i) = u(i) + a_{\text{arch}(m)}(i) + p_{\text{family}(m)}(i) + \varepsilon_m(i),$$

1134 where u is task-inherent item difficulty, a is a zero-mean architecture-family random effect, p is a
 1135 zero-mean pretraining-family random effect nested within architecture, and ε_m is an FM-specific
 1136 adaptation noise. Assume u, a, p, ε are mutually uncorrelated with variances $\sigma_u^2, \sigma_a^2, \sigma_p^2, \sigma_\varepsilon^2$. Then

the pairwise covariance of $s_m, s_{m'}$ takes a nested form:

$$\text{Cov}(s_m, s_{m'}) = \begin{cases} \sigma_u^2 + \sigma_a^2 + \sigma_p^2, & (\text{same family}) \\ \sigma_u^2 + \sigma_a^2, & (\text{same architecture, different pretraining}) \\ \sigma_u^2, & (\text{different architecture}). \end{cases}$$

Diagnostic expectation. Under this second-moment description, and absent parameter cancellations, one expects the ordering

$$\rho_{\text{within-FM}} \geq \rho_{\text{same-arch-diff-pretraining}} \geq \rho_{\text{cross-arch}}$$

as a qualitative diagnostic rather than a theorem. Empirically on 5 same-architecture-family ViT variants (Table 3) we observe the ordering on Cup, and on ACDC after applying the functional filter: $0.988 \geq 0.807 \geq 0.639$; the all-5 ACDC row is a sensitivity to including nonfunctional CLIP-B. This is consistent with the predicted ordering and with trajectory/same-basin evidence from deep ensembles and model soups [Fort et al., 2019, Neyshabur et al., 2020, Wortsman et al., 2022]; Abe et al. [2024] is better read here as motivation that predictive diversity can remain constrained in high-capacity ensembles, not as a trajectory-bound theorem.

Implication for M_{eff} . When the binary failure thresholds preserve the same broad ordering as the continuous Dice scores, M_{eff} is monotone decreasing in the observed ρ_{fail} by definition. The empirical diagnostics therefore suggest the descriptive ordering $M_{\text{eff}}^{\text{within-FM}} \leq M_{\text{eff}}^{\text{same-arch}} \leq M_{\text{eff}}^{\text{cross-arch}}$ in the tested regime, but the paper estimates this quantity from released per-case traces rather than inferring it from labels alone. Same-bundle seed pools are same-FM by construction and sit near the lower end of the measured M_{eff} range; diverse pools that cross both architecture and pretraining bundles sit closer to the upper end.

Caveats. This is a second-moment model at the per-case-score level, not a first-principles claim about SGD dynamics. The random-effects representation is descriptive; it gives a language for the observed ordering and a hook for bounding pool redundancy, not a proof that all shared-pipeline pools must have the same structure. The item-difficulty residualisation of Appendix A26 removes u and leaves $a + p + \varepsilon$; our held-out null is consistent with $a + p$ contributing after u is removed.

A29 Architecture-confound diagnostics

The main-text Table 3 reports the released 5-variant same-architecture-family ViT / different-pretraining diagnostic on Cup and ACDC. A secondary Kvasir-SEG replication used the same 5 ViT-B variants, but its aggregate JSON is outside the released artifact, so it is not tabulated as released evidence here.

Table 28: **Appendix detail for the same-architecture ViT diagnostic.** Within-FM is the decoder-seed floor; same-architecture-family ViT cross-pretraining holds the broad architecture family fixed while varying pretraining. This expands Table 3 by showing the all-5 ACDC sensitivity and the functional-filtered readout. Source: released aggregate diagnostic results/_merged/diagnostics/r7_arch_confound_analysis.json; raw NPZ prediction caches are outside the released artifact.

Task	ViT subset	Used	Seed floor	ViT cross-pretrain	Primary cross	Gap
RIGA Cup	all 5 ViTs	5/5	0.968	0.680	0.361	0.288
ACDC LV	all 5 ViTs	5/5	0.965	0.631	0.639	0.334
ACDC LV [†]	CLIP-B dropped	4/5	0.988	0.807	0.639	0.182

The same-architecture cross-pretraining separation is 0.18–0.29 across the released Cup/ACDC diagnostics under the paper’s functional filter ([†] ACDC row drops the non-functional CLIP-B at Dice=0.21<0.30, which otherwise compresses same-arch cross $\bar{r}$ and inflates the separation to 0.33). The direction is preserved, but the claim is mechanism-diagnostic rather than causal. Results are consistent with the random-effects representation in Appendix A28.

A30 Same-family cross-checkpoint decomposition

The main text explicitly cautions that the primary “same” column is dominated by same-bundle different-decoder-seed pairs. To test whether the signal collapses to that seed floor, we recompute

the released same-task per-case traces, including additional same-family checkpoint variants beyond the primary pool, under a three-stratum decomposition: (i) same-FM decoder-seed pairs, (ii) distinct checkpoints within a broad upstream family (DINOv2-B/L/S, OpenAI CLIP B/16-B/32-L/14, ConvNeXt-T/S/B, and ResNet-18/34/50/101 where functional), and (iii) cross-family pairs. This is a diagnostic broad-family grouping rather than a causal architecture/pretraining decomposition. The distinct-checkpoint same-family gap remains positive on all five primary rows, but it is substantially below the same-bundle seed floor.

Table 29: **Same-family cross-checkpoint diagnostic.** Same-FM seed pairs form the upper dependence floor; distinct-checkpoint same-family pairs remain above the diagnostic cross-family stratum but below same-bundle seeds. The Diagnostic cross column is not the Table 1 cross column. Source: `results/_merged/diagnostics/cross_checkpoint_family.json`.

Task	Same-bundle	Same-family checkpoint	Diagnostic cross	Checkpoint gap	95% CI
Kvasir	0.987	0.777	0.697	0.080	[0.043, 0.122]
ACDC LV	0.986	0.750	0.616	0.134	[0.113, 0.155]
BraTS foreground	0.977	0.682	0.589	0.093	[0.085, 0.100]
RIGA Cup	0.963	0.537	0.271	0.267	[0.204, 0.332]
RIGA Disc	0.971	0.421	0.194	0.227	[0.149, 0.293]

CI's are case-bootstrap intervals with $B=2,000$. The cross column is the diagnostic broad-family cross stratum from this expanded checkpoint grid, not the Table 1 analytic-pool cross-family column. For the slice-derived rows, subject-aggregated checkpoint-family gaps remain positive: ACDC 0.166 [0.091, 0.262] and BraTS 0.098 [0.078, 0.119]. BiomedCLIP is kept separate from OpenAI CLIP because it changes both pretraining corpus and objective. The OpenAI CLIP same-family row mixes CLIP-B/16 QuickGELU-style checkpoints with the released CLIP-L/14 `open_clip` standard-GELU dense wrapper, so that family-specific readout also absorbs this loader/activation difference.

A31 Additional mechanism and scope diagnostics

This section reports seven additional scope analyses that close axes around the main result: decoder capacity, random-initialized CNNs, loss/augmentation, fold-reslicing, BraTS train-size, BraTS multimodal input, and MC dropout. The released artifact contains the corresponding per-case Dice traces and the CPU-only summary scripts used for these appendix tables. All numbers are computed directly from the per-case traces; per-FM means use seed averaging where multiple seeds are present, and Pearson $\bar{r}$ is symmetric multi-seed.

A31.1 Decoder-capacity sweep

We rerun the Section 6 (C1) decoder sweep at four explicit head designs spanning 769 to 3.94M trainable parameters: `linear_1x1`, `mlp_2layer`, `unet_lite`, and `transformer_decoder`. Two FMs (DINOv2-B, BiomedCLIP) $\times$ two seeds {42, 43} on RIGA Cup. Across the $\sim 5,000 \times$ capacity range, pool-mean Dice ranges over 0.725–0.848 but cross-decoder within-FM Pearson $\bar{r}$ remains 0.69 on average ($n=48$ pairs across the four head designs and the two FMs)—above the RIGA Cup functional-pool cross-family reference of 0.361 (Table 1). Head capacity changes accuracy and some correlations, but does not eliminate same-backbone co-failure in this diagnostic.

Table 30: **Head-capacity sweep on RIGA Cup** (2 FMs $\times$ 2 seeds, 30-epoch frozen probe). Pool-mean Dice and trainable-parameter count per head design; the diagnostic's within-FM cross-head correlation summary is $\bar{r}=0.690$ (range 0.510–0.896), above the Table 1 RIGA Cup cross reference of 0.361.

Head design	Pool-mean Dice	Trainable params
<code>linear_1x1</code>	0.725	769
<code>mlp_2layer</code>	0.810	197,121
<code>unet_lite</code>	0.848	3,319,105
<code>transformer_decoder</code>	0.805	3,940,865

A31.2 Random-init CNN decorrelation

To check whether within-vs-cross structure persists in the absence of pretraining, we rerun the released random-init diagnostic on two CNN backbones (ConvNeXt-Tiny, ResNet-50) trained from

scratch, 8 seeds {42–49}, on ACDC LV and RIGA Cup. Even with no shared pretraining, within-FM (cross-seed) $\bar{r}$ stays clearly above cross-FM (different-architecture) $\bar{r}$: 0.916 vs. 0.830 on ACDC LV ($\Delta=0.086$), 0.841 vs. 0.720 on RIGA Cup ($\Delta=0.121$). The gap is much smaller than the frozen-FM Table 1 numbers (0.834/0.361 on Cup, $\Delta=0.473$), but it is still positive, consistent with architecture and shared recipe contributing to within-vs-cross structure in this diagnostic.

Table 31: **Random-init CNN same/cross diagnostic.** Two CNN backbones $\times$ eight seeds; 30 epochs from scratch; same 1×1 probe protocol. Positive but smaller gaps suggest architecture/recipe dependence even without shared pretrained weights. Diagnostic only.

Task	within-FM $\bar{r}$	cross-FM $\bar{r}$	Δ
ACDC LV	0.916	0.830	0.086
RIGA Cup	0.841	0.720	0.121

1202 A31.3 Loss/augmentation sensitivity

1203 The Section 4 protocol uses BCE loss, 224 inputs, no augmentation. To check whether the gap
1204 survives a closer-to-standard frozen-head loss/augmentation recipe, we evaluate the seven-bundle
1205 subset {DINOv2-B/S, CLIP-L, ConvNeXt-T, EfficientNet-B0, ResNet-50/18} $\times$ 4 seeds with com-
1206 bined soft Dice+BCE loss and paired {hflip, rot $\pm 15^\circ$ } augmentation. All runs still keep the en-
1207 coder frozen, use the same 1×1 -conv decoder, Adam 10^{-3} , and encoder-native 224×224 input;
1208 this is not a full clinical segmentation pipeline. The original RIGA Cup row remains the primary
1209 loss/augmentation diagnostic in Appendix A18; we also report three task-level diagnostic exten-
1210 sions (RIGA Disc, ACDC LV, and an ISIC-2018 Task 1 fixed 400-image held-out subset) to check
1211 whether the positive gap is an isolated Cup result. ISIC here is a supplementary per-case Dice di-
1212 agnostic on a fixed 400-image subset and is not the official ISIC challenge test/server protocol or
1213 threshold-Jaccard leaderboard score.

Table 32: **Loss/augmentation sensitivity on RIGA Cup, 7 FMs $\times$ 4 seeds.** This is a closer-to-standard frozen-head recipe using Dice+BCE and simple augmentation, not a full clinical segmentation pipeline; all rows remain frozen 1×1 heads at 224×224 . Per-FM seed-averaged Dice; pool $\bar{r}$ at the bottom. CLIP-L uses the same caveated OpenCLIP ViT-L-14/openai dense-token wrapper as the primary bundle.

FM	mean Dice	σ (seed)
DINOv2-B	0.758	0.005
DINOv2-S	0.733	0.003
CLIP-L	0.720	0.006
EfficientNet-B0	0.626	0.002
ResNet-50	0.618	0.010
ConvNeXt-T	0.583	0.004
ResNet-18	0.526	0.027
Pool: within $\bar{r}=0.903$, cross $\bar{r}=0.556$, $\Delta=0.347$ [0.275, 0.428].		

Table 33: **Task-level loss/augmentation extensions.** Same seven-bundle, four-seed frozen-head Dice+BCE+augmentation recipe as Table 32. These rows are not primary rows. Dice columns summarise per-bundle seed-averaged Dice across the seven bundles. The ISIC row is a fixed 400-image-subset held-out diagnostic, not the official ISIC test/server evaluation.

Task	N_{test}	FM-mean Dice	FM min	FM max	Δ	95% CI
RIGA Cup	95	0.652	0.526	0.758	0.347	[0.275, 0.428]
RIGA Disc	95	0.853	0.756	0.922	0.617	[0.535, 0.693]
ACDC LV	361	0.565	0.389	0.740	0.326	[0.291, 0.366]
ISIC-2018 staged	80	0.864	0.801	0.918	0.473	[0.375, 0.562]

1214 Across all four loss/augmentation rows, the Dice+BCE+augmentation recipe leaves a positive
1215 within-vs-cross gap ($\Delta = 0.326\text{--}0.617$). The magnitude varies by task: ACDC remains lower
1216 accuracy and lower-gap than the RIGA/ISIC rows, while RIGA Disc shows the largest extension
1217 gap. These rows are treated as scope checks rather than primary evidence because they use a seven-

1218 FM subset and, for ISIC, a fixed subset/split; they indicate that the simplified BCE/no-augmentation
1219 protocol is not the sole source of the co-failure signal.

1220 A31.4 Five-fold CV robustness

1221 The fold-reslicing analysis is reported with the nnU-Net complements in Appendix A17 and Ta-
1222 ble 17. The key scope point is that the released artifact has one decoder seed per fold; it supports
1223 split-stability of cross-bundle $\bar{r}$, not a within-fold vs cross-fold same-backbone estimator.

1224 A31.5 BraTS foreground train-size sweep

1225 This train-size diagnostic covers the five BraTS foreground bundles with completed feature caches
1226 (DINOv2-B, BiomedCLIP, CLIP-L/14, EfficientNet-B0, ResNet-50). It uses one decoder seed (42)
1227 and fractions $\{0.10, 0.25, 0.50, 1.00\}$ of the same cached training set, so it is an *accuracy/saturation*
1228 check rather than a within-vs-cross dependence estimator. The result is non-monotone: the 5-bundle
1229 mean Dice rises from 0.579 at 10% to 0.621–0.622 at 25–50%, then stays flat/slightly lower at full
1230 data (0.617). Thus, in this cached five-bundle subset, adding more 2D FLAIR slices improves the
1231 weakest low-data setting but does not by itself create a new high-performing, decorrelated pool; the
1232 correlation claim still has to be audited directly rather than inferred from sample size.

Table 34: **BraTS train-size accuracy/saturation diagnostic, one seed only** (5 cached bundles, seed 42). Mean/min/max Dice are across bundles at each train fraction; this table does not estimate within/cross dependence because each bundle has one seed.

Train fraction	Cells	Mean Dice	Min Dice	Max Dice
0.10	5	0.579	0.418	0.721
0.25	5	0.621	0.455	0.750
0.50	5	0.622	0.436	0.751
1.00	5	0.617	0.458	0.737

1233 A31.6 BraTS multimodal RGB derivative

1234 This analysis checks whether the BraTS foreground co-failure signal is an artifact of using a uni-
1235 modal FLAIR derivative. We build a deliberately simple three-channel 2D derivative with the same
1236 held-out subjects and top-10 foreground-area slice rule as Table 1, but pack three MRI modalities
1237 into RGB channels: R=T1c, G=T2w, B=T2-FLAIR. This is *not* the official 3D BraTS challenge
1238 protocol; it is a controlled frozen-feature stress test of the input-channel choice.

1239 The evaluated grid contains 20 cells (5 FMs $\times$ 4 seeds). Multimodal input narrows the dependence
1240 gap relative to the unimodal Table 1 BraTS row, but does not remove it: case-level within $\bar{r}=0.904$,
1241 cross $\bar{r}=0.654$, $\Delta=0.249$ [0.238, 0.262]. Aggregating each subject’s slices before correlation gives
1242 within $\bar{r}=0.946$, cross $\bar{r}=0.772$, $\Delta=0.174$. The conservative interpretation is therefore narrower
1243 than “BraTS input modality does not matter”: adding T1c/T2w raises cross-family agreement and
1244 shrinks the gap, yet same-backbone/family dependence remains positive under the same frozen-head
1245 audit.

Table 35: **Simple three-channel BraTS 2D derivative** (R=T1c, G=T2w, B=T2-FLAIR; 5 FMs $\times$ 4 seeds). This is the same 2D frozen-head audit, not the official 3D BraTS multimodal challenge protocol. Dice is seed-averaged by FM; dependence uses the same Pearson estimator as Table 1.

FM	Mean Dice	Min Dice	Max Dice	Seeds
BiomedCLIP	0.713	0.669	0.739	4
DINOv2-B	0.703	0.675	0.721	4
DINOv2-S	0.702	0.690	0.714	4
CLIP-L/14	0.686	0.682	0.692	4
ResNet-50	0.461	0.439	0.473	4
Case-level pool: within $\bar{r}=0.904$, cross $\bar{r}=0.654$, $\Delta=0.249$ [0.238, 0.262].				
Subject-aggregated pool: within $\bar{r}=0.946$, cross $\bar{r}=0.772$, $\Delta=0.174$.				

A31.7 MC dropout as a within-family decorrelator

A natural pool-design question is whether MC dropout [Gal and Ghahramani, 2016] can act as a lightweight within-family decorrelation lever. We rerun DINOv2-B on RIGA Cup with three dropout rates $p \in \{0.1, 0.2, 0.4\}$, 4 seeds, 10 MC samples per case. The deterministic within-FM cross-seed $\bar{r} \approx 0.97\text{--}0.99$; under MC averaging, $\bar{r}$ drops to 0.95 ($p=0.1$), 0.93 ($p=0.2$), 0.89 ($p=0.4$). The largest tested drop is ≈ 0.10 and remains materially smaller than the same-architecture-family ViT-vs-primary-cross separation seen on Cup (Table 3). MC dropout is therefore a weak within-family lever in this diagnostic; consistent with Appendix A13, it does not approach the architecture-family separation.

Higher- p and any-draw scope. A separate, non-released MC-dropout probe at higher dropout rates suggested lower any-draw correlations than the averaged readout above, but its raw traces are outside the released artifact and its scorer differs from the deterministic-vs-MC-averaged readout used in this section. We therefore do not tabulate those values; the released MC-dropout evidence is the table above, which uses the bundled per-case traces.

Table 36: **MC-dropout diagnostic.** DINOv2-B/RIGA Cup with four seeds and 10 MC samples per case; correlations are computed on MC-averaged per-case Dice. MC dropout provides modest decorrelation but remains far above the Table 1 RIGA Cup cross-family reference of 0.361. Diagnostic only.

Dropout p	MC mean Dice	Det. within $\bar{r}$	MC within $\bar{r}$	H_{pred}
0.1	0.666	0.972	0.947	0.062
0.2	0.657	0.972	0.925	0.067
0.4	0.633	0.987	0.893	0.076

A32 Release manifest and artifact metadata

The artifact bundle includes the primary-task result JSONs, source code, metadata, selected figure / appendix inputs, and selected aggregate diagnostic JSONs needed to audit the numerical claims from saved artifacts where redistribution is permitted. The primary rows are stored under `results/_merged`: the per-seed `per_case_dice` JSONs carry `test_ids`, `per_case_dice`, FM name, seed, and checksum fields; `within_cross/*.json`, `m_eff/*.json`, and `paper_table.json` store the primary summaries. The same directory also includes `subject_level/*.json` for the ACDC/BraTS cluster checks, `calibration_split/*.json` as split-half sensitivity (not the primary protocol), `provenance_summary.json` for source counts and case-unit metadata, and `diagnostics/*.json` for the pixel-error, matched-lift, architecture, quality-controlled pair-regression, item-difficulty, and held-out aggregate diagnostics used only as secondary evidence. The bundle does *not* include raw medical images, raw labels, full prediction NPZs, feature caches, or held-out raw prediction arrays; full split manifests beyond the case IDs embedded in the result JSONs are not included in this artifact release.

ID and source-label scope. Kvasir rows expose only image-level file identifiers in the released `test_ids`; no patient, video, or near-duplicate grouping identifier is available in the released trace bundle, so we do not claim patient-level duplicate control for that row. RIGA Cup/Disc rows retain deterministic hashed alignment identifiers and non-identifying source/task labels only to make the $n=95$ Magrabia-source test split auditable. Raw upstream folder-name components are not needed for this analysis and are not redistributed as case identifiers; the released identifiers are not validated demographic annotations and must not be used for subgroup, gender, or individual-level analysis.

Encoder-loader boundary. The released traces should be read as frozen-feature probes through the bundled loader wrappers, not as official vendor inference runs. In particular, the `clip_vitl14` traces use the version-pinned `ViT-L-14/openai` path exposed by the installed `open_clip` package; `open_clip` also exposes a `ViT-L-14-quickgelu/openai` model definition and warns that OpenAI weights are associated with QuickGELU. We therefore treat CLIP-L/14 as the bundle’s documented dense-feature wrapper and do not use it to make a claim about official OpenAI CLIP inference accuracy. The code keeps this loader unchanged so that the released per-case traces remain reproducible.

Compute and environment scope. The artifact-verification path is CPU-only: the bundled verification script and the individual `compute_*` scripts read released JSON traces and finish in minutes on a typical workstation. Raw-image reproduction is heavier and requires separate upstream data/checkpoint access. The reported training runs used a shared GPU cluster with NVIDIA A100-80GB, A6000, and L40S devices; feature extraction time depends on backbone and task size, and cached frozen-head training is lightweight relative to feature extraction. We therefore document resource classes and released trace-level commands rather than claiming a single exact GPU-hour total because raw-image retraining is outside the submitted artifact’s reproduction scope.

Table 37: Compute scope for reproducibility. The submission supports trace-level recomputation directly; raw-image training resources are provided as reproduction guidance because upstream data, checkpoints, and feature caches are not redistributed.

Reproduction target	Required compute	Notes
Primary and diagnostic summary JSONs	CPU, minutes-scale	Run <code>PYTHONPATH=code/src python3 code/scripts/verify_artifact.py</code> ; no raw images, checkpoints, Torch, or GPU required.
Frozen-feature primary traces from raw data	A100-80GB/A6000/L40S-class GPU plus storage for upstream data and features	Feature extraction is the main cost; cached 1×1 decoder-head training uses Adam for 30 epochs and is comparatively lightweight.
Appendix architecture, adaptation, and nnU-Net complements	Same cluster GPU classes; optional MONAI/nnU-Net environments where applicable	These are diagnostic or limited-scope runs. The bundled trace-level artifact is sufficient to recompute reported released summaries without rerunning them.

Table 38: Trace-level claim-to-artifact map. The artifact supports recomputation of primary audit statistics from derived per-case Dice traces. It does not reproduce raw-image training unless users separately obtain upstream datasets and checkpoints. Aggregate-only diagnostics are marked as such rather than promoted to primary evidence; exact commands and checksums are in `README.md`, `metadata/MANIFEST.md`, and `metadata/SHA256SUMS`.

Claim/readout	Released source	Recompute / limitation
Primary same-vs-cross gaps, CIs, and filtered functional pools	<code>paper_table.json</code> , <code>within_cross/*.json</code> , <code>per_case_dice/*.json</code>	<code>PYTHONPATH=code/src python3 code/scripts/verify_artifact.py</code> ; case IDs and per-case Dice are included, raw images and masks are not; trace-level recomputation only.
Subject, floor, ICC, permutation, event, and quality-control robustness	<code>subject_level/*.json</code> , <code>audit_sensitivity.json</code> , <code>failure_event_summary.json</code> , <code>quality_controlled_pairs.json</code>	The verification script calls <code>compute_all.py</code> , <code>compute_audit_sensitivity.py</code> , <code>compute_failure_event_summary.py</code> , <code>compute_quality_controlled_pairs.py</code> ; these are estimator checks, not new training runs.
Item-difficulty and cross-checkpoint diagnostics	<code>item_difficulty_robustness.json</code> , <code>cross_checkpoint_family.json</code> , <code>r7_arch_confound_analysis.json</code>	<code>compute_item_difficulty_robustness.py</code> and <code>compute_cross_checkpoint_family.py</code> ; diagnostic evidence for alternative explanations, not causal identification.
Pixel/lift, same-backbone, held-out, and nnU-Net scope rows	<code>diagnostics/pixel_error/*.json</code> , <code>diagnostics/matched_lift/*.json</code> , appendix trace files, <code>nnunet/*.json</code> , <code>nnunet_case_riga_2d_100ep/*.json</code>	Released summaries plus the RIGA case-identical nnU-Net per-model traces; other full-pipeline comparators remain aggregate or task-level only.

Table 39: Trace inventory and analytic-pool provenance. Merged counts all completed (bundle, decoder-seed) JSON entries under `results/_merged/per_case_dice`; Table 1 uses the original 44-trace primary source subset and applies the functional-floor pool shown in the final column. Same/cross pair counts for the analytic pool are reported in Table 1. The release key `brats_wt` denotes the 2D BraTS `seg>0` foreground derivative used in this paper; it is not the official 3D BraTS WT challenge target.

Task	<i>N</i>	Merged bundles	Merged members	Table 1 source	Analytic pool
<code>kvasir</code>	120	17	68	44	11 bundles
<code>acdc_lv</code>	361	17	68	44	10 bundles
<code>riga_cup</code>	95	17	68	44	9 bundles
<code>riga_disc</code>	95	17	68	44	11 bundles
<code>brats_wt</code>	3228	17	68	44	10 bundles / 40 members

Functional-floor members are retained in the source files. The paper applies a $\text{Dice} \geq 0.30$ functional floor downstream in the estimator, not at result ingestion. Thus the merged `per_case_dice` tree can list more completed JSONs than the analytic pool used by Table 1; the extra cross-checkpoint members support Appendix A30. For example, ACDC LV has 68 merged source JSONs

1301 but Table 1 retains the 44-trace primary source subset and 10 bundles after the floor. The au-
1302 thoritative row-level provenance for the table is `within_cross/*.json`, especially the `pool` and
1303 `fms_kept_after_floor` fields.

1304 **What else is in the bundle.** In addition to the primary `per_case_dice` tree, the bundle includes
1305 released trace sets for decoder capacity, UNet-skip decoding, loss/augmentation sensitivity, fold-
1306 reslicing, BraTS train-size, MC dropout, BraTS multimodal input, random-init adapted heads, and
1307 a small set of ACDC DINOv2-B full-fine-tuning provenance traces. It also includes selected aggregate
1308 diagnostic JSONs under `results/_merged/diagnostics` plus code sufficient to inspect and
1309 regenerate the included primary summaries and appendix summaries. Appendix-only diagnostics
1310 not listed here, raw prediction arrays, and full pixel-mask caches are not included in the supplement
1311 and are treated as reported diagnostics rather than fully bundled reproduction targets.

1312 **What is not in the bundle and why.** Raw medical images, raw labels, feature caches, and full
1313 prediction NPZ caches are not included. For restricted datasets such as BraTS, the release is limited
1314 to derived per-case Dice and aggregate dependence readouts; users must obtain the raw data from
1315 the upstream provider. Larger derived caches and more detailed split manifests are not included in
1316 the released artifact.

1317 NeurIPS Paper Checklist

1318 1. Claims

1319 Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contribu-
1320 tions and scope?

1321 Answer: [Yes]

1322 Justification: The abstract and introduction claim that *in the frozen and light-adaptation regime of med-*
1323 *ical segmentation pipelines*, same frozen-checkpoint/recipe pools show higher per-case Dice-score corre-
1324 lation, used as a failure-dependence proxy, than the primary-table cross-bundle reference pools. The pri-
1325 mary table uses task-specific functional pools after Dice-floor filtering: Kvasir 11 FMs, ACDC LV 10,
1326 BraTS foreground 10, RIGA Cup 9, and RIGA Disc 11. The scope is explicitly restricted to retrospective
1327 public-benchmark dependence audits; secondary pixel/lift readouts, released nnU-Net complements, and
1328 non-primary mechanism diagnostics are diagnostic or scope context rather than causal clinical-outcome
1329 evidence.

1330 2. Limitations

1331 Question: Does the paper discuss the limitations of the work performed by the authors?

1332 Answer: [Yes]

1333 Justification: The main paper gives high-level limitations in Section 7, and the appendix expands them
1334 with scoped protocol notes: frozen/light-adaptation only; LoRA and RIGA full-fine-tuning probes with-
1335 out released trace artifacts are retained only as scope notes; released CNN random-init and small ACDC
1336 full-fine-tuning provenance traces are diagnostic, not primary evidence; UNet-with-skip and nnU-Net v2
1337 complements reported, including one case-identical RIGA 2D held-out scope check, but not dense 3D
1338 FM/full-pipeline baselines for every primary row; the 3D architecture pilot is a scope-boundary diagnostic;
1339 2D axial slices at 224×224 for the primary FM audit; MedSAM used only as a frozen vision-encoder probe;
1340 broader cross-scanner / cross-demographic shifts remain untested; clinical Dice thresholds are operational
1341 conventions and many lift thresholds are support-limited.

1342 3. Theory assumptions and proofs

1343 Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and
1344 correct) proof?

1345 Answer: [N/A]

1346 Justification: The paper is empirical and benchmark-oriented; it does not include theoretical results or proofs.
1347 Section 3 introduces estimand notation but the $\bar{r}$, Δ , and lift statistics are descriptive.

1348 4. Experimental result reproducibility

1349 Question: Does the paper fully disclose all the information needed to reproduce the main experimental
1350 results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless
1351 of whether the code and data are provided or not)?

1352 Answer: [Yes]

1353 Justification: Section 4, Appendix A13, and Appendix A32 specify the public FM checkpoints, decoder
1354 architecture (1×1 -conv default plus stated sweeps), training protocol (Adam 10^{-3} , 30 epochs for the frozen-
1355 head primary runs, batch 16, BCE for the binary primary rows), seed set ($\{42, 43, 44, 45\}$ with σ -seeded de-
1356 terminism), and the actual primary split protocols: Kvasir 880/120, RIGA train = BinRushed+MESSIDOR

with Magrabia (Magrabi Eye Center) source held out ($n=95$), ACDC public patients 001–100 with patient-level fold-0 LV-positive slices ($n=361$), and the BraTS 2024 2D FLAIR `seg>0` foreground derivative: per held-out subject, up to the top-10 axial slices by foreground area with at least 64 positive pixels, subject-level 80/20 split using NumPy seed 42 ($n=3228$). This BraTS target includes the GLI resection-cavity label and is not the official 3D BraTS WT challenge target. The supplementary artifact releases primary-task per-case Dice JSONs, merged summaries, selected aggregate diagnostic JSONs, code, command documentation, hashes, and a data card for trace-level reproduction; raw-data retraining requires separately obtaining upstream datasets/checkpoints, and full raw data, full prediction caches, raw held-out prediction arrays, and more detailed split manifests are outside the supplementary artifact.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All source datasets used in the reported evidence are public or access-controlled research resources from their upstream providers, subject to their licences, research-use terms, or DUAs: RIGA+, ACDC, Kvasir-SEG, BraTS 2024, MSD tasks, and ISIC-2018. All backbones are publicly available or documented as gated research checkpoints. The supplementary artifact releases benchmark code, estimator implementations, primary-task per-case Dice traces, merged summaries, selected appendix result JSONs, README/MANIFEST commands, hashes, and a data card; authored code and metadata use the bundle licence, while derived scalar traces remain subject to upstream dataset/model terms. Raw medical images, masks, full prediction caches, some held-out-task artifacts, and complete split manifests are not redistributed in the supplementary artifact, so open access is provided for trace-level reproduction rather than raw-data end-to-end retraining.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 3 gives the formal estimands (within-/cross-family $\bar{r}$, spatial Δ_{pix} , threshold lift L) and the symmetric 4×4 seed-combo averaging rule; Section 4 gives the datasets (with patient-level splits), foundation-model list, and protocol (image size, encoder-specific upstream normalisation, decoder, optimiser, loss, epochs, batch, seed control).

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Table 1 reports paired case-level and hierarchical family/seed bootstrap 95% CIs ($B=2000$). The BraTS foreground derivative is analysed at the slice level with an explicit within-subject dependence caveat, Appendix A11 reports subject-level sensitivity checks, and Appendix A12 explains why sparse-class NCR filtering diagnostics are not used as numeric evidence.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix A32 reports the compute/resource scope. For trace-level verification, the artifact is directly reproducible on CPU in minutes through the artifact verification script. Raw-image reproduction used NVIDIA A100-80GB, A6000, and L40S GPUs across a shared cluster. Feature extraction per FM per domain ranges from minutes for CNN backbones to longer ViT/BraTS jobs; cached 1×1 -head training per seed is lightweight. For raw-image retraining, we report resource classes and per-job elapsed times where recorded, but do not claim a single exact total GPU-hour count because raw retraining is outside the submitted artifact’s reproduction scope.

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research uses only de-identified public or access-controlled research medical imaging datasets (RIGA, ACDC, Kvasir-SEG, BraTS 2024) under their original research licences, terms, or DUAs. No additional patient data are collected. The paper audits retrospective failure-correlation properties of benchmark segmentation pipelines and does not propose new clinical tools.

10. Broader impacts

1416 Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of
1417 the work performed?

1418 Answer: [\[Yes\]](#)

1419 Justification: Section 7 discusses deployment implications: shared-pipeline model pools (same FM, same ar-
1420 chitecture family, same recipe) in frozen/light-adaptation medical AI deployments may be over-estimated as
1421 providing independent safety margins. The paper quantifies a retrospective evaluation assumption (independ-
1422 ence among FM-pool members) and does not measure clinician decisions, diagnoses, or patient outcomes.
1423 Backbone-family diversification is observed to have lower dependence in the tested comparisons, not to be
1424 a complete clinical mitigation.

1425 **11. Safeguards**

1426 Question: Does the paper describe safeguards that have been put in place for responsible release of data or
1427 models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped
1428 datasets)?

1429 Answer: [\[Yes\]](#)

1430 Justification: The supplementary artifact does not release raw medical images, voxel labels, full prediction
1431 caches, upstream checkpoints, trained clinical services, or patient-facing models. Released artifacts are
1432 derived per-case scalar traces, aggregate summaries, metadata, hashes, and scorer code under upstream
1433 dataset/model constraints. The paper states that the audit is not a clinical safety guarantee, does not support
1434 re-identification or patient profiling, and should not be used as a standalone model-selection or deployment
1435 rule.

1436 **12. Licenses for existing assets**

1437 Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly
1438 credited and are the license and terms of use explicitly mentioned and properly respected?

1439 Answer: [\[Yes\]](#)

1440 Justification: RIGA+ (CC BY-NC 4.0 via Deep Blue; non-commercial use only; raw images and masks
1441 not redistributed), ACDC (research licence via challenge website), Kvasir-SEG (research/educational use
1442 via Simula; raw images not redistributed), BraTS 2024 (research licence/DUA via Synapse; no raw images
1443 or voxel-level labels redistributed), MSD tasks (research challenge resources; raw images and labels not
1444 redistributed), and ISIC-2018 (research challenge resource; raw images and masks not redistributed) are
1445 used only through derived traces or aggregate diagnostics in the submitted artifact. Backbones/checkpoints:
1446 DINOv2 (Apache-2.0), MAE ViT-B/16 (CC BY-NC 4.0), DeiT-B/16 (Apache-2.0), OpenAI CLIP ViT-
1447 B/16 and ViT-L/14 weights (OpenAI CLIP license; open_clip code MIT), BiomedCLIP checkpoint under
1448 its model-card terms rather than the open_clip library license, ConvNeXt-T/EfficientNet-B0/ResNet-
1449 50/18 loaded via torchvision weights/code under PyTorch/torchvision terms, RETFound checkpoint (gated;
1450 CC BY-NC 4.0/research-use), MedSAM/SAM (Apache-2.0). No upstream checkpoints are redistributed.
1451 **BraTS redistribution compliance:** in the supplementary artifact we release only derived scalar traces
1452 and aggregate statistics needed to reproduce the estimators; we do *not* redistribute raw BraTS images, raw
1453 voxel labels, or full prediction caches. All assets are cited, and the artifact/data-card documentation in
1454 Appendix A32 states intended use and non-use.

1455 **13. New assets**

1456 Question: Are new assets introduced in the paper well documented and is the documentation provided
1457 alongside the assets?

1458 Answer: [\[Yes\]](#)

1459 Justification: The paper releases (i) a benchmark protocol (estimands + symmetric multi-seed + bootstrap
1460 estimators); (ii) primary-task per-case Dice JSONs and merged summary JSONs for the released rows;
1461 (iii) selected appendix/ablation result JSONs and reproduction scripts; and (iv) artifact documentation in-
1462 cluding README/MANIFEST files, SHA-256 hashes, a derived-trace data card, and Croissant metadata
1463 submitted for the trace/code artifact. Raw medical images, prediction masks/caches, and full split manifests
1464 are not part of the supplementary artifact because upstream redistribution rights and artifact size constraints
1465 vary across datasets.

1466 **14. Crowdsourcing and research with human subjects**

1467 Question: For crowdsourcing experiments and research with human subjects, does the paper include the full
1468 text of instructions given to participants and screenshots, if applicable, as well as details about compensation
1469 (if any)?

1470 Answer: [\[N/A\]](#)

1471 Justification: No crowdsourcing; no new human-subjects research. We use existing public or access-
1472 controlled research datasets whose upstream governance, approvals, or access terms are handled by the
1473 dataset providers.

1474 **15. Institutional review board (IRB) approvals or equivalent for research with human subjects**

1475 Question: Does the paper describe potential risks incurred by study participants, whether such risks were
1476 disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent ap-
1477 proval/review based on the requirements of your country or institution) were obtained?
1478 Answer: [N/A]
1479 Justification: No new human-subjects data were collected. Upstream datasets carry their own governance,
1480 approvals, or access terms (RIGA, ACDC, Kvasir-SEG, BraTS 2024).
1481 16. **Declaration of LLM usage**
1482 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard com-
1483 ponent of the core methods in this research? Note that if the LLM is used only for writing, editing, or for-
1484 matting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research,
1485 declaration is not required.
1486 Answer: [N/A]
1487 Justification: LLM use was limited to writing, editing, formatting, and programming-assistant support dur-
1488 ing manuscript and artifact preparation. It is not an important, original, or non-standard component of the
1489 benchmark methodology, estimator, training pipeline, or evaluation protocol; all reported numbers come
1490 from the released scripts, logs, and JSON artifacts, and all manuscript text, code, citations, and claims were
1491 reviewed and approved by the authors.