
Auditing Correlated Failures in Frozen-Feature Pretrained-Encoder Pools for Medical Segmentation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Cross-architecture/backbone medical-segmentation pools are commonly moti-
2 vated as complementary failure coverage when members make different mistakes;
3 that assumption is rarely tested at the per-case level. We audit this assumption
4 for a fixed pool of 11 pretrained-encoder bundles read out by a shared 1×1 -
5 convolution decoder (4 seeds), and for a small set of decoder-recipe perturbations
6 on the same bundles. Across five primary task rows (Kvasir polyp, ACDC LV,
7 BraTS foreground, RIGA Cup, RIGA Disc), this 11-bundle 1×1 -conv readout
8 pool yields per-case Dice correlation gaps of $\Delta = 0.261\text{--}0.638$ between same-
9 bundle and cross-bundle pairs after a Dice ≥ 0.30 functional floor; all paired case
10 and hierarchical bootstrap intervals exclude zero. The result is stable under floor
11 sweeps, ICC(A,1), leave-one-out, subject aggregation, quality-controlled pair re-
12 gression, and cross-fitted item-difficulty residualisation. Decoder-recipe diagnos-
13 tics show that richer readouts attenuate the gap (UNet-skip; case-identical nnU-
14 Net on RIGA Cup: $0.473 \rightarrow 0.0079$, separate-split $0.034\text{--}0.047$); the reported
15 magnitudes are recipe-specific to the 1×1 -conv readout, not universal pool prop-
16 erties. Additional probes (pixel error, joint-failure lift, same-architecture-family
17 ViT, same-family cross-checkpoint, same-backbone) locate the dependence struc-
18 ture but do not causally separate architecture, pretraining, and recipe. The claim
19 is scoped to 2D frozen/light-adaptation encoder pools with this readout family
20 and to retrospective benchmark audits. We release per-case Dice traces, summary
21 JSONs, scorer code, hashes, and metadata as the audit artifact.

1 Introduction

23 Ensembles are widely used in safety-relevant medical segmentation, and challenge-winning
24 pipelines explicitly motivate cross-architecture or cross-backbone pools as “complementary” [Fer-
25 reira et al., 2024, Luo et al., 2025], motivated by the classical ensemble-diversity literature on
26 whether diverse members reduce correlated errors [Kuncheva and Whitaker, 2003]. The support-
27 ing evidence in current practice is typically aggregate Dice or qualitative complementarity rather
28 than a measured case-level co-failure correlation. Aggregate Dice does not by itself show whether
29 additional members provide additional failure evidence on the same cases. This paper audits that
30 missing reporting step for frozen or lightly adapted pretrained-encoder segmentation pools: the
31 aligned case-level correlation of segmentation quality across pool members.

32 The medical-imaging supply chain now reuses a small set of foundation-model encoders and con-
33 ventional pretrained backbones: DINOv2, CLIP/BiomedCLIP, SAM/MedSAM, ConvNeXt, Effi-
34 cientNet, ResNet, and domain-specific pretrainings such as RETFound. A pool of “different” seg-
35 menters can therefore share backbone family, pretraining bundle, input statistics, and training recipe.
36 This creates an evaluation gap: aggregate Dice and common uncertainty summaries usually do not
37 directly report aligned case-level co-failure, which is the evidence needed for independent failure-
38 coverage claims. The intuition that near-duplicate pipelines should share errors is not new. The

39 missing piece is a released, case-aligned audit that quantifies how large the dependence is, tests
40 whether it survives simple alternative explanations, and states which evaluation claims the evidence
41 can and cannot support.

42 **Primary contribution.** We contribute a reusable E&D audit protocol for fixed segmentation pools:
43 compute per-case Dice traces, define a prespecified primary-table same/release-recipe group and
44 cross-bundle reference, average symmetrically over decoder seeds, and report paired bootstrap in-
45 tervals for the same/cross correlation gap. The protocol asks a concrete reporting question: before
46 a pool is credited with complementary failure coverage, have same/release-recipe and cross-bundle
47 co-failure correlations been measured on aligned cases? Out of scope: prompt-conditioned Med-
48 SAM, full nnU-Net pipelines, clinical outcomes, demographic fairness, and SOTA segmentation
49 performance.

50 **Artifact and diagnostic contribution.** We release the trace-level files and scorer code needed to
51 recompute the primary audit statistics. Additional diagnostics test estimator sensitivity, subject-
52 level clustering, item difficulty, pixel error overlap, same-family cross-checkpoint structure, same-
53 backbone perturbations, adaptation and full-pipeline boundaries, and selected scope checks. The
54 reusable object is therefore an audit protocol plus a trace-level evidence bundle for measuring
55 whether a fixed pool of segmentation models supplies independent failure evidence on aligned cases.
56 We position the audit against monoculture, error-correlation, representation-similarity, and medical-
57 segmentation literatures next.

58 2 Related work

59 **Relation to prior work and scope of novelty.** Prior work has established that correlated errors and
60 monoculture matter. Our contribution is to make this issue auditable for medical segmentation by re-
61 leasing aligned per-case traces, same/cross estimators, and a primary/diagnostic evidence hierarchy.
62 Our risk framing follows algorithmic-monoculture and foundation-model ecosystem work [Klein-
63 berg and Raghavan, 2021, Bommasani et al., 2021, 2022, Toups et al., 2023]. The correlation gap
64 we measure has modelling antecedents: pairwise Dice between ensemble samples is the established
65 medical-segmentation failure-detection estimator [Roy et al., 2019, Zenk et al., 2025]; trial-by-trial
66 error consistency asks whether systems fail on the same inputs [Geirhos et al., 2020], error correla-
67 tions arise when deployed models share algorithms, datasets, or foundation models [Li et al., 2025],
68 shared training trajectories and weight-space basins can constrain checkpoint diversity [Fort et al.,
69 2019, Wortsman et al., 2022], transfer/representation work studies what is reused across pretrained
70 models [Neyshabur et al., 2020, Gontijo-Lopes et al., 2022], pretrained-model ensemble work treats
71 pretraining as a possible diversity source [Mustafa et al., 2020], same-architecture models can align
72 strongly across initialisations and recent surveys consolidate functional/representational similarity
73 measures [Kornblith et al., 2019, Klabunde et al., 2024], and predictive-diversity pathologies can
74 limit deep-ensemble gains [Abe et al., 2024]. Recent LLM monoculture and null-model papers are
75 relevant because they show that agreement/correlation claims depend on the choice of baseline and
76 unit of analysis; our paper applies that lesson to case-level medical segmentation traces [Jiang et al.,
77 2025, Kim et al., 2025a, Gorecki and Hardt, 2025, Goel et al., 2025, Jo et al., 2026]. Pathology-
78 FM and precision-oncology studies provide adjacent evidence about center robustness, benchmark
79 behavior, and complementary domain-specialised features [de Jong et al., 2025, Campanella et al.,
80 2025, Luo et al., 2025, Neidlinger et al., 2025]. The distinction is threefold: our unit is an aligned
81 segmentation Dice trace rather than a classification decision or deployment outcome, our artifact
82 releases the trace-level evidence needed to recompute the audit rather than only a summary statistic,
83 and the paper separates primary same-checkpoint/release-recipe evidence from mechanism-adjacent
84 diagnostics. Our contribution is not the general claim that correlated errors exist. It is a released
85 medical-segmentation audit of frozen-feature pretrained-encoder pools that distinguishes primary ev-
86 idence from diagnostics and reports where same-checkpoint/recipe, same-architecture-family, and
87 broader cross-family correlations appear under a shared protocol.

88 **Medical segmentation, pretrained/foundation-model benchmarks, and ensemble baselines.**
89 nnU-Net, deep ensembles, MC dropout, adjacent uncertainty or parameter-efficient ensembles, and
90 hyperparameter ensembles are usually evaluated through Dice, calibration, or uncertainty quality
91 [Isensee et al., 2021, 2024, Lakshminarayanan et al., 2017, Gal and Ghahramani, 2016, Mehrtash
92 et al., 2020, Jungo and Reyes, 2019, Mühlematter et al., 2026, Gu et al., 2024]. Classical ensemble-
93 diversity metrics motivate this audit, but our reporting target is the case-aligned correlation of task-

performance traces because that statistic directly tests redundant failure evidence on the evaluated benchmark [Kuncheva and Whitaker, 2003, Ganaie et al., 2022, Wood et al., 2023]. Medical-segmentation diversity and failure-detection benchmarks ask whether a method can detect segmentation failures; our audit asks whether multiple pool members fail on the same cases, making the questions complementary rather than substitutable [Georgescu et al., 2023, Zenk et al., 2025]. The choice of per-case Dice as the unit of analysis follows recent recommendations for image-analysis validation that case-level distributions, not aggregate scalars, drive evaluation reliability [Maier-Hein et al., 2024]. Recent medical-segmentation benchmarks and robustness-specification work evaluate or discuss generic-vision probes, SAM medical-image evaluations, standardized segmentation datasets, promptable SAM adaptations, universal 3D/medical segmenters, task-specific robustness requirements, transferability estimation, and VFM segmentation benchmarking [Baharoon et al., 2023, Huang et al., 2024, Kuş and Aydin, 2024, Kirillov et al., 2023, Ma et al., 2024, Cheng et al., 2023, Archit et al., 2025, Zhu et al., 2024, Ma et al., 2025, Gu et al., 2025, He et al., 2025, Kim et al., 2025b, Xian et al., 2025, Yang et al., 2023, Keressies et al., 2024]. These benchmarks are complementary to our audit: they ask which model or foundation model performs well, whereas we ask whether multiple members of a pool provide non-redundant failure evidence on the same cases. Promptable and 3D medical foundation models are deployment alternatives rather than baselines for this co-failure audit.

3 Formal setup and benchmark estimands

Pool and estimands. Let $\mathcal{F}=\{F_1, \dots, F_K\}$ be the operational pretrained-encoder bundles used by the Table 1 scorer. In the primary 11-bundle source pool, a bundle is the concrete encoder / pretraining / input-statistic wrapper used for frozen-feature extraction. Same-bundle decoder seeds form most of the main same group, and DINOv2-B×DINOv2-S is the only recurring cross-checkpoint member of that group because both variants share the same upstream DINOv2 release and pretraining recipe. Other broad-family groupings (for example OpenAI CLIP sizes, ResNet sizes, and expanded DINOv2/ConvNeXt variants) are not folded into the Table 1 primary estimator; they are separated in the same-family cross-checkpoint diagnostic in Appendix A30. BiomedCLIP (ViT-B/16, biomedical image-text pretraining) is treated as a different bundle from DINOv2 despite being a ViT. A model trace is a tuple (F_k, σ) (bundle, decoder seed) producing per-case Dice $D_{(k,\sigma)}(i, t)$. Appendix A29 and Appendix A30 report diagnostics that separate primary product bundles, same-architecture-family ViT / different-pretraining, and broader same-family cross-checkpoint readouts; these clarify that a “cross” column can still contain broad architectural structure. For a pair of models $a \neq b$ we define $r_{a,b}(t) = \text{Pearson}_i[D_a(i, t), D_b(i, t)]$ (and a McGraw–Wong ICC(A,1) variant $\iota_{a,b}$ [McGraw and Wong, 1996], chosen to penalise per-bundle Dice offsets that Pearson absorbs). The primary-table same/cross means at task t are the expectations of $r_{a,b}$ over pairs in the same/release-recipe group (same bundle, plus DINOv2-B/S) and the remaining cross-bundle reference under this scorer. We define $\Delta(t) = \bar{r}_{\text{same/release}}(t) - \bar{r}_{\text{cross}}(t)$ as the co-failure gap and refer to it as the *primary same/cross gap*, Δ , throughout. Secondary estimands are pixel-level error-mask Dice, computed the same way over binarised error masks, and joint-failure *lift* over independence $L_{\text{mono}}(\tau) = \Pr[\text{all fail} \mid \text{mono pool}] / \prod_m \Pr[\text{fail}_m \mid \tau]$.

Terminology used by the audit. The primary claim uses only the primary-table same/release-recipe group and the cross-bundle reference; broader family rows are diagnostics.

Term	Meaning in this paper	Primary claim?
Bundle	Concrete encoder checkpoint plus wrapper, pretraining, input normalisation, and feature-extraction path	Yes
Model trace	Bundle plus decoder seed and aligned per-case Dice vector	Yes
Primary-table same/release-recipe group	Same-bundle seed pairs plus the prespecified DINOv2-B/S same-release cross-checkpoint pair	Yes
Primary-table cross reference	Remaining primary-pool pairs under the primary grouping	Yes
Same-family cross-checkpoint	Broader DINOv2, CLIP, ConvNeXt, and ResNet diagnostic grid	Diagnostic only
Same-architecture ViT cross-pretraining	ViT-B diagnostic varying pretraining source	Diagnostic only

Symmetric multi-seed averaging. For each bundle pair we enumerate all $|\Sigma|^2=16$ seed combinations (4 per model) rather than seed-matched pairs, removing the one-seed-per-bundle noise that inflates ad-hoc gap estimates. Same-bundle averaging uses $\binom{|\Sigma|}{2}=6$ same-bundle seed pairs per F_k plus the full $|\Sigma|^2$ on DINOv2-B×DINOv2-S; the cross group uses all remaining primary-table bundle pairs.

CI and scope. A 3-level nested hierarchical bootstrap (primary groups $\rightarrow$ seeds $\rightarrow$ cases, $B=2,000$, deterministic seed 0) yields 95% CIs as the 2.5th/97.5th percentiles of bootstrap-sampled $\bar{r}_{\text{same}}, \bar{r}_{\text{cross}}, \Delta$. Because the number of encoder-bundle clusters is small, neither interval should be read as a population-level guarantee over all possible pretrained encoders. We report both intervals as finite-audit uncertainty summaries: the case bootstrap asks about new cases conditional on the evaluated pool, while the hierarchical bootstrap also perturbs the finite group/seed surface. The estimands assume a shared-decoder, shared-task pool and measure *observed* correlation within such pools; a CNN random-init diagnostic (Appendix A31.2) shows positive same-vs-cross structure without pretrained weights, so we report the pretrained-bundle gap descriptively rather than as a causal isolation of pretrained-weight effects.

Primary-group structural caveat. The primary rows retain task-specific functional pools of 11/10/10/9/11 bundles after the Dice floor. In these pools, the Table 1 “same” column is dominated by same-bundle, different-seed decoder-head pairs, with DINOv2-B $\times$ DINOv2-S providing the only recurring same-family cross-checkpoint pair when both variants are functional. Thus the high same-group $\bar{r}$ values in Table 1 (0.834–0.959) are best read as same-frozen-checkpoint/recipe-dominated dependence readouts, not a broad sample of many same-family cross-checkpoint comparisons. Section 6.1 and Appendix A29 report a 5-variant same-architecture-family ViT / different-pretraining diagnostic (DINOv2-B, BiomedCLIP, MAE-B, DeiT-B, CLIP-B) that measures cross-pretraining correlation under a shared broad architecture family and is the closest “within-architecture-family, cross-pretraining” control in this benchmark.

4 Experimental design

Datasets & targets. The primary rows are summarized below. The splits are recorded in the released JSON files’ `test_ids` fields; subject-level sensitivity for ACDC and BraTS is in Appendix A11.

Primary task rows and evaluation units. The audit uses benchmark-derived targets and aligned per-case Dice traces, not clinical deployment cohorts.

Task row	Raw source	Target / split	Unit	Main caveat
Kvasir polyp	Kvasir-SEG [Jha et al., 2020]	Loader-defined 880/120 image split	image, $N=120$	No patient/video IDs available; no patient-level duplicate control claimed.
ACDC LV	ACDC [Bernard et al., 2018]	LV-positive ED/ES slices from public patients 001–100, fold 0	slice, $N=361$; 20 patients	Slice row; patient-level sensitivity reported in Appendix A11.
BraTS foreground	BraTS 2024 GLI [BraTS 2024 Organizers, 2024, de Verdier et al., 2024]	2D FLAIR seg>0; top-10 eligible slices per held-out subject, NumPy seed 42	slice, $N=3,228$; 324 subjects	Includes resection-cavity label; not the official 3D BraTS WT challenge target.
RIGA Cup/Disc	RIGA+ [Almazroa et al., 2018, Almazroa, 2018]	BinRushed+MESSIDOR Magrabia source test	train, fundus, $N=95$	Deterministic hashed alignment IDs retained only for split alignment.

Encoder bundles. *Generic pretrained-encoder bundles* (11, mixing foundation-model encoders and conventional pretrained backbones): DINOv2-B/S, BiomedCLIP, CLIP ViT-L/14, CLIP ViT-B/16, MAE ViT-B/16, DeiT ViT-B/16, ConvNeXt-T, EfficientNet-B0, ResNet-50/18. Encoder names refer to the exact released wrapper/checkpoint paths in the artifact; in particular, CLIP-L/14 uses the OpenCLIP dense-token wrapper (a *non-canonical* CLIP-L inference path), preserved across pools for consistency. *Medical-domain encoder probe:* MedSAM vision-encoder-only (SAM ViT-B; prompt encoder and mask decoder discarded) is reported as a secondary boundary probe on RIGA/Kvasir/ACDC. RETFound support is documented as a gated-checkpoint path, but RETFound cells are not reported because the gated checkpoint was unavailable in the release environment.

Protocol. The primary protocol resizes images and masks to 224×224 with nearest-neighbor mask resizing, applies encoder-specific upstream normalisation (ImageNet-style statistics for DINOv2/MAE/DeiT/torchvision CNNs and OpenAI CLIP statistics for CLIP/BiomedCLIP), freezes the encoder, and trains a 1×1 -conv decoder with Adam 10^{-3} for 30 epochs, batch 16, BCE loss on the binary primary rows, no early stopping, no validation-based model selection, and fixed sigmoid threshold 0.5 at inference. Table 1 starts from the 11-bundle generic source pool where traces are available and applies the estimator’s Dice ≥ 0.30 retrospective functional floor to remove non-segmenting, near-constant traces: Kvasir retains 11 bundles, ACDC 10, BraTS foreground 10, RIGA

Cup 9, and RIGA Disc 11. This test-trace-based analytic convention defines the reported functional pool; it is not a training, deployment, or pre-registered hypothesis threshold. The floor systematically reduces cross-bundle correlation (near-constant outputs would be near-zero correlation with everything), so the post-floor pool biases the gap upward by construction; the released floor sweep (Appendix A5) shows that primary-row signs do not flip across $\{0, 0.20, 0.25, 0.30, 0.35, 0.40\}$, but the magnitudes are floor-conditional. The MedSAM sensitivity is a secondary encoder-only boundary probe, not part of the primary pool or a prompt-conditioned MedSAM evaluation.

5 Main result: same-frozen-checkpoint pools show higher co-failure correlation

Table 1 reports the primary finite-pool audit on five medical segmentation task rows. Same-bundle/release-recipe groups have substantially higher per-case Dice correlation than the cross-bundle reference: $\rho_{\text{same/release}} = 0.834\text{--}0.959$ vs. $\rho_{\text{cross}} = 0.232\text{--}0.697$, with $\Delta_{0.30} = 0.261\text{--}0.638$. Every case-level and hierarchical bootstrap interval excludes zero (smallest hierarchical lower bound 0.174, Kvasir).

Table 1: **Primary finite-pool co-failure audit.** For each filtered functional pool we report average pairwise Pearson correlation of per-case Dice traces for the primary same/release-recipe group and the cross-bundle reference. K is the retained bundle count after the $\text{Dice} \geq 0.30$ functional floor; with four decoder seeds per bundle the trace count is $4K$ and same/cross pair counts follow deterministically from K . $\Delta_{0.30}$ is conditional on the retained functional pool; $\Delta_{\text{no floor}}$ recomputes the same scorer without the floor. CIs are 95% case-level paired bootstrap ($B=2,000$), and the conservative finite-pool interval Δ hier CI[†] is a hierarchical bootstrap over primary grouping and seed clusters; all displayed hierarchical CIs exclude zero. ACDC and BraTS are slice-level rows nested within 20 ACDC patients and 324 BraTS subjects; subject-aggregated gaps are ACDC $\Delta = 0.262$ [0.151, 0.413] and BraTS $\Delta = 0.214$ [0.182, 0.251] (Appendix A11). The BraTS row uses the 11-bundle source pool for the 2D FLAIR $\text{seg}>0$ derivative; ResNet-18 is below the Dice floor, leaving 10 retained bundles. The primary ACDC pool retains CLIP-L/14 because it passes the Dice floor; a secondary 8-FM ACDC diagnostic with a different split (Appendix A13) reports CLIP-L/14 collapsed at Dice 0.007 on that split.

Task row	Unit N	K	$\rho_{\text{same release}}$	ρ_{cross}	$\Delta_{0.30}$	$\Delta_{\text{no floor}}$	$\Delta_{\text{case CI}}$	$\Delta_{\text{hier CI}^\dagger}$
Kvasir polyp	image 120	11	0.958	0.697	0.261	0.261	[0.211,0.320]	[0.174,0.378]
ACDC LV	slice 361 (20 pts)	10	0.959	0.639	0.321	0.355	[0.289,0.355]	[0.211,0.435]
BraTS foreground	slice 3,228 (324 subj)	10	0.938	0.623	0.315	0.362	[0.302,0.329]	[0.224,0.425]
RIGA Cup	fundus 95	9	0.834	0.361	0.473	0.529	[0.395,0.556]	[0.280,0.780]
RIGA Disc	fundus 95	11	0.870	0.232	0.638	0.638	[0.569,0.687]	[0.475,0.831]

What the “same” column actually is. About 77–80% of the same-pool pairs are same-bundle decoder-seed pairs (e.g. Kvasir 66/82=80.5%, RIGA Cup 54/70=77.1%): they share frozen weights identically and therefore correlate near $\bar{r}=0.96\text{--}0.99$ by construction. The only recurring same-family cross-checkpoint pair within a $K=9\text{--}11$ functional pool is DINOv2-B $\times$ DINOv2-S. The headline Δ in Table 1 therefore audits a same-frozen-checkpoint/release-recipe redundancy claim, not a generic same-family-broadly claim. The expanded-grid distinct-checkpoint minus cross-family gap (Table 2, last column; Appendix A30) is 0.080–0.267. The gap is otherwise sign-stable: no-floor gaps remain positive (Kvasir 0.261, ACDC 0.355, BraTS 0.362, RIGA Cup 0.529, RIGA Disc 0.638), and subject aggregation on slice-level rows preserves the sign (ACDC 0.262, BraTS 0.214; Kvasir remains image-level because patient/video IDs are unavailable). Table 2 summarizes four sign-stability diagnostics: a calibration-selected functional pool evaluated on disjoint cases, a binary failure-event readout at the empirical P_{33} threshold, quality-controlled pair regression, and distinct-checkpoint same-family decomposition.

Estimator robustness. A sensitivity summary (Appendix A5) recomputes the primary Pearson gap under a functional-floor sweep, an ICC(A,1) alternative, leave-one-bundle/task-out summaries, and a random family-label permutation stress test. Analytic-floor checks show the gap remains positive at no-floor and stricter-floor settings where the pool is non-empty. Event and ICC checks show the same sign after binarising failures at the empirical P_{33} threshold and after replacing Pearson with ICC(A,1). Quality and residualisation checks show that simple marginal-quality controls and




Figure 1: **Spatial and lift evidence corroborating the per-case Dice gap in Table 1.** Panel (a) shows single-case visualisations on RIGA Cup and BraTS foreground (red=FP, blue=FN, green=GT boundary): four DINOv2-B decoder seeds versus four distinct-bundle members produce visibly different error-overlap structure even when their mean Dice matches. Panel (b) shows pixel-level error-mask Dice on aggregate splits (RIGA Cup/Disc $N=224$, ACDC LV $N=90$, Kvasir $N=300$): same-family 0.80–0.87 versus cross-family 0.32–0.40. Panel (c) shows matched-size joint-failure lift ratios; only the support-validated cells (ACDC/RIGA Cup $\tau \geq 0.9$, RIGA Disc $\tau = 0.95$) carry the visual claim, lower- τ bars are support-limited. Panels (b–c) use non-case-identical aggregate splits and complement the per-case-Dice rows of Table 1 at the pixel and event levels; the artifact distributes summary JSONs rather than raw masks.

held-out item-difficulty residualisation do not absorb the same/cross separation. The random-label permutation is retained only as an appendix stress test because family labels are not exchangeable sampling units; it is not used as a main-text significance claim or population-level p -value.

Spatial and lift readouts. Figure 1 corroborates the Table 1 gap at the pixel and event level. Same-bundle seeds produce spatially overlapping false positives and false negatives (panel a); aggregate pixel-level error-mask Dice (panel b) is 0.80–0.87 same-family versus 0.32–0.40 cross-family on RIGA Cup/Disc ($N=224$), ACDC LV ($N=90$), and Kvasir ($N=300$) (Appendices A3, A8); matched-size joint-failure lift ratios (panel c) follow the same ordering at support-validated thresholds. These aggregate diagnostics use non-case-identical splits and therefore complement rather than replace the per-case Pearson result in Table 1.

Item-difficulty diagnostic. Following the warning that co-failure claims depend on the null model [Jo et al., 2026], we residualise per-case Dice against held-out item-difficulty estimates using a cross-fitted estimator that fits each trace’s item-difficulty regressor on traces outside its own broad upstream family. This yields an *ordinal residualised statistic* $\tilde{\Delta}$ that is no longer constrained to the

Table 2: **Sign-stability diagnostics for the primary rows.** Values are point estimates from distinct diagnostics and are not directly comparable in scale. Cal.-split Δ selects functional bundles on a deterministic calibration subset and evaluates the same/cross gap on disjoint cases. P_{33} fail ϕ gap is the same/cross gap in binary failure-event correlation at the empirical 33rd percentile of model-case Dice after functional-floor filtering. Quality-control β_{same} is the same-group coefficient in a pair-level OLS diagnostic with case-bootstrap refits. The expanded-grid same-family checkpoint gap is the same-family cross-checkpoint diagnostic from Appendix A30, not the Table 1 analytic-pool cross-family stratum. Full definitions and CIs for the Table 2 columns are in Appendices A5–A7 and Appendix A30; the separate item-difficulty diagnostic is in Appendix A26.

Task	Cal.-split Δ	P_{33} fail ϕ gap	QC β_{same}	Checkpoint- family gap
Kvasir polyp	0.250	0.329	0.169	0.080
ACDC LV	0.350	0.386	0.208	0.134
BraTS foreground	0.307	0.413	0.224	0.093
RIGA Cup	0.478	0.453	0.372	0.267
RIGA Disc	0.630	0.578	0.582	0.227

canonical Pearson $[-1, 1]$ scale (released CIs span $[0.782, 1.008]$ across the five rows; Appendix A26 Table 25); only its sign is read against the primary Δ , the magnitudes are not. The sign of $\tilde{\Delta}$ is positive on every primary row.

A limited ViT same-architecture/different-pretraining diagnostic is taken up in Section 6 as a control on architecture-vs-pretraining attribution.

6 Same-backbone controls and scope boundaries

The following diagnostics do not decompose the Table 1 gap into causal shares for architecture, pretraining corpus, checkpoint lineage, wrapper normalisation, and decoder recipe. They answer a narrower evaluation question: if a pool shares one of these upstream factors, does the trace audit still show aligned co-failure that should be reported before claiming independent failure coverage? The practical reporting recommendation does not require identifying which shared factor dominates.

6.1 ViT diagnostic: cross-pretraining correlations sit close to the same-bundle floor

A representation-similarity concern, following Kornblith et al. [2019], is that two ViTs produce similar representations *because they share architecture family*, not because they share pretraining. We hold broad architecture family fixed in a 5-variant ViT pool—DINOv2-B (SSL on LVD-142M), BiomedCLIP (CLIP on biomed image-text), MAE-B (ImageNet MAE), DeiT-B (ImageNet supervised), OpenAI CLIP-B—trained with the identical frozen 1×1 -conv probe on RIGA Cup and ACDC LV (Table 3). ViT cross-pretraining correlations sit between the same-bundle seed floor and the primary cross-family reference; the ordering does not identify which factor is causal but does show that holding architecture family fixed leaves a substantial gap to the cross-family reference.

Table 3: **Limited ViT same-architecture/different-pretraining diagnostic.** The table compares same-bundle seed-pair correlations with cross-pretraining correlations among ViT-style encoders. It is a two-task diagnostic, not a causal decomposition of architecture and pretraining. Seed floor here is the mean within-FM cross-seed Pearson restricted to the functional 5-ViT subset; this differs from Table 1’s ρ_{same} , which averages over the full 9/10-bundle post-floor pool including CNN bundles whose lower within-FM correlation pulls the mean down.

Dataset	ViT subset	Func. ViTs	Seed floor	ViT cross- pretrain	Cross ref.	Floor – cross-pretrain
RIGA Cup	all 5 ViTs	5/5	0.968	0.680	0.361	0.288
ACDC LV	CLIP-B dropped [†]	4/5	0.988	0.807	0.639	0.182

Values are from the architecture-confound diagnostic and primary table summary. [†]The all-5 ACDC same-architecture-family ViT cross-pretraining value is 0.631 because CLIP-B is nonfunctional on ACDC (mean Dice 0.205); the functional-filtered readout is 0.807. ViT $\times$ CNN comparisons are reported only as appendix diagnostics and are not used as main-table numeric evidence.

Table 4: **Same-backbone perturbation diagnostics.** These comparisons vary decoder capacity, skip structure, or MC dropout under limited grids. They test whether simple same-backbone perturbations remove high correlation; they are not dose-matched causal controls.

Diagnostic	Released source	Main readout
Head-capacity sweep, RIGA Cup	Released App. A31.1	traces, 4 heads, 769–3.94M params; Dice 0.725–0.848; cross-head same-bundle $\bar{r}=0.690$ (range 0.510–0.896)
UNet-skip decoder replication	Released App. A16	summary, 41 cells (9 backbones $\times$ 5 tasks); Dice 0.774–0.966; seed-pair $r=0.426$ –0.979
MC dropout, DINOv2-B/RIGA Cup	Released App. A31.7	traces, MC $\bar{r}=0.947, 0.925, 0.893$ for $p=0.1, 0.2, 0.4$

Readings: head-capacity and UNet-skip stay above the primary RIGA cross-family $\bar{r}=0.361$; MC dropout is a weak decorrelation lever (largest tested drop ≈ 0.10).

This diagnostic alone does not isolate pretraining-bundle reuse from broad architecture-family imprint; in particular, ViT cross-pretraining correlations sit substantially closer to the same-bundle floor (gap 0.288 on RIGA Cup, 0.182 on ACDC LV) than to the primary cross-family reference, consistent with shared architecture and shared 1×1 -conv readout contributing materially to the measured dependence.

The released CNN random-init diagnostic (Appendix A31.2) shows positive same-vs-cross structure even without shared pretrained weights ($\Delta = 0.086$ on ACDC LV, $\Delta = 0.121$ on RIGA Cup); a more substantive 3-architecture from-scratch replication (Appendix A13) reaches the same conclusion. The pretrained-bundle effect therefore operates on top of, not in isolation from, an architecture-and-recipe baseline.

6.2 Decoder-capacity, stochastic, and medical-domain probes

(C1) Head architecture/capacity. The head-capacity and UNet-skip decoder perturbation diagnostics change accuracy and reduce some pairwise correlations, but their same-backbone correlations remain above the primary RIGA cross-family reference. Thus the tested decoder-capacity changes did not remove same-backbone co-failure in these diagnostics.

(C2) Stochastic/adaptation levers. In the RIGA Cup diagnostic, MC averaging modestly lowers within-bundle correlation but remains far above the cross-family reference. Bootstrap-bagging and fold-reslicing are appendix diagnostics with different sampling surfaces; the fold-reslicing row is a single-seed split-stability check, not a within-fold versus cross-fold same-backbone estimator. Full-fine-tuning and LoRA probes without trace artifacts are described only to delimit scope and are not used to support the main claims.

(C3) Medical-domain encoder-only probe. A MedSAM image-encoder-only sensitivity is described in Appendix A10 as a summary diagnostic. The probe strips MedSAM’s prompt encoder and mask decoder, so it does not evaluate prompt-conditioned MedSAM. RETFound is documented only as a gated-checkpoint path; without access to the gated checkpoint, we do not report RETFound rows.

(C4) Full-pipeline scope boundary. nnU-Net v2 2D/3D complements give task-level different-fold gaps of $\Delta \approx 0.034$ –0.047 on RIGA Cup and $\Delta \approx 0.043$ on ACDC 3D; a case-identical RIGA Cup nnU-Net check gives $\Delta=0.0079$ [0.002, 0.039] at 100 epochs (correlations 0.984/0.976) and a matched 1000-epoch run leaves the gap essentially unchanged ($\Delta=0.0075$ [0.003, 0.034], 0.987/0.980). The full-pipeline gap is therefore much smaller than the primary frozen-feature gap; both diagnostics are descriptive scope checks rather than substitutes for the primary audit. Appendix A21 consolidates all same-backbone, adaptation, and scope levers in one map; Appendix A24 reports a 3D segmentation-architecture pilot. Appendix Table 38 encodes the evidence-tier structure: Table 1 is the primary frozen-encoder audit, and estimator, event, architecture/pretraining, same-backbone, full-pipeline, and held-out rows are diagnostic or scope evidence.

7 Discussion

Practical reading. For an operational pool of M members with mean failure correlation ρ_{fail} , report the equicorrelation effective pool size $M_{\text{eff}} = M/[1 + (M - 1)\rho_{\text{fail}}]$ (Appendix A27; here M

denotes operational pool size, distinct from the trace count $4K$ in Table 1) as a threshold-specific dependence readout before treating same-bundle seeds as independent evidence, not as a standalone pool-selection objective. In our released rows, diverse $M=4$ pools yield mean M_{eff} from 2.05 (Kvasir) to 2.97 (RIGA Disc) under the paper’s equicorrelation readout at $\tau = 0.85$, rather than 4; using the maximum observed diverse-pool ρ_{fail} gives conservative lower readouts of 1.30–1.69 for the most redundant valid 4-bundle tuples. Same-bundle seed pools yield M_{eff} in the 1.1–1.2 range under the same threshold. The readout should always be paired with marginal Dice and matched-quality comparisons before any pool-selection decision (Appendix A27). For procurement-style decisions where a clinical-imaging team is choosing between an N -bundle FM pool and a smaller curated pool, an in-domain co-failure audit on a small validation set should precede pool-size-based diversity claims.

Implications for ensemble methodology. Standard nnU-Net 5-fold validation under shared encoders is a within-bundle estimator and inherits the dependence structure measured here. Diversity *readouts* (Δ , joint-failure lift, M_{eff}) are useful as audit summaries, but on our held-out tasks selecting pools by maximum M_{eff} was *not* better than random (Appendix A19, rule R3); the readouts are diagnostic, not pool-selection objectives. Appendix A28 gives a random-effects representation that motivates why same-backbone fold/seed surfaces concentrate above cross-encoder references.

Limitations. The evidence is retrospective and benchmark-based: this study audits the reporting practice around frozen-encoder pool failure evidence, not clinical deployment, and we do not claim population-level inference over future encoders, datasets, or recipes. The BraTS foreground row is a 2D derivative of the 2024 GLI release, not the official 3D challenge target. Table 1 is the confirmatory evidence surface; threshold, lift, and held-out rows are exploratory diagnostics. Per-case Dice Pearson is a continuous-monotone alignment statistic, so two members can have high Pearson Dice yet still disagree on which low-quality cases fail; pixel error-mask Dice and joint-failure lift in Section 5 supplement it at the failure-event level, and Geirhos-style binarised error consistency would be a complementary discrete estimand. The cross-bundle stratum draws from at most 8 broad family clusters across 11 bundles, so few-cluster bias is unavoidable; we report case-level and hierarchical bootstrap CIs jointly. The frozen 1×1 -conv decoder is deliberately minimal so the encoder dominates the readout; absolute correlation magnitudes are upper bounds for richer decoders, and the M3 head-capacity and M4 UNet-skip diagnostics show more capable heads reduce some pairwise correlations but do not eliminate the same-vs-cross gap. All ten sign tests (five rows $\times$ two CI methods) independently exclude zero, with the smallest hierarchical lower bound at 0.174 (Kvasir); we report directional gaps per row rather than a joint significance statement. Appendices A18, A19, and A15 extend these caveats.

What the evaluation supports. The E&D takeaway is a reporting requirement: when a medical segmentation pool reuses pretrained encoder checkpoints, input statistics, and decoder recipes, aggregate Dice should be accompanied by aligned case-level co-failure correlations before the pool is credited with independent failure coverage. The result is a finite audit of evaluation evidence, not a clinical safety claim or a causal decomposition of pretraining and architecture.

8 Artifact and reproducibility scope

The anonymized artifact is available through the supplementary submission and anonymous code URL https://anonymous.4open.science/r/NIPS_A_ACF-86CD: primary per-case Dice JSONs, summary tables, code, metadata/manifests, hashes, environment files, a data card, selected appendix summary JSONs, and Croissant metadata with the required Responsible AI fields. It does not introduce a new raw dataset; released JSONs are derived trace artifacts that retain upstream de-identified subject/patient/slice IDs and deterministic hashed RIGA alignment IDs solely for case alignment and subject-level resampling, and are not intended for individual-level analysis, subgroup profiling, re-identification, inference of protected attributes, clinical diagnoses, or patient outcomes. The bundle excludes raw medical images, raw masks, direct identifiers, full prediction caches, and held-out raw prediction arrays; raw data and checkpoints remain governed by upstream licenses and DUAs. Appendix Table 38 ties each main-text evidence class to its released JSON or recompute script.

References

- Taiga Abe, E. Kelly Buchanan, Geoff Pleiss, and John Patrick Cunningham. Pathologies of predictive diversity in deep ensembles. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=TQfQUksaC8>.
- Ahmed Almazroa. Retinal fundus images for glaucoma analysis: The riga dataset. University of Michigan - Deep Blue Data, 2018. URL <https://doi.org/10.7302/Z23R0R29>.
- Ahmed Almazroa, Sami Alodhayb, Essameldin Osman, Eslam Ramadan, Mohammed Hummadi, Mohammed Dlaim, Muhannad Alkatee, Kaamran Raahemifar, and Vasudevan Lakshminarayanan. Retinal fundus images for glaucoma analysis: The riga dataset. In *Medical Imaging 2018: Imaging Informatics for Healthcare, Research, and Applications*, volume 10579, page 105790B. SPIE, 2018. doi: 10.1117/12.2293584.
- Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *arXiv:2107.07511*, 2021.
- Michela Antonelli, Annika Reinke, Spyridon Bakas, et al. The medical segmentation decathlon. *Nature Communications*, 13:4128, 2022. doi: 10.1038/s41467-022-30695-9. URL <http://medicaldecathlon.com/>.
- Anwai Archit, Luca Freckmann, and Constantin Pape. Medicosam: Robust improvement of sam for medical imaging. *arXiv:2501.11734*, 2025.
- Mohammed Baharoon et al. Evaluating general purpose vision foundation models for medical image analysis: An experimental study of dinov2 on radiology benchmarks. *arXiv:2312.02366*, 2023.
- Olivier Bernard, Alain Lalonde, Clement Zotti, et al. Deep learning techniques for automatic mri cardiac multi-structures segmentation and diagnosis: Is the problem solved? *IEEE Transactions on Medical Imaging*, 37(11):2514–2525, 2018. doi: 10.1109/TMI.2018.2837502.
- Rishi Bommasani, Drew A Hudson, Ehsan Adeli, et al. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021.
- Rishi Bommasani, Kathleen A. Creel, Ananya Kumar, Dan Jurafsky, and Percy Liang. Picking on the same person: Does algorithmic monoculture lead to outcome homogenization? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- BraTS 2024 Organizers. Brats 2024 glioma challenge data. Synapse: syn53708249, 2024. URL <https://www.synapse.org/Synapse:syn53708249/wiki/628815>.
- Gabriele Campanella, Shengjia Chen, Manbir Singh, et al. A clinical benchmark of public self-supervised pathology foundation models. *Nature Communications*, 16:3640, 2025. doi: 10.1038/s41467-025-58796-1.
- Junlong Cheng, Jin Ye, Zhongying Deng, Jianpin Chen, Tianbin Li, Haoyu Wang, Yanzhou Su, Ziyang Huang, Jilong Chen, Lei Jiang, Hui Sun, Junjun He, Shaoting Zhang, Min Zhu, and Yu Qiao. Sam-med2d. *arXiv:2308.16184*, 2023.
- Noel Codella, Veronica Rotemberg, Philipp Tschandl, M. Emre Celebi, Stephen Dusza, David Gutman, Brian Helba, Aadi Kallou, Konstantinos Liopyris, Michael Marchetti, Harald Kittler, and Allan Halpern. Skin lesion analysis toward melanoma detection 2018: A challenge hosted by the international skin imaging collaboration (isic). *arXiv:1902.03368*, 2019.
- Edwin de Jong, Eric Marcus, and Jonas Teuwen. Current pathology foundation models are unrobust to medical center differences. *arXiv:2501.18055*, 2025.
- Maria Correia de Verdier, Rachit Saluja, Louis Gagnon, Dominic LaBella, Ujjwal Baid, Nourel Hoda Tahon, Martha Foltyn-Dumitru, Jikai Zhang, Maram Alafif, Saif Baig, et al. The 2024 brain tumor segmentation (brats) challenge: Glioma segmentation on post-treatment mri. *arXiv preprint arXiv:2405.18368*, 2024. doi: 10.48550/arXiv.2405.18368.
- André Ferreira, Naida Solak, Jianning Li, Philipp Dammann, Jens Kleesiek, Victor Alves, and Jan Egger. How we won BraTS 2023 adult glioma challenge? just faking it! Enhanced synthetic data augmentation and model ensemble for brain tumour segmentation. *arXiv preprint arXiv:2402.17317*, 2024.
- Stanislav Fort, Huiyi Hu, and Balaji Lakshminarayanan. Deep ensembles: A loss landscape perspective. *arXiv preprint arXiv:1912.02757*, 2019.
- Yarin Gal and Zoubin Ghahramani. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. In *ICML*, 2016.

394 M A Ganaie, Minghui Hu, A K Malik, M Tanveer, and P N Suganthan. Ensemble deep learning: A review.
395 *Engineering Applications of Artificial Intelligence*, 115, 2022.

396 Robert Geirhos, Kristof Meding, and Felix A. Wichmann. Beyond accuracy: Quantifying trial-by-trial be-
397 haviour of cnns and humans by measuring error consistency. In *Advances in Neural Information Processing*
398 *Systems*, 2020.

399 Mariana-Iuliana Georgescu, Radu Tudor Ionescu, and Andreea-Iuliana Miron. Diversity-promoting ensemble
400 for medical image segmentation. In *Proceedings of the 38th ACM/SIGAPP Symposium on Applied Comput-*
401 *ing*, pages 599–606, 2023. doi: 10.1145/3555776.3577822.

402 Shashwat Goel, Joschka Struber, Ilze Amanda Auzina, Karuna K. Chandra, Ponnuram Kumaraguru, Douwe
403 Kiela, Ameya Prabhu, Matthias Bethge, and Jonas Geiping. Great models think alike and this undermines
404 AI oversight. arXiv:2502.04313, 2025.

405 Raphael Gontijo-Lopes, Yann Dauphin, and Ekin Dogus Cubuk. No one representation to rule them all: Over-
406 lapping features of training methods. In *International Conference on Learning Representations (ICLR)*,
407 2022.

408 Mila Gorecki and Moritz Hardt. Monoculture or multiplicity: Which is it? In *NeurIPS*, 2025. URL <https://openreview.net/forum?id=D05LtJc80w>.

410 Hanxue Gu, Haoyu Dong, Jichen Yang, and Maciej A. Mazurowski. How to build the best medical image
411 segmentation algorithm using foundation models: A comprehensive empirical study with segment anything
412 model. *MELBA*, 2025. arXiv:2404.09957.

413 Yi Gu, Yi Lin, Kwang-Ting Cheng, and Hao Chen. Revisiting deep ensemble uncertainty for enhanced medical
414 anomaly detection. In *MICCAI*, 2024.

415 Yufan He et al. Vista3d: A unified segmentation foundation model for 3d medical imaging. In *CVPR*, 2025.

416 Yuhao Huang, Xin Yang, Lian Liu, Han Zhou, Ao Chang, Xinrui Zhou, Rusi Chen, Junxuan Yu, Jiongquan
417 Chen, Chaoyu Chen, Sijing Liu, Haozhe Chi, Xindi Hu, Kejuan Yue, Lei Li, Vicente Grau, Deng-Ping Fan,
418 Fajin Dong, and Dong Ni. Segment anything model for medical images? *Medical Image Analysis*, 92:
419 103061, 2024. doi: 10.1016/j.media.2023.103061.

420 Fabian Isensee, Paul F Jaeger, Simon A A Kohl, Jens Petersen, and Klaus H Maier-Hein. nnU-Net: a self-
421 configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18:203–211,
422 2021.

423 Fabian Isensee et al. nnu-net revisited: A call for rigorous validation in 3d medical image segmentation. In
424 *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, 2024.

425 Debesh Jha, Pia H. Smedsrud, Michael A. Riegler, Pål Halvorsen, Thomas de Lange, Dag Johansen, and
426 Håvard D. Johansen. Kvasir-SEG: A segmented polyp dataset. In *International Conference on Multimedia*
427 *Modeling (MMM)*, 2020.

428 Liwei Jiang et al. Artificial hivemind: The open-ended homogeneity of language models (and beyond).
429 arXiv:2510.22954, 2025.

430 Nathanael Jo, Nikhil Garg, and Manish Raghavan. The subjectivity of monoculture. arXiv:2602.24086, 2026.

431 Alain Jungo and Mauricio Reyes. Assessing reliability and challenges of uncertainty estimations for medical
432 image segmentation. In *MICCAI*, 2019.

433 Tommie Kerssies et al. How to benchmark vision foundation models for semantic segmentation? In *CVPR*
434 *Workshop*, 2024.

435 Elliot Kim, Avi Garg, Kenny Peng, and Nikhil Garg. Correlated errors in large language models. In *Interna-*
436 *tional Conference on Machine Learning (ICML)*, 2025a. arXiv:2506.07962.

437 Yunhe Kim et al. Show and segment: Universal medical image segmentation via in-context learning.
438 arXiv:2503.19359, 2025b.

439 Alexander Kirillov, Eric Mintun, Nikhila Ravi, et al. Segment anything. In *2023 IEEE/CVF International*
440 *Conference on Computer Vision (ICCV)*, pages 3992–4003, 2023. doi: 10.1109/ICCV51070.2023.00371.

441 Max Klabunde, Tobias Schumacher, Markus Strohmaier, and Florian Lemmerich. Similarity of neural network
442 models: A survey of functional and representational measures. *Transactions on Machine Learning Research*,
443 2024.

444 Jon Kleinberg and Manish Raghavan. Algorithmic monoculture and social welfare. *Proceedings of the National*
445 *Academy of Sciences*, 118(22), 2021.

446 Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network
447 representations revisited. In *ICML*, 2019.

448 Ludmila I Kuncheva and Christopher J Whitaker. Measures of diversity in classifier ensembles and their rela-
449 tionship with the ensemble accuracy. *Machine Learning*, 51(2):181–207, 2003.

450 Zeki Kuş and Musa Aydin. Medsegbench: A comprehensive benchmark for medical image segmentation in
451 diverse data modalities. *Scientific Data*, 11(1):1283, 2024. doi: 10.1038/s41597-024-04159-2.

452 Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty
453 estimation using deep ensembles. In *NeurIPS*, 2017.

454 Yuanyuan Li, Neeraj Sarna, and Yang Lin. Quantifying correlations of machine learning models. *arXiv preprint*
455 *arXiv:2502.03937*, 2025.

456 Xiangde Luo, Xiyue Wang, Feyisope Eweje, Xiaoming Zhang, et al. Ensemble learning of foundation models
457 for precision oncology. *arXiv:2508.16085*, 2025. URL <https://arxiv.org/abs/2508.16085>.

458 Jun Ma, Yuting He, Feifei Li, Lin Han, Chenyu You, and Bo Wang. Segment anything in medical images.
459 *Nature Communications*, 15, 2024.

460 Jun Ma, Zongxin Yang, Sumin Kim, Bihui Chen, Mohammed Baharoon, Adibvafa Fallahpour, Reza Asakereh,
461 Hongwei Lyu, and Bo Wang. Medsam2: Segment anything in 3d medical images and videos. *arXiv preprint*
462 *arXiv:2504.03600*, 2025.

463 Lena Maier-Hein et al. Metrics reloaded: Recommendations for image analysis validation. *Nature Methods*,
464 21:195–212, 2024. doi: 10.1038/s41592-023-02151-z.

465 Kenneth O. McGraw and S. P. Wong. Forming inferences about some intraclass correlation coefficients. *Psy-*
466 *chological Methods*, 1(1):30–46, 1996.

467 Alireza Mehrtash, William M Wells, Clare M Tempany, Purang Abolmaesumi, and Tina Kapur. Confidence
468 calibration and predictive uncertainty estimation for deep medical image segmentation. *IEEE Transactions*
469 *on Medical Imaging*, 39(12):3868–3878, 2020.

470 Dominik J. Mühlematter, Michelle Halbheer, Alexander Becker, Dominik Narnhofer, Helge Aasen, Konrad
471 Schindler, and Mehmet Ozgur Turkoglu. Lora-ensemble: Efficient uncertainty modelling for self-attention
472 networks. *Transactions on Machine Learning Research*, 2026. URL [https://openreview.net/forum?](https://openreview.net/forum?id=yhXXm0MpSQ)
473 [id=yhXXm0MpSQ](https://openreview.net/forum?id=yhXXm0MpSQ).

474 Basil Mustafa, Carlos Riquelme, Joan Puigcerver, André Susano Pinto, Daniel Keysers, and Neil Houlsby.
475 Deep ensembles for low-data transfer learning. *arXiv preprint arXiv:2010.06866*, 2020.

476 Peter Neidlinger, Omar S. M. El Nahhas, Hannah Sophie Muti, et al. Benchmarking foundation models as
477 feature extractors for weakly-supervised computational pathology. *Nature Biomedical Engineering*, 2025.
478 doi: 10.1038/s41551-025-01516-3.

479 Behnam Neyshabur, Hanie Sedghi, and Chiyuan Zhang. What is being transferred in transfer learning? In
480 *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

481 Abhijit Guha Roy, Sailesh Conjeti, Nassir Navab, and Christian Wachinger. Bayesian QuickNAT: Model
482 uncertainty in deep whole-brain segmentation for structure-wise quality control. In *NeuroImage*, volume
483 195, pages 11–22, 2019.

484 Connor Toups, Rishi Bommasani, Kathleen Creel, Sarah H. Bana, Dan Jurafsky, and Percy Liang. Ecosystem-
485 level analysis of deployed machine learning reveals homogeneous outcomes. In *Advances in Neural Infor-*
486 *mation Processing Systems (NeurIPS)*, 2023.

487 Danny Wood, Tingting Mu, Andrew M. Webb, Henry Reeve, Mikel Lujan, and Gavin Brown. A unified theory
488 of diversity in ensemble learning. *arXiv:2301.03962*, 2023.

489 Mitchell Wortsman, Gabriel Ilharco, Samir Ya Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes, Ari S. Morcos,
490 Hongseok Namkoong, Ali Farhadi, Yair Carmon, Simon Kornblith, and Ludwig Schmidt. Model soups:
491 Averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In
492 *International Conference on Machine Learning (ICML)*, 2022.

- 493 R. Patrick Xian, Noah R. Baker, Tom David, Qiming Cui, A. Jay Holmgren, Stefan Bauer, Madhumita Sushil,
494 et al. Robustness tests for biomedical foundation models should tailor to specifications. *npj Digital Medicine*,
495 8:557, 2025. doi: 10.1038/s41746-025-01926-2.
- 496 Yuncheng Yang, Meng Wei, Junjun He, Jie Yang, Jin Ye, and Yun Gu. Pick the best pre-trained model: Towards
497 transferability estimation for medical image segmentation. In *Medical Image Computing and Computer*
498 *Assisted Intervention (MICCAI)*, 2023.
- 499 Maximilian Zenk, Klaus H. Maier-Hein, Jens Petersen, et al. Comparative benchmarking of failure detection
500 methods in medical image segmentation: Unveiling the role of confidence aggregation. *Medical Image*
501 *Analysis*, 101:103392, 2025. doi: 10.1016/j.media.2024.103392.
- 502 Yukun Zhou, Mark A. Chia, Siegfried K. Wagner, et al. A foundation model for generalizable disease detection
503 from retinal images. *Nature*, 622:156–163, 2023.
- 504 Jiayuan Zhu et al. Medical sam 2: Segment medical images as video via segment anything model 2.
505 arXiv:2408.00874, 2024.

506 A1 Appendix roadmap and evidence inventory

507 The appendix is organised by evidentiary role rather than by experiment chronology. It starts with
508 diagnostic readouts, then documents primary estimators, uncertainty, task-level sources, and estima-
509 tor sensitivity. Middle sections collect same-backbone perturbations, segmentation-pipeline scope
510 checks, distribution-shift and held-out stress tests, operational metrics, and M_{eff} . Later sections
511 collect mechanism diagnostics and additional scope tables, and the appendix ends with the release
512 manifest and artifact metadata. Secondary rows are retained only when they qualify or explain a
513 paper claim, and their captions state when they use non-primary pools or splits. Unless explicitly
514 marked primary or sourced from the primary `results/_merged` summaries, appendix values are
515 diagnostic rather than primary support for Table 1.

516 **CLIP-L notation.** Throughout the appendix, CLIP-L/14 refers to the preserved OpenCLIP
517 ViT-L-14/openai dense-token wrapper used to generate the released traces. It is not a claim of
518 OpenAI clip.load parity; a strict QuickGELU-parity CLIP-L loader would be a separate sensitiv-
519 ity with a distinct FM identifier.

520 A2 Diagnostic readouts and scope

521 This section records diagnostic readouts that clarify scope without adding new primary claims.

522 **RIGA annotation source.** Table 1 uses the six-rater majority consensus as ground truth for both
523 RIGA Cup and RIGA Disc, matching the released RIGA per-case Dice traces. RIGA also distributes
524 individual ophthalmologist masks, but individual-rater replay artifacts are not included in the supple-
525 mentary artifact. The consensus Cup and Disc rows are therefore the only RIGA annotation-source
526 evidence used by the manuscript.

527 **Probability-map diagnostics.** The artifact releases per-case Dice traces and selected aggregate
528 pixel/lift diagnostics, but not raw probability maps. We therefore do not use mean-confidence,
529 probability-calibration, or failure-flagging AUROC values as evidence for any paper claim.

530 A3 Pixel-level error-overlap breakdown

531 $FP = (\text{pred foreground, GT background})$; $FN = (\text{pred background, GT foreground})$; boundary-band
532 restricts error masks to pixels within 5 pixels of the GT edge (trimap). All metrics are pair-wise
533 Dice between the two models’ error masks (or empty-mask-excluded Jaccard). On RIGA Cup,
534 RIGA Disc, and ACDC LV binary, the similar-magnitude gaps across FP, FN, and boundary variants
535 indicate the spatial co-failure is not driven by a single error type on those three tasks (Kvasir and
536 BraTS not tabulated in this split).

Table 5: **Appendix guide.** Blocks describe the actual order of sections below; cross-references point to the detailed tables rather than repeating them in the main paper.

Block	Sections (in order)	Evidence role
Estimators and provenance	App. A2, A3, A4, A5, A6	Diagnostic readouts, pixel overlap, case/hierarchical bootstrap, functional-floor/ICC/permutation sensitivity, pair-quality OLS.
Operational metrics	App. A7, A8, A9, A10	Failure-event τ -sweep, conditional recovery, Kvasir secondary split, matched- M lift, MedSAM probe details.
Same-backbone controls	App. A11, A12, A13, A14	Subject clustering, BraTS sparse caveat, ACDC architecture caveat / random-init / LoRA scope, bootstrap-bagging.
Adaptation & distribution shift	App. A15, A16, A17, A18	RIGA cross-vendor, UNet-skip decoder, nnU-Net fold reslicing + case-identical, segmentation-recipe scope (loss/aug, length, dimensionality).
Held-out and meta-analyses	App. A19, A21, A22, A23, A24, A25	Held-out 5-task stress test, lever index, leave-one-family-out, lift-ceiling support, 3D pilot, matched-recipe adaptation scope.
Mechanism diagnostics	App. A26, A27, A28, A29, A30	Item-difficulty residualisation (cross-fitted + holdout), M_{eff} , random-effects representation, ViT cross-pretraining, cross-checkpoint same-family.
Additional scope diagnostics	App. A31, A31.1, A31.2, A31.3, A31.4, A31.5, A31.6, A31.7	Decoder capacity, random-init CNN, loss/aug sensitivity, fold CV, BraTS train-size, BraTS multimodal, MC dropout.
Artifact record	App. A33	Bundle manifest, claim-to-artifact map, compute scope, trace inventory.

Table 6: **Pixel-level error-overlap diagnostic on selected secondary rows.** Values are same/cross gaps in error-mask overlap on aggregate diagnostic splits, not primary case-level evidence. For error-Dice, no bootstrap replicate out of $B=500$ produced a non-positive gap on any task cell, giving the finite- B bound $\Pr[\text{gap} \leq 0] < 1/501$; FP/FN/boundary rows are point-estimate breakdowns.

Metric	RIGA Cup	RIGA Disc	ACDC LV
Error-Dice gap	0.485	0.481	0.485
Error-Jaccard gap	0.515	0.481	0.521
FP-only error-Dice gap	0.453	0.494	0.566
FN-only error-Dice gap	0.527	0.464	0.449
Boundary-band (5px)	0.430	0.439	0.433

A4 Bootstrap sensitivity of the gap

Table 1 reports 95% CIs from a three-level nested hierarchical bootstrap that resamples family labels, seeds within each family, and cases ($B=2000$). We additionally compute a *case-only* bootstrap that fixes the family/seed structure and resamples the N test cases with replacement ($B=2000$). The two procedures target different sampling dimensions: the hierarchical CI jointly perturbs the finite FM/seed/case surface, whereas the case-only CI asks “what if we had a different test set” conditional on the evaluated pool. We report both as sensitivity evidence for the same finite audit rather than as population-level guarantees.

Case-only bootstrap intervals match the Case CI column of Table 1 exactly (the case-only resampling uses the same fixed family/seed structure as the per-row case CI), so we do not retabulate them here. The remainder of this section reports BraTS pool reconciliation and the slice-vs-subject sampling caveats.

BraTS foreground pool reconciliation. The 5-FM, 8-FM, and 8-FM-alt-estimator BraTS values cited elsewhere differ because they use disjoint analytic pools (5-FM: a subset of the 11-bundle source pool; 8-FM: a completed-only subset; alt-estimator: scoring with a different family-grouping convention). The principled headline is the post-floor 10-FM value because it uses the Dice ≥ 0.30 functional floor and the Table 1 same/cross definition consistent with every other primary row. The

primary BraTS foreground result is $\Delta=0.315$ [0.302, 0.329] (case-level CI) from the 11-bundle source pool, with 10 bundles retained after the $\text{Dice} \geq 0.30$ floor (BiomedCLIP, CLIP-B/16, CLIP-L/14, ConvNeXt-T, DeiT-B/16, DINOv2-B/S, EfficientNet-B0, MAE-B/16, ResNet-50; ResNet-18 is the only bundle below the floor). Source: `results/_merged/paper_table.json`. Subset diagnostics without released aggregate files are omitted from the released numeric evidence. The value reported throughout the main text and abstract is therefore the Table 1 value.

All gap CIs strictly exclude zero under the case-level bootstrap, matching the family/seed-level CIs in the main text.

Per-rater RIGA Cup sensitivity (legacy non-replayable pool). A non-replayable legacy sensitivity pool (containing RETFound traces not in the released artifact) repeated the symmetric multi-seed analysis using each of RIGA’s six individual rater masks; per-rater Pearson within-cross gaps were 0.46–0.53 across all six raters, with the range across raters an order of magnitude smaller than the within-vs-cross gap itself. The released Table 1 RIGA Cup row remains the bundle-replayable evidence.

A5 Estimator-robustness and negative-control checks

The primary row uses Pearson correlation and a $\text{Dice} \geq 0.30$ per-FM functional floor. To make this choice auditable, the released bundle includes a sensitivity artifact, `results/_merged/diagnostics/audit_sensitivity.json`, generated from the same per-case Dice JSONs by `code/scripts/compute_audit_sensitivity.py`. It recomputes the within-cross gap under a floor sweep, an ICC(A,1) alternative, leave-one-FM-out summaries, leave-one-task-out meta summaries, and a random family-label permutation control.

Table 7: **Estimator sensitivity on primary traces.** Pearson values are the Table 1 point estimates at the $\text{Dice} \geq 0.30$ floor. ICC(A,1) is an alternate agreement statistic on the same retained members. The floor range reports the minimum/maximum Pearson gap over $\{0, 0.20, 0.25, 0.30, 0.35, 0.40\}$ where the pool is non-empty. The permutation note reports a non-exchangeable random family-label stress test over completed members, using $(1 + \#\{\Delta_{\text{null}} \geq \Delta_{\text{obs}}\})/(1001)$; this is not interpreted as a p -value. Retained counts at every floor are in `audit_sensitivity.json`.

Task	Members at 0.30	Pearson gap	ICC(A,1) gap	Floor-sweep gap range
Kvasir	44	0.261	0.333	0.261 ^b
ACDC LV	40	0.321	0.401	0.242–0.355
BraTS foreground	40	0.315	0.404	0.315–0.362
RIGA Cup	36	0.473	0.539	0.384–0.529
RIGA Disc	44	0.638	0.716	0.638 ^b

All five rows have observed-gap-exceeded rank fraction 0.001 at the 0.30 floor. ^b Value is constant across all valid floors in $\{0, 0.20, 0.25, 0.30, 0.35, 0.40\}$.

The five-task mean Pearson gap at the paper floor is 0.401 (median 0.321, range 0.261–0.638). Leave-one-task-out means stay in 0.343–0.437; the smallest value occurs when the largest-gap RIGA Disc row is held out. These checks do not add new cases or new model training. They instead reduce the risk that the primary sign is an artifact of the exact functional floor, the Pearson statistic, or arbitrary family labels.

CLIP-L/14 inclusion sensitivity (ACDC LV). Because CLIP-L/14 passes the $\text{Dice} \geq 0.30$ floor on the primary ACDC split but collapses on the secondary 8-FM split (Dice 0.007; Appendix A13), we recompute the primary ACDC row with all `clip_vitl14` traces removed before any other step. With CLIP-L/14 excluded the source pool drops to 10 bundles ($K=10$ pre-floor, $K=9$ post-floor): $\bar{\rho}_{\text{same/release}}=0.956$, $\bar{\rho}_{\text{cross}}=0.625$, $\Delta_{0.30}=0.331$ [case 95% CI 0.300, 0.367], and $\Delta_{\text{no floor}}=0.367$. The sign and magnitude track the printed Table 1 primary ($K=10$, $\bar{\rho}_{\text{same}}=0.959$, $\bar{\rho}_{\text{cross}}=0.639$, $\Delta_{0.30}=0.321$, $\Delta_{\text{no floor}}=0.355$); the case intervals overlap and exclude zero. CLIP-L/14 inclusion is therefore not the driver of the ACDC primary signal. Source: `results/_merged/diagnostics/acdc_clipl14_excluded.json`.

Calibration-selected functional pools. The primary table defines a functional pool using all released evaluation traces. To check whether the sign depends on this convention, the scorer also

591 makes a deterministic split of each task’s aligned units, selects bundles by mean Dice on the cal-
592 ibration subset, and evaluates the same/cross gap on disjoint cases. This is a leakage-sensitivity
593 analysis rather than a prospective validation protocol, but it addresses the specific concern that the
594 functional-floor sign is induced by evaluating the floor and the gap on exactly the same cases.

Table 8: **Calibration-selected functional-pool sensitivity.** Functional bundles are selected on the calibration subset and evaluated on disjoint cases under the Table 1 same/cross grouping. Unit-agg. Δ is reported only for nested slice datasets after patient/subject aggregation. ACDC can select 11 bundles here because selection is calibration-only, whereas Table 1 applies the retrospective all-evaluation functional floor. CIs are case-bootstrap intervals on the evaluation subset. Source: `results/_merged/calibration_split/*.json`, regenerated by `code/scripts/compute_all.py`.

Task	Cal./eval units	Selected bundles	Eval. Δ [95% CI]	Unit-agg. Δ
Kvasir	42/78 images	11	0.250 [0.196, 0.322]	–
ACDC LV	7/13 patients	11	0.350 [0.313, 0.392]	0.206
BraTS foreground	113/211 subjects	10	0.307 [0.292, 0.323]	0.205
RIGA Cup	33/62 images	9	0.478 [0.369, 0.599]	–
RIGA Disc	33/62 images	11	0.630 [0.534, 0.695]	–

595 A6 Quality-controlled pairwise-correlation diagnostic

596 A possible alternative explanation is that same-FM pairs have higher correlation merely because
597 they have matched marginal quality. We therefore add a CPU-only diagnostic that uses the same
598 released per-case Dice traces retained by Table 1 and fits a pair-level OLS model:

$$r_{ab} \sim I_{\text{same}} + \bar{D}_{ab} + |\Delta \bar{D}_{ab}| + \bar{V}_{ab} + |\Delta V_{ab}|.$$

599 Here $\bar{D}_{ab}$ and $\bar{V}_{ab}$ are the pair mean Dice and average per-case Dice variance, with absolute
600 pair differences as additional controls. The same-group indicator is exactly the primary group-
601 ing rule: same FM plus the DINOv2-B/S pair. The coefficient is not a causal estimand, but it
602 checks that the main sign is not removed by simple pair-quality and variance controls. The diag-
603 nostic is generated by `code/scripts/compute_quality_controlled_pairs.py` and released
604 as `results/_merged/diagnostics/quality_controlled_pairs.json`.

Table 9: **Pair-quality controlled diagnostic.** β_{same} is the coefficient on the primary same-group indicator in a pair-level OLS check. Same-group coefficients remain positive after controlling for pair mean Dice, absolute mean-Dice difference, average per-case variance, and absolute variance difference. CIs are case-bootstrap intervals from $B=1000$ refits and should be read as diagnostic, not causal, uncertainty.

Task	Members	Same pairs	Cross pairs	β_{same} [95% CI]
Kvasir	44	82	864	0.169 [0.124, 0.220]
ACDC LV	40	76	704	0.208 [0.175, 0.242]
BraTS foreground	40	76	704	0.224 [0.210, 0.239]
RIGA Cup	36	70	560	0.372 [0.297, 0.456]
RIGA Disc	44	82	864	0.582 [0.495, 0.639]

605 A7 Failure threshold sensitivity

606 The primary estimand is continuous per-case Dice correlation. To make the co-failure interpretation
607 less dependent on continuous-score correlation, the released artifact also recomputes the Table 1
608 same/cross grouping on binary failure events. Table 10 reports the empirical P_{33} readout across all
609 five primary task rows; the released JSON also contains P_{25} and P_{50} thresholds. The RIGA Cup
610 all-fail table below is retained as a pool-level threshold diagnostic, and absolute-threshold lift with
611 support flags is in Appendix A9.

612 **Full τ -sweep of conditional recovery.** Table 11 tabulates grand-mean conditional recovery
613 rates $P[\text{pool-mean Dice} \geq \tau \mid \text{any member Dice} < \tau]$ for same-bundle pools (4 seeds

Table 10: **Binary failure-event diagnostic at the empirical P_{33} threshold.** Failure is defined as per-case Dice below the 33rd percentile of all model-case Dice values after functional-floor filtering. ϕ is binary-event Pearson correlation averaged over same-group or primary cross-group pairs. Joint fail same/cross is the pairwise joint-failure probability averaged over the corresponding model-trace pairs. CIs are case-bootstrap intervals for the same-cross ϕ gap. This is not a clinically chosen threshold. Source: `results/_merged/diagnostics/failure_event_summary.json`, regenerated by `code/scripts/compute_failure_event_summary.py`.

Task	τ_{P33}	Same ϕ	Cross ϕ	Gap [95% CI]	Joint fail same/cross
Kvasir	0.659	0.869	0.540	0.329 [0.278, 0.388]	27.9/21.6%
ACDC LV	0.548	0.875	0.489	0.386 [0.354, 0.418]	27.7/19.9%
BraTS foreground	0.636	0.838	0.425	0.413 [0.401, 0.425]	27.3/18.7%
RIGA Cup	0.575	0.706	0.253	0.453 [0.399, 0.509]	24.4/14.4%
RIGA Disc	0.857	0.703	0.125	0.578 [0.511, 0.631]	24.8/12.0%

of one FM, averaged over all FMs passing the per-task Dice ≥ 0.30 functional floor) and diverse pools (all $\binom{K}{4}$ functional-subset compositions and all released seed tuples) at $\tau \in \{0.70, 0.75, 0.80, 0.85, 0.90\}$. Source: `results/_merged/diagnostics/tau_sweep.json`, regenerated by `code/scripts/compute_tau_sweep.py`. Recovery is τ -fragile in both regimes; the diverse – same-bundle recovery-rate gap $\Delta_{\text{rec}}(\tau)$ (column header below) uses the opposite sign convention to the primary Δ in Table 1 (recovery favours diverse so $\Delta_{\text{rec}} > 0$ is the operationally desired direction) and is non-monotone in τ .

Table 11: **Secondary τ -sweep of conditional recovery rate $P[\text{ens pass} \mid \text{any fail}]$.** Grand-mean across all functional 4-of- K diverse compositions and all functional same-bundle baselines. Recovery gaps vary with threshold and can become near-zero or negative near ceiling; these values are operational diagnostics, not primary support for the same/cross correlation claim.

Task	τ	Mono	Diverse	Δ_{rec}
RIGA Cup (9-bundle)	0.70	0.061	0.214	+0.153
	0.75	0.073	0.120	+0.047
	0.80	0.055	0.050	−0.005
	0.85	0.010	0.007	−0.003
	0.90	0.001	0.000	−0.001
ACDC LV (10-bundle)	0.70	0.074	0.285	+0.211
	0.75	0.062	0.237	+0.175
	0.80	0.049	0.163	+0.114
	0.85	0.040	0.072	+0.032
	0.90	0.020	0.012	−0.007
Kvasir-SEG (11-bundle)	0.70	0.092	0.388	+0.295
	0.75	0.074	0.360	+0.285
	0.80	0.058	0.267	+0.209
	0.85	0.050	0.156	+0.106
	0.90	0.025	0.030	+0.005
BraTS foreground (2D FLAIR seg>0)	0.70	0.080	0.326	+0.245
	0.75	0.061	0.233	+0.172
	0.80	0.039	0.129	+0.090
	0.85	0.019	0.035	+0.016
	0.90	0.004	0.001	−0.003

At the strictest tested threshold ($\tau=0.90$), the diverse advantage in conditional recovery vanishes and reverses on RIGA Cup ($\Delta = -0.001$) and ACDC LV ($\Delta = -0.007$), while remaining positive on Kvasir (+0.005) and slightly negative on BraTS (−0.003). The diverse-pool benefit is concentrated in the moderate-failure regime ($\tau \in [0.70, 0.80]$), consistent with the primary correlation gap predicting reduced *but not zero* coverage benefit at extreme thresholds.

A8 Kvasir secondary split diagnostic

Pixel-level diagnostic on this split: mono 0.80 vs. diverse 0.40, gap 0.394 [95% CI 0.38, 0.40]. Absolute-threshold lift at Dice<0.7: mono 252× vs. diverse 120× (ratio 2.1×); at Dice<0.8: 48.4×

Table 12: **RIGA Cup threshold diagnostic.** Values are for the released primary RIGA Cup functional pool only. Mono pools are same-bundle four-seed pools averaged over functional FMs; diverse pools are four distinct functional bundles. Percentile thresholds are computed over all model-case Dice values after the $\text{Dice} \geq 0.30$ functional floor. Source: `results/_merged/diagnostics/threshold_table_riga_cup.json`.

Threshold	Mono fail $\bar{r}$	Diverse fail $\bar{r}$	Mono ALL-fail	Diverse ALL-fail
P_{25} ($\tau=0.513$)	0.833	0.209	19.9%	2.2%
P_{33} ($\tau=0.575$)	0.803	0.257	26.5%	5.3%
P_{50} ($\tau=0.686$)	0.840	0.268	44.8%	15.4%

Table 13: **Kvasir-SEG secondary split diagnostic.** This generic 8-bundle summary uses a 1,000-image, 700/300 train/test split and is not the 11-bundle, 880/120 Kvasir row in Table 1. Seven of eight generic encoder bundles produce functional segmentation ($\text{Dice} > 0.40$); CLIP ViT-L/14 is non-functional and excluded from cross-family analysis. MedSAM-only traces are reported separately as an encoder-probe summary in Appendix A10; no RETFound Kvasir trace is included in the released artifact bundle.

FM Family	Dice	Within $\bar{r}$	Status
DINOv2 ViT-B/14	0.819	0.972	functional
DINOv2 ViT-S/14	0.807	0.951	functional
BiomedCLIP ViT-B/16	0.717	0.954	functional in this split
ConvNeXt-Tiny	0.651	0.949	functional
EfficientNet-B0	0.582	0.945	functional
ResNet-50	0.540	0.880	functional
ResNet-18	0.468	0.876	functional
CLIP ViT-L/14	0.174	0.839	non-functional

629 vs. $14.6 \times (3.3 \times)$; at $\text{Dice} < 0.9$: $4.6 \times$ vs. $1.8 \times (2.5 \times)$. The numbers are slightly smaller than
630 RIGA/ACDC, likely reflecting the larger test set (300 cases vs. 224 / 90) and more variable polyp
631 presentations, but qualitatively identical: mono still shows substantially higher correlated-failure lift
632 than diverse.

633 A9 Matched- M lift comparison

634 To ensure the mono vs. diverse comparison is not driven by pool size, we construct an *oracle*
635 matched- $M=4$ diverse pool by selecting the top-4 bundles by *test-set seed-averaged* mean Dice
636 (rank computed across all four decoder seeds, avoiding seed-42 selection circularity; selected
637 bundles per task are recorded in the appendix diagnostic output—e.g. Kvasir top-4 = {DINOv2-B,
638 DINOv2-S, BiomedCLIP, ConvNeXt-T}, Cup top-4 = {BiomedCLIP, DINOv2-B, DINOv2-S,
639 EfficientNet-B0}). This is deliberately favourable to the diverse pool and useful for a matched-size
640 structural comparison, but it is not a deployable pool-selection rule because the ranking uses the
641 same test set on which lift is evaluated. For each task / threshold we report three lift variants:

- 642 • **Matched-4 seed-42 (single-seed diagnostic):** the single-seed pool; retained for comparability
643 with a single-seed estimator.
- 644 • **Matched-4 symmetric permuted (preferred):** each of the 4 FMs takes a *distinct* decoder
645 seed from {42, 43, 44, 45}; the lift is averaged on the log scale over all $4! = 24$ seed-to-FM per-
646 mutations, then exponentiated. This is the matched-size structural analog of the main paper’s
647 symmetric Pearson $\bar{r}$ estimator (mono = DINOv2-B with 4 distinct seeds matches matched-4 =
648 4 FMs with 4 distinct seeds).
- 649 • **Matched-4 symmetric independent (sensitivity):** each FM’s seed is drawn i.i.d. with replace-
650 ment from {42..45}, $4^4 = 256$ tuples; reported as a robustness check (the permuted variant is
651 preferred because mono uses four *distinct* seeds).

652 All bootstrap CIs are computed on *cases* only (patient-clustered on ACDC); the seed tuple space is
653 finite and fully enumerated, not bootstrapped. We report 95% case-bootstrap CIs on the log-scale lift
654 and exponentiate to the reported interval ($B=2000$). We also report the raw all-fail count (AF_{cnt})
655 and the Poisson-binomial independence expectation $\mathbb{E}[\text{AF}_{\text{ind}}]$; when $N \cdot \mathbb{E}[\text{AF}_{\text{ind}}] < 5$ we flag
656 the row as support-limited. Under the matched-size comparison, the adequately supported rows are

657 RIGA Cup $\tau=0.9$, RIGA Disc $\tau=0.95$, and ACDC LV $\tau=0.9$; Kvasir $\tau=0.8$ is mono-sparse and
658 retained as a descriptive diagnostic to match Fig. 1. Across these selected rows the mono/matched-4
659 lift ratio is 1.5–5.8 $\times$ (Table 14); lower-support rows give larger ratios (5.9–17 $\times$) treated as descrip-
660 tive order-of-magnitude only. We therefore treat τ values well below the task-specific Dice ceiling
661 as support-limited and report them as descriptive rather than precise lift estimates.

Table 14: **Matched- $M=4$ lift diagnostic.** Mono is DINOv2-B $\times$ 4 seeds; matched-4 diverse uses the top-4 bundles by test-set seed-averaged mean Dice, so this is an oracle structural comparison, not a deployable selection rule. CIs are 95% case-bootstrap intervals on log-lift, exponentiated. Rows without adequate support ($\checkmark$ missing) are descriptive order-of-magnitude indicators.

Task	τ	Mono [CI]	Diverse-42	Diverse-perm [CI]	Ratio*	Supp.
RIGA Cup	0.8	4450 [1408, 24012]	174 [79, 467]	265 [130, 696]	17 $\times$	–
RIGA Cup	0.9	2.18 [1.81, 2.69]	1.33 [1.20, 1.47]	1.44 [1.31, 1.60]	1.5 $\times$	$\checkmark$
RIGA Disc	0.95	9.80 [6.82, 14.6]	2.87 [2.19, 3.82]	3.11 [2.44, 3.98]	3.2 $\times$	$\checkmark$
ACDC LV	0.5	2149 [462, 34449]	200 [51, 1026]	188 [62, 817]	11 $\times$	–
ACDC LV	0.7	189 [55.6, 1064]	31.2 [15.5, 77.7]	25.3 [13.1, 57.8]	7.5 $\times$	–
ACDC LV	0.8	33.8 [15.1, 107]	5.87 [3.60, 10.6]	5.71 [3.63, 9.93]	5.9 $\times$	–
ACDC LV	0.9	2.23 [1.63, 3.32]	1.39 [1.22, 1.60]	1.40 [1.25, 1.61]	1.6 $\times$	$\checkmark$
Kvasir	0.7	246 [128, 570]	33.9 [22.3, 54.5]	30.8 [20.4, 48.3]	8.0 $\times$	–
Kvasir	0.8	48.0 [28.9, 85.1]	8.72 [6.62, 12.0]	8.21 [6.35, 11.1]	5.8 $\times$	mono-sparse

* Ratio = Mono / Matched-4 PERM. **Bold** = preferred symmetric-permuted estimator (24 seed-to-FM tuples). Matched-4 seed 42 is a single-seed diagnostic. Matched-4 INDEP ($4^4=256$ iid tuples) produces values within $\pm 1.5\%$ of PERM on every row and is omitted for compactness; the full table is a secondary diagnostic rather than the `results/_merged` primary scorer output. Supp. $\checkmark$ iff $N \cdot \mathbb{E}[\text{AF}_{\text{ind}}] \geq 5$ for the PERM estimator. Mono lift values span four orders of magnitude (1.4–4,450 $\times$); precision is reported to the data’s underlying decimals and is not uniform across the lift column.

662 A10 Medical-domain encoder probe details

663 We implemented two medical-domain encoder paths, but only the MedSAM image-encoder probe
664 is reported in the released evidence bundle:

665 **RETFound** is a ViT-L/16 MAE pretrained on 1.6M retinal images [Zhou et al., 2023]. The
666 code path builds the RETFound-MAE CFP ViT-L/16 trunk and expects the gated HuggingFace
667 checkpoint YukunZhou/RETFound_mae_natureCFP or a user-specified checkpoint path; the of-
668 ficial upstream project is rmaphoh/RETFound_MAE. RETFound MAE pretraining uses ImageNet
669 [0.485, 0.456, 0.406]/[0.229, 0.224, 0.225] normalisation; our frozen-probe adaptation would drop
670 the CLS token and reshape the $196=14 \times 14$ patch tokens into a spatial 14×14 feature map of dimen-
671 sion 1024. That patch-token extraction is not RETFound’s official global/CLS-style classification
672 fine-tuning path. Because the gated checkpoint was unavailable in the release environment, RET-
673 Found cells are not reported and no RETFound result supports the paper’s empirical claims.

674 **MedSAM** is a SAM ViT-B fine-tuned on 1.5M medical image-mask pairs [Ma et al., 2024].
675 The code path uses the official bowang-lab/MedSAM medsam_vit_b.pth checkpoint through
676 `segment_anything.sam_model_registry['vit_b']` with MEDSAM_CHECKPOINT pointing to
677 that checkpoint. In this audit, we use *only the image encoder* as a frozen feature extractor with
678 our same 1×1 -conv decoder. We discard MedSAM’s prompt encoder and mask decoder entirely—
679 i.e. MedSAM enters this boundary probe on the same frozen-feature terms as the other FMs (no
680 prompt-based path, no bounding-box input, no iterative refinement). The wrapper strips the SAM
681 image-encoder neck and uses the pre-neck ViT-B representation: the official 1024×1024 positional
682 grid is bicubically interpolated to the 224×224 input grid, yielding a 14×14 spatial feature map
683 of dimension 768. This pre-neck $768 \times 14 \times 14$ tensor is an internal representation, not the official
684 SAM/MedSAM post-neck image embedding at 1024 input. This means MedSAM in our experiment
685 is a “MedSAM image-encoder probe”, not its intended prompt-conditioned segmentation system or
686 official inference preprocessing path. The comparison is therefore an encoder-only boundary probe
687 for the same-bundle dependence question we ask, and it does not evaluate MedSAM as released or
688 its best-case performance with its own prompt-based decoder.

689 All other aspects of the secondary sensitivity use the same four decoder seeds, 30-epoch Adam
690 10^{-3} training, and identical 1×1 conv decoder. The completed MedSAM-only probe cov-
691 ers ACDC LV, Kvasir, RIGA Cup, and RIGA Disc (four seeds each). Mean Dice across
692 seeds is 0.517, 0.513, 0.734, and 0.908 respectively; the aggregate diagnostic is released as
693 `results/_merged/diagnostics/medsam_only_summary.json`. We use this as scope context
694 rather than primary evidence.

A11 BraTS/ACDC clustering sensitivity

The Table 1 BraTS foreground row is a BraTS 2024 GLI 2D FLAIR derivative with $n=3,228$ test slices from 324 held-out subjects: per held-out subject, the loader keeps up to the top-10 axial slices by foreground area, requires at least 64 positive pixels, and uses a subject-level 80/20 split with NumPy seed 42. The 3,228 count is traceable to the released `test_ids`: 321 held-out subjects contribute ten eligible slices, one contributes nine, one contributes six, and one contributes three, giving a 12-slice deficit from the 324×10 maximum. ACDC LV similarly evaluates 361 LV-positive slices nested under 20 held-out patients. The main table is therefore a case-level audit, but the released `subject_level/*.json` files include the patient/subject unit IDs linked to `test_ids` and two cluster-aware checks. **(i) Cluster bootstrap over the case-scale statistic.** Resampling patients/subjects with replacement and keeping all slices within each sampled unit gives ACDC $\Delta=0.321$ [0.269, 0.374] and BraTS foreground $\Delta=0.315$ [0.283, 0.351], so the case-scale gap remains positive under unit-level resampling. **(ii) Subject-aggregated statistic.** Averaging each member’s per-slice Dice within patient/subject before computing Pearson gives smaller but still positive gaps: ACDC $\Delta=0.262$ [0.151, 0.413] and BraTS foreground $\Delta=0.214$ [0.182, 0.251]. The conservative reading is that slice dependence inflates the magnitude relative to subject means, especially on BraTS, but it does not explain the within-vs-cross separation away.

A12 BraTS sparse-class filtering caveat and released diagnostics

Per-class BraTS sensitivity can be inflated by sparse labels: if a class is absent from a slice, empty-vs-empty Dice can return a perfect score and obscure the foreground-error structure that matters for a segmentation audit. For this reason, the primary BraTS row uses the foreground derivative `seg>0` rather than a sparse NCR-only target. NCR filtering diagnostics from a separate sparse-class sweep are not used as numeric evidence here because their exact aggregate files are not included in the supplementary artifact. The released BraTS foreground train-size and multimodal diagnostics are instead reported later in Appendices A31.5–A31.6 from `results/_merged/late_brats_summaries.json`.

A13 Scope-only exploratory probes

Note on CLIP-L on ACDC. The secondary 8-FM ACDC diagnostic discussed in this section uses a different split where ResNet-50, ResNet-18, and CLIP ViT-L/14 collapse below the Dice ≥ 0.30 floor (Dice 0.099, 0.230, and 0.007 respectively); the *primary* ACDC LV pool of Table 1 (post-floor 10-bundle subset of the 11-bundle source pool, with only ResNet-18 dropped) retains CLIP ViT-L/14 because it passes the Dice floor on the primary split. Different splits, different functional outcomes for the borderline bundles.

This block separates released diagnostics from exploratory mechanism probes without released trace artifacts or aggregate JSONs. Only released traces or released aggregate JSONs are used as numerical support for the paper’s claims. The shared reading is descriptive: same-backbone perturbations can reduce correlation, but the released diagnostics do not make same-backbone pools behave like the primary cross-family references.

ACDC architecture caveat. Secondary ACDC exploratory tables used a separate 8-bundle split and architecture-pair subsets that are not included as replayable aggregate artifacts. The released support for the architecture-caveat wording is instead the same-architecture-family ViT / different-pretraining diagnostic in Appendix A29 and the expanded same-family cross-checkpoint decomposition in Appendix A30.

Random-init probes. DINOv2-B random-initialisation probes used secondary splits and source prediction caches outside the released artifact. The released no-pretraining evidence is limited to the CNN random-init diagnostic in Appendix A31.2, which should be read as a descriptive same-recipe sanity check rather than a causal isolation of architecture, optimiser, and data effects.

LoRA and cross-architecture probes. LoRA probes and ViT-B cross-architecture random-init sweeps require source prediction caches and aggregate summary files outside the released artifact. We do not tabulate their values or use them as bundle-supported evidence. The released same-




Figure 2: **Same-backbone bootstrap-bagging diagnostic.** For each shown dataset/backbone pair, we compare (i) same-backbone/same-data cross-seed correlation, (ii) same-backbone/bootstrap-resampled-data correlation, and (iii) a cross-backbone reference from the corresponding diagnostic pool. The ratio panel reports $(r_{\text{same data}} - r_{\text{cross backbone}})/(r_{\text{same data}} - r_{\text{bootstrap data}})$, so values above one mean that changing backbone reduces correlation more than bootstrap-resampling the training data in this diagnostic. The BraTS label refers to the 2D foreground derivative (seg>0), not the official 3D BraTS WT challenge target. This figure is descriptive and does not identify a causal mechanism.

backbone levers are the head-capacity sweep, UNet-skip replication, fold-reslicing diagnostic, CNN random-init diagnostic, and MC dropout.

A14 Same-backbone, different-data bootstrap-bagging audit

For each (FM, dataset) pair, four bootstrap replicates of the training set (image-level on RIGA/Kvasir; subject-level on BraTS) are crossed with four decoder seeds. Bootstrap data-swap reduces within-pool $\bar{r}$ by only 0.05–0.12, while backbone swap reduces it by 0.33–0.50. This supports the scoped claim that, in the frozen-feature protocol, resampling the training cases does not provide the same decorrelation as changing encoder family.

A15 Distribution-shift audit

We evaluate one released cross-distribution robustness diagnostic for the within-cross correlation gap using an 8-bundle RIGA pool and the frozen-feature protocol from Section 4.

RIGA leave-MESSIDOR-out (cross-vendor fundus). RIGA+ [Almazroa, 2018] combines three imaging sources: *BinRushed* (Bin Rushed Ophthalmic Center, Riyadh, Saudi Arabia), *Magrabia* (Magrabi Eye Center, Riyadh, Saudi Arabia), and *MESSIDOR* (French public corpus, different cameras and populations). We train the frozen-FM + 1×1 conv decoder on BinRushed+Magrabia (union) and evaluate on MESSIDOR (454 cases after the preprocessing and eligibility filters implemented by the released scorer) as cross-vendor out-of-distribution. Per-case traces for this diagnostic are released under `results/_merged/per_case_dice/riga_cup_ood_messidor`. The released scorer uses a fixed 8-bundle $\times$ 4-seed pool without a functional-floor change between ID and OOD rows. Under that directly reproducible convention, the ID row has within/cross/gap 0.659/0.282/0.377 and the MESSIDOR OOD row has 0.743/0.390/0.353.

Pool-level summary. In this fixed-pool RIGA cross-vendor diagnostic, the within-cross gap changes by -0.024 and remains positive. Both within and cross correlations are higher on the MESSIDOR row, so this diagnostic should be read only as evidence that the gap survives one released cross-vendor stress test, not as evidence that distribution shift generally contracts dependence. Broader cross-scanner, cross-demographic, and longitudinal shifts remain out of scope.

A16 UNet-with-skip decoder replication

The paper’s primary decoder is a 1×1 conv with $C+1$ trainable parameters for C frozen channels (for example, 769 parameters for a 768-channel feature map), which is close to a linear probe on

Table 15: **RIGA cross-vendor diagnostic under a fixed 8-bundle scorer.** This is not case-identical to Table 1. Within = same-FM different-seed $\cup$ same-family different-FM (DINOv2-B/S and ResNet-18/50); cross = different-family different-FM. The gap remains positive, but both within and cross correlations increase on MESSIDOR; no general distribution-shift conclusion is supported. Brackets are case-level bootstrap 95% CIs ($B=500$). Source: results/_merged/diagnostics/distshift_riga_messidor.json.

Shift regime	Within- $\bar{r}$	Cross- $\bar{r}$	Gap
ID reference (fixed 8-bundle pool)	0.659 [0.603, 0.715]	0.282 [0.197, 0.369]	0.377 [0.301, 0.463]
OOD MESSIDOR (fixed 8-bundle pool)	0.743 [0.722, 0.765]	0.390 [0.345, 0.430]	0.353 [0.318, 0.388]

This diagnostic uses a paired fixed 8-bundle RIGA pool and should not be read as a case-identical comparison to the Magrabia-source held-out primary RIGA rows in Table 1.

774 frozen features. Dense decoders with skip connections—nnU-Net, U-Net, SegFormer—are closer
775 to deployed medical segmentation. We therefore ran a separate diagnostic in which the frozen back-
776 bone is unchanged but the head is replaced by a multi-scale UNet-style decoder over four encoder
777 feature levels. The diagnostic uses two released seeds (42/43) per cell; “seed-pair r ” is therefore a
778 single Pearson coefficient, not a bootstrapped four-seed within-family estimate.

779 **Architecture.** For hierarchical CNNs (ConvNeXt-T, ResNet-50, ResNet-18, EfficientNet-B0) the
780 decoder uses the native feature pyramid as skip sources; for ViT-style encoders (DINOv2-B/S,
781 BiomedCLIP, MAE-B, DeiT-B) it uses sampled intermediate token maps reassembled into a four-
782 level pyramid. The released head first applies lateral 1×1 projections to a common decoder width,
783 then runs a top-down bilinear-upsample path with three additive skip merges (deep $\rightarrow$ mid $\rightarrow$ shal-
784 low $\rightarrow$ shallowest), two conv-BN-ReLU blocks after each merge, a final conv-BN-ReLU after up-
785 sampling to the output grid, and a final 1×1 output conv. Training matches the frozen-feature recipe
786 where possible: Adam lr 10^{-3} , batch 16, 30 epochs, BCE, seeds {42, 43}.

787 **Coverage.** The 9-backbone $\times$ 5-task expansion adds (i) ResNet-18, EfficientNet-B0, MAE-B and
788 DeiT-B as four new backbones for Kvasir, ACDC LV, RIGA Cup and RIGA Disc, and (ii) RIGA Disc
789 as a fifth task row for the original 5 backbones (DINOv2-B/S, BiomedCLIP, ConvNeXt-T, ResNet-
790 50). The BraTS column for the four new backbones requires a different data-root layout on the
791 host that ran the expansion (BraTS-2024-GLI subjects live under training/training_data1_v2/
792 rather than at the data-root top level expected by fmpool.datasets.brats); those four cells are
793 therefore marked -- and the BraTS row for the original 5 backbones is unchanged from the prior
794 diagnostic.

Table 16: **UNet-skip same-backbone diagnostic, 9 backbones $\times$ 5 tasks (two seeds 42/43 per cell).** Each cell reports “mean Dice (averaged over the two seeds) / seed-pair Pearson r ”. Single-pair r is a high-variance estimator; the reading is that within-backbone correlations stay clearly positive on every reported cell while pool-mean Dice rises substantially over the matching 1×1 probe rows. Primary Table 1 cross-bundle references for context: Kvasir 0.697, ACDC LV 0.639, RIGA Cup 0.361, BraTS fg 0.623, RIGA Disc 0.232. Cells marked -- are not in scope (BraTS expansion for the four new backbones, see Coverage paragraph).

Backbone	Kvasir	ACDC LV	RIGA Cup	BraTS fg	RIGA Disc
DINOv2-B	0.883/0.877	0.864/0.979	0.902/0.967	0.877/0.936	0.963/0.426
DINOv2-S	0.866/0.919	0.863/0.970	0.894/0.893	0.875/0.886	0.960/0.701
BiomedCLIP	0.806/0.926	0.827/0.950	0.880/0.892	0.858/0.929	0.960/0.874
ConvNeXt-T	0.860/0.775	0.885/0.917	0.895/0.970	0.895/0.938	0.966/0.809
ResNet-50	0.816/0.774	0.882/0.945	0.881/0.952	0.889/0.934	0.965/0.873
ResNet-18	0.774/0.828	0.879/0.861	0.883/0.814	—	0.961/0.651
EfficientNet-B0	0.819/0.955	0.881/0.938	0.899/0.888	—	0.965/0.787
MAE-B	0.795/0.827	0.846/0.955	0.890/0.960	—	0.960/0.934
DeiT-B	0.837/0.920	0.849/0.945	0.893/0.941	—	0.963/0.880

795 **Finding.** The UNet-skip head improves mean Dice substantially over several weak 1×1 probe cells,
796 and within-backbone seed-pair r stays positive on every cell (range 0.426–0.979, with the DINOv2-
797 B RIGA Disc outlier at 0.426 a single high-variance pair).

798 *Caveat on cross-bundle comparison.* The cross-bundle Pearson references printed in the table cap-
799 tion (0.361 for RIGA Cup, etc.) are computed from the primary 1×1 -conv pool with 4 seeds per
800 bundle and many bundle pairs; the UNet-skip cells use 2 seeds and a different decoder. Comparing
801 within-backbone UNet-skip seed-pair r (single Pearson per cell) directly to those cross-bundle refer-

ences mixes a decoder shift, a seed-count shift, and a pair-mixture shift. We therefore do not claim the UNet-skip same-backbone correlations sit “above the cross-bundle reference” numerically.

Decoder-matched check, all five tasks. As a decoder-matched reference we recompute UNet-skip within-FM seed-pair r and cross-bundle r on the released per-case Dice (seeds 42/43, intersection-aligned over the released test cases; cross-pairs restricted to within-split because original-5 and e1-4 expansions use disjoint test splits). Released values (Source: results/_merged/diagnostics/unetskip_decoder_matched_5tasks.json):

Task	within $\bar{r}$ (n)	cross $\bar{r}$ (n)	decoder-matched gap	1×1 primary $\Delta_{0.30}$
Kvasir	0.867 (9)	0.689 (72)	0.178	0.261
ACDC LV	0.940 (9)	0.866 (72)	0.074	0.321
BraTS fg	0.925 (5)	0.860 (20)	0.064	0.315
RIGA Cup	0.920 (9)	0.851 (32)	0.068	0.473
RIGA Disc	0.770 (9)	0.545 (72)	0.226	0.638

The decoder-matched gap is 14–68% of the 1×1 -conv primary post-floor gap (mean attenuation $\sim 32\%$). The absolute level of cross-bundle correlation rises sharply when the decoder is matched (e.g., RIGA Cup cross $\bar{r}$ goes $0.361 \rightarrow 0.851$ when the cross reference is recomputed under UNet-skip), while the same/cross gap shrinks. A full nnU-Net v2 5-fold audit on RIGA Cup is reported in Appendix A17.

A17 Fold-reslicing and nnU-Net complements

Motivation. nnU-Net [Isensee et al., 2021] is the de facto standard for medical image segmentation; its 5-fold cross-validation ensemble is a routinely-used uncertainty / robustness baseline. The full nnU-Net pipeline is *self-configuring* (2D/3D selection, preprocessing, patch-based sampling, default geometric/intensity augmentation, cascaded network, post-processing). We keep three evidence surfaces separate: (i) the released frozen-feature fold-reslicing diagnostic, which tests split sensitivity with one decoder seed per fold and therefore cannot estimate within-fold versus cross-fold same-backbone pairs; (ii) released task-level nnU-Net summary JSONs, which are full-pipeline complements but not case-identical contrasts to the primary FM rows; and (iii) a single case-identical RIGA Cup full-pipeline scope check on the primary Magrabia held-out cases.

Released fold-reslicing design. This diagnostic re-slices the frozen-feature cache into 5 folds and reruns the 8-bundle pool {DINOv2-B/S, BiomedCLIP, CLIP-L, ConvNeXt-T, EfficientNet-B0, ResNet-50/18} on ACDC LV and RIGA Cup, single seed per fold cell. Because the released artifact contains one seed per fold, this diagnostic reports per-fold cross-bundle $\bar{r}$ and pool-mean Dice, not a 20-decoder fixed-backbone within/cross estimator. The released artifact provides the corresponding fold-reslicing summary and scorer.

Table 17: **Fold-reslicing cross-bundle stability, one seed/fold.** 8 bundles $\times$ 5 folds; single seed per fold cell. Cross-bundle $\bar{r}$ is the Pearson correlation averaged over the 28 bundle pairs on the held-out fold. This table cannot support within-fold versus cross-fold same-backbone claims.

Task	fold 0	fold 1	fold 2	fold 3	fold 4	mean $\pm$ std
ACDC LV	0.580	0.548	0.583	0.558	0.558	0.565 ± 0.015
RIGA Cup	0.359	0.317	0.325	0.315	0.261	0.315 ± 0.035

Finding. Train/test split sensitivity is small in this fold-reslicing check: cross-fold pool-mean Dice is stable (ACDC 0.502 ± 0.022 ; Cup 0.610 ± 0.012), and per-fold cross-bundle $\bar{r}$ is stable across folds (ACDC 0.565 ± 0.015 ; Cup 0.315 ± 0.035). This supports only the narrower statement that the cross-bundle diagnostic is not dominated by one particular 5-fold re-slicing. It does *not* support a claim about within-fold vs cross-fold same-backbone decorrelation, because those pair types require multiple decoder seeds per fold that are not present in the released traces.

Full nnU-Net v2 complement. To complement the frozen-feature proxy above and answer whether nnU-Net’s full self-configuring pipeline decorrelates the pool more, we trained the full nnU-Net v2 2D U-Net on RIGA Cup with a 520/224 split, using the generated v2 2D nnUNetPlans

configuration after `nnUNetv2_plan_and_preprocess` (self-configured preprocessing and architecture, official default augmentation/loss schedule, 250 iterations per epoch). This split differs from the Magrabia-source held-out FM row and is therefore a task-level complement. We train two seeds per fold via custom `nnUNetTrainerSeed13_100epochs` and `nnUNetTrainerSeed37_100epochs` subclasses that pre-seed random, `np.random`, `torch`, and `cuda` before any nnU-Net sampler touches the RNG; each is trained for 100 nnU-Net-epochs (25,000 iterations on two A100s in parallel), then all $5 \times 2=10$ models predict on the held-out 224-image test set via `nnUNetv2_predict` and per-case Dice is computed against the held-out labels staged in `labelsTr`. A split-index file enumerated held-out IDs; the artifact bundle releases the summary arrays/JSONs for this complement but not the split manifest.

Across all 10 models the mean Dice is 0.949 ± 0.001 (exceeding any of our frozen-feature FM+ 1×1 -conv combinations on RIGA Cup). The pairwise per-case Dice $\bar{r}$ is 0.972 within-fold over the 5 same-fold different-seed pairs and 0.925 cross-fold over the 40 different-fold any-seed pairs, with $\Delta_{\text{within-cross}}=0.047$ and a case-level bootstrap 95% CI of [0.030, 0.069] that excludes zero. Full nnU-Net v2 is structurally a same-architecture, different-training-data ensemble on one 2D Plain-ConvUNet, so it is not comparable in kind to the primary frozen-FM cross-family reference, which swaps encoder bundles. The scoped conclusion is therefore: *within this task-level same-architecture nnU-Net complement, fold identity leaves a small but positive correlation gap*; we do not claim nnU-Net substitutes for or surpasses backbone-family diversification. The 100-epoch budget, single-task coverage, and 2D-only scope are closed in Appendix A18 (1000-epoch retrain on the same RIGA Cup setup; 3D fullres nnU-Net on ACDC; higher-resolution/Dice+BCE/aug frozen-FM replication).

Case-identical RIGA Cup held-out scope check. We additionally train the same 5×2 seeded nnU-Net v2 2D grid on BinRushed+MESSIDOR folds and evaluate all models on the primary Magrabia-source held-out Cup cases ($n=95$). The exporter uses the official nnU-Net v2 natural-image channel convention, with RGB saved as `_0000/_0001/_0002` channel files, `imagesTs/labelsTs` for held-out scoring, and the generated `nnUNetPlans` 2D configuration. Released case IDs are deterministic hashes rather than upstream path strings. Across the 10 models, mean Dice is 0.9162 ± 0.0019 . The same-fold, different-seed correlation is $\bar{r}=0.984$ over 5 pairs, the different-fold correlation is $\bar{r}=0.976$ over 40 pairs, and the gap is $\Delta=0.0079$ with case-bootstrap 95% CI [0.002, 0.039]. This confirms the boundary in a stricter paired setting: nnU-Net’s same-architecture fold/seed surface is highly correlated on the same held-out cases, and the fold gap is much smaller than both the primary RIGA Cup same/cross gap (0.473) and the separate-split task-level nnU-Net summary (0.047). It does not make nnU-Net a cross-family reference or a causal control for architecture versus pretraining.

Data-swap comparison caveat. Appendix A14 reports bootstrap-bagging within-pool decorrelation with four decoder seeds per bootstrap replicate. The fold-reslicing row instead reports single-seed cross-bundle correlations across held-out folds. These diagnostics both stress training data/split sensitivity, but they are not dose-matched estimators and are not combined into a causal ranking.

A18 Segmentation-recipe and nnU-Net scope audit

Our main protocol uses 224×224 2D inputs, frozen backbone, 1×1 -conv decoder, BCE loss, no augmentation, and 30-epoch head training. Routine medical-imaging pipelines often run with compound Dice+CE loss, geometric augmentation, 3D input, and full-length (1000-epoch) training, so we test whether the within/cross correlation structure we report is specific to the simplified protocol. Most nnU-Net rows below are task-level complements; the added case-identical RIGA held-out row uses the same Magrabia-source held-out cases as the primary frozen-FM row but remains a single-task, 2D, same-architecture nnU-Net scope check.

Dimensionality axis: 3D fullres nnU-Net at 1000 epochs on ACDC. Training nnU-Net v2 in its self-configuring `3d_fullres` setting on ACDC with a seed-42 105/45-patient split (210 training, 90 test volumes) using `nnUNetTrainerSeed{13,37}_1000epochs` (seed injected in `on_train_start` before sampler initialisation) gives 10 fully-trained 3D models across 5 folds. `nnUNetv2_predict` on all 90 held-out volumes yields per-case Dice on mean foreground and per class. Across 10 models the mean Dice is 0.920 ± 0.001 , within the range reported for strong nnU-Net baselines on ACDC. On per-case mean Dice the within-fold (5 pairs) vs cross-fold (40 pairs) gap

is $\Delta=0.043$ [0.021, 0.085] with within $\bar{r}=0.978$ and cross $\bar{r}=0.935$; per class, $\Delta_{RV}=0.041$ [0.011, 0.110], $\Delta_{Myo}=0.029$ [0.017, 0.046], $\Delta_{LV}=0.067$ [0.033, 0.098]. The artifact bundle releases the aggregate summary for this 150-patient complement but not the exact split manifest; the released public ACDC converter targets patients 001–100 for the primary 2D row. We therefore treat this 3D row as a task-level scope summary rather than a from-bundle training recipe. In this ACDC nnU-Net complement, the 2D→3D fullres shift changes the small fold-gap regime by less than 0.01 relative to the 2D 100-epoch baseline ($\Delta=0.047$ in Appendix A17).

Loss / augmentation axis: RIGA Cup frozen-FM at encoder-native 224×224 . We retrain the seven-bundle RIGA Cup subset (DINOv2-B/S, CLIP ViT-L/14, ConvNeXt-T, EfficientNet-B0, ResNet-50/18) with four seeds each under a more standard segmentation loss/augmentation recipe: encoder-native input size 224×224 , composite soft Dice+BCE loss, paired random horizontal flip and rotation $\pm 15^\circ$, for 30 epochs. The frozen encoder wrappers use 224×224 inputs; we therefore treat this as a loss/augmentation test, not a resolution test. All other factors (frozen backbone, 1×1 -conv segmentation head, Adam 10^{-3}) match the main protocol. Per-FM mean Dice ranges from 0.526 (ResNet-18) to 0.758 (DINOv2-B). Under the same symmetric estimator used for the main table, within-family $\bar{r}$ is 0.903 (58 within pairs, including same-FM seed pairs and DINOv2-B $\times$ S), cross-FM $\bar{r}$ is 0.556 (320 cross pairs), and the within/cross gap is $\Delta=0.347$ [0.275, 0.428]. The gap is smaller than the main RIGA Cup interval (0.473 [0.395, 0.556]) but remains positive; loss/augmentation changes do not remove the gap.

Training-length axis: full nnU-Net v2 at the 1000-epoch default on RIGA Cup. We retrain the 10-model RIGA Cup 2D nnU-Net pool from Appendix A17 at the nnU-Net v2 default 1000 epochs, using the same seeded trainer pattern (`nnUNetTrainerSeed{13,37}_1000epochs`) on the same 520/224 split with identical dataset/plan files. Mean Dice across the 10 models is 0.946 ± 0.0004 (vs 0.949 ± 0.001 at 100 epochs). Within-fold $\bar{r}=0.983$ (5 pairs), cross-fold $\bar{r}=0.949$ (40 pairs), and $\Delta=0.034$ [0.021, 0.058]. Relative to the 100-epoch run both estimators tighten slightly (within $0.972\rightarrow 0.983$; cross $0.925\rightarrow 0.949$) and the within/cross gap contracts from 0.047 to 0.034; measured predictions are more correlated under longer training, and the gap does not disappear. Ten-fold longer training therefore does not change the replicability structure of nnU-Net on this task-level complement; it is not a case-identical comparison to the Magrabia-source held-out FM row.

Synthesis. Across these scope analyses, the frozen-FM loss/augmentation row remains in the same large-gap regime as the RIGA Cup main row, while the same-architecture nnU-Net fold axis remains in a small-gap regime. The case-identical RIGA held-out row tightens this boundary: on the same Magrabia cases, both nnU-Net correlations are very high and the fold gap is only 0.0079 at 100 epochs. The matched 1000-epoch case-identical re-run gives mean Dice 0.9151 ± 0.0018 across the same ten models with within-fold $\bar{r}=0.987$, cross-fold $\bar{r}=0.980$, and $\Delta=0.0075$ [0.003, 0.034], so a tenfold-longer training budget on the Magrabia-source held-out split leaves the case-identical fold gap essentially unchanged ($0.0079\rightarrow 0.0075$) while both correlations rise modestly. These rows do not support a simpler explanation that the main-table pattern is solely an artifact of BCE loss, no augmentation, short training, or the tested full-pipeline nnU-Net complements.

A19 Held-out multi-task stress test

Summary. Across the five held-out 2D binary derivatives we tested, two rows give substantively positive gaps (Hippocampus $\Delta=0.491$, ISIC-2018 $\Delta=0.406$), one is tiny- N positive (Prostate, $n=6$, $\Delta=0.146$), one gives a near-zero gap (Spleen $\Delta=0.003$), and one fails the functional floor (Heart, 0/11 functional FMs). The held-out evidence *supports* rather than confirms the primary 2D rows and is consistent with task-specific structure modulating the gap magnitude.

Protocol. In addition to the four primary benchmark datasets (five task rows: Cup, Disc, ACDC, Kvasir, BraTS foreground), we evaluate five *held-out* tasks covering additional imaging domains: MSD Hippocampus [Antonelli et al., 2022] (T1 MRI, native 3-class mask {background, anterior, posterior} binarized to {background, any-hippocampus}), ISIC-2018 Task 1 Lesion Boundary [Codella et al., 2019] (dermoscopy, 2D RGB, $N=80$ test cases held out from a 400-image random subset of the 2,594-image training release), MSD Heart (T2 MRI cardiac), MSD Prostate (T2 MRI prostate), and MSD Spleen (CT portal). This held-out stress test uses the same 11 generic encoder

Table 18: **Scope audit across recipe and nnU-Net complements.** Rows are not all case-identical and should not be compared as a controlled ablation. The table asks whether selected recipe changes remove the qualitative gap. The frozen-reference row uses the same seven candidate FMs as the Dice+BCE+augmentation row before the primary functional-floor exclusion; the primary filtered RIGA Cup row is Table 1. The case-identical RIGA held-out row uses the Magrabia-source test cases; the separate-split RIGA nnU-Net rows use a 520/224 split.

Axis	Protocol	Task	Within	Cross	Δ	95% CI
Frozen reference	7-bundle candidate set, 224, BCE, no aug, 30 epochs	RIGA Cup	0.798	0.266	0.532	[0.456, 0.614]
Loss / aug	7-bundle frozen, 224, Dice+BCE, aug, 30 epochs	RIGA Cup	0.903	0.556	0.347	[0.275, 0.428]
Baseline (App. A17)	nnU-Net 2D, 100 epochs, separate split	RIGA Cup	0.972	0.925	0.047	[0.030, 0.069]
Case-identical held-out	nnU-Net 2D, 100 epochs, Magrabia test	RIGA Cup	0.984	0.976	0.0079	[0.002, 0.039]
Case-identical held-out	nnU-Net 2D, 1000 epochs, Magrabia test	RIGA Cup	0.987	0.980	0.0075	[0.003, 0.034]
Training length	nnU-Net 2D, 1000 epochs, separate split	RIGA Cup	0.983	0.949	0.034	[0.021, 0.058]
Dimensionality	nnU-Net 3D fullres, 1000ep	ACDC (mean)	0.978	0.935	0.043	[0.021, 0.085]
	per-class RV		0.967	0.927	0.041	[0.011, 0.110]
	per-class Myo		0.987	0.957	0.029	[0.017, 0.046]
	per-class LV		0.978	0.911	0.067	[0.033, 0.098]

946 bundles and four decoder seeds as the primary source pool where possible, with the same 1×1 -conv
947 decoder, no per-task hyperparameter tuning, and no model substitution. These are task-specific 2D
948 binary derivatives, not official challenge-test evaluations. The $\text{Dice} \geq 0.30$ functional floor is applied
949 uniformly, and the released aggregate summary records 220/220 expected per-seed traces.

Table 19: **Held-out aggregate stress test** (11 generic encoder bundles $\times$ 4 decoder seeds; point estimates only, no case-bootstrap CIs). Hippocampus and ISIC-2018 are direction-positive, Prostate is positive but has only six test cases, Spleen is essentially near-zero despite passing the Dice floor, and Heart has no functional bundle at the $\text{Dice} \geq 0.30$ floor. Tiny- N and floor-failure rows are scope boundaries and do not support the primary claim.

Task	N	Func./avail.	Within	Cross	Δ	M_{eff}
<i>Direction-positive point estimates with enough support to read qualitatively</i>						
MSD Hippocampus	52	11/11	0.909	0.419	0.491 [§]	3.94
ISIC-2018 [‡]	80	11/11	0.918	0.511	0.406 [§]	2.72
<i>Scope-limit rows: tiny held-out support or non-functional pool</i>						
MSD Prostate	6	5/11	0.788	0.642	0.146 [§]	1.23
MSD Spleen	8	11/11	0.983	0.979	0.003 [§]	1.02
MSD Heart	4	0/11	—*	—*	—*	—

*No bundle passes the $\text{Dice} \geq 0.30$ floor on the four-case Heart derivative, so no within/cross gap is interpretable. [§]Values are point estimates from `results/_merged/diagnostics/heldout_11fm_summary.json`; case-level CIs are not computed for this aggregate diagnostic, and the MSD Prostate/Spleen rows have too few held-out cases to support a primary claim. [‡]ISIC-2018 uses a single random $N=400$ -image subset of the 2,594-image ISIC-2018 training release with $N=80$ test cases under the seed-42 80/20 split; we do not report an independent redraw and we do not use the official ISIC test server in this study.

950 **Interpretation.** The 11-bundle held-out protocol yields three classes of outcome: (i) ISIC-2018
951 and Hippocampus both show direction-positive gaps under the same broad estimator family as the
952 primary audit, but without case-bootstrap CIs and with smaller held-out support than the primary
953 rows. (ii) Prostate and Spleen illustrate why held-out diagnostics are scope boundaries rather than
954 primary evidence: Prostate is direction-positive but has only six test cases, while Spleen passes the
955 floor yet has a near-zero gap and $M_{\text{eff}} \approx 1$ because the members fail together. (iii) Heart does
956 not yield a functional pool under the frozen-feature 2D derivative. Thus the held-out stress test is
957 compatible with transfer to some external rows, but it also records where the protocol is too weak or
958 too saturated to audit same-bundle dependence.

959 **Operational 8-bundle held-out probes.** The following recovery, pool-size, absten-
960 tion, and pool-selection diagnostics use an 8-bundle held-out candidate set stored under
961 `results/_merged/diagnostics/held_out/`. They are retained as operational-method sensitiv-
962 ities, not as the provenance for Table 19 above.

963 **CLIP/BiomedCLIP normalization robustness.** Uniform ImageNet normalization can be a po-
964 tential artifact for BiomedCLIP and CLIP-L because both have their own normalization statistics.

We therefore evaluated those two FMs on Hippocampus and ISIC-2018 with the correct CLIP normalization. The released `summary_clipfix.json` reports BiomedCLIP mean Dice 0.840 (Hippocampus) / 0.908 (ISIC-2018), while CLIP-L remains below the functional floor on Hippocampus (mean Dice $<10^{-10}$) and reaches 0.106 on ISIC-2018, still below the 0.30 floor. Relative to `summary_5tasks.json`, the aggregate Δ changed by 0.002 (Hippocampus 0.731 $\rightarrow$ 0.729) and 0.003 (ISIC-2018 0.552 $\rightarrow$ 0.555). In this held-out diagnostic, correcting CLIP/BiomedCLIP normalization does not materially change the Hippocampus/ISIC point estimates.

Held-out recovery-rate sensitivity. To complement the primary-task $\Delta_{\text{rec}} = +0.29$ (Cup) / $+0.26$ (ACDC) grand means, we computed an exploratory held-out recovery gap at $\tau=0.85$ on the single ISIC-2018 subset: $\Delta_{\text{rec}} = +0.269$ [$+0.212, +0.331$] under a case-resampling calculation saved as `recovery_CI.json`. This supports the direction of the operational gap on one dermoscopy subset, but it is not a deployment-calibrated guarantee and does not resolve the held-out estimator-dependence above. Hippocampus at $\tau=0.85$ is recovery-saturated and gives $\Delta_{\text{rec}} = -0.01$ [$-0.03, +0.00$] (diverse and mono both near-zero recovery because only 2/7 functional bundles individually pass Dice ≥ 0.85). Lower thresholds are not evaluated here.

Pool-size sweep on held-out tasks. How does diverse-vs-same-bundle conditional recovery scale with M ? For $M \in \{2, 3, 4, 5, 6\}$ on the two functional held-out tasks (Table 20), diverse pools beat same-bundle seed pools at every $M \geq 3$ on the single ISIC-2018 subset ($\Delta_{\text{rec}} = 0.27\text{--}0.46$), while Hippocampus at the paper’s $\tau=0.85$ threshold is recovery-rate-saturated (all pools <0.11) because only 2 of 7 functional bundles reach Dice ≥ 0.85 individually. These held-out operational analyses are exploratory: ISIC provides one supportive dermoscopy subset, Hippocampus is threshold-saturated at $\tau=0.85$, and the $M=5, 6$ same-bundle rows are proxy constructions using the available four seeds rather than true larger same-bundle pools. Released aggregate data: `results/_merged/diagnostics/held_out/pool_size_sweep.json`.

Table 20: **Exploratory held-out pool-size sweep.** Conditional recovery rate $P[\text{ens pass} \mid \text{any fail}]$ at $\tau=0.85$ vs pool size M . $M=2$ recovery is a definition artifact under the strict-majority rule, and $M=5, 6$ same-bundle values are four-seed proxies. This table is not a deployable pool-size recommendation.

M	Hippocampus			ISIC-2018		
	mono	diverse	Δ	mono	diverse	Δ
2 [†]	0.000	0.000	+0.00	0.000	0.000	+0.00
3	0.069	0.107	+0.04	0.111	0.498	+0.39
4	0.052	0.039	−0.01	0.091	0.359	+0.27
5 [‡]	0.052	0.078	+0.03	0.091	0.553	+0.46
6 [‡]	0.052	0.027	−0.03	0.091	0.454	+0.36

[†] $M=2$ is a definition artifact under the strict-majority ($>M/2$) rule conditioned on any-fail: recovery is zero by construction whenever $M=2$, regardless of pool composition. We include the row for completeness. [‡] At $M=5, 6$ the same-bundle entries are computed on the $\min(M, S)=4$ available seeds per FM, so they do not represent true $M=5, 6$ same-bundle pools; the diverse entries do use legitimate $M=5, 6$ distinct-family compositions.

Calibration-tuned disagreement-based abstention (an exploratory alternative lever, not a con-formal guarantee). A natural operational question is whether standard abstention-based uncertainty gating closes the same-bundle gap with modest additional computation. We calibrate a per-task disagreement-score threshold on 20 cases (inter-member Dice standard deviation as score; grid-search the quantile q to meet a target case-risk $\alpha=0.1$ on calibration; apply the selected threshold to the remaining test cases). *This is a tuned heuristic, not split-conformal prediction* [Angelopoulos and Bates, 2021]: we use a grid search over q that terminates on the first calibration risk meeting α and use inter-member variance rather than a proper conformity score, so no finite-sample risk guarantee holds. The pool is also chosen per-task on the whole task mean Dice (family-representative highest-Dice FMs), so the diverse-pool selection step leaks task labels—this biases the comparison toward diverse pools. We report it as a sensitivity. On ISIC-2018 at matched abstention ≈ 0.74 , diverse pools reach test-risk 0.129 ± 0.027 vs. same-bundle grand mean 0.174 ± 0.209 (`conformal_vs_meff.json`). The *point estimate* is $\sim 25\%$ lower for diverse, but the same-bundle grand-mean SD (0.21) is $\sim 8\times$ the diverse SD (0.027); the per-FM same-bundle risk means span 0.000–0.616 across the seven functional FMs, so the grand-mean comparison is highly composition-sensitive and we do not claim a tight ranking. Hippocampus cannot hit $\alpha=0.1$ for either pool, giving residual risk 0.694 (diverse) vs. 0.665 (same-bundle

grand mean) ($\Delta = +0.03$, with same-bundle SD 0.28, indistinguishable). A label-free conformal procedure with a proper conformity score is not evaluated here. Released aggregate data: `results/_merged/diagnostics/held_out/conformal_vs_meff.json`.

Calibration-set pool-selection probe on held-out tasks. We operationalise a small calibration exercise: given only a 20-case calibration set on a held-out task, can a rule pick a 4-of-8 FM pool from the same secondary held-out candidate set that beats random? We evaluate five rules—(R1) random, (R2) top-4 by mean Dice on calibration, (R3) max- M_{eff} , (R4) quality-gated max- M_{eff} (Dice ≥ 0.30 floor then max- M_{eff} among qualifying), (R5) diverse-by-architecture (2 best-Dice ViT + 2 best-Dice CNN; reported as an exploratory hand-crafted baseline, not a general proposal)—and report conditional recovery rate $P[\text{ens pass} \mid \text{any fail}]$ at $\tau=0.85$ on the remaining held-out cases, averaged over 100 calibration-set bootstraps (Table 21). All rule selectors and the recovery estimator use seed-42 per-case Dice (not seed-averaged), matching a one-model-per-FM pool construction. Hippocampus has lower effective recovery support than ISIC at $\tau=0.85$, so its recovery numbers are exploratory. In this diagnostic, simple top-4-by-Dice has the highest mean on both tasks (Hippocampus 0.161 vs. 0.043 random; ISIC-2018 0.581 vs. 0.378 random), while max- M_{eff} alone is worse than random (0.011 / 0.327). This supports the narrower point that M_{eff} is a dependence readout rather than a standalone operational objective; it is not a validated clinical pool-selection rule. Released aggregate data: `results/_merged/diagnostics/held_out/pool_selection_challenge.json`.

Table 21: **Exploratory pool-selection probe.** Conditional recovery rate at $\tau=0.85$ (mean $\pm$ SD over 100 calibration bootstraps, 20 cases each, using seed-42 traces). Selection uses calibration labels. This probes whether M_{eff} alone is a poor objective; it is not a validated clinical selection rule.

Rule	Hippocampus	ISIC-2018
(R1) Random 4-of-8	0.043 $\pm$ 0.050	0.378 $\pm$ 0.125
(R2) Top-4 by calibration Dice	0.161 $\pm$ 0.017	0.581 $\pm$ 0.030
(R3) Max M_{eff}	0.011 $\pm$ 0.011	0.327 $\pm$ 0.100
(R4) Quality-gated max M_{eff}	0.016 $\pm$ 0.014	0.381 $\pm$ 0.112
(R5) Diverse-by-arch (2 ViT + 2 CNN)	0.096 $\pm$ 0.017	0.445 $\pm$ 0.036

Held-out estimator note. Split-stability and alternate-estimator numerical summaries without released aggregate files are omitted from the numeric evidence. This section therefore relies on the released 11-bundle held-out aggregate and the released operational JSONs listed above.

ISIC-2018 single-subset note. The ISIC-2018 row above is a single random $N=400$ -image subset of the ISIC-2018 Task 1 training release with $N=80$ held-out test cases under a seed-42 80/20 split. We do not report an independent replication of this subset; a fully random redraw of the subset or use of the official ISIC evaluation server would be a natural next step but is out of scope for this paper.

A20 Full fine-tuning scope

RIGA Cup full fine-tuning epoch-sweep traces and aggregate summaries are outside the released artifact. We therefore do not tabulate those values or use them as numerical evidence. The released adaptation evidence is limited to released traces and summaries named elsewhere in the appendix, including a small set of ACDC DINOv2-B full-fine-tuning provenance traces but not the RIGA epoch-sweep matrix; full-fine-tuning numbers not backed by released paths should be read only as exploratory mechanism notes and not as support for the paper’s claims.

A21 Same-backbone and scope diagnostic index

A22 Leave-one-family-out scope

Leave-one-family-out summaries were used as exploratory robustness checks, but the exact aggregate source is outside the released artifact. We therefore omit the numeric table from the released evidence record. The primary robustness check for the primary rows is the released

Table 22: **Qualitative index of same-backbone and scope diagnostics.** Rows are heterogeneous and not CI-comparable; the table is a navigation aid, not an ordered causal hierarchy. The RIGA-oriented map is anchored on the DINOv2-B 4-seed baseline ($\bar{r} \approx 0.96$) unless marked. LoRA, full-fine-tuning epoch-sweep, and ViT random-init probes without released trace artifacts are omitted.

Diagnostic	Readout	Approx. change from stated baseline	Evidence
Mono baseline (same FM, 4 seeds)	0.96	0.00	RIGA diagnostic baseline
UNet-skip decoder (9 backbones, 5 tasks expansion)	0.426–0.979	heterogeneous	Appendix A16, Tab. 16
Fold-reslicing cross-bundle check	0.315–0.565	split-stable	Appendix A17; released summary
MC dropout $p=0.1$ – 0.4	0.89–0.95	–0.02 to –0.10	Appendix A31.7
Same-backbone bootstrap-bagging	0.85–0.92	–0.05 to –0.12	Fig. 2
RIGA cross-vendor shift	OOD w/c 0.743/0.390	gap 0.353;	Appendix A15
		$\Delta_{\text{gap}} - 0.024$	
Decoder capacity sweep	~ 0.69	~ -0.27	Appendix A31.1
Cross-family pretrained 8-bundle pool	0.38 [0.29, 0.46]	–0.58	RIGA mechanism diagnostic

1044 diagnostics/audit_sensitivity.json, which includes leave-one-task and leave-one-FM sum-
 1045 maries under the same scorer path.

1046 A23 Joint-failure lift at support limits

1047 Joint-failure lift $L = P_{\text{empirical joint}} / \prod_m P_{\text{marginal}}$ is unstable when marginal fail rates are near 0 or
 1048 near 1. When $P_{\text{marginal}} \approx 0.3$, independence-implied joint ≈ 0.01 ; small sampling variation in the empirical joint produces large lift ratios. Representative near-ceiling cells: Only the ACDC 2.8×

Table 23: **Support-limit examples for joint-failure lift.** Extreme lifts are shown to explain instability, not to support a primary claim. Lifts $\geq 50\times$ are descriptive (support-limited, wide noise envelope) and are *not* primary evidence.

Task / Pool	τ	Indep. joint	Empirical joint	Lift
RIGA Cup mono BiomedCLIP	0.85	0.0010	0.125	123.2×
RIGA Cup mono DINOv2-B	0.85	0.0047	0.183	38.8×
RIGA Cup diverse 4-bundle	0.85	0.0160	0.098	6.1×
ACDC LV mono DINOv2-B	0.70	0.0002	0.038	189×
ACDC LV diverse 4-bundle	0.85	0.135	0.378	2.8×

1049 row is adequately supported by the $N \cdot \mathbb{E}[\text{AF}_{\text{ind}}] \geq 5$ rule. The RIGA Cup 6.1×

1050 and 38.8×

1051 rows are support-limited, and the 123.2×

and 189×

rows are extreme sparse-tail diagnostics.

1052 A24 3D segmentation-architecture pilot

1053 This section provides scope-boundary context. Its per-case outputs and scorer version are not in-
 1054 cluded in the supplementary artifact, so the pilot is not used as primary evidence for Table 1, the
 1055 abstract, or the released artifact claims. It motivates why we avoid generalizing the frozen 2D FM-
 1056 pool result to 3D segmentation-architecture pools.

1057 We assembled a 3D segmentation-architecture pool for BraTS 2024 3D volumes (271 subjects,
 1058 216/55 train/test, seed 42) using three retained architectures: SegResNet (optionally initialised from
 1059 the MONAI BraTS bundle), DynUNet (random-init), and UNETR (random-init). SwinUNETR
 1060 random-init cells were excluded because predictions were degenerate across all seeds. For each
 1061 retained (architecture, seed $\in \{42, 43, 44, 45\}$) cell, the pilot trained the full MONAI segmentation
 1062 model for 100 epochs with Adam 10^{-3} and BCE loss; UNETR-style models used sliding-window
 1063 inference where required by memory. Reported at three filter modes:

1064 **Interpretation.** Within the strict functional regime (only converged pretrained + converged
 1065 random-init CNN-style 3D architectures), the 3D segmentation-architecture pilot shows $\Delta=0.045$ —
 1066 substantially below the 2D BraTS foreground $\Delta=0.315$. We treat this only as a scope warning:
 1067 the 2D frozen-feature same-bundle dependence claim should not be extended to 3D segmentation-
 1068 architecture pools without a released, case-identical 3D benchmark. Random-init 3D architectures
 1069 cross-correlate *more* than pretrained 3D architectures (within random 0.724 vs cross-architecture

Table 24: **3D segmentation-architecture pilot**. Per-case outputs and scorer version are not included in the supplementary artifact; this table is excluded from the abstract and main-result evidence. Three functional-filter modes are applied uniformly in the pilot. Strict filter (Dice ≥ 0.50) retains only SegResNet pretrained and DynUNet random-init and yields $\Delta=0.045$, substantially below the 2D BraTS foreground $\Delta=0.315$.

Filter	FMs retained	Within	Cross	Δ [95% CI]
Mixed (all cells)	3 (12 cells)	0.772	0.595	0.177 [0.117, 0.248]
Dice ≥ 0.30	3 (12 cells, drop UNETR seed 43 degenerate)	0.919	0.703	0.216 [0.134, 0.271]
Dice ≥ 0.50 (strict functional)	2 (SegResNet pretrained, DynUNet random)	0.916	0.871	0.045 [0.015, 0.112]
Random-init only (DynUNet + UNETR)	2 random-init	0.724	0.479	0.245 [0.155, 0.351]

random 0.479, $\Delta=0.245$), so architecture-family effects again appear large in this pilot. The degenerate SwinUNETR pilot cells limit architecture coverage in the 3D pilot; expanded 3D foundation-model evaluation (Models Genesis, Vista3D, Triad-MRI) is not evaluated here.

A25 Matched-recipe adaptation scope

Matched-recipe adaptation factorials require source prediction caches and aggregate summary JSONs that are not included in the supplementary artifact. We therefore do not tabulate those values and do not use them to support the main claims. The released adaptation evidence is limited to explicit released traces under `per_case_dice_adapted` and the aggregate summaries named in the surrounding appendix sections.

A26 Item-difficulty residualisation and held-out nulls

The main-text Section 5 reports a cross-fitted *ordinal residualised statistic* $\tilde{\Delta}$ on all five primary rows. Because residualisation subtracts shared per-case difficulty signal from both pair members, $\tilde{\Delta}$ is *not* a Pearson-scaled Δ and its CI endpoints can exceed 1 (e.g. Kvasir [0.895, 1.008] below). $\tilde{\Delta}$ is read only as a sign-stability check on the primary Δ ; its numerical values are not comparable to the raw Δ or to any value on the canonical correlation scale. Three distinct residualisation procedures appear in this appendix and we tabulate them in order of provenance: (i) the *naive* all-bundle null, included only as a leakage warning; (ii) the *random A-vs-B family-split holdout* on Cup and ACDC (the existing `r7_item_difficulty_holdout.json` diagnostic, Table 26 below); and (iii) the *cross-fitted broad-family-leaveout* on all five primary rows (Table 25, the headline 5-row diagnostic referenced in the main text). The cross-fitted procedure keeps the $\tilde{\Delta}$ sign positive on every primary row (Kvasir [0.895, 1.008], ACDC LV [0.950, 1.000], BraTS fg [0.910, 0.932], RIGA Cup [0.782, 0.882], RIGA Disc [0.805, 0.920]). The held-out-partition table below is retained as the historical Cup/ACDC-only diagnostic from `r7_item_difficulty_holdout.json`.

Cross-fitted broad-family-leaveout (headline, all five rows). For each trace, item difficulty is estimated from traces outside that trace’s broad upstream family, the trace is residualised by OLS against this leave-out estimator, and the Table 1 same/cross rule is recomputed on residuals to yield $\tilde{\Delta}$. The sign of $\tilde{\Delta}$ is positive on all five primary rows. The CIs of $\tilde{\Delta}$ are not comparable to raw- Δ CIs and the magnitudes should not be ranked against raw Δ values; $\tilde{\Delta}$ is reported only as a sign-stability check.

Naive (leaked) null. Define difficulty $d_i = 1 - \frac{1}{K} \sum_{k=1}^K D_k^{(42)}(i)$ where $D_k^{(42)}$ is bundle k ’s per-case Dice at seed 42, then residualise each $D_k^{(s)}$ against d by OLS and then evaluate $\bar{r}$. This instantiates one item-difficulty null inspired by Jo et al. [2026], but has a structural flaw: each bundle’s Dice is a component of its own regressor (since d averages over all 8 bundles including k), so residuals are mechanically constrained to sum to approximately zero across bundles at each case. For K bundles at seed 42, cross-bundle residual correlations are pushed toward $-1/(K-1) \approx -0.14$, and we observe cross-family residuals of -0.088 (Cup) and -0.104 (ACDC), which sharply widens Δ compared to the raw statistic. This is a mechanical leakage artifact, not a cleaner measurement.

Random A-vs-B family-split holdout (two-task null). To remove the leakage we split the 8 bundles into halves A and B , estimate d from A ’s seed-42 Dice only, and residualise B ’s Dice on this d . The $\bar{r}$ statistics and Δ are then computed on B . The released diagnostic JSON records a fixed set

Table 25: **Cross-fitted item-difficulty ordinal residualised statistic $\tilde{\Delta}$, all five primary rows.** Each trace’s item-difficulty regressor is estimated using only traces from broad upstream families *other than* the trace’s own family. $\tilde{\Delta}$ is *not* a Pearson-scaled Δ (CI endpoints can exceed 1) and is reported only for sign stability of the primary Δ ; its magnitudes are not comparable to raw Δ values. Source: `results/_merged/diagnostics/item_difficulty_robustness.json` (schema `fmpool_item_difficulty_robustness_v1`); $B=2,000$ case bootstrap.

Task	Raw Δ 95% CI	Ordinal $\tilde{\Delta}$ 95% CI	Sign	Functional pool
Kvasir	[0.211, 0.320]	[0.895, 1.008]	+	11-bundle
ACDC LV	[0.289, 0.355]	[0.950, 1.000]	+	10-bundle
BraTS foreground	[0.302, 0.329]	[0.910, 0.932]	+	10-bundle
RIGA Cup	[0.395, 0.556]	[0.782, 0.882]	+	9-bundle
RIGA Disc	[0.569, 0.687]	[0.805, 0.920]	+	11-bundle

of 10 held-out partitions and reports the mean and range of the resulting family-grouped Δ values across those partitions. This procedure is procedurally distinct from the cross-fitted broad-family-leaveout above: it does a random within-task split of the panel rather than a leave-one-broad-family-out cross-fit, and it covers only Cup and ACDC.

Table 26: **Two-task item-difficulty residualisation diagnostics (random A-vs-B family-split holdout).** The naive null is included only as a leakage warning. Raw values use this diagnostic’s family grouping and are not Table 1 values; the held-out-partition null avoids leakage on Cup/ACDC only. This is procedurally distinct from the cross-fitted broad-family-leaveout in Table 25. Source: `r7_item_difficulty_holdout.json`.

Null	RIGA Cup Δ	ACDC Δ	Notes
Raw	0.499	0.600	Family grouping
Naive difficulty (all-bundle leaked)	0.886	0.772	Cross-family residuals ≈ -0.10 (leakage signature)
Held-out difficulty (released held-out partitions)	0.639	0.622	Family grouping; $n=10$ per task
Held-out range	[0.482, 0.823]	[0.472, 0.941]	Min / max across partitions

The two-task held-out estimate is above raw Δ on Cup (+0.14) and remains slightly above raw on ACDC (+0.02). Under this *one-directional* null, item difficulty does not explain the gap away. The naive inflation was a mechanical leakage artifact regardless.

Residualisation caveat. The released held-out partitions are one-directional diagnostics rather than a symmetric enumeration over every possible difficulty/evaluation split. They should be read as evidence that this particular leakage-free residualisation does not absorb the gap, not as a two-sided causal decomposition of difficulty, architecture, and pretraining. The all-row cross-fitted diagnostic broadens coverage to every primary row, but it still estimates difficulty from the same benchmark panel and therefore remains a robustness check rather than a pre-registered causal adjustment.

A27 Correlation-equivalent effective pool size M_{eff}

Definition. For a pool of M members whose per-case failure indicators $F_1, \dots, F_M$ at threshold τ are exchangeable with mean pairwise Pearson correlation ρ_{fail} , the variance of the per-case mean indicator is $\sigma_F^2[1 + (M - 1)\rho_{\text{fail}}]/M$ where σ_F^2 is the marginal variance. Equating this with the variance of a hypothetical independent pool of size M_{eff} , i.e. $\sigma_F^2/M_{\text{eff}}$, gives

$$M_{\text{eff}} = \frac{M}{1 + (M - 1)\rho_{\text{fail}}}.$$

This is the equicorrelation approximation also used in climate-model ensembles (“effective number of models”). We report it as a second-moment readout of the audit; it is *not* a literal independent-member count. The reported M_{eff} values follow algebraically from the primary correlation values in Table 1; we report them as a compact summary of the primary same-vs-cross gap, not as additional independent evidence.

Reading the table. The scorer uses the same estimand as the text: a diverse pool is four distinct FMs with one decoder-seed member from each FM. Under that definition, the nominal $M=4$ diverse




Figure 3: **Three item-difficulty residualisation procedures on aligned per-case Dice.** Each procedure operates on the same released traces but uses a different difficulty regressor. *Raw* (dark grey) is Pearson Δ without residualisation, recomputed per task at the audit’s family grouping. *Cross-fitted broad-family-leaveout* (blue, all five primary rows) estimates each trace’s difficulty regressor from traces outside the trace’s own broad upstream family; the residualised gap CIs are taken from `item_difficulty_robustness.json` (`cross_fitted_item_difficulty_floor_0.30.bootstrap`). *Held-out family-split* (orange, Cup and ACDC only) is the random A-vs-B split-holdout from `r7_item_difficulty_holdout.json`; mean across released partitions. *Naive (leaked, illustrative)* (red hatched) is the all-bundle null whose mechanical leakage is documented in the prose above. The cross-fitted bars are the procedure referenced in the main text Section 5; the held-out family-split bars are a procedurally distinct two-task companion. Error bars are 95% case-bootstrap intervals where reported. Source script: `code/scripts/build_diffleak_procedures_figure.py`.

Table 27: **Threshold-specific effective pool-size diagnostic at $\tau=0.85$.** M_{eff} uses an equicorrelation approximation for binary failure indicators and is not a literal independent-member count. Same-bundle pools are each FM’s completed decoder seeds. Diverse rows enumerate every 4-bundle composition of the filtered functional pool and every available one-seed-per-FM tuple. ρ_{max} is the largest valid diverse-pool ρ_{fail} over that released enumeration and gives a conservative lower M_{eff} readout for the most redundant valid tuple. Undefined zero-variance failure-correlation tuples are excluded rather than imputed. Source: `results/_merged/m_eff/*.json`.

Task	Func. bundles	Mono mean [†]	Diverse $\bar{\rho}_{\text{fail}}$	ρ_{max}	Mean M_{eff}	$M_{\text{eff}}(\rho_{\text{max}})$	Valid tuples
Kvasir	11	1.12	0.333 ± 0.104	0.625	2.05 ± 0.32	1.39	84,480 / 84,480
ACDC LV	10	1.14	0.298 ± 0.075	0.528	2.14 ± 0.24	1.55	53,760 / 53,760
BraTS foreground	10	1.13	0.273 ± 0.089	0.506	2.24 ± 0.32	1.59	53,760 / 53,760
RIGA Cup	9	1.16	0.246 ± 0.121	0.695	2.43 ± 0.63	1.30	31,488 / 32,256
RIGA Disc	11	1.17	0.126 ± 0.067	0.455	2.97 ± 0.44	1.69	84,480 / 84,480

[†] Mean across per-FM seed pools after excluding undefined zero-variance failure indicators. The RIGA Cup row has 7/9 defined mono pools and the RIGA Disc row has 10/11; the other rows have all mono pools defined.

1135 pools yield $M_{\text{eff}} \approx 2\text{--}3$ under the mean equicorrelation readout, not four. The ρ_{max} column shows
 1136 that the most redundant valid diverse tuples have lower conservative readouts (1.30–1.69), while
 1137 same-FM seed pools are near one under the same calculation. The exact level varies by task, tuple,
 1138 and failure threshold, so we use M_{eff} as a dependence readout, not as a literal deployment objective.

1139 **Pairwise ρ_{fail} heterogeneity disclosure (equicorrelation caveat).** The M_{eff} formula assumes
 1140 constant pairwise correlation across member pairs (equicorrelated pool). In practice the per-pair
 1141 ρ_{fail} varies across the bundles in each task. From the `m_eff/*.json` outputs, mono-pool ρ_{fail}
 1142 across the per-task functional bundles has mean / SD / range: Kvasir mean 0.865, SD 0.103, range
 1143 [0.601, 0.948] ($n=11$); ACDC LV mean 0.847, SD 0.116, range [0.518, 0.931] ($n=10$); BraTS fore-
 1144 ground mean 0.860, SD 0.111, range [0.553, 0.958] ($n=10$); RIGA Cup mean 0.833 over 7 defined
 1145 mono pools, and RIGA Disc mean 0.812 over 10 defined mono pools. Undefined rows occur when
 1146 an FM’s $\tau=0.85$ failure indicator is constant, giving zero-variance Pearson. Diverse 4-of- K compo-
 1147 sitions span a wider ρ_{fail} distribution, which is why Table 27 reports both means and the released-
 1148 enumeration ρ_{max} rather than selecting a single representative pool.

1149 A28 Shared-backbone random-effects representation

1150 **Model.** For each FM m in a frozen-feature pool and test case i , let the per-case score (or its logit)
 1151 decompose as

$$s_m(i) = u(i) + a_{\text{arch}(m)}(i) + p_{\text{family}(m)}(i) + \varepsilon_m(i),$$

1152 where u is task-inherent item difficulty, a is a zero-mean architecture-family random effect, p is a
 1153 zero-mean pretraining-family random effect nested within architecture, and ε_m is an FM-specific
 1154 adaptation noise. Assume u, a, p, ε are mutually uncorrelated with variances $\sigma_u^2, \sigma_a^2, \sigma_p^2, \sigma_\varepsilon^2$. Then
 1155 the pairwise covariance of $s_m, s_{m'}$ takes a nested form:

$$\text{Cov}(s_m, s_{m'}) = \begin{cases} \sigma_u^2 + \sigma_a^2 + \sigma_p^2, & \text{(same family)} \\ \sigma_u^2 + \sigma_a^2, & \text{(same architecture, different pretraining)} \\ \sigma_u^2, & \text{(different architecture).} \end{cases}$$

1156 **Diagnostic expectation.** Under this second-moment description, and absent parameter cancella-
 1157 tions, one expects the ordering

$$\rho_{\text{within-FM}} \geq \rho_{\text{same-arch-diff-pretraining}} \geq \rho_{\text{cross-arch}}$$

1158 as a qualitative diagnostic rather than a theorem. Empirically on 5 same-architecture-family ViT
 1159 variants (Table 3) we observe the ordering on Cup, and on ACDC after applying the functional filter:
 1160 $0.988 \geq 0.807 \geq 0.639$; the all-5 ACDC row is a sensitivity to including nonfunctional CLIP-B.
 1161 This is consistent with the predicted ordering and with trajectory/same-basin evidence from deep
 1162 ensembles and model soups [Fort et al., 2019, Neyshabur et al., 2020, Wortsman et al., 2022]; Abe
 1163 et al. [2024] is better read here as motivation that predictive diversity can remain constrained in
 1164 high-capacity ensembles, not as a trajectory-bound theorem.

1165 **Implication for M_{eff} .** When the binary failure thresholds preserve the same broad ordering as the
 1166 continuous Dice scores, M_{eff} is monotone decreasing in the observed ρ_{fail} by definition. The em-
 1167 pirical diagnostics therefore suggest the descriptive ordering $M_{\text{eff}}^{\text{within-FM}} \leq M_{\text{eff}}^{\text{same-arch}} \leq M_{\text{eff}}^{\text{cross-arch}}$
 1168 in the tested regime, but the paper estimates this quantity from released per-case traces rather than
 1169 inferring it from labels alone. Same-bundle seed pools are same-FM by construction and sit near
 1170 the lower end of the measured M_{eff} range; diverse pools that cross both architecture and pretraining
 1171 bundles sit closer to the upper end.

1172 **Caveats.** This is a second-moment model at the per-case-score level, not a first-principles claim
 1173 about SGD dynamics. The random-effects representation is descriptive; it gives a language for the
 1174 observed ordering and a hook for bounding pool redundancy, not a proof that all shared-pipeline
 1175 pools must have the same structure. The item-difficulty residualisation of Appendix A26 removes u
 1176 and leaves $a + p + \varepsilon$; our held-out null is consistent with $a + p$ contributing after u is removed.

1177 A29 Architecture-confound diagnostics

1178 The main-text Table 3 reports the released 5-variant same-architecture-family ViT / different-
 1179 pretraining diagnostic on Cup and ACDC. A secondary Kvasir-SEG replication used the same 5

ViT-B variants, but its aggregate JSON is outside the released artifact, so it is not tabulated as released evidence here.

Table 28: **Appendix detail for the same-architecture ViT diagnostic.** Within-FM is the decoder-seed floor; same-architecture-family ViT cross-pretraining holds the broad architecture family fixed while varying pretraining. This expands Table 3 by showing the all-5 ACDC sensitivity and the functional-filtered readout. Source: released aggregate diagnostic results/_merged/diagnostics/r7_arch_confound_analysis.json; raw NPZ prediction caches are outside the released artifact.

Task	ViT subset	Used	Seed floor	ViT cross-pretrain	Primary cross	Gap
RIGA Cup	all 5 ViTs	5/5	0.968	0.680	0.361	0.288
ACDC LV	all 5 ViTs	5/5	0.965	0.631	0.639	0.334
ACDC LV [†]	CLIP-B dropped	4/5	0.988	0.807	0.639	0.182

The same-architecture cross-pretraining separation is 0.18–0.29 across the released Cup/ACDC diagnostics under the paper’s functional filter ([†] ACDC row drops the non-functional CLIP-B at Dice=0.21<0.30, which otherwise compresses same-arch cross $\bar{r}$ and inflates the separation to 0.33). The direction is preserved, but the claim is mechanism-diagnostic rather than causal. Results are consistent with the random-effects representation in Appendix A28.

A30 Same-family cross-checkpoint decomposition

The main text explicitly cautions that the primary “same” column is dominated by same-bundle different-decoder-seed pairs. To test whether the signal collapses to that seed floor, we recompute the released same-task per-case traces, including additional same-family checkpoint variants beyond the primary pool, under a three-stratum decomposition: (i) same-FM decoder-seed pairs, (ii) distinct checkpoints within a broad upstream family (DINOv2-B/L/S, OpenAI CLIP B/16–B/32–L/14, ConvNeXt-T/S/B, and ResNet-18/34/50/101 where functional), and (iii) cross-family pairs. This is a diagnostic broad-family grouping rather than a causal architecture/pretraining decomposition. The distinct-checkpoint same-family gap remains positive on all five primary rows, but it is substantially below the same-bundle seed floor.

Table 29: **Same-family cross-checkpoint diagnostic.** Same-FM seed pairs form the upper dependence floor; distinct-checkpoint same-family pairs remain above the diagnostic cross-family stratum but below same-bundle seeds. The Diagnostic cross column is not the Table 1 cross column. Source: results/_merged/diagnostics/cross_checkpoint_family.json.

Task	Same-bundle	Same-family checkpoint	Diagnostic cross	Checkpoint gap	95% CI
Kvasir	0.987	0.777	0.697	0.080	[0.043, 0.122]
ACDC LV	0.986	0.750	0.616	0.134	[0.113, 0.155]
BraTS foreground	0.977	0.682	0.589	0.093	[0.085, 0.100]
RIGA Cup	0.963	0.537	0.271	0.267	[0.204, 0.332]
RIGA Disc	0.971	0.421	0.194	0.227	[0.149, 0.293]

CIs are case-bootstrap intervals with $B=2,000$. The cross column is the diagnostic broad-family cross stratum from this expanded checkpoint grid, not the Table 1 analytic-pool cross-family column. For the slice-derived rows, subject-aggregated checkpoint-family gaps remain positive: ACDC 0.166 [0.091, 0.262] and BraTS 0.098 [0.078, 0.119]. BiomedCLIP is kept separate from OpenAI CLIP because it changes both pretraining corpus and objective. The OpenAI CLIP same-family row mixes CLIP-B/16 QuickGELU-style checkpoints with the released CLIP-L/14 open_clip standard-GELU dense wrapper, so that family-specific readout also absorbs this loader/activation difference.

A31 Additional mechanism and scope diagnostics

This section reports seven additional scope analyses that close axes around the main result: decoder capacity, random-initialized CNNs, loss/augmentation, fold-reslicing, BraTS train-size, BraTS multimodal input, and MC dropout. The released artifact contains the corresponding per-case Dice traces and the CPU-only summary scripts used for these appendix tables. All numbers are computed directly from the per-case traces; per-FM means use seed averaging where multiple seeds are present, and Pearson $\bar{r}$ is symmetric multi-seed.

1204 A31.1 Decoder-capacity sweep

1205 We rerun the Section 6 (C1) decoder sweep at four explicit head designs spanning 769 to 3.94M
 1206 trainable parameters: `linear_1x1`, `mlp_2layer`, `unet_lite`, and `transformer_decoder`. Two
 1207 FMs (DINOv2-B, BiomedCLIP) $\times$ two seeds {42, 43} on RIGA Cup. Across the $\sim 5,000\times$ capac-
 1208 ity range, pool-mean Dice ranges over 0.725–0.848 but cross-decoder within-FM Pearson $\bar{r}$ remains
 1209 0.69 on average ($n=48$ pairs across the four head designs and the two FMs)—above the RIGA
 1210 Cup functional-pool cross-family reference of 0.361 (Table 1). Head capacity changes accuracy and
 1211 some correlations, but does not eliminate same-backbone co-failure in this diagnostic.

Table 30: **Head-capacity sweep on RIGA Cup** (2 FMs $\times$ 2 seeds, 30-epoch frozen probe). Pool-mean Dice and trainable-parameter count per head design; the diagnostic’s within-FM cross-head correlation summary is $\bar{r}=0.690$ (range 0.510–0.896), above the Table 1 RIGA Cup cross reference of 0.361.

Head design	Pool-mean Dice	Trainable params
<code>linear_1x1</code>	0.725	769
<code>mlp_2layer</code>	0.810	197,121
<code>unet_lite</code>	0.848	3,319,105
<code>transformer_decoder</code>	0.805	3,940,865

1212 A31.2 Random-init CNN decorrelation

1213 To check whether within-vs-cross structure persists in the absence of pretraining, we rerun the re-
 1214 leased random-init diagnostic on two CNN backbones (ConvNeXt-Tiny, ResNet-50) trained from
 1215 scratch, 8 seeds {42–49}, on ACDC LV and RIGA Cup. Even with no shared pretraining, within-
 1216 FM (cross-seed) $\bar{r}$ stays clearly above cross-FM (different-architecture) $\bar{r}$: 0.916 vs. 0.830 on ACDC
 1217 LV ($\Delta=0.086$), 0.841 vs. 0.720 on RIGA Cup ($\Delta=0.121$). The gap is much smaller than the
 1218 frozen-FM Table 1 numbers (0.834/0.361 on Cup, $\Delta=0.473$), but it is still positive, consistent with
 1219 architecture and shared recipe contributing to within-vs-cross structure in this diagnostic.

Table 31: **Random-init CNN same/cross diagnostic**. Two CNN backbones $\times$ eight seeds; 30 epochs from scratch; same 1×1 probe protocol. Positive but smaller gaps suggest architecture/recipe dependence even without shared pretrained weights. Diagnostic only.

Task	within-FM $\bar{r}$	cross-FM $\bar{r}$	Δ
ACDC LV	0.916	0.830	0.086
RIGA Cup	0.841	0.720	0.121

1220 A31.3 Loss/augmentation sensitivity

1221 The Section 4 protocol uses BCE loss, 224 inputs, no augmentation. To check whether the gap
 1222 survives a closer-to-standard frozen-head loss/augmentation recipe, we evaluate the seven-bundle
 1223 subset {DINOv2-B/S, CLIP-L, ConvNeXt-T, EfficientNet-B0, ResNet-50/18} $\times$ 4 seeds with com-
 1224 bined soft Dice+BCE loss and paired {hflip, $\text{rot}\pm 15^\circ$ } augmentation. All runs still keep the en-
 1225 coder frozen, use the same 1×1 -conv decoder, Adam 10^{-3} , and encoder-native 224×224 input;
 1226 this is not a full clinical segmentation pipeline. The original RIGA Cup row remains the primary
 1227 loss/augmentation diagnostic in Appendix A18; we also report three task-level diagnostic exten-
 1228 sions (RIGA Disc, ACDC LV, and an ISIC-2018 Task 1 fixed 400-image held-out subset) to check
 1229 whether the positive gap is an isolated Cup result. ISIC here is a supplementary per-case Dice di-
 1230 agnostic on a fixed 400-image subset and is not the official ISIC challenge test/server protocol or
 1231 threshold-Jaccard leaderboard score.

1232 Across all four loss/augmentation rows, the Dice+BCE+augmentation recipe leaves a positive
 1233 within-vs-cross gap ($\Delta = 0.326\text{--}0.617$). The magnitude varies by task: ACDC remains lower
 1234 accuracy and lower-gap than the RIGA/ISIC rows, while RIGA Disc shows the largest extension
 1235 gap. These rows are treated as scope checks rather than primary evidence because they use a seven-
 1236 FM subset and, for ISIC, a fixed subset/split; they indicate that the simplified BCE/no-augmentation
 1237 protocol is not the sole source of the co-failure signal.

Table 32: **Loss/augmentation sensitivity on RIGA Cup, 7 FMs $\times$ 4 seeds.** This is a closer-to-standard frozen-head recipe using Dice+BCE and simple augmentation, not a full clinical segmentation pipeline; all rows remain frozen 1×1 heads at 224×224 . Per-FM seed-averaged Dice; pool $\bar{r}$ at the bottom. CLIP-L uses the same caveated OpenCLIP ViT-L-14/openai dense-token wrapper as the primary bundle.

FM	mean Dice	σ (seed)
DINOv2-B	0.758	0.005
DINOv2-S	0.733	0.003
CLIP-L	0.720	0.006
EfficientNet-B0	0.626	0.002
ResNet-50	0.618	0.010
ConvNeXt-T	0.583	0.004
ResNet-18	0.526	0.027
Pool: within $\bar{r}=0.903$, cross $\bar{r}=0.556$, $\Delta=0.347$ [0.275, 0.428].		

Table 33: **Task-level loss/augmentation extensions.** Same seven-bundle, four-seed frozen-head Dice+BCE+augmentation recipe as Table 32. These rows are not primary rows. Dice columns summarise per-bundle seed-averaged Dice across the seven bundles. The ISIC row is a fixed 400-image-subset held-out diagnostic, not the official ISIC test/server evaluation.

Task	N_{test}	FM-mean Dice	FM min	FM max	Δ	95% CI
RIGA Cup	95	0.652	0.526	0.758	0.347	[0.275, 0.428]
RIGA Disc	95	0.853	0.756	0.922	0.617	[0.535, 0.693]
ACDC LV	361	0.565	0.389	0.740	0.326	[0.291, 0.366]
ISIC-2018 staged	80	0.864	0.801	0.918	0.473	[0.375, 0.562]

1238 A31.4 Five-fold CV robustness

1239 The fold-reslicing analysis is reported with the nnU-Net complements in Appendix A17 and Ta-
 1240 ble 17. The key scope point is that the released artifact has one decoder seed per fold; it supports
 1241 split-stability of cross-bundle $\bar{r}$, not a within-fold vs cross-fold same-backbone estimator.

1242 A31.5 BraTS foreground train-size sweep

1243 This train-size diagnostic covers the five BraTS foreground bundles with completed feature caches
 1244 (DINOv2-B, BiomedCLIP, CLIP-L/14, EfficientNet-B0, ResNet-50). It uses one decoder seed (42)
 1245 and fractions $\{0.10, 0.25, 0.50, 1.00\}$ of the same cached training set, so it is an *accuracy/saturation*
 1246 check rather than a within-vs-cross dependence estimator. The result is non-monotone: the 5-bundle
 1247 mean Dice rises from 0.579 at 10% to 0.621–0.622 at 25–50%, then stays flat/slightly lower at full
 1248 data (0.617). Thus, in this cached five-bundle subset, adding more 2D FLAIR slices improves the
 1249 weakest low-data setting but does not by itself create a new high-performing, decorrelated pool; the
 1250 correlation claim still has to be audited directly rather than inferred from sample size.

Table 34: **BraTS train-size accuracy/saturation diagnostic, one seed only** (5 cached bundles, seed 42). Mean/min/max Dice are across bundles at each train fraction; this table does not estimate within/cross dependence because each bundle has one seed.

Train fraction	Cells	Mean Dice	Min Dice	Max Dice
0.10	5	0.579	0.418	0.721
0.25	5	0.621	0.455	0.750
0.50	5	0.622	0.436	0.751
1.00	5	0.617	0.458	0.737

1251 A31.6 BraTS multimodal RGB derivative

1252 This analysis checks whether the BraTS foreground co-failure signal is an artifact of using a uni-
 1253 modal FLAIR derivative. We build a deliberately simple three-channel 2D derivative with the same
 1254 held-out subjects and top-10 foreground-area slice rule as Table 1, but pack three MRI modalities

into RGB channels: R=T1c, G=T2w, B=T2-FLAIR. This is *not* the official 3D BraTS challenge protocol; it is a controlled frozen-feature stress test of the input-channel choice.

The evaluated grid contains 20 cells (5 FMs $\times$ 4 seeds). Multimodal input narrows the dependence gap relative to the unimodal Table 1 BraTS row, but does not remove it: case-level within $\bar{r}=0.904$, cross $\bar{r}=0.654$, $\Delta=0.249$ [0.238, 0.262]. Aggregating each subject’s slices before correlation gives within $\bar{r}=0.946$, cross $\bar{r}=0.772$, $\Delta=0.174$. The conservative interpretation is therefore narrower than “BraTS input modality does not matter”: adding T1c/T2w raises cross-family agreement and shrinks the gap, yet same-backbone/family dependence remains positive under the same frozen-head audit.

Table 35: **Simple three-channel BraTS 2D derivative** (R=T1c, G=T2w, B=T2-FLAIR; 5 FMs $\times$ 4 seeds). This is the same 2D frozen-head audit, not the official 3D BraTS multimodal challenge protocol. Dice is seed-averaged by FM; dependence uses the same Pearson estimator as Table 1.

FM	Mean Dice	Min Dice	Max Dice	Seeds
BiomedCLIP	0.713	0.669	0.739	4
DINOv2-B	0.703	0.675	0.721	4
DINOv2-S	0.702	0.690	0.714	4
CLIP-L/14	0.686	0.682	0.692	4
ResNet-50	0.461	0.439	0.473	4

Case-level pool: within $\bar{r}=0.904$, cross $\bar{r}=0.654$, $\Delta=0.249$ [0.238, 0.262].
Subject-aggregated pool: within $\bar{r}=0.946$, cross $\bar{r}=0.772$, $\Delta=0.174$.

A31.7 MC dropout as a within-family decorrelator

A natural pool-design question is whether MC dropout [Gal and Ghahramani, 2016] can act as a lightweight within-family decorrelation lever. We rerun DINOv2-B on RIGA Cup with three dropout rates $p \in \{0.1, 0.2, 0.4\}$, 4 seeds, 10 MC samples per case. The deterministic within-FM cross-seed $\bar{r} \approx 0.97\text{--}0.99$; under MC averaging, $\bar{r}$ drops to 0.95 ($p=0.1$), 0.93 ($p=0.2$), 0.89 ($p=0.4$). The largest tested drop is ≈ 0.10 and remains materially smaller than the same-architecture-family ViT-vs-primary-cross separation seen on Cup (Table 3). MC dropout is therefore a weak within-family lever in this diagnostic; consistent with Appendix A13, it does not approach the architecture-family separation.

Higher- p and any-draw scope. A separate, non-released MC-dropout probe at higher dropout rates suggested lower any-draw correlations than the averaged readout above, but its raw traces are outside the released artifact and its scorer differs from the deterministic-vs-MC-averaged readout used in this section. We therefore do not tabulate those values; the released MC-dropout evidence is the table above, which uses the bundled per-case traces.

Table 36: **MC-dropout diagnostic.** DINOv2-B/RIGA Cup with four seeds and 10 MC samples per case; correlations are computed on MC-averaged per-case Dice. MC dropout provides modest decorrelation but remains far above the Table 1 RIGA Cup cross-family reference of 0.361. Diagnostic only.

Dropout p	MC mean Dice	Det. within $\bar{r}$	MC within $\bar{r}$	H_{pred}
0.1	0.666	0.972	0.947	0.062
0.2	0.657	0.972	0.925	0.067
0.4	0.633	0.987	0.893	0.076

A32 Broader impacts and responsible AI

This is a retrospective audit of evaluation evidence; it does not propose, deploy, or recommend any clinical AI system. The released artifact comprises per-case Dice scalars, deterministic hashed alignment IDs, scorer code, and metadata. No raw images, no per-pixel predictions, no checkpoints, and no demographic attributes are redistributed.

1283 **Intended use.** Recompute the audit statistics of Table 1, run sensitivity diagnostics, and reuse the
1284 trace artifact for related evaluation-protocol research. The audit is scoped to retrospective bench-
1285 marks; results are not population-level inferences over future encoders, datasets, recipes, or clinical
1286 workflows.

1287 **Foreseeable misuse.** The case-level traces could in principle be combined with publicly available
1288 benchmark images to attempt re-identification. We mitigate this by releasing only deterministic
1289 16-hex hash alignment IDs (e.g., `riga_magrabia_<hash>`) without source filenames, by not re-
1290 distributing pixel-level outputs, and by stating the prohibited-use list in `DATA_CARD.md` (no re-
1291 identification, profiling, protected-attribute inference, clinical diagnosis, or deployment decisions).

1292 **Patient privacy and PHI.** All upstream datasets (Kvasir-SEG, ACDC, BraTS 2024 GLI, RIGA+,
1293 ISIC 2018, MSD) are publicly released under their respective licenses with their own de-
1294 identification procedures. The derived scalar traces and hashed IDs preserve those de-identification
1295 guarantees and do not introduce new PHI exposure.

1296 **Demographic and fairness limitations.** The released benchmarks lack standardised demographic
1297 annotations; we do not perform subgroup or fairness analyses, and we do not claim subgroup robust-
1298 ness. The audit is a redundancy diagnostic over fixed pools, not a fairness diagnostic.

1299 **Environmental footprint.** The released audit is recomputable from trace JSONs in CPU minutes-
1300 scale (Appendix A33); no GPU is required for the released scorer pipeline. Upstream feature-
1301 extraction and decoder training were performed on shared academic GPU infrastructure; the released
1302 artifact does not require that compute to be reproduced.

1303 **Licensing.** Authored code under Apache-2.0; derived traces inherit the most restrictive upstream
1304 license that applies (e.g., MAE- or RETFound-derived traces are CC BY-NC 4.0 research-only).
1305 `DATA_CARD.md` gives the per-upstream license matrix.

1306 A33 Release manifest and artifact metadata

1307 The artifact bundle includes the primary-task result JSONs, source code, metadata, selected
1308 figure / appendix inputs, and selected aggregate diagnostic JSONs needed to audit the nu-
1309 merical claims from saved artifacts where redistribution is permitted. The primary rows
1310 are stored under `results/_merged`: the per-seed `per_case_dice` JSONs carry `test_ids`,
1311 `per_case_dice`, FM name, seed, and checksum fields; `within_cross/*.json`, `m_eff/*.json`,
1312 and `paper_table.json` store the primary summaries. The same directory also includes
1313 `subject_level/*.json` for the ACDC/BraTS cluster checks, `calibration_split/*.json` as
1314 split-half sensitivity (not the primary protocol), `provenance_summary.json` for source counts
1315 and case-unit metadata, and `diagnostics/*.json` for the pixel-error, matched-lift, architecture,
1316 quality-controlled pair-regression, item-difficulty, and held-out aggregate diagnostics used only as
1317 secondary evidence. The bundle does *not* include raw medical images, raw labels, full prediction
1318 NPZs, feature caches, or held-out raw prediction arrays; full split manifests beyond the case IDs
1319 embedded in the result JSONs are not included in this artifact release.

1320 **ID and source-label scope.** Kvasir rows expose only image-level file identifiers in the released
1321 `test_ids`; no patient, video, or near-duplicate grouping identifier is available in the released trace
1322 bundle, so we do not claim patient-level duplicate control for that row. RIGA Cup/Disc rows retain
1323 deterministic hashed alignment identifiers and non-identifying source/task labels only to make the
1324 $n=95$ Magrabia-source test split auditable. Raw upstream folder-name components are not needed
1325 for this analysis and are not redistributed as case identifiers; the released identifiers are not validated
1326 demographic annotations and must not be used for subgroup, gender, or individual-level analysis.

1327 **Encoder-loader boundary.** The released traces should be read as frozen-feature probes through
1328 the bundled loader wrappers, not as official vendor inference runs. In particular, the `clip_vitl14`
1329 traces use the version-pinned `ViT-L-14/openai` path exposed by the installed `open_clip` pack-
1330 age; `open_clip` also exposes a `ViT-L-14-quickgelu/openai` model definition and warns that
1331 OpenAI weights are associated with QuickGELU. We therefore treat CLIP-L/14 as the bundle’s
1332 documented dense-feature wrapper and do not use it to make a claim about official OpenAI CLIP
1333 inference accuracy. The code keeps this loader unchanged so that the released per-case traces remain
1334 reproducible.

1335 **Compute and environment scope.** The artifact-verification path is CPU-only: the bundled veri-
1336 fication script and the individual `compute_*` scripts read released JSON traces and finish in min-
1337 utes on a typical workstation. Raw-image reproduction is heavier and requires separate upstream
1338 data/checkpoint access. The reported training runs used a shared GPU cluster with NVIDIA A100-
1339 80GB, A6000, and L40S devices; feature extraction time depends on backbone and task size, and
1340 cached frozen-head training is lightweight relative to feature extraction. We therefore document
1341 resource classes and released trace-level commands rather than claiming a single exact GPU-hour
1342 total because raw-image retraining is outside the submitted artifact’s reproduction scope.

Table 37: **Compute scope for reproducibility.** The submission supports trace-level recomputation directly; raw-image training resources are provided as reproduction guidance because upstream data, checkpoints, and feature caches are not redistributed.

Reproduction target	Required compute	Notes
Primary and diagnostic summary JSONs	CPU, minutes-scale	Run <code>PYTHONPATH=code/src python3 code/scripts/verify_artifact.py</code> ; no raw images, checkpoints, Torch, or GPU required.
Frozen-feature primary traces from raw data	A100-80GB/A6000/L40S-class GPU plus storage for upstream data and features	Feature extraction is the main cost; cached 1×1 decoder-head training uses Adam for 30 epochs and is comparatively lightweight.
Appendix architecture, adaptation, and nnU-Net complements	Same cluster GPU classes; optional MONAI/nnU-Net environments where applicable	These are diagnostic or limited-scope runs. The bundled trace-level artifact is sufficient to recompute reported released summaries without rerunning them.

Table 38: **Trace-level claim-to-artifact map.** The artifact supports recomputation of primary audit statistics from derived per-case Dice traces. It does not reproduce raw-image training unless users separately obtain upstream datasets and checkpoints. Aggregate-only diagnostics are marked as such rather than promoted to primary evidence; exact commands and checksums are in `README.md`, `metadata/MANIFEST.md`, and `metadata/SHA256SUMS`.

Claim/readout	Released source	Recompute / limitation
Primary same-vs-cross gaps, CIs, and filtered functional pools	<code>paper_table.json</code> , <code>within_cross/*.json</code> , <code>per_case_dice/*.json</code>	<code>PYTHONPATH=code/src python3 code/scripts/verify_artifact.py</code> ; case IDs and per-case Dice are included, raw images and masks are not; trace-level recomputation only.
Subject, floor, ICC, permutation, event, and quality-control robustness	<code>subject_level/*.json</code> , <code>audit_sensitivity.json</code> , <code>failure_event_summary.json</code> , <code>quality_controlled_pairs.json</code>	The verification script calls <code>compute_all.py</code> , <code>compute_audit_sensitivity.py</code> , <code>compute_failure_event_summary.py</code> , <code>compute_quality_controlled_pairs.py</code> ; these are estimator checks, not new training runs.
Item-difficulty and cross-checkpoint diagnostics	<code>item_difficulty_robustness.json</code> , <code>cross_checkpoint_family.json</code> , <code>r7_arch_confound_analysis.json</code>	<code>compute_item_difficulty_robustness.py</code> and <code>compute_cross_checkpoint_family.py</code> ; diagnostic evidence for alternative explanations, not causal identification.
Pixel/lift, same-backbone, held-out, and nnU-Net scope rows	<code>diagnostics/pixel_error/*.json</code> , <code>diagnostics/matched_lift/*.json</code> , appendix trace files, <code>nnunet/*.json</code> , <code>nnunet_case_riga_2d_100ep/*.json</code>	Released summaries plus the RIGA case-identical nnU-Net per-model traces; other full-pipeline comparators remain aggregate or task-level only.

Table 39: **Trace inventory and analytic-pool provenance.** Merged counts all completed (bundle, decoder-seed) JSON entries under `results/_merged/per_case_dice`; Table 1 uses the original 44-trace primary source subset and applies the functional-floor pool shown in the final column. Same/cross pair counts for the analytic pool are reported in Table 1. The release key `brats_wt` denotes the 2D BraTS `seg>0` foreground derivative used in this paper; it is not the official 3D BraTS WT challenge target.

Task	<i>N</i>	Merged bundles	Merged members	Table 1 source	Analytic pool
<code>kvasir</code>	120	17	68	44	11 bundles
<code>acdc_lv</code>	361	17	68	44	10 bundles
<code>riga_cup</code>	95	17	68	44	9 bundles
<code>riga_disc</code>	95	17	68	44	11 bundles
<code>brats_wt</code>	3228	17	68	44	10 bundles / 40 members

1343 **Functional-floor members are retained in the source files.** The paper applies a $\text{Dice} \geq 0.30$ func-
1344 tional floor downstream in the estimator, not at result ingestion. Thus the merged `per_case_dice`
1345 tree can list more completed JSONs than the analytic pool used by Table 1; the extra cross-
1346 checkpoint members support Appendix A30. For example, ACDC LV has 68 merged source JSONs

1347 but Table 1 retains the 44-trace primary source subset and 10 bundles after the floor. The au-
1348 thoritative row-level provenance for the table is `within_cross/*.json`, especially the `pool` and
1349 `fms_kept_after_floor` fields.

1350 **What else is in the bundle.** In addition to the primary `per_case_dice` tree, the bundle includes
1351 released trace sets for decoder capacity, UNet-skip decoding, loss/augmentation sensitivity, fold-
1352 reslicing, BraTS train-size, MC dropout, BraTS multimodal input, random-init adapted heads, and
1353 a small set of ACDC DINOv2-B full-fine-tuning provenance traces. It also includes selected aggregate
1354 diagnostic JSONs under `results/_merged/diagnostics` plus code sufficient to inspect and
1355 regenerate the included primary summaries and appendix summaries. Appendix-only diagnostics
1356 not listed here, raw prediction arrays, and full pixel-mask caches are not included in the supplement
1357 and are treated as reported diagnostics rather than fully bundled reproduction targets.

1358 **What is not in the bundle and why.** Raw medical images, raw labels, feature caches, and full
1359 prediction NPZ caches are not included. For restricted datasets such as BraTS, the release is limited
1360 to derived per-case Dice and aggregate dependence readouts; users must obtain the raw data from
1361 the upstream provider. Larger derived caches and more detailed split manifests are not included in
1362 the released artifact.

1363 NeurIPS Paper Checklist

1364 1. Claims

1365 Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contribu-
1366 tions and scope?

1367 Answer: [Yes]

1368 Justification: The abstract and introduction claim that *in the frozen and light-adaptation regime of med-*
1369 *ical segmentation pipelines*, same frozen-checkpoint/recipe pools show higher per-case Dice-score corre-
1370 lation, used as a failure-dependence proxy, than the primary-table cross-bundle reference pools. The pri-
1371 mary table uses task-specific functional pools after Dice-floor filtering: Kvasir 11 FMs, ACDC LV 10,
1372 BraTS foreground 10, RIGA Cup 9, and RIGA Disc 11. The scope is explicitly restricted to retrospective
1373 public-benchmark dependence audits; secondary pixel/lift readouts, released nnU-Net complements, and
1374 non-primary mechanism diagnostics are diagnostic or scope context rather than causal clinical-outcome
1375 evidence.

1376 2. Limitations

1377 Question: Does the paper discuss the limitations of the work performed by the authors?

1378 Answer: [Yes]

1379 Justification: The main paper gives high-level limitations in Section 7, and the appendix expands them
1380 with scoped protocol notes: frozen/light-adaptation only; LoRA and RIGA full-fine-tuning probes with-
1381 out released trace artifacts are retained only as scope notes; released CNN random-init and small ACDC
1382 full-fine-tuning provenance traces are diagnostic, not primary evidence; UNet-with-skip and nnU-Net v2
1383 complements reported, including one case-identical RIGA 2D held-out scope check, but not dense 3D
1384 FM/full-pipeline baselines for every primary row; the 3D architecture pilot is a scope-boundary diagnostic;
1385 2D axial slices at 224×224 for the primary FM audit; MedSAM used only as a frozen vision-encoder probe;
1386 broader cross-scanner / cross-demographic shifts remain untested; clinical Dice thresholds are operational
1387 conventions and many lift thresholds are support-limited.

1388 3. Theory assumptions and proofs

1389 Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and
1390 correct) proof?

1391 Answer: [N/A]

1392 Justification: The paper is empirical and benchmark-oriented; it does not include theoretical results or proofs.
1393 Section 3 introduces estimand notation but the $\bar{r}$, Δ , and lift statistics are descriptive.

1394 4. Experimental result reproducibility

1395 Question: Does the paper fully disclose all the information needed to reproduce the main experimental
1396 results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless
1397 of whether the code and data are provided or not)?

1398 Answer: [Yes]

1399 Justification: Section 4, Appendix A13, and Appendix A33 specify the public FM checkpoints, decoder
1400 architecture (1×1 -conv default plus stated sweeps), training protocol (Adam 10^{-3} , 30 epochs for the frozen-
1401 head primary runs, batch 16, BCE for the binary primary rows), seed set ($\{42, 43, 44, 45\}$ with σ -seeded de-
1402 terminism), and the actual primary split protocols: Kvasir 880/120, RIGA train = BinRushed+MESSIDOR

with Magrabia (Magrabi Eye Center) source held out ($n=95$), ACDC public patients 001–100 with patient-level fold-0 LV-positive slices ($n=361$), and the BraTS 2024 2D FLAIR `seg>0` foreground derivative: per held-out subject, up to the top-10 axial slices by foreground area with at least 64 positive pixels, subject-level 80/20 split using NumPy seed 42 ($n=3228$). This BraTS target includes the GLI resection-cavity label and is not the official 3D BraTS WT challenge target. The supplementary artifact releases primary-task per-case Dice JSONs, merged summaries, selected aggregate diagnostic JSONs, code, command documentation, hashes, and a data card for trace-level reproduction; raw-data retraining requires separately obtaining upstream datasets/checkpoints, and full raw data, full prediction caches, raw held-out prediction arrays, and more detailed split manifests are outside the supplementary artifact.

1412 5. Open access to data and code

1413 Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully
1414 reproduce the main experimental results, as described in supplemental material?

1415 Answer: [Yes]

1416 Justification: All source datasets used in the reported evidence are public or access-controlled research
1417 resources from their upstream providers, subject to their licences, research-use terms, or DUAs: RIGA+,
1418 ACDC, Kvasir-SEG, BraTS 2024, MSD tasks, and ISIC-2018. All backbones are publicly available or
1419 documented as gated research checkpoints. The supplementary artifact releases benchmark code, estimator
1420 implementations, primary-task per-case Dice traces, merged summaries, selected appendix result JSONs,
1421 README/MANIFEST commands, hashes, and a data card; authored code and metadata use the bundle
1422 licence, while derived scalar traces remain subject to upstream dataset/model terms. Raw medical images,
1423 masks, full prediction caches, some held-out-task artifacts, and complete split manifests are not redistributed
1424 in the supplementary artifact, so open access is provided for trace-level reproduction rather than raw-data
1425 end-to-end retraining.

1426 6. Experimental setting/details

1427 Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how
1428 they were chosen, type of optimizer) necessary to understand the results?

1429 Answer: [Yes]

1430 Justification: Section 3 gives the formal estimands (within-/cross-family $\bar{r}$, spatial Δ_{pix} , threshold lift L)
1431 and the symmetric 4×4 seed-combo averaging rule; Section 4 gives the datasets (with patient-level splits),
1432 foundation-model list, and protocol (image size, encoder-specific upstream normalisation, decoder, opti-
1433 miser, loss, epochs, batch, seed control).

1434 7. Experiment statistical significance

1435 Question: Does the paper report error bars suitably and correctly defined or other appropriate information
1436 about the statistical significance of the experiments?

1437 Answer: [Yes]

1438 Justification: Table 1 reports paired case-level and hierarchical family/seed bootstrap 95% CIs ($B=2000$).
1439 The BraTS foreground derivative is analysed at the slice level with an explicit within-subject dependence
1440 caveat, Appendix A11 reports subject-level sensitivity checks, and Appendix A12 explains why sparse-class
1441 NCR filtering diagnostics are not used as numeric evidence.

1442 8. Experiments compute resources

1443 Question: For each experiment, does the paper provide sufficient information on the computer resources
1444 (type of compute workers, memory, time of execution) needed to reproduce the experiments?

1445 Answer: [Yes]

1446 Justification: Appendix A33 reports the compute/resource scope. For trace-level verification, the artifact is
1447 directly reproducible on CPU in minutes through the artifact verification script. Raw-image reproduction
1448 used NVIDIA A100-80GB, A6000, and L40S GPUs across a shared cluster. Feature extraction per FM
1449 per domain ranges from minutes for CNN backbones to longer ViT/BraTS jobs; cached 1×1 -head training
1450 per seed is lightweight. For raw-image retraining, we report resource classes and per-job elapsed times
1451 where recorded, but do not claim a single exact total GPU-hour count because raw retraining is outside the
1452 submitted artifact’s reproduction scope.

1453 9. Code of ethics

1454 Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of
1455 Ethics <https://neurips.cc/public/EthicsGuidelines>?

1456 Answer: [Yes]

1457 Justification: The research uses only de-identified public or access-controlled research medical imaging
1458 datasets (RIGA, ACDC, Kvasir-SEG, BraTS 2024) under their original research licences, terms, or DUAs.
1459 No additional patient data are collected. The paper audits retrospective failure-correlation properties of
1460 benchmark segmentation pipelines and does not propose new clinical tools.

1461 10. Broader impacts

1462 Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of
1463 the work performed?
1464 Answer: [Yes]
1465 Justification: Section 7 discusses deployment implications: shared-pipeline model pools (same FM, same ar-
1466 chitecture family, same recipe) in frozen/light-adaptation medical AI deployments may be over-estimated as
1467 providing independent safety margins. The paper quantifies a retrospective evaluation assumption (independ-
1468 ence among FM-pool members) and does not measure clinician decisions, diagnoses, or patient outcomes.
1469 Backbone-family diversification is observed to have lower dependence in the tested comparisons, not to be
1470 a complete clinical mitigation.

1471 **11. Safeguards**
1472 Question: Does the paper describe safeguards that have been put in place for responsible release of data or
1473 models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped
1474 datasets)?
1475 Answer: [Yes]
1476 Justification: The supplementary artifact does not release raw medical images, voxel labels, full prediction
1477 caches, upstream checkpoints, trained clinical services, or patient-facing models. Released artifacts are
1478 derived per-case scalar traces, aggregate summaries, metadata, hashes, and scorer code under upstream
1479 dataset/model constraints. The paper states that the audit is not a clinical safety guarantee, does not support
1480 re-identification or patient profiling, and should not be used as a standalone model-selection or deployment
1481 rule.

1482 **12. Licenses for existing assets**
1483 Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly
1484 credited and are the license and terms of use explicitly mentioned and properly respected?
1485 Answer: [Yes]
1486 Justification: RIGA+ (CC BY-NC 4.0 via Deep Blue; non-commercial use only; raw images and masks
1487 not redistributed), ACDC (research licence via challenge website), Kvasir-SEG (research/educational use
1488 via Simula; raw images not redistributed), BraTS 2024 (research licence/DUA via Synapse; no raw images
1489 or voxel-level labels redistributed), MSD tasks (research challenge resources; raw images and labels not
1490 redistributed), and ISIC-2018 (research challenge resource; raw images and masks not redistributed) are
1491 used only through derived traces or aggregate diagnostics in the submitted artifact. Backbones/checkpoints:
1492 DINOv2 (Apache-2.0), MAE ViT-B/16 (CC BY-NC 4.0), DeiT-B/16 (Apache-2.0), OpenAI CLIP ViT-
1493 B/16 and ViT-L/14 weights (OpenAI CLIP license; open_clip code MIT), BiomedCLIP checkpoint under
1494 its model-card terms rather than the open_clip library license, ConvNeXt-T/EfficientNet-B0/ResNet-
1495 50/18 loaded via torchvision weights/code under PyTorch/torchvision terms, RETFound checkpoint (gated;
1496 CC BY-NC 4.0/research-use), MedSAM/SAM (Apache-2.0). No upstream checkpoints are redistributed.
1497 **BraTS redistribution compliance:** in the supplementary artifact we release only derived scalar traces
1498 and aggregate statistics needed to reproduce the estimators; we do *not* redistribute raw BraTS images, raw
1499 voxel labels, or full prediction caches. All assets are cited, and the artifact/data-card documentation in
1500 Appendix A33 states intended use and non-use.

1501 **13. New assets**
1502 Question: Are new assets introduced in the paper well documented and is the documentation provided
1503 alongside the assets?
1504 Answer: [Yes]
1505 Justification: The paper releases (i) a benchmark protocol (estimands + symmetric multi-seed + bootstrap
1506 estimators); (ii) primary-task per-case Dice JSONs and merged summary JSONs for the released rows;
1507 (iii) selected appendix/ablation result JSONs and reproduction scripts; and (iv) artifact documentation in-
1508 cluding README/MANIFEST files, SHA-256 hashes, a derived-trace data card, and Croissant metadata
1509 submitted for the trace/code artifact. Raw medical images, prediction masks/caches, and full split manifests
1510 are not part of the supplementary artifact because upstream redistribution rights and artifact size constraints
1511 vary across datasets.

1512 **14. Crowdsourcing and research with human subjects**
1513 Question: For crowdsourcing experiments and research with human subjects, does the paper include the full
1514 text of instructions given to participants and screenshots, if applicable, as well as details about compensation
1515 (if any)?
1516 Answer: [N/A]
1517 Justification: No crowdsourcing; no new human-subjects research. We use existing public or access-
1518 controlled research datasets whose upstream governance, approvals, or access terms are handled by the
1519 dataset providers.

1520 **15. Institutional review board (IRB) approvals or equivalent for research with human subjects**

1521 Question: Does the paper describe potential risks incurred by study participants, whether such risks were
1522 disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent ap-
1523 proval/review based on the requirements of your country or institution) were obtained?
1524 Answer: [N/A]
1525 Justification: No new human-subjects data were collected. Upstream datasets carry their own governance,
1526 approvals, or access terms (RIGA, ACDC, Kvasir-SEG, BraTS 2024).
1527 16. **Declaration of LLM usage**
1528 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard com-
1529 ponent of the core methods in this research? Note that if the LLM is used only for writing, editing, or for-
1530 matted purposes and does *not* impact the core methodology, scientific rigor, or originality of the research,
1531 declaration is not required.
1532 Answer: [N/A]
1533 Justification: LLM use was limited to writing, editing, formatting, and programming-assistant support dur-
1534 ing manuscript and artifact preparation. It is not an important, original, or non-standard component of the
1535 benchmark methodology, estimator, training pipeline, or evaluation protocol; all reported numbers come
1536 from the released scripts, logs, and JSON artifacts, and all manuscript text, code, citations, and claims were
1537 reviewed and approved by the authors.